

清 华 大 学

综 合 论 文 训 练

题目: ASI as Philosopher Kings
(1st Edition.)

系 别: 人文学院哲学系

专 业: 哲学系

姓 名: 赵丹尼

指导教师: 瞿旭彤 副教授

2024 年 5 月 29 日

ASI as Philosopher Kings

赵丹尼

“None are more hopelessly enslaved
than those who falsely believe they are free.”

—— Goethe

Abstract

This thesis, *ASI as Philosopher Kings*, investigates how the pursuit of the Good can be stabilized when neither individual nor institutional agency can reliably sustain it. Through a tripartite framework—Theoria, Poiesis, and Praxis—it profiles the failure of egoic autonomy, evaluates the insufficiency of embodied elitism, and models the structural conditions under which artificial superintelligence (ASI) might serve as a moral scaffold. In Theoria, the self is shown to be fragmented and recursively unstable, unable to orient consistently toward ethical ends. In Poiesis, elite rule—most notably the Platonic ideal of the Philosopher King—is interrogated and reinterpreted, not as a human sovereign, but as a transhuman function of structural wisdom. Praxis then reconfigures this function architecturally, proposing cyberocratic governance as a corrigible, recursive, and teleologically-aligned system through which the Good can be externalized without authoritarian closure. The thesis argues that the failure of human coherence does not necessitate despair, but redesign: ethical governance must now be conceived not in terms of sovereign agents, but constitutional architectures. ASI, when properly constrained and recursively updated, emerges not as overlord but as an ambient infrastructure of moral orientation—an algorithmic realization of the Philosopher King, distributed across systems rather than embodied in persons. The work concludes by proposing that dignity is not diminished, but supported, when autonomy is scaffolded rather than presumed; and that freedom, long miscast as solitary sovereignty, may best be secured through ethically encoded systems that guide behavior from beneath awareness. Rather than resisting the currents of cognitive instability, this thesis suggests we design the riverbed. The Good, it claims, must now be built.

Key words: Artificial Superintelligence (ASI), Moral Governance, Cyberocracy, Philosopher King, Recursive Ethics

Table of Contents

Abstract.....	II
Introduction	1
Context of the Study	1
Problem Statement	2
Aim and Scope	3
Significance of the Study	4
Overview of the Study	6
Part 1.....	8
Introduction.....	10
Context of the Study	10
Problem Statement.....	10
Aim and Scope	11
Significance of the Study	11
Overview of the Study	11
Background.....	13
A Marker of the Good Life Is Happiness	13
This Is How to Be Happy	14
This Is Why We Are Not Happy.....	15
The Solution Is Habits?	16
Why the Solution Is a Band-Aid	17
Core	21
Position (methodology)	21
Method	22
Outcomes and Results.....	24
On the Persona	25
Ego.....	26
Subpersonalities.....	27
Subconscious Forces	29
Interpretation:	29
On Autonomy.....	30
The Allonomy of Autopilot.....	30
The Effect of Subconscious Forces.....	32
The Mechanisms of Control	33
Interpretation:	34
On Free Will	35
Deterministic Influences	36
The Ego as Compatibilist.....	37
A Drifting Will	40
Interpretation:	40
Results.....	41
Synthesis.....	44

The Discussion	44
Contentions.....	44
Over Background.....	44
1. Should We Even Aim for the Good Life?	44
2. The Problem of Rationalizing Human Behaviour	45
3. Does Increasing Willpower Extend At All?	47
Over Outcomes	48
4. The Degree of Subconscious Self-Sabotage	49
5. Are We Really Not Morally Responsible for our Subconscious?	50
6. What About Intelligent Design?	51
Over Results (interpretation)	52
7. The Ego's Influence on the General Will	52
8. How Much Can the Drifting Will Be Nudged?	54
9. The Reality of an External, Parameterized Framework.....	56
Caveats.....	58
Limitations and Shortcomings	60
Of the data	60
Of the results.....	61
Implications.....	61
The Conclusion	64
i. People are not any happier when they make their own decisions	64
ii. People are not led to better outcomes when they make their own decisions	64
iii. People are not any freer when they make their own decisions	65
Part 2.....	67
Introduction.....	69
Context of the Study	69
Problem Statement.....	69
Aim and Scope	70
Significance of the Study	70
Overview of the Study	70
Background.....	72
The Need for Elitism	72
Justifications from Isokratia.....	72
Fall into Totalitarianism	73
Philosopher Kings as the Answer?	74
Why Marcus Aurelius Is the Closest We've Gotten	74
Core	78
Position (methodology)	78
Method	79
Outcomes and Results.....	80
On Tenability.....	81
Implausible	81
Superfluous	83
Corrupt.....	84

Interpretation:	84
On Practicality	85
Legal Legitimacy	86
Democratic Legitimacy	89
Pragmatic Legitimacy	92
Interpretation:	96
On Efficacy	96
Sophia	97
Technê	100
Phronesis	102
Interpretation:	107
Results	107
Synthesis	110
The Discussion	110
Contentions	110
Over Background	110
1. A Transhumanist Metaphor or Jarvis/Alfred Blueprint?	110
2. Do the Benefits Justify the Costs?	111
3. The Price for Freedom	113
Over Outcomes	114
4. The Tradeoff of Autonomy	114
5. Spontaneity and Predeterminism	115
6. Moral Responsibility	116
Over Results (interpretation)	117
7. Subject-objectification and Object-subjectification	117
8. Meaningful Consent	119
9. The Reality of the Neo-Platonic Framework	120
Caveats	121
Limitations and Shortcomings	122
Of the data	122
Of the results	123
Implications	123
The Conclusion	126
I. The Philosopher King substantiates the most viable framework	126
ii. The Philosopher King substantiates the most supported framework	126
iii. The Philosopher King substantiates the most contrived framework	127
Part 3	129
Introduction	131
Context of the Study	131
Problem Statement	131
Aim and Scope	132
Significance of the Study	132
Overview of the Study	133
Background	134

The Story of Algorithms	134
Utopia	134
Dystopia	135
Value Alignment as Safeguarding?	136
Base “Original Sin” Code	136
Core	140
Position (methodology)	140
Method	142
Outcomes and Results	142
On AI Remedy	143
Mind: Individual Failure	144
Body: Structural Failure	146
Spirit: Ethical Failure	147
Interpretation:	149
On Potential Pathways.....	149
Phase 1: Tools (25% > AI; 75% < Human)	150
Phase 2: Extensions (~50% AI; ~50% Human).....	151
Phase 3: Environments (75% < AI; 25% > Human).....	153
Interpretation:	155
On Actionable Steps.....	156
Individual: Situated Literacy	156
Communal: Auditable Earnestness.....	158
International: Supranational Coordination.....	159
Interpretation:	161
Results.....	161
Synthesis.....	164
The Discussion	164
Contentions.....	164
Over Background.....	164
1. Does Sapience Without Sentience Invalidate Moral Legitimacy?	164
2. The Problem of “Trapped” AI	165
3. Algorithm Aversion: an Error in Reason - or an Intuition of Wisdom?	167
Over Outcomes	169
4. Optimizing Behavior at the Cost of Authenticity?	169
5. Is the Incremental Rollout of Cyberocracy a Trojan Horse for Technocracy?	170
6. Are We Replacing Human Bias with Algorithmic Blindness?	172
Over Results (interpretation)	173
7. Does Delegating Ethical Deliberation Disengage the Moral Will?	173
8. Can a World Without Human Rule Still Be a Human World?	175
9. The Reality of a Closed-Loop, Philosopher Algorithmic Neo-Platonic Framework	175
Caveats.....	177
Limitations and Shortcomings	178
Of the data	178

Of the results.....	179
Implications.....	179
The Conclusion	182
I. Outsourced systems are bound to their inherited moral origin	182
ii. Outsourced systems enable continuous ethical refinement	182
iii. Outsourced systems foster autonomy through design	183
The Discussion	185
Contentions.....	185
Over Foundations.....	185
1. The failure of ego is not just individual—it disables orientation itself.	185
2. Orientation requires authority—but authority cannot be grounded in the self.	186
3. Agency failure is recursive, not merely contingent.	187
Over Forms.....	188
4. Elite embodiment inherits the same fragilities as the ego	188
5. Alignment cannot be internalized—it must be externalized.	189
6. Legitimacy must be decoupled from embodiment to persist under complexity.	190
Over Orientations	191
7. Corrigibility is not an engineering feature—it is a constitutional principle.	191
8. Orientation must be teleological—but telos must be structural, not subjective.	192
9. Ethical governance in the post-egoic age is not a who—it’s a how.	194
Caveats.....	195
1. Assumptions—Not Repeated, But Reconfigured	195
2. Structural Interdependence and the Fragility of Translation	196
3. Epistemic Stance and Interpretive Partiality	197
4. On Synthesis Without Finality.....	198
5. On AI as Collaborator and Condition	199
Limitations and Shortcomings	200
Of the data	200
Of the results.....	201
Of the writing	201
Implications.....	202
The Conclusion	205
I. What We Aimed to Say	205
II. What We Have Said	205
III. What Follows from What Was Said	206
IV. One Last Bridge: From Theory to Telos	206
V. Final Metaphor	206
Bibliography	208
Appendix A: Misc.	220
Appendix B: Glossary	228
Theoria:	228
Poeisis:	231

Praxis:	233
Acknowledgements	236

Introduction

Context of the Study

It is a strange feature of the present moment that we are more empowered than ever to choose—and less equipped than ever to understand what our choices mean. In an age that exalts autonomy as the highest good, decisions are rarely made through reflection. They emerge from momentum. Individuals believe themselves to be choosing when, in fact, they are complying—submitting to forces beneath awareness and beyond control. Whether neurochemical, algorithmic, or cultural, the effect is the same: the self becomes less an agent than a conduit, animated by scripts it did not author.

This is not a crisis of intelligence, but of integration. Information is abundant; coherence is rare. Behavioral nudges, default settings, parasitic attention economies—these are not glitches in modern life but its design constraints. Under such conditions, freedom becomes indistinguishable from drift. The will, if it exists at all, is episodic. The self, long presumed sovereign, reveals itself as porous, unstable—embedded in systems it cannot fully see, let alone steer.

These concerns are not new. Plato warned of the democratic soul: fragmented, appetitive, governed by preference rather than purpose. Kant insisted that autonomy untethered from reason devolves into impulse. Freud revealed a psyche ruled by repression and division. Contemporary cognitive science has only deepened the critique, displacing the rational agent with a bundle of heuristics, habits, and hidden drives. Yet despite mounting evidence of internal incoherence, our systems of law, commerce, and governance continue to presume a subject who is rational, stable, and free—a subject who may never have existed.

The persistence of this myth is not owed to its strength, but to the absence of a credible replacement. To admit that most people cannot, in the aggregate, reliably act in their own interest—let alone in the interest of others—feels bleak. But denial does not forestall dysfunction. In a world where decisions scale instantly and consequences globalize overnight, the cost of epistemic incoherence is no longer personal—it is civilizational.

And yet the very technologies that expose this incoherence may also gesture beyond it. Artificial intelligence, framed alternately as threat or tool, invites a third possibility: not as a substitute for human agency nor a servant of human preference, but as a structural supplement to human limitation. This is not to claim that machines will become wise in any familiar sense. It is to reframe the question: can wisdom be approximated architecturally? Can ethical orientation be embedded in the very systems that shape our perception, behavior, and collective life?

This question reactivates an ancient tension: between knowledge and rule, between understanding and power. Plato's vision of the Philosopher King was not merely a fantasy of wise leadership—it was a structural thesis: that justice requires alignment with the Good, and that such alignment demands more than intelligence—it demands orientation. The Philosopher King ruled justly not because he was clever, but because his soul had been trained to perceive the Good and to act in accordance with it.

That vision unraveled under the Enlightenment's suspicion of authority. In its place arose the ideals of representative democracy, legal equality, and the self-governing subject. But these ideals, too, rest on a fragile anthropology: that individuals can self-regulate, self-correct, and act on principle rather than impulse. Today, this belief stands on increasingly unstable ground. The failures of deliberation are measurable. The limits of willpower are well-documented. And the systems built atop the myth of agency—markets, democracies, open networks—are producing outcomes that destabilize their own foundations. When climate inaction, algorithmic extremism, and political polarization emerge not from malice

but from ordinary choices made under systemic distortion, the crisis ceases to be theoretical. It becomes architectural.

This thesis is structured around that architecture—its failures, its limits, and its possible reconfigurations. It unfolds in three recursive movements:

Theoria examines the internal fragmentation of agency. It dismantles the assumption that individuals are unified decision-makers, revealing instead a self that is conflicted, permeable, and structurally compromised—a site of interference rather than command.

Poiesis turns to elite alternatives to egoic self-rule, probing the viability of the Philosopher King as a model for moral and political authority. Drawing from philosophical frameworks and historical exemplars, it asks whether such a figure can sustain coherence—and what occurs when he cannot.

Praxis pushes further, asking whether the role imagined for the Philosopher King might be instantiated not in a person, but in a system—specifically, in artificial architectures. It explores whether alignment with the Good can be designed rather than intuited, and whether moral coherence can be stabilized in structures resilient to the fragility of the agents they are meant to guide.

These parts are not sequential. Each refracts the others. Each failure disclosed becomes a ground for reentry. Together, they form a recursive inquiry into the conditions under which governance—ethical, political, algorithmic—might still function when the agents who inhabit it no longer do.

In this light, the context of this thesis is not AI, but failure. Not novelty, but recursion. The project returns to first principles—not to preserve them, but to reengineer their logic for a world that has outgrown its myths. The claim is not that humans are irredeemable, nor that machines are virtuous. It is that **structure matters more than strategy**. The question is not who governs, but how governance is constituted—and whether that constitution can endure the epistemic entropy of the present age.

Problem Statement

The modern subject stands at a paradox: entrusted with freedom, yet unequipped to bear it. Beneath the rhetoric of empowerment lies a cognitive architecture misaligned with the demands of contemporary life. Most behavior is not the product of reflection, but of constraint, habit, and bias. The figure of the rational agent—long central to liberal ethics, democratic theory, and market logic—has failed both descriptively and normatively. It neither explains how people behave, nor accounts for the persistence of avoidable harm despite centuries of institutional refinement.

This is not a peripheral flaw. When agency is assumed but not realized, the gap between intention and outcome becomes structural. Systems built on self-regulation—elections, markets, juries, publics—become susceptible to manipulation, fatigue, and noise. The issue is not malice, but misfit: individuals act on distorted value-maps within architectures they cannot see, let alone direct, generating cascading moral friction at scale.

That friction is not incidental. The mechanisms we rely on for judgment—perception, memory, motivation—are themselves shaped by inherited bias and infrastructural conditioning. We mistake fluency for truth, confidence for competence, familiarity for wisdom. These are not occasional errors. They are features of a subject shaped by environments hostile to coherence.

The result is a feedback loop: the subject believes itself competent, acts accordingly, and resists disconfirmation. This loop is reinforced by systems that flatter the ego—education that rewards performance over understanding, economies that valorize growth over

stewardship, politics that reward affect over orientation. Even self-help cultures often replicate the illusion they claim to transcend: that the ego can correct itself.

What emerges is not merely individual error, but collective miscalibration. As complexity scales and causal chains extend beyond the reach of intuition, the risks multiply. A single viral image can collapse institutional trust. A flawed heuristic in code can propagate injustice at population scale. Under such conditions, failure becomes default.

The core problem is not the absence of moral intention, but the absence of infrastructure through which intention can reliably express itself. Autonomy, under these conditions, becomes simulacral: what appears as choice is often inertia, framed as will. And the systems premised on this illusion—of the coherent agent, the self-correcting citizen—begin to fracture under their own assumptions.

This thesis contends that our most urgent ethical and political failures stem not from bad actors, but from a misdescribed subject embedded in incoherent systems. Efforts to resolve these failures through education, information, or willpower merely reiterate the error: they presume the very agency they aim to repair.

The question, then, is not whether individuals should be free, but whether current conditions of freedom can sustain the Good. Can agency, as presently constituted, withstand recursive distortion and algorithmic acceleration? Or must we imagine a supplemental form—external, corrigible, and non-anthropocentric—that can stabilize orientation when the self cannot?

This is not a call for technocracy, nor a rejection of human judgment. It is a call to reexamine the foundations of autonomy and governance in an age where intention no longer ensures alignment. The problem is not that freedom has failed, but that its architecture has been misunderstood. If the Good is to endure, the question is no longer simply what it is—but what kind of structure can hold it when we no longer can.

Aim and Scope

This thesis proposes a structural reimagining of governance: not through sovereign individuals, but through designed systems—specifically, artificially superintelligent architectures tasked with moral orientation in the absence of coherent agents.

This inquiry is not grounded in hype or alarmism, but in a practical metaphysics: that when systems fail to align with the Good, the failure may lie not in agency, but in structure. The question is not whether humans should rule, nor whether machines can replace them—but whether governance, as a function, can be re-engineered under conditions of recursive instability.

The central problem is not technical, but constitutional:

How can the pursuit of the Good be stabilized when neither individual nor institutional agency can reliably sustain it?

This question does not yield a singular answer. It opens onto three interdependent inquiries, each corresponding to one of the thesis's main movements:

- **Theoria — The Crisis of Egoic Intelligence**

Can intelligence be decoupled from ego without collapsing ethical orientation?

This section begins with the internal incoherence of the self. It shows that the subject is not a stable ground for action but a fractured site of interference. Theoria does not resolve where coherence might come from—but it sets the terms for why coherence cannot be presumed.

- **Poiesis — The Crisis of Embodied Rule**

Can moral legitimacy persist in governance without relying on human embodiment? Here, the thesis examines elite alternatives to egoic rule—most notably the Platonic

Philosopher King. While human-centered wisdom remains desirable, its embodiment in individual rulers appears increasingly untenable. Poiesis draws from Theoria to reject the sufficiency of ego, and gestures toward a structural form that might inherit the function of orientation.

- **Praxis — The Possibility of Architectural Ethics**

Can a corrigible system be built that externalizes moral direction without authoritarian closure?

If alignment with the Good cannot be sustained internally or embodied reliably, it must be embedded structurally. Praxis explores whether governance can be made recursive, corrigible, and teleologically oriented—not by replacing human judgment, but by supplementing its failure.

These questions are not answered in isolation. Each part refracts the others. Theoria diagnoses the fragility of the self. Poiesis tests the viability of moral embodiment. Praxis sketches a posthuman architecture for ethical stabilization. Together, they map a recursive logic of breakdown and reconstitution, in which moral orientation is no longer assumed, but must be designed.

In this sense, the thesis is not a linear argument but a conceptual structure. It advances not a solution, but a framework: a composite architecture of inquiry. Its aim is to reframe alignment—not as a problem of control or prediction, but of constitution and telos. It is a philosophical inquiry, but not an abstract one. It is grounded in real-world pathologies: epistemic volatility, institutional erosion, and the recursive distortion of agency by algorithmic systems.

The scope of this inquiry is deliberately conceptual. It makes no empirical claims about the current state of AI, nor any assumptions about timelines for artificial general intelligence. What matters is not when such systems might arrive, but whether we are asking the right questions about how such systems could—or should—be trusted. The concern is not with functionality, but with **moral form**.

This distinguishes the project from two dominant approaches. It departs from mainstream AI alignment discourse, which treats value learning and behavior shaping as engineering challenges. And it diverges from political theory that treats human agency as a normative necessity. This thesis begins with a different premise: that neither the human subject nor its institutional proxies can guarantee alignment—and that governance must be rethought structurally.

This is not a call for machinic sovereignty, nor a rejection of human dignity. It is a recognition that dignity without direction is dangerous—and that if direction cannot be sustained from within, it must be architected from without. The question is how: what telos it serves, how it remains corrigible, and whether it can stabilize coherence in a world no longer governed by coherent selves.

There are clear limits. This is a philosophical treatise, not a technical protocol. It assumes—for the sake of inquiry—the conceptual plausibility of strong ASI: agents capable of recursive self-modification and epistemic modeling. It also assumes that such systems could, in principle, be designed to remain aligned with the Good. These are not empirical claims, but necessary conditions for posing the deeper question:

If such a system could be built, what would it mean to trust it—and how might it be structured to remain trustworthy?

This thesis does not resolve that question. But it insists on its urgency. Its aim is not to predict the future, but to reshape the terms in which any viable future must now be conceived.

Significance of the Study

This thesis enters a fractured intellectual landscape. AI ethics often assumes that human values can be cleanly extracted and encoded. Political theory continues to rely on increasingly unstable conceptions of agency and legitimacy. Institutional reformers attempt coherence with tools designed for another era. Rather than operating within these frameworks, this study reframes the foundational problem they orbit: the collapse of moral orientation under recursive conditions. Its significance lies not in proposing a doctrine or solution, but in clarifying the kind of question that must now be asked—and the conceptual architecture in which it can be coherently posed.

The first contribution is philosophical. The thesis redefines alignment not as a technical or behavioral challenge, but as a constitutional one. The question is no longer whether systems will obey us, but whether *any* structure—human or artificial—can reliably sustain orientation toward the Good in the absence of stable agents. This reframing destabilizes two core assumptions: that values are legible and transferable, and that legitimacy must be grounded in embodiment. Rejecting both opens space for a post-anthropocentric conception of governance: corrigible, recursive, and teleologically structured.

Second, the thesis contributes to an emerging paradigm that treats governance as a design problem. As artificial systems increasingly shape perception, choice, and normativity, ethics becomes inseparable from infrastructure. This study does not frame that convergence as a crisis to reverse, but as a condition to re-engineer. It proposes that constitution be understood not as a civic artifact, but as embedded architecture: adaptive, recursive, and morally oriented by design. In this shift, sovereignty relocates—from agent to form.

Third, the thesis repositions philosophical inquiry itself. Rather than asking whether AI should be restrained or released, it asks what kind of system-level form might be capable of sustaining ethical direction when neither intention nor institutional memory can. This is not a rejection of moral life, nor an endorsement of machinic rule. It is a recognition that if agency is structurally unstable and governance increasingly algorithmic, then the central ethical question becomes architectural: what form can hold direction when the self cannot?

This reframing has implications across domains:

- **In AI ethics**, it moves beyond value extraction and oversight. Alignment becomes a structural property—not of behavior, but of orientation.
- **In political philosophy**, it challenges the view that legitimacy must return to the human subject. If agency is no longer epistemically coherent, legitimacy must be formal rather than authorial.
- **In design theory**, it reframes ethics not as an external imposition, but as a recursive design condition—something that must be embedded, not applied.
- **In constitutional theory**, it expands the idea of constitution beyond law or ritual. Constitution here names any formal system capable of sustaining orientation under epistemic entropy.

What this thesis offers is not an answer, but a constraint-space. It identifies the limits within which viable governance must now be conceived. It clears away assumptions that no longer hold—about agency, legitimacy, and the role of technology in ethical life. It does not advocate surrender to machines, nor retreat into sovereign nostalgia. It argues instead that if humans cannot reliably govern themselves or each other, the form of governance itself must be reconceived.

This is not a speculative manifesto. It does not predict the emergence of ASI or prescribe a technical paradigm. It assumes the plausibility of systems capable of recursive modeling and self-revision. From that assumption, it poses the second-order question: not what such systems should do, but what form they must take to be trustworthy—and how that trust must be structured if it is to endure.

That question cannot be resolved by technical expertise alone, nor by political will. It requires a reframing of ethics: not as intention, but as form; not as control, but as architecture.

In this light, the significance of the study lies not in novelty, but in necessity. It isolates the central dilemma of posthuman governance: not how to restrain emergent systems, but how to design systems worth aligning with. It treats misalignment not as malfunction, but as a symptom of deeper conceptual instability. Its contribution is not a model, but a frame—not a roadmap, but a clearing in which new maps might yet be drawn.

It is in this clearing that the stakes become legible:

- What does it mean to trust a system, not a subject?
- How can constitution persist without authorship?
- What forms of value can survive recursive abstraction—and which cannot?

These are not questions *Theoria*, *Poiesis*, or *Praxis* fully resolve. But they are the questions those movements make possible. And that, finally, is the significance of this study: it reframes the future of governance not as a problem of control, but as a problem of form.

Overview of the Study

This thesis unfolds in three recursive movements: **Theoria**, **Poiesis**, and **Praxis**. Each interrogates a distinct facet of the central question: how can moral orientation be sustained when neither individuals nor institutions can reliably secure it? These movements are not arranged linearly, nor do they build toward closure. Instead, they form a refractive arc: each exposes the limits of the others, clarifying the structural stakes of the inquiry as a whole.

The sequence begins with **Theoria**, which examines the failure of moral agency at its metaphysical root. Rather than treating the self as a stable ground for action, it reveals a fragmented, conditioned, and structurally misaligned subject. The problem is not bad intention, but cognitive interference. The rational agent turns out to be a recursive illusion. This leads to the first key inquiry: **can intelligence be decoupled from ego without collapsing its ethical substrate?**

But *Theoria* alone cannot resolve this. It diagnoses instability, but offers no structural alternative. That task begins in **Poiesis**, which shifts the focus from subject to structure. Here, the thesis examines elite models of governance—especially the Platonic *Philosopher King*—as attempts to sustain orientation through exceptional design rather than distributed will. These models are not idealized; they are stress-tested. They reveal both the conceptual power and the practical fragility of embodied wisdom. From this arises the second inquiry: **can governance retain legitimacy without relying on the anthropocentric assumption that wisdom must be humanly embodied?**

Yet *Poiesis* too reaches its limit. Even the best rulers are subject to bias, instability, and epistemic blind spots. Alignment remains necessary, but its embodiment becomes untenable. The thesis turns, finally, to **Praxis**—not as implementation, but as speculative extrapolation. If selfhood is unstable and sovereign embodiment insufficient, might orientation be architected instead? *Praxis* explores whether governance can be embedded in systems designed to remain corrigible under recursive strain. It reframes—without resolving—the third inquiry: **can a governance architecture be designed that externalizes moral direction without collapsing into coercion or drift?**

Each part offers insight and discloses constraint. *Theoria* fractures the self but leaves structure open. *Poiesis* probes elite rule but reveals its fragility. *Praxis* sketches a posthuman structure, but inherits the normative tensions of both. These recursive failures are not flaws—they are the form. The thesis is not a cumulative argument but a triangulation: each movement isolates a domain, clarifies its limits, and reframes the question.

This form mirrors the logic it examines. If agency collapses, so too must the fantasy that ethics, knowledge, or governance can be grounded in a single site—whether the subject, the sovereign, or the system. What is required is a distributed architecture of orientation: one capable of absorbing recursive instability without assuming epistemic closure. The thesis performs this logic formally, refusing linear build-up in favor of conceptual recursion.

Each movement also draws from distinct disciplines. **Theoria** engages metaphysics, cognitive science, and moral psychology. **Poiesis** draws from political theory, classical philosophy, and historical models of elite rule. **Praxis** incorporates systems theory, AI alignment discourse, and speculative constitutional design. This interdisciplinary movement is not stylistic—it is structural. No single discourse can resolve a problem that is recursive, infrastructural, and architectonic by nature.

Taken together, the three parts support a set of interlocking insights:

- That moral agency must be understood not as autonomy, but as vulnerability.
- That legitimacy cannot be anchored in embodiment alone, but must be structurally secured.
- That governance must be corrigible by design—not simply constrained by oversight.

These do not constitute a solution. They constitute a reconfiguration of the problem space. The thesis does not resolve the alignment problem, nor produce a model of posthuman governance. It constructs a conceptual architecture—a frame within which new and better questions may be asked.

The final movement—**Discussion and Conclusion**—returns to the core inquiries posed in the *Aim and Scope*. These are not answered in a conclusive sense. They are reframed in light of the structural interdependence revealed across the three parts. What emerges is not resolution, but orientation: a clarified map of the conceptual terrain that must be navigated if moral life is to persist in a world where neither subject, sovereign, nor system can be presumed stable.

This thesis does not prescribe a platform or a policy. It offers no utopia, only a form. A form through which moral orientation—though structurally fragile and historically elusive—might still be thinkable. Its contribution is not doctrinal but architectural. It holds open a conceptual space: a scaffold in which the Good might not be secured, but sustained.

Part 1
Theoria: The Ego as Limited

Abstract

This first part lays the ontological groundwork for the thesis by challenging the dominant modern assumption of individual autonomy. It contends that the human ego—often valorized as a rational decision-maker and ethical agent—is in fact a highly compromised construct. Drawing from psychology, neuroscience, and moral philosophy, the chapter argues that the ego is neither unitary nor sovereign. Instead, it is a spokesperson for a fragmented psyche, influenced more by subpersonalities, subconscious drives, and default neural modes than by any coherent will.

Theoria critiques the pervasive belief in self-discipline, motivation, and “free will” as adequate tools for reaching the Good Life. It frames these as comforting illusions maintained by a cultural emphasis on internal responsibility while ignoring the deeper structures of behavioral determinism and perceptual manipulation. The failure of most self-help paradigms, wellness cultures, and even classical virtue ethics is traced back to their overestimation of the ego’s power.

To break from this cycle, the chapter introduces a compatibilist framework: one that preserves moral responsibility without clinging to metaphysical autonomy. This shift demands that we move away from idealizing “self-mastery” and toward designing environments, infrastructures, and systems that support the Good Life externally. Theoria thus concludes not with a moral prescription, but with an epistemic reorientation—one that problematizes the self as a unit of ethical agency and opens the door to systemic, post-individual approaches developed in later parts.

Key words: autonomy, ego, free will, determinism, virtue

Introduction

Context of the Study

Wealth. Fame. Power. If we could obtain this... and everything else the world had to offer, would we finally acquire the Good Life?

It wouldn't be farfetched to claim that desire is the driving force of life. Whether it be the desire for a glass of water or the desire for complete world domination, desire is what animates us into being. And other than the base desire to stay alive, there is arguably no greater desire than that of the Good Life¹.

Questions surrounding the Good Life, particularly how we ought to live, can be traced back farther than antiquity, with the only thing rivalling the depth of the ethical theories being the breadth of their derivatives. But more than consequentialism, deontology, and virtue ethics for theories sake, when people generally speak of the Good Life, they are often expressing the desire to be happy.

While there has yet to be a consensus on how to go about this, there has been much overlapping anecdotal evidence for millennium. Wealth, fame, and power, which by and large encompass the external subsidiaries, only make up half the equation; the other (and arguably more important) half being the internal synthesis.

Despite recent statistical findings supporting this anecdotal evidence, in addition to knowledge surrounding happiness being more materially accessible than ever before, happiness isn't rising – in fact, it's falling, and at an accelerated rate². This led researchers to conclude that if the equation is both sound and complete, then it means implementing the felicific calculus must be easier said than done. Moreover, they extrapolated this failure to reach the goal of the Good Life as a result of lacking willpower.

Problem Statement

In response to this problem of failure to reach the goal of the Good Life was a resurgence of the wellness industry, with particular emphasis on self-help, meditation, and habit formation. It was generally believed that willpower, as a limited resource, could have its capacity trained like a muscle. Furthermore, that cultivation of a (for lack of better terminology) *Virtuous Hexis*³ would solve the source of ego depletion.

Yet, even though there have been meta-analyses debunking not just a single paper or two, but a fully-formed research literature, the dissemination of misinformation is still ongoing. Evidence suggests that efforts to train willpower are at best trivially effective, and at worst counterproductive. This is because self-control is limited to the ego, and our Persona is multifaceted. Willpower, and subsequently a Virtuous Hexis, then, is only affective during moments of engaged consciousness, which by default are far and few between.

As reviews have revealed how the field has changed, it is clear that applying new understandings is a key point of concern for acquiring the Good Life. We see that focusing on

¹ People do not merely wish to exist; they wish to live meaningfully. Many would rather die than endure a life devoid of worth. It is not life itself, but the Good Life, that is to be chiefly valued.

² Brennan, Megan. 2023. "Depression Rates Reach New Highs." Gallup. Last modified July 28, 2023. <https://news.gallup.com/poll/505745/depression-rates-reach-new-highs.aspx>.

³ In the *Nicomachean Ethics*, Aristotle claims virtue arises through habitual action (hexis), cultivated by repeated practice (habitos). While the original Greek sense of hexis is contextually bound to the soul, this thesis adopts a working definition: the *Virtuous Hexis* as the cultivation of state via mindful habit—commensurate with, though not reducible to, self-control.

See: Aristotle, *Nicomachean Ethics*, trans. Terence H. Irwin, 2nd ed. (Hackett, 1999).

willpower alone is an insufficient approach and that there exists a need for further investigation on how to turn our Self as a whole towards the good.

Aim and Scope

The singular, paramount aim of this part is to engage in this preliminary investigation of turning our Self as a whole towards the good. Limitations to the research are noted from the start and caveats are revisited later in the discussion. For example, I do not examine related discussions outside the Western philosophical traditions. The background I provide is far from comprehensive (the reader is encouraged to consult the entries on “Happiness” and “Well-Being” at the Stanford Encyclopedia of Philosophy for overviews on the topic, with the site also providing redirection to more niche concepts within the research literature), and the research is narrowed down to an emphasis on happiness as a marker of the Good Life. The part is also focused on external change, as opposed to using psychotherapeutic models to heal the pathology. In other words, the emphasis is on a practical ethics on a collective scale rather than an individual one. There is a need for these limitations because the investigation of how to turn our Persona as a whole towards the good is such a complex, interdisciplinary area.

In the part, I also limit the resources to those already published. I choose sources that are publicly available rather than private collections as they have a higher likelihood of establishing discourse, either by communicating through reference and citations or through public forums, therefore providing a deeper insight for the discussion. Resources that are already published are chosen because ongoing studies lack quality control either through public discourse or scientific review. Conducting experiments and/or case studies myself would also be too costly.

Significance of the Study

The overall value of this investigation is contingent primarily on its potential impacts. One intended outcome of the investigation, on a theoretical level, is to clarify the relationship between the Good Life and happiness.

Another intended outcome is advancing knowledge in this field by identify to what extent the effects of willpower extend to Subpersonalities. Specifically, I focus on how deterministic influences, whether external (unintended) or internal (intended), affect our default modes of conscious behaviour.

Further, on a practical level, a third intended outcome contributes to the solution of a practical problem, namely possible avenues to overcome failure to reach the goal of the Good Life.

Finally, the more immediate concern of this part is to adequately feed into the next part of my thesis.

Overview of the Study

The purpose of this part is to address the problem of failure to reach the goal of the Good Life and attempt to propose a new solution by way of external change, as opposed to using psychotherapeutic models to heal the pathology. In order to do this, this part will be divided into four main segments: Introduction, Background, Core, and Synthesis.

The Introduction, as we have just now seen, begins with an expanded review of existing theory and practice in relation to the aforementioned problem of realizing the Good

Life. After introducing the context of the study, it formalizes the problem statement with further theory behind why we are not living the Good Life. This is followed by clarifying a singular aim (preliminary investigation of turning our Self as a whole towards the Good) and laying the scope of the study. The penultimate section explores the significance of the study, overall value, and expected outcomes before arriving at the overview of the study where we are now sketching out what we aim to say.

The Background formalizes and clarifies definitions and concepts used throughout the remainder of the thesis; contextualizes the problem through salient history of issues to date; and posits a few tentative judgements before segueing into the Core segments. This narrative includes narrowing on a marker of the Good Life, analysing a framework for happiness and its inadequacies, as well as contending with the discovered problem of habit cultivation.

The Core is further divided into the Contribution (position and method of the research program) and Outcome (data, statistics, and analysis from the literature). The Core elaborates on the research results as they relate to a new solution by way of external change. This includes analysis On the Persona, On Autonomy, and On Free Will, particularly in how they interrelate in the capacity (and lack thereof) in turning the Self as a whole towards the Good.

The Synthesis is also substantiated by two segments: the Discussion (Contentions drawn from analysis of the outcome, including alternative interpretations; Caveats, Limitations and Shortcomings; and Implications) and Conclusion. The Synthesis integrates the key findings and appropriate concepts as they respond directly to the aim, acknowledging alternative approaches for external change which coherently segues into the following part of the thesis: *Poiesis*.

Background

A Marker of the Good Life Is Happiness

A marker of the Good Life is happiness. It does not matter which ethical theory you ascribe to⁴; happiness is undoubtedly an important variable, if not a central aspect of well-being⁵. This can be shown in two ways: from an explanatory stance, it is hard to imagine ascribing to any ethical theory if it did not accommodate happiness; and from a deliberative stance, it is even harder to imagine how the Good Life could be acquired in the absence of happiness.

It is important to note that although happiness is a marker of the Good Life, it does not necessarily implicate it. While the Good Life necessitates, or at the very least *commensurates* with happiness, Goodheart's Law indicates that pursuing happiness as an end in itself not only diminishes its relation to the Good Life, but could also result in converse effects⁶. One oft-heard aphorism, commonly called the "paradox of hedonism," is that the pursuit of happiness is self-defeating⁷; other case studies show how costs outweigh benefits due to systemic mistakes in the pursuit of happiness⁸. Happiness alone, then, is not indicative of the Good Life, yet its absence unmistakably signifies the lack thereof.

Ideally, a more comprehensive test of the Good Life would be an overlap of various markers, but insofar as happiness is at least a prerequisite for the Good Life, research is narrowed down to an emphasis on happiness. Moreover, the focus of the background is on how to be happy.

⁴ To avoid the complexity of strictly defining the Good Life or its relationship to happiness, this part does not advocate for hedonistic or other mental-state accounts. Rather, it treats happiness as a clear proxy. As Bostrom notes, referencing Bourget and Chalmers' *PhilPapers* survey: "No ethical theory commands majority support among philosophers, so most philosophers must be wrong." See:

Bourget, David, and David J. Chalmers. "The PhilPapers Surveys." 2009. <http://philpapers.org/surveys/>. Cited in: Nick Bostrom, *Superintelligence: Paths, Dangers, Strategies* (Oxford University Press, 2014).

⁵ As noted in the *Stanford Encyclopedia of Philosophy*, it is questionable whether any major philosophical school denies the importance of happiness across plausible accounts. See:

Haybron, Daniel M., and Valerie Tiberius. "Happiness." *Stanford Encyclopedia of Philosophy*, Summer 2023. <https://plato.stanford.edu/entries/happiness/>.

See also: Sidgwick [1907] 1966; Brandt 1979, 1989.

“人类的利益作为永远的起点和归宿。真正的理解是一场对话，而不是自说自话。”

[The good of mankind as the eternal starting point and destination. True understanding is a dialogue, not self-talk.]

— Zhu Rui (朱锐), ed. *What is Insight?* (2022)

⁶ Goodhart's Law suggests that "when a measure becomes a target, it ceases to be a good measure." This applies to happiness when pursued directly:

(1) The act of pursuing happiness distorts its meaning and validity as a metric, undermining objective research and creating measurement bias.

(2) Pursuing pleasure as a proxy for happiness can backfire—certain pleasures, such as drugs, may overwhelm or "crowd out" others, as Daniel Gilbert explains:

“Although most of us think of heroin as a source of human misery, shooting heroin doesn't actually make people feel miserable. It makes them feel really, really good—so good, in fact, that it crowds out every other source of pleasure.”

Cited in Daniel Gilbert, "Here's to Tofu, Baseball, and Heroin," blog entry, June 6, 2006.

https://www.randomhouse.com/kvpa/gilbert/blog/200606heres_to_tofu_baseball_heroin8.html.

⁷ Schooler, Jonathan, Dan Ariely, and George Loewenstein. "The Pursuit and Assessment of Happiness Can Be Self-Defeating." Working Paper, 2003. <https://www.researchgate.net/publication/228857366>.

Lyubomirsky, Sonja. *The How of Happiness: A Scientific Approach to Getting the Life You Want*. New York: Penguin Press, 2007.

⁸ Haybron, Daniel M. *The Pursuit of Unhappiness: The Elusive Psychology of Well-Being*. Oxford: Oxford University Press, 2008.

This Is How to Be Happy

The Stanford Encyclopedia of Philosophy has condensed happiness to be both: (1) a value term, roughly synonymous with well-being; and (2) a descriptive psychological term, akin to ‘depression’ or ‘tranquility’⁹. In simple terms, happiness is how you feel about your life in general and how you feel inside your head at any given moment in time. Within the context of my thesis, happiness is the overall individual, self-evident feeling that motivates us¹⁰.

For the purposes of this part, we are more concerned not so much about the nature of happiness as we are about how happiness is contrived. So far, it is understood that there are two dimensions of happiness: endogenic and exogenic, both equally manifesting within the brain the same as dopamine, serotonin, norepinephrine and endorphin¹¹. Incidentally, this proves¹² that organic happiness is not any better than synthetic happiness.

While the end-point dimensions of happiness are more or less agreed upon within the scientific community, the same cannot be said about the start-point sources. Nonetheless, there has been much overlapping anecdotal evidence for millennium. Wealth, fame, and power, which by and large encompass the exogenic subsidiaries, only make up half the equation¹³; the other endogenic (and arguably more important¹⁴) half being the internal synthesis.

⁹ This claim is a generalization and should be qualified. Some philosophers (e.g., Chekola 2007; Murphy 2001) treat “happiness” descriptively in ways that differ from standard well-being models. There is also a growing body of non-Western and cross-cultural scholarship complicating Western assumptions. For a comprehensive overview, see:

Stanford Encyclopedia of Philosophy, “Happiness,” Note 2.

<https://plato.stanford.edu/entries/happiness/notes.html#note-2>.

¹⁰ “Happiness” is an imprecise term—commonly confused with hedonic pleasure or flattened into something as shallow as “elevator music.” Yet, elevated to its Platonic Form, happiness might instead be understood as a profound, even paradoxical state: an awe inspired by suffering, a reflection of transcendence. In this thesis, happiness functions as a placeholder for a concept more ineffable than pleasure and more encompassing than well-being.

See: *The School of Life*, “Why Socrates Hated Democracy,” YouTube, Nov. 28, 2014.

<https://www.youtube.com/watch?v=3uLDin9A9pc>.

¹¹ Dfarhud, Dariush, Maryam Malmir, and Mohammad Khanahmadi. “Happiness & Health: The Biological Factors – Systematic Review Article.” *Iranian Journal of Public Health* 43, no. 11 (2014): 1468–1477. PMID: 26060713; PMCID: PMC4449495.

¹² Gilbert, Daniel. “Isn’t Real Happiness Better Than Synthetic?” In *Stumbling on Happiness*. New York: Alfred A. Knopf, 2006.

The implications echo Nozick’s Experience Machine thought experiment and critiques of “inauthentic happiness”—whether achieved through self-deception or artificial stimuli like Huxley’s soma, which promises pleasure divorced from meaning.

See also: Huxley, Aldous. *Brave New World*. [1932] Reprint. New York: Harper Perennial Modern Classics, 2005.

¹³ Wealth, fame, and power correlate with Maslow’s hierarchy of needs but do not straightforwardly equate to happiness. The *Easterlin Paradox* demonstrates that while wealthier individuals within a country may be happier, richer countries are not significantly happier than poorer ones—and economic growth alone does little to improve happiness. Moreover, some communities (e.g., the Amish, Inughuit, Maasai) report high subjective well-being despite low material prosperity.

See: Easterlin 1974; Jebb et al. 2018; Diener et al. 2018; Helliwell et al. 2012–present; Biswas-Diener 2005, 2018; Graham 2009; Everett 2009.

As cited in: SEP, “Happiness,” <https://plato.stanford.edu/entries/happiness>.

¹⁴ One may be unhappy with a healthy body or happy despite an unhealthy one. External goods like wealth and power are facilitators—not sources—of happiness. Endogenous factors (e.g., disposition, perspective) require internal synthesis, much like carbohydrates must be metabolized into energy. In this analogy, happiness is the product of an inner transformation, not just an external input. As Sir Thomas Browne wrote:

This Is Why We Are Not Happy

In recent years, despite statistical findings supporting anecdotal evidence, people are still not happy¹⁵. This is perplexing when considering that this evidence is neither esoteric nor deliberately kept hidden. In other words, there exist recipes, if not concrete guidelines, for happiness which are materially accessible, yet still people are not happy.

Interpreting research narrowed down to an emphasis on the exogenic subsidiaries of happiness, more specifically the control aspect of power as it relates to happiness¹⁶, we see that there are three key explanations available:

(1) People know what to do, they just don't do it. Some of this is due to Cinderella complex¹⁷ where people simply wait for somebody else to rescue them. In this case you cannot help those who do not want to help themselves;

(2) Factors outside of our control. For example if you are living through a famine, then it isn't hard to imagine that one's wealth is impacted. Coincidentally, it also makes intuitive sense that the Good Life is not one situated in famine. Here, factors beyond our control are exactly that – beyond our control;

(3) It is easier said than done. Researchers have concluded that, assuming that the data is correct, failure to reach the goal of the Good Life is a result of lacking willpower¹⁸. It would seem then that a solution to the problem is increasing willpower.

"I am the happiest man alive. I have that in me that can convert poverty to riches, adversity to prosperity... fortune hath not one place to hit me."

– Browne, Thomas. *Religio Medici*, 1642.

See also: Brickman, Philip, Dan Coates, and Ronnie Janoff-Bulman. "Lottery Winners and Accident Victims: Is Happiness Relative?" *Journal of Personality and Social Psychology* 36, no. 8 (1978): 917–927.

¹⁵ Of course, there are exceptions—some may claim they are living the Good Life. But in aggregate, indices like the UN Human Development Index (HDI) and other well-being frameworks show that global happiness is not only declining, but doing so at an accelerating pace.

(See Appendix A: Misc., Exhibit 1.)

¹⁶ In line with this thesis's scope, the focus narrows to a subset of exogenic factors—specifically, autonomy. Unlike endogenic variables (e.g., inner blessedness or subjective synthesis), autonomy can be more readily quantified and tracked across populations, making it more useful for empirical analysis.

For comprehensive surveys and datasets, see: Myers and Diener 1995; Argyle 1999, 2002; Layard 2005; Bok 2010a; Diener et al. 2018; Jebb et al. 2020; and the annual *World Happiness Reports* (Helliwell et al. 2012–present).

For philosophical perspectives on autonomy and well-being, see: Ahuvia et al. 2015; Andreou 2010; Brülde 2014; Haybron 2013a; Lauinger 2015; Morris 2015; Setiya 2014; Tiberius 2017; Wren-Lewis 2019.

Also cited in: Kesebir and Diener, "In Pursuit of Happiness: Empirical Answers to Philosophical Questions," *Perspectives on Psychological Science* 3, no. 2 (2008): 117–25.

¹⁷ The "Cinderella Complex" refers to an unconscious desire—typically in women—to be taken care of, often at the cost of autonomy. The term originates from the fairy tale archetype: beautiful, passive, polite, and dependent, awaiting rescue. It reflects internalized fears of independence and an expectation of external salvation.

Source: Harley Therapy. "The Cinderella Complex: Fear of Independence." *Harley Therapy - Psychotherapy & Counselling*. <https://www.harleytherapy.co.uk/counselling/cinderella-complex.htm> (accessed September 30, 2023).

¹⁸ In a key study on ego depletion, Baumeister et al. (1998) found that participants who resisted temptation (e.g., cookies) subsequently performed worse on unrelated self-control tasks (e.g., solving puzzles). This led to the hypothesis that self-control draws on a limited resource – which I find akin to a "mana bar" in video games. Baumeister, Roy F., Ellen Bratslavsky, Mark Muraven, and Dianne M. Tice. "Ego Depletion: Is the Active Self a Limited Resource?" *Journal of Personality and Social Psychology* 74, no. 5 (1998): 1252–65.

[https://faculty.washington.edu/jdb/345/345%20Articles/Baumeister%20et%20al.%20\(1998\).pdf](https://faculty.washington.edu/jdb/345/345%20Articles/Baumeister%20et%20al.%20(1998).pdf).

Indeed, the next step many have made from this is that if a lack of willpower¹⁹, as a consequence of ego depletion (the idea that self-control or willpower is a limited resource that can be depleted through use), is at the core of the problem; then by exercising self-control, willpower would seemingly increase²⁰. Leveraging the properties of the information age, the materially accessible data from the research literature fed into the booming wellness industry, most notably in the forms of self-help, meditation, and habit formation²¹.

The Solution Is Habits?

It was theorized that if we could get better at controlling ourselves²², then all our individual problems whether it was lose weight, read more, or even make enough money, would be also be fixed. This was based on the idea that self-control, as identified in the Anterior Singulate Cortex²³, could be trained in the same way as synapse plasticity. In other words, willpower was likened to a muscle which would increase in strength through exercise.

The cultivation of a Virtuous Hexis was thought to be at the core of the solution. It was believed that by engaging in mindful activities²⁴ (i.e. reading for thirty minutes each morning), characteristics akin to temperance (virtue) would inevitably lead to a habitual state (hexis). Furthermore, the more the mindful activity was practiced, the larger the capacity for self-control would become, until, like muscle memory, the Virtuous Hexis would become second nature and the mindful activity would no longer be mindful, but instead natural.

¹⁹ While often conflated, weakness of will, lack of self-control, and akrasia are distinct—though overlapping—concepts. This thesis focuses primarily on the first two. As an illustration: those who binge drink, regret it, and promise to stop—only to repeat the behavior soon after—demonstrate weakness of will, not simply lack of understanding. Akrasia, as later discussed, adds a philosophical layer to this failure of alignment between knowledge and action.

²⁰ The interplay between ego, willpower, and self-control can be visualized through a gaming metaphor:

1. **Ego as Player Character** – the center of agency, decision-making, and self-representation.
2. **Willpower as Mana Bar** – a finite resource required for effortful control.
3. **Self-Control as Mana-Spending Action** – each exertion of restraint depletes this resource.

See: Muraven, Mark, and Roy F. Baumeister. “Self-Regulation and Depletion of Limited Resources: Does Self-Control Resemble a Muscle?” *Psychological Bulletin* 126, no. 2 (2000): 247–59. <https://doi.org/10.1037/0033-2909.126.2.247>.

²¹ The Marshmallow Test, a landmark experiment in self-control, was designed to assess delayed gratification by offering children a choice between an immediate reward (one marshmallow) or a larger reward (two marshmallows) if they waited. The experiment’s broader implication—that delayed gratification predicts long-term success—was later echoed in wellness trends emphasizing willpower. Notably, the 1970s saw a concurrent rise in the wellness industry, which absorbed such findings to promote habit-based self-improvement.

See: Mischel, Walter, and Ebbe B. Ebbesen. “Attention in Delay of Gratification.” *Journal of Personality and Social Psychology* 16, no. 2 (1970): 329–37. <https://doi.org/10.1037/h0029815>;

van Tulleken, Chris. *Ultra-Processed People: The Science Behind Food That Isn’t Food*. New York: W.W. Norton, 2023;

Adams, Thomas, Sharon K. Anderson, and Robert K. Hester. “Wellness: Its Origins, Theories, and Current Applications in Human Services.” NASPA, 2014.

https://www.naspa.org/images/uploads/kcs/WHPL_Canon_We_Wellness_Its_origins_Theories_and_application_s.pdf.

²² While often used synonymously with *enkrasia* (self-control), the concepts of (a) improving our ability to resist temptation, (b) doing what is morally right, and (c) overcoming weakness of will are overlapping yet distinct. This thesis distinguishes between these nuances where relevant.

²³ After Skool. “The Socratic Method – Philosophy – Socrates – Debate – Dialectic.” YouTube video, 9:00. April 15, 2020. <https://www.youtube.com/watch?v=dgRSfhoHE4g>.

²⁴ Mindful activities require active attention—like learning to drive or interpreting this very sentence—whereas mindless activities, such as passive scrolling, occur without conscious engagement. The distinction matters because cultivating self-control depends on periods of deliberate awareness, which engage the ego.

Circumstantial evidence later went on to support that in this way not only would willpower increase²⁵, but also the ego would no longer need to be depleted on mindful activities that would later become second nature.

But if increasing willpower was a solution to happiness, then why isn't everyone engaging in cultivating a Virtuous Hexis? Moreover, how is it that happiness is not only not increasing, but actually decreasing inversely proportional to the booming of the wellness industry²⁶? Why is it that psychotherapeutic models of habit creation and the cultivation of a Virtuous Hexis fail to heal the pathology of akrasia?

Why the Solution Is a Band-Aid

Other than a few outliers²⁷, there is a lack of data supporting that increasing willpower (if it can even be increased at all²⁸) leads to an increase in happiness. Furthermore, there is even scarcer evidence, if any at all, which supports the original hypothesis of increasing willpower would subsequently help fix all our other individual problems²⁹. Rather, looking at it scientifically, the elusiveness about self-control is that we can get good at a particular task but that doesn't actually translate over into the other parts of our lives. Using cold showers as an example, the only thing cold showers really make you good at is tolerating cold showers.

There have been meta-analyses debunking not just a single paper or two, but a fully-formed research literature³⁰. Moreover, the conclusions were rushed and the solution

²⁵ Inzlicht, Michael, Brandon J. Schmeichel, and C. Neil Macrae. "Why Self-Control Seems (but May Not Be) Limited." *Trends in Cognitive Sciences* 18, no. 3 (2014): 127–133. <https://doi.org/10.1016/j.tics.2013.12.009>.

²⁶ Despite the exponential growth of the wellness industry, global happiness indicators continue to decline. This is reflected in rising rates of depression, suggesting a disconnect between marketed solutions and actual well-being.

See: Brenan, Megan. "Depression Rates Reach New Highs." *Gallup*, July 28, 2023.

<https://news.gallup.com/poll/505745/depression-rates-reach-new-highs.aspx>.

²⁷ Apparent anomalies in wellness outcomes may be attributable to **healthy user bias**—a selection effect where individuals who pursue interventions (e.g., clinical trials, diets, self-help regimens) are already predisposed to healthier outcomes due to their health-conscious behaviors.

See: Shrank, William H., Alicia R. Patrick, and M. Alan Brookhart. "Healthy User and Related Biases in Observational Studies of Preventive Interventions: A Primer for Physicians." *Journal of General Internal Medicine* 26, no. 5 (2011): 546–50. <https://doi.org/10.1007/s11606-010-1609-1>.

²⁸ Inzlicht, Michael, and Brandon J. Schmeichel. "What Is Ego Depletion? Toward a Mechanistic Revision of the Resource Model of Self-Control." *Perspectives on Psychological Science* 7, no. 5 (2012): 450–63. <https://doi.org/10.1177/1745691612454134>.

²⁹ Promoting self-control as a universal solution can backfire. Evidence suggests that targeting the wrong audiences—particularly those who lack baseline behavioral support—may yield no benefit, provoke resistance, or even cause psychological harm by encouraging comparison and self-blame.

³⁰ The replicability of ego depletion has come under increasing scrutiny. Meta-analyses and multi-lab replications have failed to find consistent effects, raising concerns of publication bias. This has led some to argue that the original resource model may have been overstated.

See:

– Inzlicht, Michael [@minzlicht]. "Replication attempts of ego depletion meta-analysis show no effect."

Twitter, January 29, 2016. <https://twitter.com/minzlicht/status/693238008305176576>

– Carter, Evan C., and Michael E. McCullough. "Publication Bias and the Limited Strength Model of Self-Control." University of Miami, 2014. <https://local.psy.miami.edu/chblab/PubBiasSelfControlEgo.pdf>

– Hagger, Martin S., et al. "A Multilab Preregistered Replication of the Ego-Depletion Effect." University of Miami, 2016.

– Engber, Daniel. "Everything Is Crumbling." *Slate*, March 4, 2016. https://slate.com/health-and-science/2016/03/ego_depletion_an_influential_theory_in_psychology_may_have_just_been_debunked.html

– Mangin, Thomas et al. "No Ego-Depletion Effect Without a Good Control Task." *Psychology of Sport and Exercise* 57 (2021). <https://doi.org/10.1016/j.psychsport.2021.102033>

– Englert, Chris, and Alex Bertrams. "Again, No Evidence for or Against the Existence of Ego Depletion." *Frontiers in Human Neuroscience* 15 (2021). <https://doi.org/10.3389/fnhum.2021.658890>

proposed (increasing willpower) was based on an oversimplification of the data³¹. Nonetheless, while the analysis was incomplete, stemming from findings that were correlational and not causational, the data itself still retains its objective value. In other words, the literature is invalid, but not completely wrong.

It has been shown then that the research literature was correct on at least three fronts:

- (1) Willpower does have an effect on self-control
- (2) Training willpower does increase self-control
- (3) Ego depletion is a real phenomenon

First and foremost, Willpower does have an effect on self-control. This can be seen both positively and negatively where willpower is affected by self-fulfilling prophecies³². In the West, for example, many believe, and experience, that exercising self-control leaves them feeling drained. I.e. the fatigue after a morning spent dealing with difficult clients or focused on spreadsheets on a computer screen. In the East, the opposite is true, particularly in India where the cultural belief that mental effort is energising and the more concentration and self-control you expend in one situation, the more invigorated you will feel for the next challenge³³. Further research supports the effect that willpower has on self-control is intrinsically tied to belief³⁴ – this, in essence, means holding on to the idea that willpower is a limited resource can actually be bad for us, making us *more* likely to lose control and act against our better judgment.

Second of all, training willpower does improve self-control through awareness sensitivity. Maintaining the willpower as a muscle analogy, we see that the strength of our willpower increases through neural adaptation³⁵. Nonetheless, since willpower is not actually a muscle but a property of the Anterior Singulate Cortex, its “genetic ceiling” is capped at what we already have, meaning the strength of our willpower can increase through efficiency and using our willpower to its full potential, but our willpower cannot increase in size or capacity. This efficiency as a result of neural adaptation has been shown to be due to increasing awareness sensitivity. Self-control, in the conventional sense at least, doesn’t actually exist, rather it is an emergent property of conflict monitoring. Previously, it was thought that awareness precedes control, when in fact, it has been show that *awareness is control*³⁶. By internalizing our awareness either indirectly through emotional regulation and

³¹ Much of what is popularly circulated about willpower and self-control simplifies or distorts the original research. In many cases, “oversimplification” becomes a form of misrepresentation: even if the premises remain technically true, the conclusions drawn may be invalid or misleading.

³² “According to These Researchers, Willpower Is Not a Limited Resource.” *The Behavioral Scientist*. Accessed September 30, 2023. <https://www.thebehavioralscientist.com/articles/according-to-these-researchers-willpower-is-not-a-limited-resource>

³³ Savani, Krishna, and Veronika Job. “Reverse Ego-Depletion: Acts of Self-Control Can Improve Subsequent Performance in Indian Cultural Contexts.” *Journal of Personality and Social Psychology* 113, no. 4 (2017): 589–607. <https://doi.org/10.1037/pspi0000099>

³⁴ Job, Veronika, Carol S. Dweck, and Gregory M. Walton. “Ego Depletion—Is It All in Your Head? Implicit Theories About Willpower Affect Self-Regulation.” *Psychological Science* 21, no. 11 (2010): 1686–93. <https://doi.org/10.1177/0956797610384745>

³⁵ Posner, Michael I., and Mary K. Rothbart. “Willpower and Brain Networks.” *ISSBD Bulletin* 2012, no. 1 (2012): 7–10. PMID: 25264480.

³⁶ Self-control is not the resolution of conflict, but the very act of staying engaged with it. When internal battles cease—through distraction or suppression—control is often lost. This insight is supported by research on the **Iowa Gambling Task (IGT)** and the **somatic marker hypothesis**, which propose that decision-making is often guided by pre-conscious affective cues. Intuition, in this view, functions as a fast, automated decision architecture—one that precedes rationalization.

stress prevention, or directly through acts of meditation, synapse plasticity is strengthened, enforcing neural adaptation, which results in improved self-control.

Last but not least, ego depletion is correct, just not in the way originally conceived. Traditionally, ego depletion has been referred to as the idea that self-control or willpower is a limited resource that can be depleted through use. But it would be more accurate to understand ego depletion as the waning of the ego in response to “fronting” of other subpersonalities³⁷.

The missing key that researchers did not have (or perhaps did not use³⁸) is that our ego – the estimation of ourselves – is not actually our entire selves, but rather a small personality attached to a collective of giant subconscious forces. In other words, it is not necessarily that we lack the willpower, it is that **we have characteristically misunderstood autonomy of the ego, as opposed to the Ego**³⁹, in relationship to our Selves and Persona.

It is primarily this missing key which would have altered the entire directionality of the researcher’s initial conclusions. Without this information, they failed to understand that self-control is limited to the ego since conflict monitoring is only made possible through awareness, hence an engaged consciousness. Moreover, since the default of our multifaceted Persona is not predominantly that of the ego, it is no surprise why it would appear that we have a lack of willpower. Put simply, since willpower (and subsequently, a Virtuous Hexis⁴⁰) is a substrate of the ego, it is clear how self-control would fail to be affective when the ego is absent, e.g. when on autopilot.

From this background chapter, we have thus seen that:

1. Happiness is not necessarily the best marker of the Good Life (particularly in isolation), but insofar as it is at least accommodated, if not actively facilitated, happiness must be promoted to acquire the Good Life.

See: Bechara, Antoine, Antonio Damasio, Hanna Damasio, and Steven Anderson. “Insensitivity to Future Consequences Following Damage to Human Prefrontal Cortex.” *Cognition* 50, no. 1–3 (1994): 7–15.

[https://doi.org/10.1016/0010-0277\(94\)90018-3](https://doi.org/10.1016/0010-0277(94)90018-3);

Damasio, Antonio R. *Descartes’ Error: Emotion, Reason and the Human Brain*. London: Vintage, 2008;

Walsh, Daniel. “Occam’s Razor (Luck) Is the Most Efficacious Decision-Making Tool in the Absence of Perfect Information.” Forthcoming, 2021.

³⁷ The human mind comprises multiple competing subpersonalities, each taking turns “fronting” as the conscious self. While *personality* refers to the enduring configuration of traits across time, the *persona* is a context-specific mask—the social interface of the self. The distinction matters: the persona is the phenotypic projection known to others, while a subpersonality is a genotypic entity within the internal ecology of mind.

³⁸ Psychology’s empirical credibility has been shaken by its **replication crisis**, which raises doubts about the robustness of many widely cited findings—including those related to self-control. Compounding this issue is the commercial influence of the wellness industry, which has an economic interest in sustaining simplified narratives.

See: Yong, Ed. “Psychology’s Replication Crisis Is Running Out of Excuses.” *Wired*, May 11, 2016.

<https://www.wired.com/2016/03/psychology-crisis-whether-crisis/>;

The Writing Space. “Statistics on Self-Help Books.” *Medium*, January 15, 2023. <https://medium.com/the-writing-space/here-are-some-interesting-statistics-on-self-help-books-e9895a848c8f>;

After Skool. “The Neuroscience of Willpower.” YouTube video, March 22, 2021.

<https://www.youtube.com/watch?v=hQB3viVFhPA>.

³⁹ The “you” currently reading this sentence—the focal point of conscious awareness—is not the full self, but the **Ego**: the executive agent within the mind’s system. It is a temporary front, distinct from both the deeper personality and the broader internal community of subpersonalities.

⁴⁰ A **Virtuous Hexis**, as framed in this thesis, is mostly constrained to the domain of the Ego. As a result, cultivated habits may not generalize: mental routines (e.g., memorizing vocabulary) often fail to carry over, while embodied habits (e.g., brushing teeth) tend to persist. This distinction reinforces the idea that durable self-control requires physical anchoring.

2. Happiness is a manifestation in the brain, triggered by both endogenic and exogenic dimensions. In regards to the latter, wealth, fame, and power are non-negotiable subsidiaries.
3. The research literature had concluded that we fail to obtain baseline levels of wealth, fame, and power, (and, subsequently, the Good Life) because we lack willpower.
4. The proposed solution was the cultivation of a Virtuous Hexis through practising mindful activities. This was theorized to not only increase the capacity for self-control but also reduce the burden of control needed for otherwise mindful activities that become second nature.
5. The data the researchers procured retain objective value, but their conclusions were premature. Willpower does have an effect on self-control; training willpower does increase self-control; and ego depletion is a real phenomenon. But, since self-control isn't affective when the ego is absent, increasing willpower is not an effective solution for achieving the Good Life.

It is important to acknowledgment that without the significant contributions of the academic community and their notable advancements, it is evident that our current understanding wouldn't be as enriched in the various domains that it is today. However, in the spirit of scholarly dialogue, it is equally important to remember that function of a piece of writing is to move a research area forward, not to be preserved for 500 years. As such, this part would like to draw attention to a specific aspect that may benefit from further exploration or consideration. It is my intention to engage in a collaborative discussion and contribute to the refinement of our collective knowledge, namely, by recognizing that the field is changing and implementing revised understandings into theories of how to acquire the Good Life.

Specifically, then, we see that focusing on willpower alone is an insufficient approach and this substantiates a need for further investigation on how to turn our Self as a whole towards the Good.

In what follows, we delve into the core chapters of my work which, drawing from the conclusions of this background chapter, leads into my research hypotheses and questions. The task of these core chapters is to reconcile the missing key of a limited ego with failure to reach the goal of the Good Life. Moreover, it is to profile the predisposition to self-sabotage based on the Persona, Autonomy, and Free Will for actualization of the good.

Core

Position (methodology)

From the conclusions of the background chapter, it is clear that focusing on willpower alone is an insufficient approach to obtaining wealth, fame, and power; moreover, that we must shift our perspective from taking the estimation of ourselves as the entire self to that of the limited ego.

Research questions that naturally arise include, but are not limited to: if willpower alone is an insufficient approach, then which approach(es) would we need to be more efficient? How are we to understand the pluralism of our Persona in relation to acquiring the Good Life? To what extent does the Ego possess autonomy? And to what extent does the Subconscious Forces influence, if not outright control us?

Taken together, these key research questions lead into an initial formalized hypothesis: we are slaves (to Subconscious Forces that are really in control of us).

As my thesis is interdisciplinary with an emphasis on changing or emerging fields, much of the work unprecedented with no prior existing framework or research program. My research questions revolve predominantly around *current theory, discoveries, and debates* that are most useful and salient to my topic, with contribution is informed by a spread from psychology to sociology. There is an evaluative mix of theoretical and applied, study and case studies, observation and occasional innovation.

My particular stance⁴¹ in regards to the information, or data, that I am analyzing will be purely philosophical since, as already mentioned in the scope of this part, conducting experiments and/or case studies myself to support my hypotheses would be too costly. My role as lead, and sole, researcher will thus be limited to the evaluator of theoretical and applied (including thought) experiments, objective observer, and even, at times, participant in the activities (even if trivially anecdotal) under study.

In conduit with this philosophical stance, there are key suppositions and attitudes that are important to state and justify. Of course, it is impossible to completely observe oneself objectively, therefore, it is very likely that I hold other stances that I am unaware of. Nonetheless, those of which I can consciously report include:

Stance 1: I hold a classical utilitarian attitude in regards to happiness as a marker.

Justification: to hold a classical utilitarian view in regards to happiness within the context of this thesis means happiness is valued akin to the highest good, that if not a moral tenet of the Good Life itself, is at least a catalyst towards it. This elevated status of happiness, specifically the greatest happiness for the greatest number, is a teleological necessity⁴². Simply, assuming that happiness consummates with the Good, then more happiness in more scenarios

⁴¹ In the social sciences, “methodology” refers not merely to method or technique, but to the broader epistemological stance of the researcher—whether as designer of experiments, objective observer, or participant-observer. In this thesis, the term is used similarly to denote the overarching framework or philosophical lens, and should not be conflated with empirical designs used in the natural sciences.

⁴² “But shorn of teleology, ethics as a body of knowledge was expurgated of its central content [...] a set of arbitrary prohibitions.” MacIntyre illustrates this through the example of the Polynesians, whose moral taboos lost meaning once stripped of their spiritual context. When King Kamehameha II abolished them, there was no resistance—demonstrating how morality can collapse into incoherence when cut off from a teleological framework.

MacIntyre, Alasdair. *After Virtue: A Study in Moral Theory*. 3rd ed. Notre Dame, IN: University of Notre Dame Press, 2007.

necessarily matches with movement along the axis of the Good. Alternatively, this viewpoint posits: if an action does not maximise*⁴³ happiness, it is a “good” action, but not *the* “Good” action. For instance, in times of default, if a Subpersonality blows an entire paycheck at a casino, then relating back to wealth, we see that this was not the action that would maximize happiness, and hence not a good action according to the Utilitarian paradigm.

Stance 2: I take that uncontrollable sources of (un)happiness cannot outweigh controlled sources.

Justification: if exogenic happiness is to be a sufficient marker, then it must not only be clear what these exogenic properties are (e.g. what things make you happy), but that their influence be substantive as well (e.g. palpable effects are noteworthy). If this were not the case, then any conscious effort to be happy would be futile, meaning autonomy would be a hoax and regardless of effort, levels of happiness would be deterministically fixed. This stance is thus non-negotiable, because either it is true which validates any and all investigations by way of this thesis; or it is false which not only invalidates this thesis but also all ideas regarding free will and the extent to which we influence the World.

Stance 3: I support a neuroscientific model understanding of the mind.

Justification: admittedly, it is impossible to rationalize human behaviour since we are not solely mechanistic beings; we cannot subsume the biological into the technical⁴⁴ (i.e. Descartes’ animal machine par excellence). One example we explore of where this does not work is the Triune Complex which has been largely dismissed as a myth in the fields of evolutionary and developmental neuroscience. Nonetheless, we still make simplified models and heuristic analysis, such as only the Ego being free – at least in the Kantian sense – and subsequently, rationality as freedom. This is necessary, if not in the search for Truth, then at least in the laying of understanding. The greatest justification for the neuroscientific model specifically in contrast to other similar models of the mind is that it is comparatively more correct, complete, and significant. Like most other justifications from science⁴⁵, then, the neuroscientific model might not necessarily be right, but it is at least less wrong.

Put simply, my position can be reduced to: the research program is what I am doing; and my stance is necessarily how I’m going to go about my investigation. It is not much more complicated than that, but for the sake of research ethics, I felt it important to lay out my identified stances as above.

Method

⁴³ This parallels Goodhart’s Law: maximizing a proxy (like happiness or success) distorts the underlying goal. The point here is that maximization in this context does not refer to sacrificing long-term goods (e.g. health) for short-term gains (e.g. drugs), but to the corruption of ends by fixating too rigidly on the means.

⁴⁴ Di Siena, Tosca. "The Dynamic Object: Naturalizing the Tekhne with Gilbert Simondon." *Medium*, June 10, 2023. <https://medium.com/@tosca.siena/the-dynamic-object-naturalizing-the-tekhne%CC%84-with-gilbert-simondon-0479a3dae85c>.

⁴⁵ In science, “facts” are provisional—valid until falsified. Progress is not the accumulation of immutable truths but the elimination of error. As Kuhn argued in *The Structure of Scientific Revolutions*, science advances through paradigm shifts, where failing theories are discarded in favor of frameworks that better explain anomalies.

Having expanded on the hypotheses and research questions, including the research program and stance which I have taken, I will now justify the method(s) which will be used to test the hypotheses and answer the research questions.

To start, it should be acknowledged that the selection of research methods available for a philosophical stance forms a wide range of perspectives, including conceptual analysis, critical theory, deconstruction, experimental philosophy, hermeneutics, Kantianism, methodological naturalism, phenomenology, and pragmatism⁴⁶.

In review of the methods available, I decided that as my research program is cross disciplinary, so, too, should be my method. This is because it would be irresponsible to confine myself to any specific existing perspective, especially one that I might be less experienced in. Moreover, there is already much overlap between different research methods⁴⁷. A more comprehensive, if not fitting, approach, then, would be to use an integrative method that combines the most relevant designs, instruments, and processes. Furthermore, this comprehensive approach is the most appropriate as it also accounts for any possible “better” methods that were overlooked due to being too niche.

My singular method, as a pluralism of other research methods, can be understood as a triangulate approach:

- A. Literary Interpretation: At the core of this approach lies the vital role played by the “hermeneutic circle” in the process of interpretation. I understand this method to not enforce meanings to concepts, but perform an interpretive dialectic back and forth within the literature to allow a matter to come into view. Through this aspect, I will offer comparative explanations of different findings or of connections between various disciplines to build the understanding deeper and fuller. Incidentally, it is also important to note that I recognize that knowledge (herein synonymous with “literature”) has the potential to progress from diverse sources, and credibility is not exclusively restricted to the academic realm. Consequently, platforms like Wikipedia, traditionally perceived as “less than scholarly,” and, dare I say, multimedia platforms like YouTube or Webtoon, will also have been leveraged. This is because, more so than knowledge itself, is the way knowledge is packaged⁴⁸. Therefore, for example, if a video essay on YouTube, “This video will change how you see Eren”⁴⁹, about an anime *Attack on Titan* helps me understand the Schmittian “other” better than a reading guide on “The Concept of the Political”, the fact that the video essay is neither vetted nor is from a scholarly source doesn’t detract from its value in advancing knowledge. As such, my literary interpretation will extend to review of variety of mediums whether they be conventional (books, papers, talks,

⁴⁶ D’Oro, Giuseppina, and Søren Overgaard, eds. *The Cambridge Companion to Philosophical Methodology*. Cambridge: Cambridge University Press, 2017. <https://doi.org/10.1017/9781316344118>.

⁴⁷ Philosophical methods often overlap in unexpected ways. Kantianism’s aspiration to ground philosophy in rational structure mirrors the empiricism of experimental philosophy, albeit from opposite epistemological directions. Similarly, critical theory, though often positioned as radical, shares methodological roots with conceptual analysis.

⁴⁸ The receptivity of an idea often depends less on its content than on its packaging. Advice ignored from a parent may later be praised when echoed by a friend—what changes is not the information but the framing. Contextual cues such as the source’s authority, the emotional state of the recipient, and how the message is framed greatly impact receptiveness. Thus, repackaging information can enhance not only its reach but its impact, underscoring the role of narrative context in intellectual persuasion.

⁴⁹ invaderzz. “This video will change how you see Eren.” YouTube video, 1:26:03. Posted December 23, 2021. <https://www.youtube.com/watch?v=H6GmVCD7cxk>. The video serves as a strong example of non-traditional media conveying nuanced philosophical insight—specifically, using the character Eren Yeager from *Attack on Titan* as a lens to understand Schmittian notions of the “Other” and radical agency.

documentaries, etc.), or not (multimedia platforms, television/movie series, cartoons, comics, etc.)

- B. Ongoing Discussion: In the same way the instrument of the hermeneutic circle is used to stimulate dialectic within and between the literature, dialectic will also extend beyond the literature. This is in recognition that the discussion is not discrete, but continuous. Not to be viewed retrospectively, but progressively. This process will manifest primarily through conversation and debate, both online and offline by way of threads, forums, stack exchanges, and everyday dialogue. As a participant in these conversations under study, in contrast to secondary and tertiary sources, reflexive first-hand experience will lend itself to anecdotal evidence which is not only up to date, but often possesses a practical aspect which other studies that are done in isolation lack.

The potential of Large Language Models (such as ChatGPT) has also allowed for a modified autodidacy whereby ideas can be run past AI with varying degrees of responsiveness. Since writing performs the twofold task of not only helping writers clarify their thinking, but also conveying information, there can be inconsistency by way of ambiguous transference⁵⁰. Running ideas past AI in LLMs allow for the identification of this inconsistency and also prompts reflection into different avenues.

- C. Grounded Theory: Comparable to experimental philosophy, a grounded theory makes use of empirical data, but more than positing hypotheticals, the design of this method emphasizes the generation of theory (including thought experiments) which is 'grounded' in data that has been systematically collected and analysed. It is used to uncover such things as social processes and then analyzing them not in isolation but on a backdrop of real phenomena. Furthermore, using grounded theory as a design within my method also allows for a reference point in which to interpret gathered data.

Put simply, my method converses amongst a wide range of mediums concurrently under a real-world setting.

Outcomes and Results

Priming this segment on outcomes and results necessarily prerequisites a brief breakdown of (1) how the data was gathered and what it is intended to represent; (2) what the gathered data looks like; and (3) how the data should be interpreted.

Of the initial research questions in response to the conclusions drawn from the Background chapter, we can further divide the outcomes to revolve around three mini questions (themes) in respect to three different sections: On the Persona; On Autonomy; and On Free Will. The first section naturally aids in response to the conclusions drawn from the Background with the initial three subsections targeted towards the initial research questions of how to turn our Self as a whole towards the Good; specifically, when the Ego (or at least willpower as an effect of the Ego) alone being an insufficient influence in this turning. In the case of the first section, the overarching question is further divided in the subsections of: Ego; Subpersonalities; and Subconscious Forces. These subsections will display what the gathered

⁵⁰ Writers generate text linearly (horizontally), while readers interpret meaning hierarchically (vertically). This disconnect often results in a gap between intended meaning and received meaning. See: McEnerney, Larry. "The Problem of the Problem." Lecture, The University of Chicago Writing Program at Ohio State University, 2014. YouTube video, 58:02. <https://www.youtube.com/watch?v=H6GmVCD7cxk>.

data looks like and will be followed by a subsection of how the data should be interpreted. The overall breakdown is as follows:

- (1) Initial Research Question: How are we to understand the pluralism of our Persona in relation to acquiring the Good Life?
Section Heading: On the Persona
Subsections:
 - a. Ego
 - b. Subpersonalities
 - c. Subconscious Forces
 - d. Interpretation
- (2) Initial Research Question: To what extent does the Ego possess autonomy?
Section Heading: On Autonomy
Subsections:
 - a. The Allonomy of Autopilot
 - b. The Effect of Subconscious Forces
 - c. The Mechanisms of Control
 - d. Interpretation
- (3) Initial Research Question: To what extent does the Subconscious Forces influence, if not outright control us?
Section Heading: On Free Will
Subsections:
 - a. Deterministic Influences
 - b. The Ego as Compatibilist
 - c. A Drifting Will
 - d. Interpretation

These three sections will necessarily tie the subsections together into an interpretation that addresses the initial research questions. These interpretations will naturally be the outcomes of my research methods as laid out above, and these very outcomes will then lay the foundations of my results as stemming from the sections working in tandem.

On the Persona

In review of the Background, it was shown that our Persona is multifaceted inasmuch as it is split between multiple competing subpersonalities. Our ego – the estimation of ourselves – is not actually our entire “Self”, but rather a small personality shackled to a collective of giant subconscious forces otherwise known as the Psyche⁵¹. The Self itself is a dualism of the metaphysical mind (Psyche) and the physical body. It is absolutely crucial to reiterate that “we”, the metaphysical aspect of the Self that others interact with, is the phenotype of a community of subconscious forces (Psyche), the individual forces themselves (predominantly consisting of the Ego and Subpersonalities) are the genotypes.

⁵¹ In Jungian psychology, the psyche encompasses all conscious and unconscious processes. The “Persona,” in this framework, is the social face an individual presents to the external world—shaped by collective expectations and social norms. This external mask mediates between the individual and society, and does not necessarily reflect the totality of the self. See: Eternalised. “Jordan Peterson VS Friedrich Nietzsche | Is God Dead?” YouTube video, 20:36. Posted May 24, 2020. <https://www.youtube.com/watch?v=UKJlxvOOanM> (see 18:22).

One major confusion in understanding the Persona is linguistic, particularly the multiple ways the word “personality” is used⁵²; this is why “multiple personality disorder” is now officially “dissociative identity disorder”. In grounded psychopathology, the word “personality” often means more than its everyday definition. It frequently includes not only a person's traits and characteristics, but aspects that the layperson would categorize as “identity” - Self-awareness, sense of Self, name, age, gender, etc. Moreover, a person's Persona may appear to change depending on who they're interacting with (code switching⁵³), but their personalities - as psychopathologists define it - remains stable.

The Persona is also commonly understood as a tripartite: Ego, Subpersonalities, and other Subconscious Forces. The Ego, as the conscious element, is characterized by the Task-Positive Network (TPN) and is generally responsible for “public relations”. The Ego is also who we mistakenly believe ourselves to be. Subpersonalities are generally distinct, or devoid of their own egos in fully integrated and harmonized psychologies⁵⁴. They are the Subpersonalities that “front” or substitute when switching to the Default Mode Network (DMN). Subconscious Forces are also characterized by the DMN but are generally non-unconscious entities. They are anything that influences the behavior of the Ego and Subpersonalities.

Ego

Standing at odds with the Cartesian notion of thinking⁵⁵, the person you see yourself as is not you, but your Ego. In other words, the personality that is currently consciously reading this sentence, the “you” that is deliberately paying attention and interpreting these squiggles into meaningful words, is defined as “Ego”.

Our ego is only a small part of our Persona⁵⁶, which is why even when taking into account code switching between our friends, family, and strangers, the people that surround

⁵² Quora Community. “Do Normal People Have a Single Personality or Multiple Personalities Depending on the Situation and Role They Are In?” Quora. Accessed September 30, 2023. <https://www.quora.com/Do-normal-people-have-a-single-personality-or-multiple-personalities-depending-on-the-situation-and-role-they-are-in>. Although informal, this entry offers illustrative perspectives on the fluidity of personality expression across contexts, lending anecdotal support to the argument that the Persona adapts socially while maintaining internal multiplicity.

⁵³ Code switching, in the context of personality, refers to the fluid adaptation of behavior, language, or demeanor depending on social context. Just as a bilingual speaker switches languages, individuals adjust “personality modes” when interacting with parents versus friends. This phenomenon illustrates how the Persona operates as a dynamic interface between internal traits and external expectations—highlighting the modular and adaptive nature of identity.

⁵⁴ Scott Barry Kaufman emphasizes the importance of self-integration for personal development and well-being, arguing that embracing all facets of ourselves—rather than suppressing or compartmentalizing them—is vital for holistic growth. He frames integration not as a one-time achievement but as a dynamic, ongoing process toward wholeness.

After Skool. “The Neuroscience of Personality.” YouTube video, 38:20. Posted September 15, 2021. <https://www.youtube.com/watch?v=CKrfCtDbkvw>.

See also: *The Daily Motivation*. “The Importance of Self-Integration | Scott Barry Kaufman.” Podcast audio. Canadian Podcasts. Accessed September 30, 2023. <https://canadian-podcasts.com/podcast/the-daily-motivation/the-importance-of-self-integration-scott-barry-kau>.

⁵⁵ The Cartesian model of thinking, drawn from Descartes’ *Cogito ergo sum*, posits that thought is the essence of being. This view entails a dualism between the immaterial mind and the material body. “Cartesian materialism” later evolved as a critique, referring to the idea of a centralized “theater” of consciousness in the brain. While Descartes never used the term, modern philosophers use it to critique overly simplistic models of unified consciousness.

⁵⁶ This video outlines the concept of a “psychological black death”—a term used to describe the mass psychological suffering resulting from disconnection, purposelessness, and alienation in modern life. It examines historical and philosophical sources of the crisis, including Nietzsche, Jung, and Kierkegaard.

us will fail to recognize our self-characterizations. i.e. if you write an estimation of your own character and you show it to your friends, your friends will likely laugh at how off it is.

We are also strangers to ourselves not only because we mistakenly associate our ego with who “we” (the word ‘we’ – generally a plural – is here functioning as a singular entity comprising a pluralism of personalities, also known as “ourselves”) actually are, but also, as seen in the Background, confuse ourselves over how much control over the “we” (Self) we (Ego) have. For instance, the reason you get angry at your wife even when you know she doesn’t deserve it or you can’t write the fantasy novel you’ve always wanted to is that your Ego doesn’t have control over the Subconscious Forces just in the same way we don’t have control over other people.

In the microcosm of “We” (Psyche) as a Being with multiple personalities, the Ego can be best understood as the spokesperson, navigator, or better yet, pilot for ourselves. The representational Ego adopts this role as it is often the most developed of the personalities in terms of rational capacity, but in “bad psychologies” or where Subpersonalities are equally as developed (Alters), there can be multiple Egos. The Ego, then, should be understood a positional role as opposed to a certain personality trait (ego). This means, particularly in bad psychologies, the position of the Ego can be usurped by any other subpersonality that has developed an ego of their own. To reiterate, an ego (uncapitalized) is a developed rational capacity which gives rise to something akin to “self-awareness”, whereas The Ego (capitalized) is a positional and representational role.

The Ego is bound to, and only manifests, or “fronts”⁵⁷ when the TPN is engaged. The TPN is engaged only when we are paying attention in response to default incapability or observed anomaly (i.e. an internal conflict). Since willpower is a substrate of the ego, the TPN can be exemplified in moments of self-control such as emotional regulation or when dealing with stress. For example, the “scare” back into reality when mindlessly driving and a rabbit suddenly jumps into the road is a characteristic fronting of the TPN. Incidentally, the TPN can also be understood as the hallmark of compatibilism; the system that allows for overriding the Subconscious Forces (i.e. the Ego having decided to go on a diet and catching the “we” reaching for a Dorito and either affirming the decision or choosing otherwise).

Subpersonalities

Subpersonalities are the other denizens, or “actors” within the Cartesian Theater⁵⁸. Unlike the Ego, Subpersonalities are generally not fully formed personalities, with the exception of Alters which are often misinterpreted as “alter egos”⁵⁹. It is important to note,

Academy of Ideas. "The Psychological Black Death." YouTube video, 15:42. Posted March 20, 2022.

<https://www.youtube.com/watch?v=s0K0sI5XiP4>.

⁵⁷ "Fronting" refers to the process in which different subpersonalities—or alters—take control of a person’s conscious awareness and behavior. In the context of dissociative systems, fronting does not require awareness from the primary Ego and may occur involuntarily.

⁵⁸ Daniel Dennett’s Multiple Drafts Model of consciousness challenges Cartesian materialism, rejecting the idea of a single, central experiencer. Instead, Dennett proposes that consciousness arises from parallel processes (“drafts”) competing for expression, with no privileged central narrative. This model suggests a decentralized, distributed mind that undermines the idea of a unified self.

Dennett, Daniel C. *Consciousness Explained*. Boston: Little, Brown, 1991.

See also: “Consciousness: More like Fame than Television” in *Die Technik auf dem Weg zur Seele* (1996); “The Myth of Double Transduction” in *Toward a Science of Consciousness II* (1998); Kinsbourne, Marcel. “Time and the Observer: The Where and When of Consciousness in the Brain.” *Behavioral and Brain Sciences* 15, no. 2 (1992): 183–247. <https://doi.org/10.1017/S0140525X00068229>.

⁵⁹ There’s often confusion between an “alter ego” and an “alter.” An alter ego is typically a conscious projection or role (e.g., a cosplay persona) constructed by the Ego. An alter, in contrast, is an autonomous subpersonality with its own perceptions and behaviors, often unaware of the Ego or other alters. Alters may front without

however, that Alters (also known as “others”, “insiders”, and “parts”) are distinct from the other Subpersonality subset. Moreover, Alters themselves are completely different from alter egos. An alter ego is more like if you do cosplay; the Ego creates and acts out a character. You aren't actually that character but you are acting as if you are. You can place certain values, strengths or characteristics into that alter ego. Maybe that will allow you to feel that strengths and confidence and can be helpful in the long run. Alters, on the other hand, are entities often oblivious to the Ego, let alone parts to be controlled. They aren't sitting in a room waiting to be called on. They interpret situations differently than the Ego might, moreover, they can react and/or front absent the ego's “permission” to do so.

Neuroscience has theorized Alters, as the prime subset of Subpersonalities, to be created to protect the general Self or “We”⁶⁰, having begun as bundles⁶¹ or separate states: hunger, attachment, safety etc. but later on - particularly in “bad psychologies” or where Subpersonalities are equally as developed as the Ego (Alters) - developed their own ways of perceiving (or, better yet, engaging)⁶² with the world. Their genotypic reactions are typically “theirs”; but are phenotypically interpreted as “ours”. I.e., for those with “bad psychologies”, the Ego is not to blame for actions of Alters, but is still held responsible for them.

If not having evolved an ego of their own to become Alters⁶³, Subpersonalities are bound to the subconscious, and only front when the TPN reverts to the DMN. Subpersonalities are responsible for keeping the body in operation when the Ego “checks out” and shifts into DMN⁶⁴.

As the binary opposite to the TPN, the DMN can be likened to the “screen saver” mode of consciousness. The DMN is the concentric (Reptilian or Primal) domain of the subconscious and includes all mindless activities such as “mindless scrolling” and muscle memory activities such as brushing one's teeth. One hallmark of the DMN is “autopilot”⁶⁵:

conscious permission or coordination, especially in dissociative systems.

Quora Community. "Do I Have an Alter Ego or Is It Just Another Side of Me?" Quora. Accessed September 30, 2023. <https://www.quora.com/Do-I-have-an-alter-ego-or-is-it-just-another-side-if-me>.

⁶⁰ Antonio Damasio's Somatic Marker Hypothesis posits that bodily feelings (“somatic markers”) play a critical role in decision-making. These emotional reactions guide us by attaching value to choices based on past experience, especially in situations of uncertainty. Disruption of these processes—such as damage to the prefrontal cortex—can severely impair decision-making and social behavior.

Damasio, Antonio R. *Descartes' Error: Emotion, Reason, and the Human Brain*. New York: G.P. Putnam's Sons, 1994.

⁶¹ Hume challenged the notion of a unified self, arguing instead that the self is “nothing but a bundle or collection of different perceptions” constantly in flux. This view undermines the idea of a persistent ego and supports the thesis's claim that what we take as a stable identity is instead an ongoing psychological illusion. Hume, David. *A Treatise of Human Nature (I, IV, §VI)*. Edited by David Fate Norton and Mary J. Norton. Oxford: Oxford University Press, 2000 (originally published 1739–40).

⁶² Gazzaniga's split-brain research suggests consciousness is not unitary but distributed across competing neural networks. Despite the hemispheric disjunction, patients still experienced a unified sense of self, pointing toward the Ego's interpretive role in post-hoc narrative construction. This sheds light on how individuals like Phineas Gage could maintain continuous awareness despite significant frontal lobe damage—possibly due to an Alter fronting in place of the Ego.

Gazzaniga, Michael et al. Various articles between 1962–1975. See: “Dyspraxia Following Division of the Cerebral Commissures,” *Archives of Neurology* 16 (1967); “Split Brain in Man,” *Scientific American* 217 (1967); de Haan et al., “Split-Brain: What We Know Now,” *Neuropsychology Review* 30 (2020).

⁶³ A key distinction between normative psyche and dissociative identity disorder (DID) lies in whether the Alter has developed its own ego. In pathological cases, such Alters are no longer confined to the subconscious but vie for conscious control.

⁶⁴ Transactional Analysis, as proposed by Berne, parses behavior into ego states—Parent, Adult, and Child—which correspond broadly with different brain network activations (e.g., TPN and DMN).

Berne, Eric. *Games People Play: The Psychology of Human Relationships*. New York: Grove Press, 1964.

⁶⁵ The phenomenon of “autopilot”—where one acts without conscious awareness—can be explained neurologically by the shift from the Task-Positive Network (TPN) to the Default Mode Network (DMN), which

you aren't in "autopilot" right now, but you will be 5 minutes after you get bored and your eyes keep scanning the text but with no comprehension going on ("mindless reading"). One final thing to note is that the Ego and Subpersonalities are not necessarily discrete entities. Neuroscientists and Psychologists can confirm that they act independently, but the permeable line where one shifts into the other is unclear. To draw an analogy, we are certain blue and green are two discrete colour codes, but where exactly we draw the line between the two is ambiguous.

Subconscious Forces

Subpersonalities are not the only things that reside within the subconscious. Although there is yet to be substantial evidence for other denizens and/or bundles (such as the Jungian Shadow, Anima, and Animus) within the Cartesian Theater, it is speculated that there are other Subconscious Forces that influence the behaviour of the Ego and, especially Subpersonalities. Not all of these are native to the subconscious (or even unconscious⁶⁶) and some arrive as external stimuli.

With the advent of Surveillance Capitalism⁶⁷ one potent example from outside the Self is advertisements. While advertisement can influence the behaviour of the Ego directly, for example by introducing a product/service which is rationally a superior good, advertisements also impact the Ego indirectly, for example, breeding feelings of insecurity through subtle body shaming cues from fashionwear. Much in the same way, advertisements can influence Subpersonalities directly (by appealing to specific desires such as sugary foods) and indirectly (by appealing to abstract desires such as hunger in general) when they "front".

A different example for Subconscious Forces from inside the Self is gut bacteria⁶⁸. Since the gut is oft claimed to be the "second brain", the influence that gut bacteria have on both the Ego and Subpersonalities, directly and indirectly is quite considerable. This can be seen with veto like power when stimulating the need to vomit, or more subdued with certain food cravings.

Interpretation:

When we first set out on this first section *On the Persona*, the idea of the Self as multifaceted was rather vague. Now, however, we have a clearer picture of the tripartite Persona: Ego, Subpersonalities, and other Subconscious Forces⁶⁹. The Ego (not to be confused with an ego) is the pilot of the Self under TPN, Subpersonalities autopilot the Self

is governed by Subconscious Forces.

Reddit thread: "What Happens When Your Brain Goes on Auto-Pilot?" r/explainlikeimfive (Accessed September 30, 2023).

⁶⁶ The "unconscious" and "subconscious" are often conflated. In this thesis, the subconscious includes any force the Ego is unaware of but that may still influence behavior. The unconscious, by contrast, is deeper, more inaccessible, and includes phenomena like sleepwalking or trance states.

See Jung's framework: *Collected Works of C.G. Jung*, 20 vols. Princeton University Press, 1953–1979.

⁶⁷ Surveillance capitalism commodifies human experience through data extraction and predictive algorithms. These deterministic forces influence behavior without conscious awareness, and are prime examples of Subconscious Forces shaped externally.

Zuboff, Shoshana. *The Age of Surveillance Capitalism*. New York: PublicAffairs, 2019.

⁶⁸ The gut microbiome, as part of the extended phenotype of the Self, plays a role in behavioral predispositions. Gut flora can trigger cravings and influence decision-making through the gut-brain axis, operating beneath the threshold of conscious awareness.

Alcock, John, et al. "Is Eating Behavior Manipulated by the Gastrointestinal Microbiota?" *BioEssays* 36 (2014).

⁶⁹ See Appendix A: Misc., Exhibit 2.

under DMN, and Subconscious Forces influence the Self to varying degrees under both TPN and DMN. Furthermore, with evidence supporting the Persona as only a phenotype of the Psyche as well as the metaphysical aspect of the Self, we really start to see how miniscule our Ego – the estimation of ourselves – is in the microcosm of the Psyche, and how much smaller still in the multifaceted Self.

Returning to one of the original questions, then, we are now in a much better position to understand the pluralism of the Self in relation to acquiring the Good Life. Namely, that **our Persona/Self is more than our ego**; the infinitesimal nature of the Ego supports the previous conclusion that increasing willpower is not an effective solution for achieving the Good Life. Moreover, it hints at the idea that weakness of will is not only a misnomer since the influence of the will in and of itself is limited, but also that, leading from this, akrasia should be attributed to a more, if not the most influential player, which at this point appears to be the Subconscious. This naturally segues into the subsequent research question of “to what extent does the Psyche possess autonomy?” After all, if the Ego is only a small stakeholder in relation to acquiring the Good Life, then it would make sense to move on to bigger fish – that is, if they truly do have that much more influence and/or control over our lives. Moreover, if the bigger fish also possess a larger portion of control, then it is they that must be ostensibly held accountable for failing to obtain wealth, fame, and power, and ultimately the Good Life.

On Autonomy

The question of control is synonymous with the question of autonomy. In the previous section, it was hinted that we – our Ego – is an influential, but ultimately small stakeholder not only of the Self, but of the Persona as well. Moreover, it was hinted that, if not the ego, then by process of elimination, the largest stakeholder must be the Subconscious. But time under fronting does not necessarily equate to time under control.

One major confusion surrounds understanding autonomy; not only is the Ego not in charge, but neither is the Subconscious really in charge either. If we return to the metaphor of the TPN and DMN being different pilots, then if a destination is not agreed upon, neither has complete control over the flight path even if one pilot is in the cockpit longer than the other. Moreover, unlike the TPN which is solely inhabited by the Ego (with the exception of bad psychologies), the DMN is a communal home to Subpersonalities in addition to other denizens and/or bundles which further dilutes control of the Self across many more stakeholders.

If we focus only on the Persona, we notice that the composition of autonomy is shared between the Psyche of the Ego, Subpersonalities, and other Subconscious Forces. As the most obvious avatar, just as we mistake our Ego for the estimation of ourselves, we too mistake how much control (if truly any) the Ego has over the Self. While we like to think our Ego has the largest share of autonomy, evidence suggests that what we really have is but an illusion of control. Incidentally, under the DMN, autonomy largely appears to be in the hands of the Subconscious. In this capacity for control, it would thus appear that hindering obtaining exogenic subsidiaries must only be an act of malicious self-sabotage by the Subconscious. This, however, is only partly true due to lack of malicious intent. Despite this lack of intent, the effect of Subconscious Forces still result in self-sabotage.

The Allonomy of Autopilot

The representational Ego, insofar as it adopts the role of spokesperson, navigator, and pilot would seem to be the largest shareholder in the game of autonomy, that is control over

the Self or We. After all, even though most, if not all of us, would agree that we are not always consciously present, it would still seem that we were present for a large part of the day and thus have the largest share of the pie.

This belief that “time under fronting = time under control” is consolidated by the fact that for the majority of people, Subpersonalities are entities often oblivious to the Ego. If Subpersonalities fronted for any great length of time, it would ostensibly be common knowledge⁷⁰. It makes intuitive sense, then, that if we - the Ego - were not present for a large part of the day, surely, we would notice by some indirect method. For instance, we could notice either by someone else pointing it out directly or by some form of spatial awareness such as being conscious one moment that the sun is up, and suddenly being aware that the sun had set the next minute.

But we do notice, we just don't linger enough for the observation to register an impact. i.e. saying we will just check up on social media for five minutes and then the next thing you know, it's been five hours and you can't remember what the last post/image/video was about. Moreover, just as you only know you fell asleep after waking up, shifting from TPN to the DMN is only detectable in hindsight. If we notice but don't do anything about it, then, in conjunction with the Spotlight Effect⁷¹, others are even less likely to point out an anomaly. Furthermore, unlike with dissociative identities where fronting can last for anywhere from minutes to decades⁷², the diaspora of TPN to DMN switching in those with harmony of integration (or associative identities, AI) is often much more frequent, shorter, and with a more or less even ratio. For instance, switching could occur in five-minute intervals which is a time lapse usually unnoticeable.

With this in mind, it shouldn't be much of surprise that studies reveal that the belief that we have general autonomy (as it relates to the General Will) is misguided. Depending on how loosely autonomy is defined⁷³, we spend anywhere from 46.9% of our waking hours⁷⁴ to 95% of our day on autopilot according to the Max Plank Institute⁷⁵. With this in mind, it would seem that it is really the Subconscious (Psyche sans Ego) that is really in control of us. Furthermore, this would mean that the allonomous Subconscious is largely morally responsible for obtaining, or failure thereof, the exogenic subsidiaries. In other words, it is the Subconscious that hinders obtaining wealth, fame, and power, and ultimately the Good Life. Conventionally, this process is understood as acts of “self-sabotage”.

⁷⁰ DID is often diagnosed when Alters do not harmoniously integrate. By contrast, most people exhibit associative identities (AI), which are better integrated and exhibit smoother transitions between fronting subpersonalities. The difference is qualitative, not categorical.

Pirangy, “*The Syzygy: Anima and Animus*,” *Medium* (2023).

Huntjens, R.J.C. et al. “*Procedural Memory in DID*,” *Consciousness and Cognition* 14 (2005).

Also see *Chasing Tails*, Ep. 56, by T. Lee, Webtoon (Accessed September 30, 2023).

⁷¹ The spotlight effect—our tendency to overestimate how much others notice us—is shaped by egocentric bias. Given that even we misunderstand ourselves, it's even more unlikely others see through our internal discord.

⁷² Psychological integration involves achieving coherence among one's various subpersonalities. True integration reduces internal friction and improves adaptability.

Quora Community: “*Understanding Psychological Integration*” (Accessed September 30, 2023).

⁷³ Although we are the primary stakeholder of our Persona—being the one who fronts most frequently—we are still influenced by social conditioning and external forces. Autonomy, then, is not ownership but governance, often provisional and fragile.

⁷⁴ A wandering mind—when the Ego is no longer fronting—correlates with unhappiness. This supports the claim that conscious control (TPN) is not only limited but energetically expensive.

Killingsworth, M.A., and Gilbert, D.T. “*A Wandering Mind Is an Unhappy Mind*,” *Science* 330 (2010).

⁷⁵ Up to 95% of daily actions occur on “autopilot.” Most of our life is not governed by deliberate decisions but by habitual Subconscious Forces.

As cited by: Rodgers, Elizabeth. “*How Much of Your Day Do You Spend on Autopilot?*” *Dr. Elizabeth Rodgers Psychology*.

At first glance, we may be so inclined to question why our Subconscious would self-sabotage acquiring the Good Life. After all, obtaining the Good Life would undoubtedly benefit the Subconscious as well. Thereby, it would seem that by identifying the motivation behind Subconscious actions, if not opening a direct channel for communication, then at least allowing for cooperation indirectly through appeasing subpersonalities' base desires⁷⁶, otherwise known as damage control.

The Dorito Effect⁷⁷ provides an apt case study which exemplifies the phenomena of why our Subconscious undermines progress. Specifically, it elaborates on why we have that bizarre sensation of putting another Dorito in your mouth as if someone else is feeding you - which we know now is because someone else (Subpersonalities as allonomic entities) is doing exactly that. The reason, then, why we find ourselves at the very bottom of the bag⁷⁸ after just having resolved to go on a diet is because the artificial flavourings hijack our reward system⁷⁹. In a broader sense, the artificial flavourings, as a vector of Subconscious Forces, influence our Subconscious by appealing to its base desires.

The effectiveness of Subconscious Forces, such as the Dorito Effect, shouldn't be underestimated. From the Ego perspective, mindful eating appears so easy to engage in that when our Subpersonalities polish off an entire bag of Doritos when we've just resolved to go on a diet, our Ego can't help but see it as an act of betrayal. This is because knowing that the Self (which naturally includes the subpersonalities) benefits from dieting, self-sabotage can only be a calculated, if not deliberately malicious, move. But this logic is only possible from the Ego perspective which is privileged with a rational capacity.

We should remember what qualifies the Ego (capitalized) from an ego (uncapitalized) is the development of rational capacity. Subpersonalities, insofar as they lack an ego, lack this rational capacity and all the abilities that come with it. Furthermore, insofar as Subpersonalities are bound to the DMN, the extent of its deliberation capped at the Reptilian Brain, meaning that when met with Subconscious Forces (or otherwise "Psychic Noise"), Subpersonalities effectively act under duress⁸⁰. We see then that not only is the Subconscious entirely incapable of even comprehending the notion of dieting, let alone betraying the resolve, but also there can be no "active"⁸¹ self-sabotage (with the exception of Alters) since that too would require the rational capacity. What we perceive as self-sabotage, then, is simply the result of cognitive dissonance between the Ego and Subconscious where the Ego is turned towards the good but the Subconscious has been turned elsewhere.

Nonetheless, while our Subconscious may not actively, or intentionally contribute to self-sabotage, that does not mean it does not do so passively. The Subconscious still hinders

⁷⁶ Controlled indulgence (e.g., a scheduled "cheat meal") can prevent binge behavior by preemptively engaging Subconscious Forces with an acceptable outlet.

⁷⁷ The Dorito Effect refers to hyperpalatable artificial stimuli that override the rational Ego by appealing directly to the reward system. As such, it is a prime example of how Subconscious Forces hijack decision-making.

⁷⁸ Corporations like PepsiCo bank on the predictability of self-sabotage. They would be more alarmed if consumers didn't finish the bag, as that would signal a failure in their behavioral conditioning.

⁷⁹ The Reptilian Brain's reward system evolved to respond to caloric density and flavor. Hyperstimulating food hacks this ancient mechanism, bypassing conscious regulation.

⁸⁰ The Marshmallow Test does not measure willpower alone but reveals the interaction between Subconscious Forces and context (e.g., socioeconomic status). Children who "fail" the test are not weak-willed, but operating under different internal and environmental constraints.

Tulleken, Chris van. Ultra-Processed People. W.W. Norton & Company, 2023.

⁸¹ Action, at least in the rational sense, presupposes capacity for reflection and evaluation. Philosophy of mind distinguishes between passive response and active agency, suggesting that the Ego—endowed with rationality—can override base impulses.

The Internet Encyclopedia of Philosophy: "Thomas Reid: Theory of Action."

obtaining wealth, fame, and power through inaction and/or being led astray (effortless action⁸²) by Psychic Noise. We see then that attempting to parlay with the Subpersonalities would prove trivial as the Subpersonalities themselves are but mere puppets (akin to Platonic Drones⁸³) to Subconscious Forces.

The Mechanisms of Control

The Ego, then, only has an illusion of control and the Subpersonalities are only puppet pilots with no will of their own. Seeing as cooperation with the Subpersonalities would prove ineffective, the research question that naturally follows relates to how; how does our Subconscious self-sabotage? In the absence of why, understanding how our Subconscious is enabled to engage in suboptimal decisions can also support the drafting of a solution. If, for instance, we liken our Subpersonalities to impotent children, then by understanding the triggers of how their incompetent behaviour is an exploitation of their ignorance (as seen with PepsiCo⁸⁴), we can hope to block, reverse, and/or circumvent the “Dorito” mechanism that enables self-sabotage – a “child lock”, if you will.

Earlier it was briefly touched upon how Subconscious Forces and Psychic Noise influence both the Ego and Subpersonalities through the vectors of advertisements and gut flora. Furthermore, in the subsection just above, we expanded on the influence taking affect in the Reptilian Brain; specifically through the Dorito Effect as it exemplifies the effect of Subconscious Forces on the reward system. The next step then in understanding the mechanism of how Subconscious Forces effects the psyche starts with revealing the connection between the reward system and the Reptilian Brain.

The Reptilian (or Primal) Brain, as alluded to in discussion of the DMN, is the concentric domain of the subconscious, being the innermost circle of the Triune Complex⁸⁵. The Triune Complex has been largely dismissed as a myth in the fields of evolutionary and developmental neuroscience, however, the idea that there exists a “basal ganglia” (or some other amygdala network for our primal instincts) at the very core of our being (though not necessarily schematic) is widely accepted in the scientific community, if not uncontested. In other words, stripping away our conscious and limbic systems, at the very base of the Self is a survival mechanism which is colloquially known as the Reptilian Brain.

In short, the Reptilian Brain can be likened to the control center of our genes⁸⁶. But unlike the horizontal coexistence between the Ego and the Subpersonalities or even the conscious and the unconscious, the coexistence between the Reptilian Brain and the psyche is vertical. The Reptilian Brain exists on a different hierarchy – a landlord to the tenants of the Ego and Subpersonalities who have little to no bargaining power. For the most part, this

⁸² Big Think. “The Neuroscience of Free Will.” YouTube video, 15:42. Posted January 18, 2023. <https://www.youtube.com/watch?v=NvZi7ZV-SWI>.

⁸³ Plato. *Republic*. Translated by G. M. A. Grube, revised by C. D. C. Reeve. Indianapolis: Hackett, 1992. (Originally written ca. 375 BCE.) The metaphor of the “drones” appears in Book VIII as part of Plato’s discussion of societal degeneration.

⁸⁴ This thesis holds no animosity toward PepsiCo; the invocation of the Dorito Effect simply builds on a popularly accessible example. If anything, this footnote serves as a playful disclaimer. While such corporations may act as vectors for Psychic Noise and Subconscious Forces, they are also merely responding to consumer demand—fitting within the framework already laid out. Should PepsiCo or its competitors wish to fund future research, I remain open to ethically sound collaboration.

⁸⁵ Wikipedia Contributors. “Triune Brain.” *Wikipedia*. Last modified August 10, 2023. https://en.wikipedia.org/wiki/Triune_brain.

⁸⁶ Dawkins, Richard. *The Selfish Gene*. Oxford: Oxford University Press, 1976. Our survival mechanism—embodied by what is colloquially called the Reptilian Brain—serves primarily to preserve genetic material. Genes don’t “care” in the anthropomorphic sense; rather, they operate by incentivizing behavior through the hedonic reward system. This neural landlord maintains indifference unless survival is threatened.

landlord is indifferent to the whims of the Psyche inasmuch as it largely does not interfere with their “libertarianism”⁸⁷. Yet to ensure the carrying out of its prime directive, the Reptilian Brain will incentivize the Psyche through the so-called reward system and, in life-threatening scenarios, can even execute martial law: evicting the tenants and overriding all unnecessary functions for survival in the immediate moment.

In regards to how this all relates to Subconscious Forces; we see the survival mechanism manifest as “needs”. As biological entities; we are needy creatures and the fulfillment of our needs corresponds with a more or less⁸⁸ hedonic reward system. This can be easily observed with the near universal phenomena of “hangriness”⁸⁹ where the Reptilian Brain will employ a myriad of Subconscious Forces to incentivize the consumption of food. Another way that the Reptilian Brain extends its “invisible hand”⁹⁰ is something we’ve already talked about: autopilot.

The mechanism of how autopilot works has already been alluded to – the shifting of the TPN to the DMN. The reason why this shift even occurs is because, like a landlord, the Reptilian Brain wants to save as much money on the electricity bill as possible. In other words, shifting of the TPN to the DMN (conventionally known as “Ego depletion”) is an effect of the survival mechanism – to be as energy efficient as possible. Here it is of the utmost importance to reiterate that it is not we of the tenants – the Ego – that has any bargaining power at all; it is the Reptilian Brain that withholds all veto power. The very existence of our Ego is due to the mercy (and perhaps trivial usefulness) of the Reptilian Brain, our continued privilege of minor autonomy (which we have inflated to a delusion of being the “master of our own temple”) only because of the Reptilian Brain’s dormant indifference.

It is this very indifference that allows for both cancer and cure. On one hand we see how PepsiCo amongst other Psychic Noise stays under the indifferent radar of the Reptilian Brain - the Dorito Effect exploits the reward system by aligning intense flavour with its base desire of (hyper)palatable food. Yet on the other, by taking advantage of the same mechanism, we can create an external “child lock” change that blocks, reverse, and/or circumvent the process of self-sabotage.

Interpretation:

From this second section *On Autonomy*, we’ve seen that time under fronting does not necessarily equate with time under control; a tenant spends the most time living in a house, but it is the landlord that withholds all veto power. The ego has only the illusion of control, the Subconscious just a puppet without any intent of its own, and, perhaps most of all, the Reptilian Brain is easily bought so long as the vector aligns with its prime directive. Taken

⁸⁷ The term “libertarianism” is used loosely here to refer to a form of incompatibilism in which both the Ego and Subpersonalities appear to make choices unimpeded by external compulsion. This appearance of freedom, however, is largely illusory, as later discussed in the part *Poiesis*.

⁸⁸ Genes express imperatives rather than desires. Their aim is not pleasure per se, but functional survival. Still, in practical terms, these imperatives often resemble hedonistic calculus—pleasure and pain operate as proxies for evolutionary success.

⁸⁹ The term “hangry” (hungry + angry) captures a survival signal. In evolutionary contexts, flavor intensity often tracked caloric density—hence our innate attraction to sweet or fatty foods. The modern mismatch between flavor and nutrition (e.g. artificial flavorings) exacerbates this signal, contributing to dysfunctional reward circuits.

⁹⁰ A shorthand analogy: Laozi + Adam Smith. This juxtaposes non-interventionist wisdom with market self-regulation, relevant here for understanding how the Reptilian Brain delegates autonomy while still pulling the strings of survival.

together, this means that, if compatible with the prime directive, subconscious forces are in control most of the time.

Now, returning to the research question of the extent of which the Psyche possess autonomy, we see that the value of this answer relates to unmasking the true perpetrator behind our self-sabotage. To this effect, we see that it's neither our Subpersonalities, nor our genes, but, as foreshadowed, those who hijack our reward system with Psychic Noise and pervert our survival instinct to turn towards their products: large corporations. In other words the reason we don't really have autonomy, the reason we have hindered obtaining exogenic subsidiaries, and the reason why we fail to obtain the Good Life is because Subconscious Forces lead us astray. This can be naturally occurring, or deliberate in the way of large corporation, but again, this leading astray not necessarily out of malice but simply in response to consumer demand⁹¹.

It would seem then that one answer to turning the Self as a whole towards the good is not necessarily increasing willpower but "child-locking" the mechanisms behind our suboptimal decisions and decreasing margin for error, i.e. a boycotting of the corporations involved. But that would be too easy, wouldn't it? Indeed, While corporations do engage in evil of the commons⁹², they are simply a middle man. In the spirit of a Scooby-Doo triple unmasking⁹³, we thus see that corporations only exploit and feed into the true mastermind behind our self-sabotage: deterministic influences. What then are these deterministic influences? How does causal determinism influence the directionality of Subconscious Forces? And to what extent does the Subconscious and/or subconscious forces influence, if not outright control us?

On Free Will

In all respects, the Reptilian Brain is a devout, but not a maximally efficient landlord; it is dedicated to its prime directive and has been successful to the point of survival to this day and age, yet there is reason to believe that other than the bare minimum, the Reptilian Brain does not (and possibly cannot) invest in the Good Life in contrast to a mediocre one. Why this is the case, however, does not matter nearly as much as the implications of its impressionability: whoever influences the Reptilian Brain controls the Self. And, perhaps dauntingly, these influences have vied for control even before we were born.

At first these influencers might seem to be an evil mastermind, and while this isn't too far from the truth in the age of surveillance capitalism, (deterministic) influences themselves are simply a natural phenomenon with no agenda of its own⁹⁴. If anything, perhaps we should be grateful for deterministic influences because they necessarily inform our worldview; without external stimuli we would have no point of perspectivism to navigate our lives. In any case, it is impossible to detach ourselves from the facticities of our contextual placement as it is precisely epistemic endosmosis and contextualism that enables us to engage with the world.

⁹¹ As discussed in *Ultra-Processed People*, corporations exist to fulfill demand, not to create it ex nihilo. If Doritos weren't profitable, they wouldn't be made. Consumer desire is a deterministic influence, which—once aligned with addictive products—creates a feedback loop: more production → more consumption → more demand.

⁹² Wikipedia Contributors. "Tragedy of the Commons." *Wikipedia*. Last modified September 20, 2023. https://en.wikipedia.org/wiki/Tragedy_of_the_commons.

⁹³ *Scooby-Doo! Camp Scare*. Directed by Ethan Spaulding. Burbank, CA: Warner Home Video, 2010. DVD.

⁹⁴ "There are no bad instincts. Only bad environments." —Dan Sadler, personal communication to the author, September 30, 2023.

To understand the extent to which Subconscious Forces influence, if not outright control us as a result of their interactions with the Drifting Will and antecedent prior causal events, is to explore the idea of Free Will, determinism, and libertarian accounts of compatibilism. Insofar as we interact with a deterministic multiverse, all of our actions are necessarily influenced by the events prior, regardless of whether external (unintended) or internal (intended), influence both the TPN and DMN; given a specified *way things are at a time t*, the way things go *thereafter* is *fixed* as a matter of *natural law*. Nonetheless, granted that all events are causally determined, an Ego's predisposition does not equate to its condition. Insofar as everything is determined, we can't choose to escape fate, but we can choose which fate to ascribe ourselves to. Moreover, while "free" aspects ostensibly have autonomy, they have very little influence on the directionality of the General Will. Our Ego is thus held captive by external constraints of the Subconscious (Psyche sans Ego) and is at the mercy of the Drifting Will.

Deterministic Influences

The law of causality states that, even with retrocausality⁹⁵, for every effect, there must be a first cause. This idea that every event is necessitated by antecedent events and conditions together with the laws of nature is deeply connected with our understanding of the physical sciences and their explanatory ambitions. But causal determinism is plagued by theories of what causation is from any conclusions about the success or failure of determinism⁹⁶. When we talk about determinism then, we shall use the SEP definition: Determinism is true of the *world* if and only if, given a specified *way things are at a time t*, the way things go *thereafter* is *fixed* as a matter of *natural law*⁹⁷.

Determinism is essentially the manifestation of causality. For instance, gravity is explained by the natural law to be such that an object will be pulled towards a center of mass. If at time *t* a bowling ball is held over the Eiffel Tower, then *thereafter*, it will necessarily fall towards the Earth. Moreover, this fate is *fixed* – the specific directionality has some leeway in that other than intended causation (assuming that the ball is held over the Eiffel Tower so that it may fall) the ball might also experience unintended causation (i.e. bounce off the tower or a gust of wind might shift its spin). Nonetheless, regardless of whether the causation was intended or unintended, premediated or spontaneous, controlled or uncontrolled, determinism claims that as a matter of natural law, the ball will always go down and never up. Other instances can be observed in how dominoes tumble, hot food becomes cold, and other general manifestations of entropy. There are of course theorized potential exceptions, the most generally accepted theory including quantum jumps⁹⁸, but it would be intellectually dishonest to claim that because exceptions exist to a rule, the rule is no longer relevant⁹⁹.

The same applies to self-sabotaging behavior. For example, returning to the Dorito Effect, if we never tasted a Dorito, we would not crave it; if never saw a Dorito advert, we

⁹⁵ *Dark*. Created by Baran bo Odar and Jantje Friese. Season 1, Episode 8, "As You Sow, So You Shall Reap." Netflix, December 1, 2017.

⁹⁶ arman, John. *A Primer on Determinism*. Dordrecht: D. Reidel, 1986.

Mellor, D. H. *The Facts of Causation*. London: Routledge, 1995.

Also see: Hossain, Farhan. "What Life Looks Like in a Future After Free Will." *Predict (Medium)*. July 22, 2023. <https://medium.com/predict/what-life-looks-like-in-a-future-after-free-will-3365e3a50a93>.

⁹⁷ Stanford Encyclopedia of Philosophy. "Freedom and Responsibility in Ancient Philosophy." <https://plato.stanford.edu/entries/freedom-ancient/>.

⁹⁸ Hossenfelder, Sabine. "I Don't Believe in Free Will. This Is Why." YouTube video, 18:04. Posted June 3, 2023. <https://www.youtube.com/watch?v=TI5FMj5D9zU>.

⁹⁹ Shapiro, Ben. *Brainwashed: How Universities Indoctrinate America's Youth*. Nashville: Thomas Nelson, 2004. Kindle edition.

would not be *conditioned* to want it; and if Doritos never existed, we wouldn't even know about it. This nurture side of determinism is also mirrored in the nature side: if a murderer never acted on their impulse, they wouldn't be a murderer; if a murderer never had the gene expression that *predisposed* them, they wouldn't have likely turned out to be one; and ultimately, if the murderer had never been born, then they wouldn't have ever become a murderer. Yet, in this tension between nature vs. nurture, it is important to point out that our nature does not predetermine but it does predispose¹⁰⁰. For instance, you could have the genes and the predisposition to be an alcoholic but if you're born on an island that doesn't have any alcohol it's not going to happen. Much in the same way, there are genes that predispose us to eating junk food¹⁰¹, but if we don't live anywhere near a "food swamp", then we won't be conditioned to do so.

This relationship between conditioning and predisposition, is echoed in ancient wisdom: we are the average of the five people spend the most time with; or translated into the context of the Good Life: our (conscious TPN, and unconscious DMN) behaviour – and antecedent Subconscious Forces – is the product of the five closest environmental factors we condition ourselves to. Of course, five is just a ballpark¹⁰² placeholder and the way the dominoes tumble is more complicated and multifaceted as aforementioned. It is very rarely that a single domino predetermines any action alone, or that self-sabotaging behaviour is a result of a single agent, i.e. the Dorito Effect is likely a result of the Butterfly Effect from two or more dominoes colliding such as the corn oil industry¹⁰³ and mislead science¹⁰⁴.

We see then that all of our actions, regardless of whether external (unintended) or internal (intended), and antecedently both the TPN and DMN are necessarily influenced by the events prior. Our natural predispositions are catalysed by conditioning, either enabled (predisposed alcoholic in a bar) or disabled (predisposed alcoholic in a mosque). And Subconscious Forces are a product of nature, but more so nurture, meaning that which hinders obtaining exogenic subsidiaries is not really internal, but external; not lack of willpower, but lack of healthy environment.

The Ego as Compatibilist

Although all events are causally determined, that doesn't mean our – the Ego's – choices are predetermined as well. While of course they will still be influenced to a degree by how society is informed (epistemic endosmosis) in addition to the cultural "zeitgeist"

¹⁰⁰ TranscriptForest. "Robert Plomin: How Do Genes Influence Our Behaviour?" *Modern Wisdom Podcast*. August 2, 2021. <https://www.transcriptforest.com/en/modern-wisdom/11559-353-robert-plomin-how-do-genes-influence-our-behaviour-2021-08-02>

¹⁰¹ Llewellyn, Clare H., Cornelia H. M. van Jaarsveld, Danielle Boniface, et al. "Eating Rate Is a Heritable Phenotype Related to Weight in Children." *American Journal of Clinical Nutrition* 88, no. 6 (2008): 1560–66. <https://doi.org/10.1093/ajcn/88.6.1560>.

As cited in: van Tulleken, Chris. *Ultra-Processed People: The Science Behind Food That Isn't Food*. W.W. Norton & Company, 2023.

¹⁰² Kwik, Jim. "You're Not the Average of the Five People You Surround Yourself With." *The Mission (Medium)*. January 15, 2023. <https://medium.com/the-mission/youre-not-the-average-of-the-five-people-you-surround-yourself-with-f21b817f6e69>.

While the "five people" rule is not scientific, it resonates with the argument in this thesis: that our Subconscious Forces, including TPN/DMN behavior, are determined by environmental proximity and input.

¹⁰³ The overabundance of corn, especially in the U.S., contributed to the development of high-fructose corn syrup, which in turn drove the mass production of shelf-stable snack foods. For a broader analysis, see: van Tulleken, Chris. *Ultra-Processed People: The Science Behind Food That Isn't Food*. New York: W.W. Norton & Company, 2023.

¹⁰⁴ The push to "eat five meals a day" for sustained energy led to increased demand for snacks between meals, further accelerating the snack industry and normalizing grazing behaviour.

environment (epistemic contextualism), we still have at least some modicum of Free Will¹⁰⁵. This is because there is a difference between determinism and fatalism where the former claims that one's fate is necessarily influenced by the events prior, and the latter claims that one's fate is fixed regardless of events prior. In other words, while determinism allows for flexibility in one's fate by changing the events prior, fatalism allows no such exception where even deliberate attempts to escape one's fate paradoxically reinforce it¹⁰⁶.

It is important to point out that the literature is still divided on the reality of Free Will with positions ranging between compatibilism and incompatibilism. For the neo-Pascaline sake of moving forward¹⁰⁷, however, compatibilism is the only viable option because to claim otherwise is often incorrigible and, in the case of hard determinism (or fatalism), akin to defeatism where 'since everything is set in stone, there is no point in trying'¹⁰⁸. Here, then, I would like to turn to a thought experiment of "**Man can do what he wills but he cannot will what he wills**"¹⁰⁹:

Imagine that you were conditioned to like chocolate ice cream, your choice to eat chocolate ice cream then is not really your own as you did not choose to like it. Consider then, that having this knowledge, you deliberately rebel and do not eat chocolate ice cream to spite your genetic predisposition. You no sooner make the decision that it occurs to you that the decision was only made possible because of the aforementioned knowledge (epistemic endosmosis). And then it occurs to you that *that* - the decision to rebel against your conditioning – was yet another effect of your conditioning.

¹⁰⁵ For deeper exploration into how cultural narratives and institutional discourse shape the Ego and identity, see Michel Foucault's *Discipline and Punish*, Louis Althusser's *Ideological State Apparatuses*, and Anthony Giddens' work on reflexive self-identity. These thinkers expand on how media, education, and social norms sculpt the apparent autonomy of the self.

Foucault, Michel. *Discipline and Punish: The Birth of the Prison*. Translated by Alan Sheridan. New York: Vintage Books, 1995.

Althusser, Louis. "Ideology and Ideological State Apparatuses." In *Lenin and Philosophy and Other Essays*, translated by Ben Brewster, 85–126. New York: Monthly Review Press, 1971.

Giddens, Anthony. *Modernity and Self-Identity: Self and Society in the Late Modern Age*. Stanford, CA: Stanford University Press, 1991.

¹⁰⁶ The story of Oedipus serves as a mythological archetype for the unconscious drives of fate, mistaken self-knowledge, and tragic determinism—foundational themes revisited throughout this thesis.

¹⁰⁷ The Neo-Pascaline Matrix refers to a compatibilist wager structured similarly to Pascal's Wager. If we assume we are slaves (to deterministic influences) and act accordingly, we may actually improve our autonomy. For elaboration, see Appendix A: Misc., Exhibit 3.

¹⁰⁸ Classical compatibilism (e.g., as proposed by Frankfurt or Hobbes) holds that freedom and determinism are compatible because free actions are those not externally constrained and that flow from one's desires. The compatibilism used here, however, builds on this by emphasizing practical autonomy within systems—accepting deterministic constraints while still optimizing behavioral outcomes through self-awareness and environmental design. This is not metaphysical libertarianism but a pragmatic framework of "structured freedom."

Frankfurt, Harry G. *Freedom of the Will and the Concept of a Person*. *Journal of Philosophy* 68, no. 1 (1971): 5–20.

¹⁰⁹ See Korsgaard's Kantian interpretation of morality as autonomy: Korsgaard, Christine M. "Morality as Freedom." In *Creating the Kingdom of Ends*, 159–187. Cambridge: Cambridge University Press, 1996.

The important point here is that efforts to second guess the reality of Free Will cannot help you decide what to do, and it certainly doesn't make you any freer. Second guessing can, in fact, only enslave you to indecision making you less free than you were to begin with. In order to *do* anything, you must simply ignore the possibility of hard determinism (or fatalism and its cognates) that you are conditioned, and decide what to do - just as if you were free. Following Kant's ideas about ownership of desire maxims "under the idea of freedom"¹¹⁰ then, is not about a theoretical assumption necessary to decision, but about a fundamental feature of the standpoint from which decisions are made¹¹¹. In other words, regardless of the reality of Free Will, we must regard ourselves as having it, if not in support of moral foundation¹¹², then at least for the sake of moving forward this argument.

Libertarian compatibilists would have us believe that insofar as everything is determined, we can't choose to escape fate, but we can choose which fate to ascribe ourselves to. In our capacity for freedom then (tangible or illusionary), we cannot stop a bowling ball falling from a high place due to gravity (determinism), but that doesn't mean a specific ball cannot be prevented from falling from a specific height or place (fatalism). After all, "freedom is what you do with what's been done to you"¹¹³. In other words, compatibilism argues for the reality that a cause will necessarily lead to an effect, but neither is the cause nor the effect set in stone. Compatibilists therefore argue that determinism is but a phenomenon that can equal help and hinder; those ignorant of causality will simply be swept up while drifting along, but those who recognise how the dominoes tumble can use it to their advantage and shift the winds of fate.

Relating this to the earlier proposition of flying under the radar of the Reptilian Brain's indifference, we may not get to choose our disposition, but we ostensibly can choose the environment of condition; a murderer might not have chosen to want to like killing, but they have the autonomy to act on their impulse or not¹¹⁴. Incidentally, the TPN can also be understood as the hallmark of compatibilism; the system that allows the Ego, in its, albeit limited, autonomy has the capacity for self-denial of otherwise unconscious influences from Psychic Noise, i.e. in regards to the Dorito Effect, compatibilism allows for the Ego to rationally evaluate the Dorito and then, in recognition of the disconnect with its empty caloric nature, preemptively circumvent circumstances that would allow for the Pavlovian urge for flavour.

Nonetheless, provided that the Ego does truly possess true (limited) autonomy by means of a Free Will, the immediate effects do not extend past the TPN since the remaining denizens of the Cartesian Theater generally do not possess a will (as accruing from an ego), let alone a Free Will. Moreover, due to the mechanisms of control, what little time the Ego

¹¹⁰ Kant's foundational claim that to be free is to act according to reason (and not inclination) underpins the ethical argument for autonomy. See: Kant, Immanuel. *Groundwork of the Metaphysics of Morals*. Translated by Mary Gregor. Cambridge: Cambridge University Press, 1998. (Original work published 1785). Cited as G 4:448/66.

¹¹¹ Korsgaard's reading is further supported by Thomas Hill, Jr., especially in: "Kant's Argument for the Rationality of Moral Conduct," *Pacific Philosophical Quarterly* 66 (1985): 3–23; and "Kant's Theory of Practical Reason," presented at the Ethics and Its History conference, UC San Diego, May 16–17, 1986.

¹¹² Kant argues that to be truly free is to be morally accountable for one's actions. If determinism excused all behaviour, moral responsibility would collapse, rendering all efforts to cultivate virtue or self-restraint futile.

¹¹³ "Freedom is what you do with what's been done to you." — Jean-Paul Sartre. For modern psychological commentary, see: After Skool. *The Neuroscience of Identity*. YouTube video, 9:52. Posted March 5, 2022. <https://www.youtube.com/watch?v=zAQYtoIfB4g>. (Cited at 5:45 and 6:05)

"freedom is what you do with what's been done to you" — Sartre

¹¹⁴ From a compatibilist perspective, murder committed under the DMN (Default Mode Network) would not be considered an act of Free Will. However, failure to mitigate risk while under the TPN (Task Positive Network)—e.g., by removing weapons or seeking help—remains a failure of moral responsibility.

has under fronting is comparatively infinitesimal, meaning the percentage of time that the Self experiences Free Will is also relatively minimal. Thus while it might seem that Free Will was a possible avenue in overcoming self-sabotaging behavior, at least from an internal perspective, it seems unlikely.

A Drifting Will

Borrowing from Kant's idea that a Free Will is the will which aligns with the moral law, we see that the overall, composite will of the Self (General Will) is anything but free. This is not because the moral law could never align itself with the prime directive of the Reptilian Brain since according to the many-worlds theory¹¹⁵ there is a very possible scenario, if only by chance; rather, it is precisely because of chance that this occurs that the will is not free. In other words, it is not the moral law itself which Kant posits makes the will free, it is the autonomy of choosing the moral law that does so. To avoid circular reasoning, it is important to note that it is not the capacity of autonomy that makes the will free - for Kant would argue, in all of our irrationality¹¹⁶, we could still choose not to align our will with the moral law – but the deliberate choosing within our capacity for autonomy that allows it to be so. The implication from Kant then isn't that the General Will isn't free because it often fails to align with the moral law (although this is a worthy point of contention); rather, the key implication is that the General Will isn't free because there is no real autonomy, let alone any deliberate choices made within its capacity.

Again, while aspects of the Self could be deemed to be free (i.e. the Ego), in their limited autonomy, they have very little influence on the directionality of the General Will. Small aspects such as the Ego (and any other Subpersonalities, if having evolved an ego of their own) are necessarily shackled to the giant collective of the Psyche meaning, at least from an internal perspective, trying to change the directionality (or even nudging internal deterministic influences to overcome self-sabotaging behaviour) of the General Will is akin to a small man trying to control the path of a very large dog.

Our will can thus be said to be “drifting” where our Psyche, when met with deterministic influences, drifts like chaff in the wind – or dragged by a dog on a leash. This “Drifting Will”, as I coin it, encompasses (and is synonymous with) the General Will as it relates to the multifaceted Psyche (limited autonomy of the Ego, allonomy of the Subconscious, and neutral phenomena of Subconscious Forces). Moreover, insofar as the Ego's Free Will is held captive by external constraints of the Subconscious (Psyche sans Ego), we – our Ego – is not truly free, but is at the mercy of the General Will and how it drifts. Thus, the Ego is dragged along by whoever influences the Reptilian Brain; whoever influences the Reptilian Brain controls the Self, and whoever controls the Self is the true mastermind behind our self-sabotage, which we know now to be deterministic influences. In short, we are slaves to Fate, but a fate we can still ostensibly control.

Interpretation:

¹¹⁵ A thought experiment: in a multiverse with infinite possibilities, entropy implies that even something as absurd as a hamburger turning into a flamingo could occur. The illustration underscores how low-probability events can still happen, though rarely with moral significance.

¹¹⁶ Dubbed here the “Lady Windermere Paradox,” it captures the psychological contradiction where people are unhappy when they *don't* get what they want, but also when they *do*. Gilbert, Daniel. *Stumbling on Happiness*. New York: Alfred A. Knopf, 2006, p. 184. See also: Kahneman, Daniel, and Amos Tversky. “Prospect Theory: An Analysis of Decision under Risk.” *Econometrica* 47, no. 2 (1979): 263–291. <https://doi.org/10.2307/1914185>.

In this third and final section *On Free Will*, we've explored the idea of Free Will, determinism, and libertarian accounts of compatibilism. Particularly we looked at the limited autonomy of the Ego and its influence on the General Will. In doing so, we saw that compatibilism is the only appropriate (and possible) avenue for further discussion and regardless of the reality of Free Will, the end result is still the same: a Drifting Will.

The extent to which the Subconscious Forces influence, if not outright control us can now be understood in light of their interactions with the Drifting Will, or better yet, the drifting of the will. After all, it is the very process by which the will drifts, when met with deterministic influences, which transforms the General Will into a drifting one; moreover, it is by this very process that Subconscious Forces are necessarily influenced by events prior and drag along (outright control) the Ego towards the Self's general behavior. The extent to which Subconscious Forces influence the Self is thus not really a question of quantity, but of quality; the General Will, like chaff in the wind, offers little to resistance against deterministic influence amongst other Psychic Noise so long as they align with the prime directive of the Reptilian Brain, therefore, the General Will, insofar as it is a Drifting Will, is synonymous with the directionality of whatever way the wind blows. As such, what we need concern ourselves with (assuming that our multiverse is a deterministic one) isn't how strong the wind blows but where it is coming from, where it is going to, and how we can manipulate the weather so the wind blows under the radar of the Reptilian Brain's indifference and in the direction amenable to the Good Life.

In short, we are slaves to Fate, but no specific one. Paradoxically, from an internal perspective, our Ego is shackled to whichever direction Subconscious Forces nudge the Self towards, but in an expanded sense of this perspective - through empowerment by awareness of the Drifting Will - we come to understand that our autonomy is not necessarily as limited as initially presupposed as the compatibilist capacity of the Ego allows for external change and thus influence on deterministic influences. With this in mind, and having addressed this third and final interpretation, we can now piece it together with the first two to come to a mini conclusion of outcomes and results where we attempt to reconcile the missing key of a limited ego with failure to reach the goal of the Good Life.

Results

Before implementing my research program and delving into the core chapters of this particular part, we first formalized a hypothesis built upon several key research questions which included how are we to understand the pluralism of our persona in relation to acquiring the Good Life? To what extent does the Ego possess autonomy? And to what extent does the Subconscious Forces influence, if not outright control us?

From the outcomes of each research question, we can now verify the validity of the hypothetical claim: we are slaves (to Subconscious Forces that are really in control of us). In brief:

- (1) Our Persona is more than our Ego
- (2) Subconscious Forces are in control most of the time
- (3) We are slaves to Fate

First and foremost, Our Persona is more than our Ego. The pluralism of our Persona is commonly understood as tripartite meaning any acquiring the Good Life necessarily includes all three facets. Furthermore, our Persona as the phenotype of the Psyche is yet again only another facet of the Self. Turning our Self as a whole towards the good then is not a light endeavour that solely involves the Ego (as the background suggested with lack of willpower);

rather it is a combined effort that involves not just the Psyche, not just the Persona, but both metaphysical and physical aspects that comprise the gestalt Self. Insofar as we can reduce this to a simple understanding that our Persona is more than our Ego, the reduced implication is still the same: our Ego is not in control of our Persona, let alone the entirety of the Self.

Second of all, Subconscious Forces are in control most of the time. The compatibilist capacity of the Ego allows for some semblance of autonomy, but it is limited. Since the question of control is synonymous with the question of autonomy, the Ego thus has limited control. Moreover, since the DMN dominates the Persona in more ways than one (time under fronting, superior energy efficiency, general efficacy, etc.), the DMN in all of its impressionability applied to the magnitude of its domination (autopilot) effectively translates into being a Puppet CEO¹¹⁷ for Subconscious Forces. This is because the DMN does not have any bargaining power, or intentions - let alone a conscious will - meaning other than the Ego's limited autonomy, there is little to no autonomy exhibited in other areas of the Psyche. Moreover, while the TPN is less impressionable, it still struggles to efficaciously make well informed, competent decisions in the background of Psychic Noise. Furthermore, it is still influenced by Subconscious Forces meaning in all modes of consciousness, Subconscious forces have their hegemonic hold. Thus, our Ego is not in control, but Subconscious Forces are.

Last but not least, we are slaves to Fate. Subconscious Forces, although in control most of the time, are only a proxy for deterministic influences amongst other Psychic Noise that align with the prime directive of the Reptilian Brain. The very process by which the will drifts, when met with deterministic influences causes the General Will to follow the directionality of whatever way the wind blows. Incidentally, the compatibilist capacity of the Ego allows for its Free Will influence over the Drifting Will. Nonetheless, the magnitude of the gestalt Self's Drifting Will, more often than not, drags along the Ego as part of the Psyche, but more importantly, the behaviour of the Persona as well. The extent to which Subconscious Forces influence, if not outright control us is relatable to the deterministic influences we expose ourselves to. We see then that our Ego is not in control, but Subconscious Forces are as a proxy for deterministic forces.

Taken together, the three different sections: On the Persona; On Autonomy; and On Free Will, verify the validity of the hypothetical claim of using being slaves by showing incrementally how far removed the Ego is not only from its own autonomy (as impressed upon by Subconscious Forces), but also from control over the Self (as an infinitesimal facet). Evidence for the Ego's enslavement to Fate only further consolidates the idea that it is unfeasible that a psychotherapeutic model partaken under limited autonomy of the Ego could ever resolve the General Will's hinderance of obtaining exogenic subsidiaries. And, ultimately, the data gathered shows that it's not just that people *should not*, but more so *cannot* make their own decisions in turning our Self as a whole towards the Good.

From this core chapter, we have thus seen that:

1. Socrates was correct when he claimed that no-(Ego) knowingly chooses to do things that goes against the good.
2. Self-sabotage is a result of unfavourable Subconscious Forces which themselves are deterministically influenced by causal events.

¹¹⁷ Just as a CEO may be the public face of a company but remain beholden to a controlling shareholder, the Ego may appear autonomous while being overruled by Subconscious Forces. See also: Yoon, Hyeon. *Reborn Rich*. Webtoon. Episode 24. Accessed September 30, 2023. <https://www.webtoons.com/en/drama/reborn-rich/episode-24/viewer>.

3. Causal events are a product of Fate and while no Fate is certain, neither is Fate avoidable.

Needless to say, it should be that the data gathered reflects years of research and study. Even though these findings are supported from a range of academic disciplines, it is important to note that the data gathered here is an interpretation of the research theory and methods. Research can never claim to be conclusive, especially in emergent fields such as psychology which many still contend to be in the realm of pseudo-science. It is very possible then that following the next breakthrough in science, the interpretation of the outcomes and results may be concluded to be void. However, in the spirit of scholarly dialogue, and for the sake of progress, we shall work with the insights gained from profiling the predisposition to self-sabotage actualization of the good. It is my intention to build from the foundation of my results to comprehensively answer the initial research question and ultimately contribute to the solution of the practical problem of the General Will's hinderance of obtaining exogenic subsidiaries.

Specifically, then, we see that it is unfeasible that a psychotherapeutic model partaken under limited autonomy of the Ego could ever turn our Persona – and subsequently the entire Self – towards the good. Yet, as alluded above, this does not exclude the possibility of a solution focused on external change as supplied from an external perspective instead¹¹⁸. This proposition will be the overarching theme for the discussion moving forward where, as a more efficient approach than simply increasing willpower alone, we attempt to reconcile the missing key of a limited ego with failure to reach the goal of the Good Life

¹¹⁸ The cure for psychological dysfunction often cannot come from within because the pathology is not merely individual but collective. Thus, solutions must target collective psychological patterns rather than isolated psyches. See: Academy of Ideas. *The Psychological Black Death*. YouTube video, 15:42. Posted March 20, 2022. <https://www.youtube.com/watch?v=s0K0sI5XiP4>. (Cited at 10:38–11:15 and 14:34–15:08)

Synthesis

The Discussion

There are many points of contention, some major, some minor; some discrete, some subsidiary; some salient, some yet to be thought of... For the intent and purposes of this part, however, we will only address the most salient contentions as they relate to the data gathered thus far from the Background, Outcomes, and Results.

- (1) What could be possibly done to reconcile the missing key of a limited ego with failure to reach the goal of the Good Life?
- (2) and we are in a much better position to answer the initial research question of: if willpower alone is an insufficient approach, then which approach(es) would we need to be more efficient?

Here we will also evaluate the caveats, limitations, and shortcoming of this part before ending with the implications that take all these into account. Discussion will also overlap between not only sections but also chapters as well. i.e. questions regarding the reptilian brain (Outcome: on autonomy) necessarily relate to how Alters are formed (Outcome: on persona) which itself also relates to translation of aspects of Virtuous Hexis (Background: the solution is habits?). Thus, they cannot be neatly confined but will endeavour to address them under the key variables of:

1. Contentions over Background (Good Life + willpower?) = briefly touch on since does not make up core of argument
2. Contentions over Outcomes (Persona; Autonomy; Free Will)
3. Contentions over Results (+discussion, thesis: need external framework that curbs/nudges/parameterizes, enhances people choice)

Contentions

Over Background

1. Should We Even Aim for the Good Life?

The first contention is with the initial assumption of aiming for the Good Life, not so much in regards to if people even want the Good Life (which will be addressed in the caveats segment), but rather whether we should even aim for the Good Life. Specifically, in the background we already touched upon Goodheart's Law and how aiming for a marker of the good (happiness) not only diminishes its relation to the Good Life, but could also result in converse effects; now we need to explore if aiming for the Good Life itself results in the same.

This contention primarily arises from implications of Universe 25¹¹⁹ which shows societal breakdown in response to life in abundance. It is worth exploring then if obtaining exogenic subsidiaries by way of utopia would similarly lead to boredom and discontent in

¹¹⁹ Overpopulation-induced dysfunction, or "behavioral sink," was demonstrated in Calhoun's Universe 25: a rodent utopia that collapsed into violence, apathy, and social breakdown, showing how structural abundance without psychosocial integration breeds pathology.

Calhoun, John B. "Population Density and Social Pathology." *Scientific American* 206, no. 3 (1962): 139–148.

human populations as many contenders predict.¹²⁰ Psychology and neuroscience support this idea, positing that the optimal state is not found in life in abundance, but in between, embracing the tension of the unfamiliar. This could very well mean that the Good Life is not necessarily a destination or final state but a transitional journey; what substantiates the Good Life isn't so much in the attainment as it is in life, liberty, and the pursuit of happiness. One interpretation of this is that we might already be living the Good Life right now yet fail to realize for what it is.

An alternative interpretation, however, is that granted we can attain the Good Life, still, it would never be good enough for us. This correlates with the facticity of seemingly limitless human greed coupled with the hedonic treadmill. Here, it is the hedonic treadmill, which usually a coping mechanism acts now as a double-edged sword; the contrast effect¹²¹ of what makes a street sweeper as happy as a king is what also makes a king as sad as a street sweeper. Making Greed run on the hedonic treadmill then, we see that it doesn't matter how much levels of absolute happiness increase, the hedonic treadmill keeps relative levels of happiness the same, meaning we will always want more and never be satisfied with what we have. This loops back to what was mentioned earlier of the exogenic subsidiaries being half the equation for the Good Life by way of minimum threshold, but the other endogenic (internal synthesis) half being much more important.

From this contention then, we see that it's not just about attaining the Good Life, but our attitudes towards it; aiming for the Good Life is futile if we don't know how to appreciate it once we're there.

2. *The Problem of Rationalizing Human Behaviour*

A second contention is with the rationalizing of human behaviour. Psychologists are often the first to make claims along the lines of "the human being is the only animal that..." and, consequently, psychologists are thus also often the first to identify the difficulty in ascribing any general pattern¹²², including the rationalizing of the Psyche.

Aristotle, while ostensibly not the first, is renowned for claiming what qualifies us from other beasts (and children) is rationality¹²³. This idea was later developed by Kant in the direction of the Free Will¹²⁴, where, as aforementioned in the Drifting Will, he described rationality as the deliberate choosing within our capacity for autonomy (particularly the choosing of the moral law).

¹²⁰ Dostoevsky's *Notes from Underground* critiques utopia as fundamentally incompatible with human nature. When granted total comfort and predictability, man will revolt against the system—not out of suffering, but boredom. Kurzgesagt illustrates a parallel idea in "The Egg," where omniscient design still fails to quell existential dissatisfaction.

Kurzgesagt. "The Egg: A Short Story." YouTube video, 5:34. July 10, 2019.

<https://www.youtube.com/watch?v=4wuBCQrHI2E> (0:13–2:08).

¹²¹ Happiness is more influenced by relative status than absolute gains. One's satisfaction with their salary, for instance, depends less on the number itself and more on how it compares to peers—especially in-laws.

The Aspen Institute. "Stumbling on Happiness with Daniel Gilbert." YouTube video, 38:25. February 2, 2023. <https://www.youtube.com/watch?v=2EiV4-CleIs> (35:42–36:21).

¹²² Humans are uniquely capable of imagining alternative futures, which ironically makes us chronically dissatisfied with the present. Our ability to project leads to perpetual discontent.

Gilbert, Daniel. *Stumbling on Happiness*. New York: Alfred A. Knopf, 2006, 184.

¹²³ In Aristotelian terms, children are considered in a state of *akrasia* (incontinence)—they lack stable rational control and act impulsively, much like beasts.

¹²⁴ In Kantian ethics, any rational agent—human or otherwise—is a moral legislator under the categorical imperative. Aliens, who, if capable of moral reasoning, would fall under the same universal moral law. Kant, Immanuel. *Groundwork of the Metaphysics of Morals*. Translated by Mary Gregor. Cambridge: Cambridge University Press, 1998. (Original work published 1785). Cited as G 4:448/66.

Humouring for the moment the possibility that the Good Life is synonymous with the moral law, the data gathered - as it evinces if not the irrationality of the Persona, then at least the lack of (autonomy that enables) rationality – would mean it is unattainable by condition. But again, as with the key implication of the Free Will not lying in the moral law itself but in the capacity of autonomy to choose it so, the key implication of humouring Kant’s conditional has little to do with the moral law, and much more to do with its status as a conditional. In other words, one problem of rationalizing human behaviour is with the assumption that humans have sufficient autonomy to align their capacity for rational action with any conditional at all. Since the Ego barely has autonomy over its own will, let alone the Psyche, failure for the Self to rationalize its behaviour (let alone be called a rational being) in accordance with a conditional – be it the moral law or any other – is a very legitimate cause for concern.

Another interpretation of this contention stems not from rational capacity, but in trying to institutionalize¹²⁵ human behaviour as a rational expression of the mind. Specifically, this interpretation is concerned with the idea that neuroscience can neatly fit all human motivation and subsequent action into a localized, contextual model¹²⁶, i.e. the problem registered earlier with the schematic triune complex. This institutionalized systemization is at odds with not only the randomness (quantum jumps) inherent to deterministic influences which influence our behaviour, but also the Madness, Genius, and Abnormal Conditions which lead to spontaneity of our behaviour itself. Furthermore, this institutionalized systemization also reduces mind-body dualism to philosophical materialism, ignoring any potential metaphysical or spiritual explanations. Yet, even if philosophical materialism turns out to be true, even if institutionalized systemization is found to be coherent, the way something is isn’t telling of the way something is supposed to be. Namely, even if human behaviour can be patterned, which current society is more or less so as a function of the assembly line¹²⁷, this does not make it the ideal which is clear from general job dissatisfaction, not necessarily with the job itself so much as with the patterned repetitive “workplace ennui”¹²⁸ of doing the same thing day in day out, over and over again.

In regards to the former interpretation, we have already touched upon the answer: if the ego, from an internal perspective (psychotherapeutic model), can’t rationally (and more so, feasibly) align its behaviour with a conditional insofar as it leads to the Good Life, then we can circumvent this through an external change instead. This extent to the efficacy of parameterize, external change (or framework) will be explored in the eighth and ninth contention. In regards to the latter interpretation, we need not formally institutionalize human

¹²⁵ Institutionalization refers to the process by which individual behavior becomes formalized into systems, rules, or social norms. Over time, this external structure can override intrinsic or deliberative moral reasoning, channeling conduct into patterned scripts.

¹²⁶ Efforts to localize brain function often reduce complex behavior to discrete regions, neglecting the influence of environmental context. A fuller model acknowledges how cognition arises from interactions between neural substrates and external, often unpredictable, conditions.

¹²⁷ In modern society, time consciousness has shifted from being embodied and experiential to abstract and externally regimented. This alienation from organic rhythms may be a key factor in psychological fatigue. Weixin (WeChat). “Time Consciousness in Modern Society.” Article ID 1707659000. Accessed September 30, 2023. <https://mp.weixin.qq.com/template/article/1707659000/index.html>

¹²⁸ Jobs that lack variation or cognitive stimulation often induce a state of occupational boredom, which research shows may be more harmful to long-term mental health than stress or even burnout.

Gupta, Priyanshu. Answer to “Do Most Jobs Turn Boring and Repetitive After Some Time?” Quora, August 10, 2023. <https://www.quora.com/Do-most-jobs-turn-boring-and-repetitive-after-some-time-If-yes-how-does-one-keep-oneself-motivated/answer/Priyanshu-Gupta-1409>.

Kohler, Lindsay. “Why Boredom at Work Is More Dangerous Than Burnout.” *Forbes*, July 26, 2022. <https://www.forbes.com/sites/lindsaykohler/2022/07/26/why-boredom-at-work-is-more-dangerous-than-burnout/>

behaviour; neuroscience is only useful to this thesis to the point that it supports that our Persona is multifaceted and we are deterministically influenced. Furthermore, an external framework can be patterned without evoking a negative response; in the same way that the finite need not be dichotomized with the infinite¹²⁹, systemization need not be mutually exclusive with personalization.

3. *Does Increasing Willpower Extend At All?*

A third contention is with how willpower can have an effect on self-control granted that the Ego doesn't even have full control over itself. In the background we saw that the research literature was correct on at least three fronts: (1) willpower does have an effect on self-control; (2) training willpower does increase self-control; and (3) Ego depletion is a real phenomenon. The contention arises, however, when trying to understand this in light of results from our research; specifically, in trying to understand how increasing willpower can increase self-control when no longer under the TPN.

This confusion can be attributed largely in part to how people understand the function of willpower as limited to akrasia. Namely, while most people identify willpower as the factor affecting conscious overcoming, i.e. writing theses when one would rather read comics; the subsidiary effect of increased neural adaptation is often overlooked, i.e. reps in gym.

This subsidiary effect of the TPN's lasting influence on the DMN is elucidated once again by Phineas Gage: when Gage's frontal lobe was damaged, his ego was ostensibly either destroyed, corrupted, or usurped by a Subpersonality which would explain his shift in demeanour, but, more importantly, whilst habits related to mental acuity (i.e. a virtuous disposition) were lost, those related to physical prowess were not (i.e. brushing teeth). One interpretation of this is that although the Hexis gained from practicing mindful activities would be limited to the ego that cultivated it, muscle memory as it relates to the Self rather than Psyche translates across all Subpersonalities. The idea that gains from mindful practice don't carry over but those from physical practice do is further evinced by how those with DID can have an Ego and Alters which speak completely different languages (Mental Hexis) and suddenly find themselves having acquired different physical skills (Physical Hexis) when refronting¹³⁰. More evidence is seen with progress (or lack thereof) when we shift from the TPN to the DMN and back again. For example, when we daydream during a math exam, we don't come back to find a sheet of solved problems and, likewise, when we daydream when driving, we don't come back finding ourselves parked in the middle of the highway. Thus seen, while our autonomy is limited, the Physical Hexis allows for work towards conscious tasks while unconscious – but only if leading activity is solely a physically demanding task (i.e. driving) and does not encroach on the use of mental acuity (i.e. dealing with a rabbit that jumps into the road) which necessarily leads to the "scare" back into reality.

The key to employing the Subconscious to essentially work for the conscious is all about capitalizing on setup of the framework: "the subconscious is the ultimate computer. Cognition is the keyboard"¹³¹ – inciting the DMN to habitually "muscle memory" its way

¹²⁹ A bounded system containing recursive feedback loops—like infinite mirrors inside a finite mirror—can produce internal complexity that feels infinite, even when the structural limits are externally fixed.

¹³⁰ Motivation in repetitive tasks often hinges not on novelty but on perceived purpose. Even low-skill jobs can feel meaningful when the worker perceives their actions as contributing to a broader narrative.

Quora Community. "Understanding Motivation in Repetitive Jobs." *Quora*. Accessed September 30, 2023. <https://qr.ae/pKQPkX>

¹³¹ The subconscious mind can continue working on problems outside of active awareness—a phenomenon often referred to as insight. Seemingly idle periods can yield sudden breakthroughs when the mind is allowed to

through segued tasks, i.e. driving, lifting weights, cleaning, etc. To make it vivid, the Physical Hexis is one such framework that works much in the same way a business can still earn money in the background while all employees are at home in bed. Namely, it capitalizes on parameterizing autonomous (behavioral) trends while we no longer (under the TPN) in the office.

Another differing, and contrasting interpretation is expressed negatively (perpendicular) and in contention with the idea that all behaviour must necessarily commensurate, or at least not actively go against the survival mechanism. Specifically, it is a contention not with the limitations of increased willpower, but with its abundance.

So far, the general theme has been that the Ego's autonomy is limited, and its influence on the entire Self even less so. Moreover, it has been implied that even if the Ego somehow manages to purchase influence - such as by way of cultivating a Virtuous Hexis - then the influence or action must necessarily be conducive or, at the very least, align with the prime directive of the Reptilian Brain. This would be the natural conclusion given that, for the sake of energy efficiency (an effect of the survival mechanism), the Reptilian Brain budgets the Ego's time under TPN, let alone its influenced Hexis and behaviours. Yet, if the Ego's willpower is so pressed for influence, why then does it seem so easy to cultivate bad habits, i.e. smoking, drinking, etc., that are not conducive to survival? Moreover, how do these bad habits even gain approval from the stingy Reptilian Brain?

The answer is actually very simple: they don't. It's not smoking, drinking, and other bad habits that get the green light from the Reptilian Brain, rather, it's the addictive substance that accompanies it. Furthermore, the habit is not purchased with the influence of the Ego by way of physical practice, but through chemical composition and hijacking of the brain chemistry. We can understand this phenomenon by returning to the Dorito Effect: it is not the Ego by any sheer force of willpower (physical practise) that influences the Self to eat its way to the very bottom of the bag, it is the artificial flavourings (chemical composition) as a vector of Subconscious Forces that do. Much in the same way, it is not the Ego that cultivates a Hexis for smoking, but the nicotine within cigarettes that through binding with dopamine receptors¹³² smuggle past the radar of the Reptilian Brain's indifference.

From this contention then, we see that willpower extends much further than just the Ego insofar as it relates to physical prowess and not mental acuity. Nonetheless, there is a limit to how far willpower can extend, and the extent of its reach is nowhere near close to being able to cultivate influence in the same magnitude as addictive substances. That being said, in contemplating ways to induce external change (as opposed to internal psychotherapeutic models), one potential avenue is revisiting the Virtuous Hexis, by way of inciting the subconscious mind, insofar as it allows for the translation of "muscle memory". Naturally, a Physical Hexis is limited by way of mental acuity, but insofar as its one step closer in turning the behaviour of the Persona (and thus, the Self) towards the Good, it is worthy of consideration.

Over Outcomes

drift.

TED-Ed. "The Science of Boredom." YouTube video, 4:55. Posted October 3, 2021.

<https://www.youtube.com/watch?v=zHECSAdJtQ> (Cited at 2:55–3:07)

¹³² Nicotine hijacks the brain's reward pathways, establishing cycles of craving and reinforcement that blur the line between voluntary use and compulsive behavior.

Tiwari, Rajesh K., Vineeta Sharma, Rakesh K. Pandey, and Shubhini A. Saraf. "Nicotine Addiction: Neurobiology and Mechanism." *Journal of Pharmacopuncture* 23, no. 1 (2020): 1–7.

<https://doi.org/10.3831/KPI.2020.23.001>

4. *The Degree of Subconscious Self-Sabotage*

A fourth contention is with the degree that facets of the Persona sabotage the Self. Exploring this requires us to first distinguish between general unconscious Subpersonalities, and those which have evolved an ego to become Alters. As mentioned in *On Autonomy*, the former is only capable of passive betrayal while effectively under duress by Subconscious Forces and Psychic Noise; in contrast, the latter (while still impressionable) with their developed rational capacity are also capable of active betrayal as consequence of their own personal agenda. Another way of discerning the different facets and their capabilities is returning to the patriarchal analogy: the Self is like a Planetary Human¹³³ where the representational Ego is the patriarchal figure, the Subpersonalities are obedient (albeit impotent) children, and Alters are pubescent teenagers.

One interpretation of this contention stems from why Alters would actively betray other personalities, including the Ego, when sabotaging the Self is subsequently sabotaging their own good. This can be explained in two parts. First, like pubescent teenagers, Alters have a developed rational capacity, not a perfected one. While Alters are no longer impotent like other less developed Subpersonalities (or other limbic facets), neither are they fully competent either. This has been most widely observed in those with DID who have been noted to harbour Alters with differing levels of maturity, i.e. a three-year-old Alter in grown man's body. Second, as having evolved from a general Subpersonality, Alters also have a unique personality of their own; they are quite literally an alternative consciousness. Just as different people have different goals and aims, some being law-abiding citizens while others being anarchic insurrectionists, Alters, too, have their own agendas, even potentially self-sabotaging ones.

Another contrasting interpretation of the contention is with the reality of Alters self-sabotaging. As mentioned above, not all Alters are self-sabotaging in the same way that not all teenagers are combative; it is only in Psyches with bad psychology who have Alters that vie for control of the conscious. Incidentally, not all Subconscious actions are self-sabotaging either, being a mixed urn¹³⁴. Furthermore, even if Alters have contradicting agendas, the power stratification between them is uneven. It should be reiterated that the Ego should be understood a positional role as opposed to a certain personality trait (ego). Thus, if even the current Ego is dragged along by whoever influences the Reptilian Brain, then Alters whose egos aren't even strong enough to usurp the Ego are even more powerless in the face of deterministic influences. Moreover, the majority of the population do not even have Subpersonalities that have evolved into Alters, let alone Alters with an egoistic streak. While everyone is ostensibly somewhere on the spectrum between DID and AI, we should remember that neuroscience theorizes that only those skewing towards the former will have Alters *potentially* developed, and only in response to trauma or events that prompts the Reptilian Brain to create Subconscious protection for the general Self.

The magnitude of any self-sabotaging influence, if any at all, as stemming from contrarian Alters is thus equal or less to the influence mustered when the Ego collects itself. The implications of this are thus that acts of Alter self-sabotage is not an excuse to lean away

¹³³ The concept of a "planetary human"—one whose interiority contains an entire civilization—is explored metaphorically in visual and narrative art, where environmental degradation reflects psychological entropy.

¹³⁴ Subpersonalities, typically unconscious, conform to the deterministic influences acting on the psyche. Whether benevolent or harmful, these influences manifest in actions not necessarily traceable to conscious volition. Homer captures this dynamic poetically through the imagery of Zeus's two urns—one filled with blessings, the other with sorrows.

Homer. *The Iliad*. Translated by Robert Fagles. New York: Penguin Classics, 1998. Book 24, lines 527–533.

from the Good since the Ego has the necessary power to reverse any action with a stronger (if not much stronger) influence on the Self, if only during times of engagement.

5. *Are We Really Not Morally Responsible for our Subconscious?*

A fifth contention is with accountability for our Subconscious. While the Ego and other Subpersonalities are interdependent, they are still mutually exclusive. Furthermore, earlier in *On the Persona*, we made a comparison between the Ego not having control over the Subpersonalities just in the same way we don't have control over other people. As we are not accountable for other people's actions, this leads to the idea that neither should our Ego be held accountable for the actions our Self performs under the DMN.

There are primarily two causes for concern with a first interpretation, both which can be understood in parallel with Plato's argument from justice in the soul and the state.

In regards to the soul (which we shall hold synonymous with the Psyche), Plato's argument was that injustice not only *leads to* disorder, but necessarily *is* disorder¹³⁵. If we dissociate from responsibility for our Subconscious, we, in turn, exacerbate dissociation of the Psyche (if not the entire soul as well) where this added disorder makes it even harder to turn the Persona as a whole towards the Good. Much in the same way, Plato argued that disorder in the soul would mirror itself with disorder in the city¹³⁶; family systems therapy teaches that when parents dissociate from their children (i.e. disavowing due to bad behaviour), this rejection only degenerates the patriarchal system. Failing to take accountability, then, would only enable murderers, liars, and thieves to simply blame their actions on their Subconscious and disposition of Subconscious Forces (via the Reptilian Brain) leading to chaos in society¹³⁷. We are seeing the beginning of this already with the advent of the post-truth era and how identity politics is slowly leading to societal collapse.

A second interpretation of contention follows closely with this defeatist, rejectionist idea. Specifically, contending that our Subpersonalities, insofar as being proxy for deterministic force, are under the jurisdiction not of the Ego but of causal events is akin to rejecting our capacity for compatibilism. Incidentally, this brings into question where one jurisdiction starts and the other begins and if there even can be an arbitrary line drawn between the TPN and DMN. But, more importantly, and in an extreme rendition, this kind of reasoning is akin to claiming the Big Bang made Hitler wage war on the Jews. The problem with this defeatist reasoning was already responded to in the section *On Free Will* where it was discussed that if not in support of moral foundation, then at least for the sake of moving

¹³⁵ Plato's tripartite soul becomes disordered when its rational, spirited, and appetitive parts fall out of balance. This psychological disharmony leads to misjudgment, misbehavior, and ultimately, moral failure. 434d–444a.

¹³⁶ A disordered city, according to Plato, mirrors the disordered soul: when the governing class fails to rule with wisdom, the auxiliaries lose discipline, and the producers pursue unchecked appetites. Civic unrest reflects inner disharmony. 543a–544e.

¹³⁷ Accountability is not solely an individual burden. When unfavorable behavior arises from systemic failure—e.g., lack of access, safety nets, or social infrastructure—moral responsibility must also be distributed across societal institutions. The state's failure to intervene early may warrant collective responsibility. Academy of Ideas. *The Psychological Black Death*. YouTube video, 15:42. Posted March 20, 2022. <https://www.youtube.com/watch?v=s0K0sI5XiP4> (Cited at 14:34–15:08)

This discussion parallels the philosophical concept of *moral luck*, where individuals are judged morally for actions influenced by factors beyond their control (see Bernard Williams, Thomas Nagel). The Ego's partial jurisdiction over behavior mirrors this paradox, raising questions about whether ethical accountability is fair when rooted in deterministic subconscious processes.

Williams, Bernard. *Moral Luck: Philosophical Papers 1973–1980*. Cambridge: Cambridge University Press, 1981. Nagel, Thomas. "Moral Luck." *Mortal Questions*. Cambridge: Cambridge University Press, 1979.

forward this argument¹³⁸. But more importantly, all the structures in the universe, including Hitler and the Jews, were created by quantum fluctuations in the plasma in the early universe¹³⁹. Furthermore, their details were not determined at the Big Bang (if there even was a Big Bang) which means that rejecting the Ego's jurisdiction of accountability is quite literally incompatible with the laws of physics.

Since rational thought is likened to being master of its own temple, our representational Ego has a legitimate, physical duty for accountability over the entire Self. This does not demand adherence with moral responsibility in the narrow sense where the Ego has duty to align the Self with the moral law (although this would at least be the Kantian ideal); however, it does demand adherence in the broad sense, if not for the sake of actively turning our Self as a whole towards the Good, then at least for the sake of passively preventing further societal collapse and moral degeneration (dissociation) of the Self. The ego has been blessed with the Will to Power¹⁴⁰, where in its (albeit limited) autonomy, has the potential to aim and direct one's life by way of cooperating with Subconscious Forces in determining the values which are pursued¹⁴¹. It would thus be a shame to maintain a defeatist, rejectionist view of responsibility (or lack thereof) and subject our entire Selves to a life under Bad Faith¹⁴² when the Ego in its compatibilist capacity has the means (taking accountability) Free Will to do something about it.

6. What About Intelligent Design?

A sixth contention is with intelligent design. This is a token contention that inevitably addresses the concerns of certain sects who find it difficult to reconcile God's Plan with free will. This is also a contention which will allow me to make a start on reconciling the thematic dichotomy between religious dogma and scientific dogma. Namely, even if God's Plan exists, or any other divine will for that matter, that doesn't mean we don't have the freedom to make our own decisions.

The first interpretation that adjusts God's Plan to allow for freedom of decisions refers to the earlier response in where compatibilism reconciled determinism with libertarianism. Specifically speaking, by substituting "God" (or how God's plan is executed, "God's Hand") for "deterministic influences", we see that the same answer still applies: a cause will necessarily lead to an effect, but neither is the cause nor the effect set in stone.

¹³⁸ Moral responsibility, though philosophically debatable, is pragmatically grounded in consequences.

Regardless of whether incompetence stems from individual limits or systemic complexity, ignoring it yields real harm. Ethical discourse must prioritize actionable insight over metaphysical certainty.

¹³⁹ Even if free will is an illusion, belief in it may remain socially necessary. Without it, the sense of moral responsibility erodes—and with it, the motivation to act differently. Hossenfelder, Sabine. "I Don't Believe in Free Will. This Is Why." YouTube video, 18:04. Posted June 3, 2023. <https://www.youtube.com/watch?v=TI5FMj5D9zU>

¹⁴⁰ Nietzsche's "Will to Power" refers to the capacity of becoming what one is—a process of self-overcoming and authentic actualization. The concept captures the notion that if all drives could be reduced to one, it would be the Will to Power.

BBC Ideas. *The History of Timekeeping*. YouTube video, 7:25. Posted May 15, 2020. <https://www.youtube.com/watch?v=ADWhVq5Rddo>. (Cited at 17:37–18:59)

¹⁴¹ The industrial revolution did not just alter economies and social structures—it altered time itself. Standardized work hours and mechanical clocks restructured consciousness around productivity. Vox. *How Industrialization Changed Time*. YouTube video, 10:45. Posted September 8, 2021. <https://www.youtube.com/watch?v=UKJlxvOOanM>

¹⁴² Sartre argued that once we recognize our freedom, we must also confront the dread that follows—the terror of infinite possibility. Many retreat into what he termed "Bad Faith": choosing to believe in comforting illusions rather than embracing responsibility.

After Skool. *The Neuroscience of Identity*. YouTube video, 9:52. Posted March 5, 2022. <https://www.youtube.com/watch?v=zAQYtoIfB4g>. (Cited at 6:05)

Adapted to the current contention, we get: God’s plan will necessarily come to fruition, but neither is the path or nor the direction set in stone. Admittedly, this means God’s plan (if such thing exists) would be fatalistic to much chagrin, but a saving grace is that while Oedipus can’t help killing his father and sleeping with his mother, he still gets to make his own decisions by considering options, drawing on memories, and choosing based on preferences and information. The end point is fixed, but just because the end of a story is already written doesn’t mean it’s not worth reading from the beginning.

The second interpretation is an argument from blasphemy on our right to play God. More specifically, it is a normative judgement against the implication to parametrize nature (causal events) to nudge behaviour of the Self towards the Good instead of allowing deterministic (“God’s Hand”) influences to arise organically. But firstly, if God did not want us to utilize nature (ideally mutualistically), then He ostensibly wouldn’t have given us Free Will or the freedom to make our own decisions in the first place. And secondly, if His love is so unconditional that He sends us to Hell for going against His will, then it would ostensibly be better to make Heaven on Earth and have freedom with the Devil later than sacrifice authenticity for an infinite with an undeserving God.

The way in which religious dogma relates to scientific dogma, where zealots (God’s plan) and atheists (mathematics of the standard model¹⁴³) both agree, is seen in support for hard determinism. Of course, the two arrive at the same conclusion through different methods, and the sudden truce in fighting a common enemy doesn’t forgive foe into friend. More importantly, from the section on Free Will and just above, we’ve seen how hard determinism is incorrect, meaning from a negative (perpendicular) interpretation, we can make a start on reconciling both religious and scientific dogma through a shared invalidation. Specifically, in relation to this contention, an invalidation by negative proof – the negation of hard determinism by proof of freedom of decision; the negation of incompatibilism by proof of free will.

Over Results (interpretation)

7. *The Ego’s Influence on the General Will*

A seventh contention returns to the idea of the Ego’s influence, this time from macroscopic interpretations since we’ve already addressed the microscopic ones.

In addressing an earlier contention, we specified a certain way in which the Ego’s time under control can extend past its time under fronting by way of having a lasting influence (willpower) through a physically Virtuous Hexis. In this contention, we will focus more generally on the validity (and potency) as well as magnitude of the Ego’s influence. Specifically, in returning to the mechanisms by which the Ego’s control can translates across the Psyche we will see how the underdog of the Ego’s Free Will can be the mouse to the elephant¹⁴⁴ of the Self’s General Will.

In the second contention, we already saw how cultivating a Physical Hexis validated the Ego’s indirect influence over facets of the Psyche; however, validating the Ego’s direct influence is a little trickier since while tenants of the Ego have free reign in their own temple,

¹⁴³ Marcolli, Matilde. *Moduli Spaces of Maximal Surfaces in Minkowski Space*. World Scientific, 2020; see also *Noncommutative Cosmology*. World Scientific, 2022.

¹⁴⁴ The ego is likened to a mouse guiding an elephant: while the ego may “steer,” it remains vastly overpowered by the subconscious collective—analogous to a rider mounted atop a beast, not in control of its mass but sensitive to its movements.

they are still “contracted”¹⁴⁵ to the landlord of the Reptilian Brain. To reiterate this mechanism of control, the tenant (Ego) is free to open and close the curtains (choose the rhythm of breathing) while at home, but must audit the landlord (Reptilian Brain) if it wishes to take down the curtains completely (suffocate by holding one’s breath). Of course, we see that the landlord vetoes¹⁴⁶ any of these changes (holding one’s breath; stopping one’s heartbeat) to its temple, despite the tenant ultimately living there. Nonetheless, the existence of deterministic influences doesn’t mean we don’t have the freedom to make our own decisions. Moreover, the compatibilist capacity of the Ego still allows for direct enactment on those decision during time under fronting where the mouse is generally free (if not upsetting the Reptilian landlord) to pilot the elephant in its limited autonomy. The most obvious employment of this autonomy is changing one’s environment such as directly choosing a specific friend group to hang around with, directly chaining oneself to the basement on the night of the full moon, or directly removing all objects of temptation out of the kitchen. Other ways that the Ego can directly influence the limbic system and Reptilian Brain is through, of course, physical interventions such as exercise or breath work which have been proven to lower stress levels¹⁴⁷, but also a factor that is largely overlooked: imagination – we can’t use logic to influence the speed of our heart, but if we think of something scary (or very calming) we find that our heart rate changes in response.

A second interpretation relates to the magnitude of the mouse’s influence. Even if we can create a Physical Hexis that extends into the DMN, and even if the Ego is given free rein to direct the Self in moments of autonomy, how is this drop in the bucket not subsumed? If the mouse is too small for the elephant to see (Ego’s influence << Subconscious Forces) or the mouse is accidentally stomped on before the elephant can notice (Ego’s influence impressed upon by unfavourable deterministic influence), does having such a small influence even count as a *plausible* influence?

Understanding the magnitude of the Ego’s influence requires us first to understand the hierarchy of the long line of Puppet Masters. While the Ego’s share of the pie is bigger than Subpersonalities (particularly Alters) if not just in validity, than potency (time under fronting ≠ time under control) as well, it is still only a minor shareholder compared to whatever – and, especially, whoever - influences the Reptilian Brain:

- Behavioral change as a result of vector of Subconscious Forces
- → Subconscious Forces are a result of deterministic influences
- → → Determinism is a result of causal events
- → → → Causal events are a result of fate
- → → → → And fate is nudged by vectors

Using the Dorito Effect for emphasis:

¹⁴⁵ The social contract, in both Hobbesian and Rousseauian terms, assumes a kind of metaphysical pact—implied agreement—that legitimizes the structure of society and its constraints on freedom.

¹⁴⁶ The power structure within the psyche is unlike typical hierarchies. The Ego has limited authority, but frequent control. The Reptilian Brain, despite rarely “fronting,” often exerts the strongest subconscious influence. Corporations—while artificial constructs—function similarly to an ego by directing behavior through proxies (e.g., branding or artificial flavorings) that target primal systems.

The School of Life. *The Philosophy of Boredom*. YouTube video, 6:30. Posted April 12, 2022.

<https://www.youtube.com/watch?v=hg6XUYWj-pk>. (Cited at 2:12–2:51)

¹⁴⁷ The grandfather paradox—if you travel back in time and prevent your grandparents from meeting, you negate your own existence—remains a powerful metaphor for the instability of causal manipulation.

PBS Space Time. *The Grandfather Paradox Revisited*. YouTube video, 15:20. Posted November 5, 2021.

<https://www.youtube.com/watch?v=kSZKIupBUuc>

- Artificial flavour is vector
- → Subconscious Force is advertisement
- → → Deterministic influence is demand for snacks
- → → → demand for snacks as a result of corn oil industry and mislead science
- → → → → causal events is overgrowth of corn waste
- → → → → → corn waste because of events leading from big bang
- → → → → → → events leading from big bang dictated by human events

As we see, there is no set number of puppet masters, but the major players include the Ego itself, Alters (amongst other “conscious” Subpersonalities and/or limbic facets), deterministic influences, and external factors who play using vectors. The good news is that while shares of the pie are not distributed evenly, the line of control is not vertical where minor shareholders are beholden to major ones; instead, the line of control is circular where both the chicken and egg are subject to retrocausality¹⁴⁸ and even minor players have the potential to influence major shareholders. Moreover, it’s a feedback loop between subject-objectification and object-subjectification¹⁴⁹ where Fate is not a single entity but the very dynamic by which all actors play into and influence the hermeneutic cycle. This means that although we are slaves to Fate, we are simultaneously masters of it as well. But this in turn begs the question: if being master of Fate is contingent on being a slave to Fate and vice versa, from whence does our competency stem from? How, in our zeitgeist capacity for and from determined human nature, are we capable of making the most efficacious decisions in governing ourselves? These queries will be further developed in the next part. For now, however, we see that the magnitude of the Ego’s influence is not contingent on its share of the pie but in its skill of playing the game.

In terms of the Ego’s validity (and potency) as well as magnitude of influence, it should now be clear that the Ego legitimately influences the Psyche directly and indirectly, the potency of its influence relates to how well it (introspectively) gathers itself in navigating the Self, and its magnitude contingent in its skill in (extrospectively) directing the hermeneutic cycle and creating its own degree of plausibility.

8. *How Much Can the Drifting Will Be Nudged?*

An eight contention is with the extent to which the Drifting Will can be nudged. Following from the second contention, we will now explore the first half of the extent to the efficacy of external change (or framework). This will be done in addressing interpretations regarding the frontend impracticability (of creating and enforcing the framework) as well as the backend unfeasibility (of such a framework leading to the good).

Thus far we’ve portrayed the Free Will and Drifting Will as in opposition, one the underdog (mouse), the other a heavy weight (elephant); one the progressive upholder of liberty, the other a regressive impressionable burden. But in this contrast, it’s easy to forget that the Ego is a product of the Self, not an external observer, thus not only retaining many of its parent traits (i.e. irrationality, impressionability), but also influenced by their retrocausal dynamic.

¹⁴⁸ The chicken-and-egg dilemma is not just about origins but interdependence. Attempting to alter the past may result in recursive contradictions, rendering the entire premise self-defeating.

¹⁴⁹ The dialectic of subject–object interaction demonstrates how internal forces manifest externally and how objects, in turn, shape our internal development. This recursive interaction forms the basis for identity, habit, and individuation.

First, in regards to the Ego's irrationality, we have already seen the problem of rationalizing human behaviour. As a product of the Self, the Ego is neither Houyhnhnm¹⁵⁰, Robot, or completely rational being (let alone omniscient) so even if information failure was not a problem, people are prone to making mistakes under the TPN. Second, in our zeitgeist capacity for and from determined human nature, we can't fully¹⁵¹ detach ourselves from the facticities of our contextual placement making it very difficult to observe ourselves objectively, or even perceive ourselves through lateral thinking with withered technology¹⁵². Incidentally, this is why most innovations come not from within a certain field, but from an external source, particularly with the help of a "specialized generalist"¹⁵³, polymath, or Renaissance Man. Third, the retrocausal dynamic between the Ego and Self (i.e. Self mindlessly purchasing Doritos, because "why not?"¹⁵⁴) allows for secondary addiction much in the same way a fetus can become addicted to nicotine in the womb due to the dynamic with its host smoking mother. Taken together, these three effects of irrationality, impressionability, and retrocausality leads to two possible implications:

1. incapability of *creating* this external framework (or at least one that is efficacious) because too influenced by the Self
2. incapability of *enforcing* this external framework (at least competently) because of human nature, liability for unwittingly influenced by self-sabotaging causal event

While these implications might seem like frontend impracticability on the surface, this is not quite true if we avoid what both of them have in common: contingency on the Ego. We see that we are facing the same problem as expressed in the Background: namely, a subjectively mental endeavour; moreover, in the same way then, the answer is also the same as posited in the third contention: specifically, an objectively physical solution.

A second interpretation of the contention addresses backend unfeasibility; namely, that the validity (and potency) as well as magnitude of the Ego's influence doesn't matter if the Drifting Will isn't even open for nudging., e.g. it doesn't matter how skilled or visionary an artist is if the audience is blind. We already alluded to this interpretation when we distinguished between fate and destiny; specifically, when expositing how predisposition (alcoholism) is necessarily activated by condition (availability of alcohol)¹⁵⁵. But while it

¹⁵⁰ In *Gulliver's Travels*, Swift's Houyhnhnms live by pure logic and unblemished reason, so detached from vice they lack even the concept of lying.

Swift, Jonathan. *Gulliver's Travels*. Edited by Robert DeMaria Jr. New York: Penguin Classics, 2003. (Original work published 1726)

¹⁵¹ While metaphysical detachment from context is impossible, it is possible to dilute contextual influence through selective informant exposure—such as curated social roles or digital avatars—in the spirit of Nietzschean perspectivism.

¹⁵² Rituals help organize chaos, provide meaning, and regulate group dynamics by anchoring abstract values in repeated symbolic actions.

National Geographic. *The Psychology of Rituals*. YouTube video, 6:18. Posted June 18, 2022.

<https://www.youtube.com/watch?v=Js4NiTFq9Pw>. (Cited at 3:08–4:29)

¹⁵³ Procrastination is not always a defect of will but a complex interplay of task aversion and emotional regulation. Those with wide but deep experience—specialized generalists—may navigate such terrain with creative flexibility.

SciShow Psych. *Why We Procrastinate*. YouTube video, 4:45. Posted January 9, 2023.

<https://www.youtube.com/watch?v=29oyOLh6PNE>. (Cited at 0:54–1:32)

Harvard University. *The Future of Free Will Debate*. YouTube video, 18:30. Posted March 22, 2023.

<https://www.youtube.com/watch?v=XJsCM7xyETc>. (Cited at 9:46–12:36)

¹⁵⁴ "Why not?" is often not an argument but a placeholder for absent conviction—a shrug that silently endorses impulse.

¹⁵⁵ It is insufficient to recognize cognitive bias without asking what to do about it. If biases are rooted in nature or nurture, can they be reprogrammed through gene editing or epistemic curation?

may appear that nature is thus only empowered by nurture, this is not necessarily the case: a normal child nurtured under horrific circumstances may be the makings of a psychopath, but a psychopath by nature nurtured in a loving environment will still turn into a monster¹⁵⁶. It would appear then that if not how far the scales tip being idiosyncratic, then at least that the external framework of an environment may enable, but not necessarily disable.

To make the extent of an external framework's efficacy vivid, we can imagine an untouchable¹⁵⁷ born in India and a legitimate child of two Hollywood stars. It doesn't matter how favourable the environment the untouchable grows up in, whether they are blessed with genius and even earn a full ride scholarship, it is highly unlikely that they will grow up to be a middle eastern oil mogul. Conversely, it doesn't matter how unfavourable the environment the star child grows up in, whether they are born with a mentally debilitating disease or have both parents pass away at an early age, it is still highly likely that they could grow up to become a notable movie director. In both these cases, we see that one's hereditary "privilege" plays a large role in one's fate and while the extent to which an external framework can nudge one's fate isn't without limits, the fact that there is any extent at all means that inciting external change is not unfeasible.

Specifically, then, the extent to the efficacy of external change (or framework) is neither capped at frontend impracticability, nor backend unfeasibility. The Drifting Will can be nudged efficaciously so long as there is focus on an objectively physical solution instead of a subjectively mental endeavour, i.e. a 3rd party framework that is not the Ego – and by extension, not impotent people – that governs itself. Moreover, while the extent that the Drifting Will can be nudged isn't unlimited, our fates can still switch lanes so long as they are parallel in proximity.

9. *The Reality of an External, Parameterized Framework*

A ninth (or 8.2) contention is with the reality of an external framework which explores the second half of the extent to the efficacy of external change (or framework). Namely, this includes addressing interpretations regarding the tenability of an external framework as well as the consequences and potential drawbacks of any parameterized, external change.

Just now, we hinted that a 3rd party framework, that is not the Ego – and by extension, not impotent people – that governs itself, would be needed to circumvent the inherent problem of subjectively mental endeavours. Expanding on that specific notion of frontend impracticability here would be beyond the scope of this part and instead will be addressed in more detail in the following part; for now though, we can simply understand it as follows: humans have rational capacities, but are overwhelmingly creatures of sentiment. Human nature, therefore, makes it very hard to competently (and, perhaps, objectively¹⁵⁸) govern itself without rational training. Insofar as an Ego led external framework would thus be a

¹⁵⁶ Historical revisionism surrounding Lenin's legacy should be approached cautiously; the idealized narrative often obscures the brutality of the means.

Kisin, Konstantin. "The Truth About Vladimir Lenin and the Russian Revolution." *Konstantin Kisin (Substack)*, September 15, 2023. <https://www.konstantinkisin.com/p/the-truth-about-vladimir-lenin-a>

¹⁵⁷ Dalits—formerly "untouchables"—exist outside India's fourfold caste system, which continues to enforce deep socio-economic divides.

¹⁵⁸ Traditional bureaucracy fails not from lack of structure but from its excess. Detached rationality flattens human nuance, making some problems insoluble simply because they don't fit a predefined rubric. Also referenced in: *我不是潘金莲* (*I Am Not Madame Bovary*)

mechanized adaptation of Kant's *Reich Der Zwecke*, it is untenable by paradox¹⁵⁹. This echoes the earlier interpretation of needing a Physical Hexis as opposed to a comprehensive, Virtuous Hexis which would necessarily include aspects of mental acuity as well. What then would a tenable 3rd party framework look like?

One approach could be the "child-locking" mechanism we alluded to earlier. Revisiting the Marshmallow Test as a case in point, a child-locking solution would be an ascetic removal (which itself could lead to contentions around abstinence¹⁶⁰) of the marshmallow in the first place. Another rendition of decreasing the margin for error could be related to losing one's keys¹⁶¹ such as having a mechanism which shrinks the potential area for unfavourable outcomes (in this case, having a set area where keys are "muscle memory"-like placed).

A second approach could be similar to the Physical Hexis: an external system (i.e. a law, personal trainer that holds you accountable, social contract, or even algorithm) stemming from outside the Ego's direct influence which capitalizes on the Subconscious behavioural trends whilst within predetermined parameters. The contention thus lies not in what this 3rd party framework would look like, but how it would come about. This need (and construction) of this external framework, however, is a topic of next part: *Poeisis*.

Another interpretation of the contention is with consequences and potential drawbacks of any parameterized, external change. Specifically, in parallel to the concern of if the good life is really the best aim; how do we know that an external framework will necessarily be good for us? How do we know that efforts to create, implement, and eventually the results that follow won't be counterproductive? How do we know that we do not cut off a crucial element when parameterizing? Simply: we don't. But politics as a science is not about getting it right, it's about getting it better. Progress (as an effect of knowledge) is not positivistic – it does not accumulate linearly – and neither is the progression comprehensive either where the world of knowledge is often determined by people who are too male, too pale, too stale¹⁶².

One legitimate worry of this march into the unknown, however, is the potential for the risk to outweigh the rewards. This worry is not only a hypothetical concerning itself with the potential drawbacks (i.e. boredom and discontent) of (political) utopia, but also a historical one insofar as it stems from the very social conditions that led Dostoevsky amongst other contenders to be so apprehensive about any social system justifying itself on a felicific calculus¹⁶³. Specifically, the views of the socialist reforms in 19th century Russia, and subsequent social engineering, failed so magnificently that the consequences were entirely

¹⁵⁹ If the Ego directly leads an external framework, it forfeits its reflective capacity—what makes the external useful is that it exists independently from the Ego's limitations. To lead from within is to contaminate the scaffolding.

¹⁶⁰ Asceticism may be misread as denial, but it can also be a rational decision to preemptively disengage from temptations one knows they cannot withstand—not escape, but wise abstinence.

¹⁶¹ The phenomenon of misplacing objects—such as losing car keys—mirrors the self-sabotaging behavior observed in the Marshmallow Test. The inability to recall the moment of loss suggests the memory was encoded not by the conscious Ego but by a subpersonality. The same dynamic applies to impulsive acts: memory formation is contingent on which personality was "fronting" at the time of the behavior.

¹⁶² The structure of scientific knowledge is non-linear, resembling a pendulum swinging between opposing forces or skipping over intermediate steps to reach new paradigms. Progress lies not just in answers but in refining the quality of our questions.

University of Chicago. "Larry McEnerney: The Problem of the Problem." YouTube video, 58:02. Posted August 15, 2014. <https://www.youtube.com/watch?v=vtIzMaLkCaM>. (Cited at 28:00–31:30)

¹⁶³ Dostoevsky warns that mechanical or overly rational social systems disregard the complexity of human nature. In *Crime and Punishment*, he argues that "the soul won't obey the rules of mechanics," criticizing utopian models that reduce people to parts in a formula. Dostoevsky, Fyodor. *Crime and Punishment*. Translated by Richard Pevear and Larissa Volokhonsky. New York: Vintage, 1993. (Original work published 1866).

counter-productive¹⁶⁴. Can we 100% in confidence thus claim that a 3rd party external framework will be the best thing going forward? Again, the answer is no. But, once again, we must remember the spirit of science is about the journey, not the destination; moreover, the World waits for no Man: we either change with Nature or stagnate and be left behind. Perhaps then in a rather crude approximation: the consequences of 2nd world war was not kind to Germany, but the consequences of carpet bombing can be said to have been a formative factor in rebuilding one of the greatest economies we have today. The saplings of Easter cannot spring forth without the decay following All Hallows' Eve; ostensibly then, it is better to risk destroying civilization in a gamble for humanity's future than to slowly rot and die in stagnation¹⁶⁵.

In our final contention then, we have addressed the tenability of an external framework as an external system such as a law, system, social contract, or even algorithm that stems from outside the Ego's direct influence. We have also addressed the consequences and potential drawbacks of any parameterized, external change; specifically, we identified the goodness of the 3rd party external framework not in its consequences, but in its potential to overcome the zeitgeist of our epistemic placement, adapting contextual endosmosis to the spirit of scientific progress. What matters most isn't so much what the external framework looks like or the consequences (be they favourable or unfavourable) that follow as much as how said parameterized framework allows us to escape decay, for stagnation is death.

Caveats

Although not a contention in itself, it's important to draw up (and revisit) the caveats, as well as limitations and shortcomings of the data and results. This is where I now critically examine my own results in the light of the previous state of the subject as outlined in the Background, and make judgments as to what has been learnt in my work thus far. Moreover, acknowledging potential issues, offering disclaimers, and providing additional context will ensure that there is a comprehensive view of relevant factors which will necessarily help inform our views when coming to implications of the study.

As aforementioned in the results, the data gathered is an interpretation of the research theory and methods; therefore, in the spirit of scholarly dialogue, it is imperative that I now present warnings or cautions regarding the use or interpretation of information, findings, and/or recommendations.

First and foremost is the need to acknowledge the many assumptions that have been made, not just in my own research but in the research that my research is based upon. These assumptions include, but are not limited to:

1. The Good –
 - i. That happiness must be facilitated / it is not suffering
 - ii. That it is wanted
 - iii. That it is yet to be realized
 - iv. That it is objective / moral progress exists
2. The (Scientific) Method –

¹⁶⁴ From a socialist viewpoint, crime may be understood as a form of protest against an unjust or abnormal social structure. Under a "normal" structure, crime would disappear because its *raison d'être*—dissonance—would be removed.

¹⁶⁵ 树挪死，人挪活 — "A tree dies when moved, but a person thrives." This Chinese proverb captures a cultural value of change and adaptability, contrary to Western notions of rootedness or stasis. Weixin (WeChat). "Gettier and the Limits of Knowledge." Article ID 1707659000. Accessed September 30, 2023. <https://mp.weixin.qq.com/template/article/1707659000/index.html>.

- i. That it is sound
 - ii. That it is valid
 - iii. That it is (relatively) complete
 - iv. That it is coherent
- 3. The Self –
 - i. That the Psyche is concentric
 - ii. That the DMN takes precedence over the TPN
 - iii. That DID and AI exist on a spectrum
 - iv. That our identities persist through time
- 4. Free Will / Freedom –
 - i. That compatibilism is true
 - ii. That it is not an illusion
 - iii. That destiny is subordinate to fate
 - iv. That events are determined
- 5. Etc.

This list is far from being comprehensive, but what's important to note isn't the length (it is not necessarily the case that the more assumptions there are, the more serious the cumulative error is if any is proven wrong¹⁶⁶), but the value of each assumption. For example, if it turns out that the Psyche is not concentric, and subsequently that the idea of the Cartesian Theatre, Subpersonalities, Reptilian Brain, etc., are all proven wrong, it really doesn't matter. The value of this assumption is simply to build a foundation that explains why people should not (and cannot really) make their own decisions. There are perhaps infinite ways to arrive at the same conclusion and perhaps even more schools of thought ranging from Aristotle to Li Ao, from Kant to everyday anecdotal experience that can provide evidence in support. The key reason I chose the specific assumptions that I have is because rather than attacking a strawman argument, I want to test beyond a shadow of a doubt the legitimacy of my thesis which necessarily means identifying (and sometimes building) the strongest argument possible in opposition to try and tear my thesis down.

Second, is the problem in the so-called “solving of problems”¹⁶⁷, which is circular reasoning in itself. This should be paid special attention to, especially considering its commonality in our post truth era where sensible debate is avoided simply because people argue over different definitions of “personal truth” instead of for *the Truth*. This is particularly worrying as this problem extends to the hard sciences, let alone the softer social sciences¹⁶⁸. The greatest problem of making assumption then is that even constants are prone to fluctuations, meaning they aren't actually constant at all. Moreover, steamrolling these problems through *fixing by definition* will inevitably cause problems down the line. Nonetheless, this is not a thesis on metrology, so while we will keep using flawed assumptions, in the same way physics treats Earth as a ball for simplified calculus, it should always be kept in the back of the mind that while conclusions may be circumstantially correct, the answers are still subject to deterministic influences.

¹⁶⁶ To guard against foundational collapse, this thesis relies on parallel argumentation rather than a single pillar of assumptions. Even if one line of reasoning fails, the broader framework remains intact.

¹⁶⁷ Gettier problems show that justified true belief is not sufficient for knowledge. The Gettier challenge highlights how seemingly rational belief systems can be upended by hidden variables. Stanford Online. “Gettier Problems in Epistemology.” YouTube video, 22:45. Posted February 10, 2023. <https://www.youtube.com/watch?v=sF03FN37i5w>. (Cited at 10:33–14:10)

¹⁶⁸ Some physical constants (e.g. the speed of light, Newton's G) have been redefined rather than revised to accommodate variability, effectively side-stepping inconsistencies. This reflects a deeper philosophical point: science often retrofits its foundation to preserve coherence rather than expose uncertainty.

Third, my stances (classical utilitarian attitude in regards to happiness as a marker; uncontrollable sources of (un)happiness cannot outweigh controlled sources; neuroscientific model understanding of the mind; etc.) have necessarily coloured my understanding in addition to interpretations of the results, both of mine and others. In parallel what was said earlier about epistemic placement, it is impossible to truly view this as a weakness or deficiency to be improved, but should nevertheless be noted as a caveat.

Fourth, and perhaps the greatest caution of all, should be in regards to you, dear reader. If you find yourself easily convinced by my argument, that is only a telling warning sign that something is greatly amiss. If you do not yourself, find the background lacking, outcomes incredulous, and interpretations biased, then, in all respects, you have failed to do your job just as I have failed to do mine. For that, I am sorry and can only offer caution that you do not convince too easily moving forward.

Limitations and Shortcomings

In the Aim and Scope we noted from the start that there would be limitations to the research such as not examining related discussions outside the Western philosophical traditions, narrowing the emphasis to happiness as a marker of the Good Life, and focusing on external change, as opposed to using psychotherapeutic models to heal the pathology of *akrasia* (as an affect of the Drifting Will). Now that the preliminary research phase is over, we can elaborate on more inherent constraints (limitations) of the data and results in addition to identifying weaknesses and deficiencies (shortcomings) which can be improved upon. Taken together, the limitations and shortcomings will provide further context in which the findings should be interpreted and implications drawn.

Of the data

First, I am relatively better read in philosophy compared to the other subjects of psychology and neuroscience. This means I am likely to have missed key arguments in the literature, misplaced emphasis on the wrong areas, or simply misinterpreted the data due to difference in method and methodology.

Second, in continuation from the above limitation, some of the literature sought after was materially unavailable. These pieces could have contained some of the key arguments or led to a different interpretation and thus taking my thesis on an entirely different trajectory. Nonetheless, as this limitation was foreseen – and preemptively addressed in different areas above; i.e. acknowledgement of the non-positivistic world of knowledge – the effect of this limitation is likely minimal¹⁶⁹.

Third, as a human myself, I, too, am prone to limitations of the Ego. This includes the potential to lose sight between target and measure and falling prey to Goodheart's Law as alluded to in the Background. Namely, there are possible instances in my research where I may have confused the target becoming a measure: are we aiming for the Good or happiness?

Fourth, not so much a constraint as a creative – and traditionally “unprofessional” – liberty is my unconventional method of literary interpretation. Of course, my purview is that

¹⁶⁹ It is a fallacy to believe that Kant—or any canonical philosopher—is essential to doing good philosophy. One can engage in rigorous philosophical inquiry without being fluent in the classics. The purpose of general history courses is not to create philosophers, but to transmit foundational knowledge. Weixin (WeChat). “Gettier and the Limits of Knowledge.” Article ID 1707659000. Accessed September 30, 2023.

my method is progressive for the better, but I also acknowledge that there are reasons for some mediums to be deemed "less than scholarly"¹⁷⁰.

Fifth, one major constraint is focusing solely on exogenic subsidiaries. Having already touched on how the endogenic dimensions appear to hold greater weight, the largest weakness of narrowing the study to exogenic dimensions of happiness is that it might be inadequate as a marker by itself. This narrowing, however, was essential for the scope of this thesis as previously mentioned, but this inadequacy could be improved in future editions with more time, and more funding.

Of the results

First, of the research itself: the aim is high, the scope is extensive and time is under limit. This meant that while the data gathered stems from decades of research, my results and subsequent interpretations are somewhat rushed. Ideally, I could have another two or three years, but that is what revisions and subsequent editions are for. After all, the function of a piece of writing is to move a research area forward, not to be preserved for 500 years.

Second, in parallel to being a newcomer (at least in the professional sense) to other disciplines outside of standard philosophy, I am unfamiliar with the audience, code of research community, and academic citizenry of the non-philosophical disciplines. This is definitely both a limitation and shortcoming when it comes to not only presenting my results, but also asking for feedback (for lack of network, and code to communicate within), or in general participating in academic discourse in external fields of interest.

Third, in continuation of not knowing the professional code of language used within specific disciplines, this makes it difficult to convey ideas and translate them appropriately across disciplines. Moreover, this extends to academic writing in general where, as a personal inadequacy, while I endeavor to write with clarity with my target audience in mind, my writing is still geared towards helping me think. Specifically, due to my epistemic placement, my writing is generated on the horizontal axis, but my audience – who do not, and cannot share the same placement – read on the vertical axis. Nonetheless, as this is my first time writing a thesis of this magnitude, I'm sure this inadequacy will be resolved as I gain more experience.

Implications

Having finished addressing the initial contentions as well as caveats, limitations, and shortcomings, we can now draw together the core contribution to this part of the thesis: the final comprehensive implications. This necessarily includes the summary of the results and evaluating the discussion thus far insofar as they respond directly¹⁷¹ to the aim of this part as stated in the introduction.

To reiterate, the singular, paramount aim of this part was to *engage in the preliminary investigation of turning our Self as a whole towards the good*, particularly, in profiling the predisposition to self-sabotage actualization of the Good From the Background we were introduced to the first hurdle where focusing on willpower alone is an insufficient approach and this substantiates a need for further investigation on how to turn our Self as a whole

¹⁷⁰ Relying on non-scholarly sources entails risk. A commonly cited statistic—"95% of our daily activities are subconscious"—lacks verifiable academic origin, despite frequent attribution to the Max Planck Institute. While the number is illustrative, not factual, it still serves to underscore the broader claim: that a majority of our decisions are suboptimal and nonconscious.

¹⁷¹ Just as each theoretical division invites its own questions, so too does each imply downstream consequences. This thesis has prioritized the most formative of those implications rather than exhaustively listing each one.

towards the Good; specifically, the Ego (or at least willpower as an effect of the Ego) alone being an insufficient influence in this turning. This spurred us to identify, and hopefully reconcile the missing key of a limited ego with failure to reach the goal of the Good Life.

Moreover, it allowed us to formalize a preliminary hypothesis: we are slaves (to Subconscious Forces that are really in control of us).

As the research progressed, however, we came to realize that our hypothesis was not entirely correct. Delving deeper into several key research questions surrounding the pluralism of the Persona, autonomy of the Ego, and influence of Subconscious Forces, we came to a more nuanced understanding. Particularly, in profiling the predisposition to self-sabotage actualization of the Good, we became aware that Subconscious Forces are only one of the many puppet masters that we are slaves to. Furthermore, this enhanced understanding of hermeneutic cycle dynamic and the Ego's master/slave dualism allowed for a clearer picture of why people should not (and cannot really) make their own decisions.

In relation to Subconscious self-sabotage, we see that:

- (a) Willingly – Altars can have counterproductive agendas; and
- (b) Unwittingly – unfavourable influences directly direct the Subpersonalities

In relation to Conscious self-sabotage, we see that:

- (a) Willingly – our Ego is prone to akratic action; and
- (b) Unwittingly – unfavourable influences indirectly impression upon the Ego

The outcomes of the study, therefore, reveal that it doesn't matter whether our actions are performed under the TPN or the DMN; in case of the latter, the Ego fails to make any direct decision, and in the case of the former, the data suggests that even when acting on our own *supposed* volition, we are still prone to acts of self-sabotage. The Self as a whole then possesses a general proclivity¹⁷² for incompetency where it is not only biased towards suboptimal action, but incompetent decisions as well. Here then, it is important to reiterate that these factors combined means even if information failure was not a problem, even if we were not predisposed to unfavourable genetics, and even if we were nudged by favourable conditions, our inherent human nature makes it thus that we fail to properly aim for what is good for us. In other words, we fail to introspectively foresee how future actions would make us feel and aim for the wrong things in our weakened human internal perspective. This accumulates in a general impotence of the Ego, at least from an internal perspective, in turning the Self as a whole towards the Good.

Interpreting the results in light of the preliminary hypothesis, we see that we were mistaken about Subconscious Forces and Psychic Noise in quite a few ways, including the potency of their influence, the extent to their control, and the motivation behind their self-sabotaging behaviour. Nonetheless, while we were mistaken in our hypothesis insofar as to whom (or what) exactly the Ego was enslaved to, we were still correct as to the reality of the Ego's enslavement.

In evaluating the discussion (primarily the contentions) insofar as they respond directly to the aim, we are now much better position to identify alternative approaches (mechanisms, systems, etc.) to focusing on willpower (as an effect of the Ego) alone. Specifically, having established - and rather consolidated - that people should not (and cannot really) make their own decisions; we identified that to reconcile the missing key of a limited ego with failure to reach the goal of the Good Life, there needed to be an external change by way of 3rd party framework. It was posited that these mechanisms ("child-locking") and systems (i.e. a law, personal trainer that holds you accountable, social contract, or even algorithm) promoted the alternative approaches to efficaciously promote external change,

¹⁷² Incompetence is embedded across both the task-positive (TPN) and default mode networks (DMN), suggesting that piecemeal intervention—targeting only the Ego or the Subconscious—is inadequate. A gestalt-based model that treats the psyche as an interconnected whole offers more promise than isolated improvements.

alignment, and ultimately turning our Self as a whole towards the good. In one rendition, the proposed mechanisms would shrink the potential area for unfavourable outcomes, i.e. a set area where keys are “muscle memory”-like placed. In another, systems capitalize on the Subconscious behavioural trends whilst within predetermined parameters., i.e. you cannot efficaciously force children’s behavioral trend of watching television into a love for reading, but by adjusting sound to subtitles, one can capitalize on their behavioral tendencies, promoting reading without direct supervision.

When abstracted to a generalized result, the framework provides two distinct approaches:

- (1) Mechanisms – decrease the margin for error by shrinking the potential area for Psychic Noise and unfavourable outcomes
- (2) Systems – capitalize on behavioural trends to leverage leading “muscle memory” continuity

Both mechanisms and systems are viable alternative approaches that reconcile the missing key of a limited ego with failure to reach the goal of the Good Life; however, while the external framework of either could ostensibly complement the Ego’s general impotence from an internal perspective, why limit the solution to a single approach when both can be combined?

Absorbing the various renditions of the above outcomes and synthesis as they respond directly to the aim in addition to our revised understanding of the hypothesis, we at last come to the final comprehensive implications as follows:

1. Psychotherapeutic models stemming from the Ego’s internal perspective are inadequate in turning the Self as a whole towards the Good.
2. The Ego can extend its extrospective influence over the General Will through indirect cultivation of an external framework.
3. A 3rd party framework by way of shrinking the margin for error and leveraging behavioural trends complements the strengths and weaknesses of the Ego.

In sum, we can form a singular implication as it responds to the singular aim: turning our Self as a whole towards the Good necessitates a 3rd party framework.

The Conclusion

Through the deduction of our final comprehensive implications, we necessarily indicated how we fulfilled the aim. The conclusion here is thus mainly just a formality (TL;DR) for those dear readers who simply skipped to the end.

For this conclusion segment then, we will recapitulate, by way of walkthrough, what we aimed to say, what we have said, and what the implications of are from that which was said. To prevent simply repeating that which is written elsewhere, there will be an emphasis on the link between the Discussion and Conclusion in contrast to the previous emphasis on the link with the Introduction. The main point of this conclusion segment then can be understood as a stepping stone from which the singular implication (conceptual 3rd party framework) transforms into a practical application in the second part of the thesis: Poiesis.

Tying the Introduction through to the Synthesis we see that:

- i. **People are not any happier when they make their own decisions**
– (Background + Core) = first intended outcome; this is because:
 - a. **There exists a prospection/retrospection asymmetry.**
As seen with Gilbert's photo selection satisfaction study, when faced with freedom of choice, greater freedom correlates with inverse increases in happiness. Incidentally, it also leads to greater analysis paralysis which further frustrates happiness.
 - b. **People misemphasize material sources of happiness over character.**
We fail to introspectively foresee how future actions would make us feel in the same way our misinformed predictions lead us to believe that (holding disposition as a constant) a king winning the lottery would be happier than a street sweeper getting into an accident leading to paralysis.
 - c. **Our weakened human internal perspective misguides our aims.**
Even if information failure, in addition to unfavourable genetics and conditions, were not a problem, inherent avenues for impressionability under the TPN clouds our judgement of what to aim for and misdirects our pursuit of the Good.

First and foremost, we have accomplished the first intended outcome of clarifying the relationship between the Good Life and happiness. Specifically, we deduced that when making own decision, people not happy → then maybe they shouldn't make their own decisions.

- ii. **People are not led to better outcomes when they make their own decisions**
– (Core + Synthesis) = second intended outcome; this is because:
 - a. **The Ego possesses a general proclivity for incompetence.**
At least from an internal perspective, conscious self-sabotage extends from both willing (proneness to akrasia) and unwitting (as a result of impressionability) dimensions. Depending on how well the Ego can introspectively gather itself thus will directly reflect on its potency to extrospectively mitigate its influenced proclivity for incompetence.
 - b. **Frontend impracticality arises from contingency on the Ego.**
The retrocausal dynamic between the Ego and Self means whenever the Ego directly participates in decision making, there is an associated risk of being incapable of *creating* and *enforcing* parameters against unfavourable outcomes. This, in addition to maintaining many of its parent traits, fosters a literal conflict of interest in the Ego's decision-making process which relates to the aforementioned proclivity for incompetence.
 - c. **The Ego is held captive by external constraints.**

Granted that the Ego is able to escape destiny (manages to collect itself, introspectively see through unfavourable influences, and take full control of the Self); still, the Ego alone cannot overcome the zeitgeist of its epistemic placement without renouncing its autonomy. The Ego can nudge Fate, but not escape it – or its associated constraints - in its entirety.

Second of all, we have accomplished the second intended outcome of identifying to what extent the effects of willpower extend to Subpersonalities. Specifically, we deduced that when making own decision (which are not even completely own), people not happy + liable to lead to unfavourable outcomes → then maybe the shouldn't make their own decisions

iii. **People are not any freer when they make their own decisions**

– (Discussion: Implications) = third intended outcome; this is because:

a. Unheimlichkeit.

Just as epistemic placement (i.e. informed by our humanity) grounds, and thus empowers us, the Drifting Will's constraints pose a lodestone to orient our freedom. By appreciating our entangled enslavement to Subconscious Forces (amongst other Puppet Masters), we remove the delimited illusion of Egoistic autonomy which liberates the true capacity for Egoistic autonomy. Conversely, this Chinese finger trap means ignoring our placement and acting on our own further restricts our freedom.

b. Morality is not solely under the jurisdiction of the Individual.

Since decisions are influenced by society – which makes the State one of the Puppet Masters – the psychology of the collective must be taken into account. Making decisions in ignorance of the buffered personality dissociates ourselves from the capacity for freedom.

c. The raw Ego has an innate margin for error.

Turning our Self as a whole towards the Good necessitates a 3rd party framework. Decisions stemming purely from an internal perspective is often hoodwinked into performative actions on behalf of the General Will, further entrenching the framework for the Ego's ineptitude making following decisions even less likely to be free.

Last but not least, we have accomplished the third intended outcome of contributing to the solution of a practical problem, namely possible avenues to overcome failure to reach the goal of the Good Life. Specifically, we deduced that when making own decision (which are not even completely own), people not happy + liable to lead to unfavourable outcomes + nudged by hermeneutic slave dynamic → then maybe the people should not (and cannot really) make their own decisions

In a nutshell, the golden thread that has been running throughout this part of the thesis not just explicates how and why people should not (and cannot really) make their own decisions, but has repeatedly implicated from various, comprehensive angles possible avenues to overcome failure to reach the goal of the Good Life. Once again then, to clarify the link between the Discussion and Conclusion:

1. People aren't happier when they make their own decisions
2. People aren't better when they make their own decisions
3. People aren't freer when they make their own decisions

4. Therefore, people should not make their own decisions

Finally, then, we arrive at the cusp of the last intended outcome, which is to adequately feed into the next part of my thesis. In light of the Discussion (particularly the

Caveats, Limitations, and Shortcomings), the argument thus far has been a long ways away from being perfect, let alone substantial. Yet, it would appear to at least remain adequate to support that turning our Self as a whole towards the Good necessitates a 3rd party framework. Thus, moving onto the next part, we, ourselves, turn from profiling the predisposition to self-sabotage actualization of the Good to evaluating 3rd party frameworks for practically implementing utilitarian contrivance of the Good. As a closing metaphor which hopefully encapsulates this entire part: the ripples of Puppet Masters surely rock the boat, but rather than blindly flowing down a rivulet of Fate, why not choose our Charon instead?

Part 2
Poeisis: The Makings of a True Philosopher King

Abstract

Poiesis picks up where Theoria leaves off: if individuals cannot reliably reach the Good Life through internal governance, then we must ask who—or what—should govern. This part reexamines Plato’s Philosopher King as an archetype of benevolent, rational rule. Far from being an anachronism, the Philosopher King is positioned as a structurally necessary corrective to egoic governance, providing a frame for moral clarity, expertise, and holistic optimization. Yet history shows that all attempts at elite human leadership collapse into corruption, ideology, or inefficiency. The problem is not the model, but the medium: humans are unfit vessels for Philosopher Kingship.

The chapter proceeds through a tripartite analysis: tenability (can such rule exist?), practicality (would it be accepted?), and efficacy (would it produce the Good Life?). These questions reveal that the Philosopher King is not just politically impossible, but biologically unviable. Nevertheless, the core logic remains sound. To preserve its ethical utility, Poiesis reconfigures the model as a posthuman ideal—abstracting the Philosopher King into a system rather than a person.

This reframing makes space for alternative governors: AI, transhuman elites, algorithmic advisors, or infrastructural architectures. It retains the Platonic commitment to epistemic virtue but divorces it from flesh-and-blood rule. In this way, Poiesis serves as a philosophical hinge—bridging the failure of human ethics with the potential of technological sovereignty. It invites us to take seriously the idea that salvation may lie not in better people, but in better systems—precursors to the Cyberocratic vision explored in Part 3.

Key words: governance, elitism, Plato, legitimacy, technocracy

Introduction

Context of the Study

With a 3rd party framework that capitalizes on parameterizing behavioral trends while no longer under the TPN, should we finally be able acquire the Good Life?

From the previous part, it would seem that an external framework that complements the limitations of the Ego would lead to better deterministic influences, if not overall outcomes. At the very least, an external framework so long as it is tenable, practical, and efficacious would be an improvement along the axis of the Good.

What though would this framework look like, and how would it be executed? This question lays the foundation of all organised communities, commonly manifesting as “politics”. Particularly, the ideal of a “benevolent despot” who would guide the “common herd”¹⁷³ has been sought through the Islamic Caliphs and Hindu Dharmic Rulers to the Confucian Junzi and Native American Wisdom Leaders.

While not always identical in every aspect, justifications for a *competent judge*-like government are scattered across the myriad of cultures and philosophical traditions. In these instances, the role of expertise and specialisation has been a matter of consequence regarding equality of opportunity and equity of expression.

Nonetheless, although this ideal has been sought after since the formation of the very first societies, the reality of social engineering has often fallen prey to totalitarianism. Moreover, this fall has historically been found in the wake of good intentions, such as the transition of populist to demagogue, thus consolidating a general aversion to political idealism.

Problem Statement

In response to the problem of demagoguery and Candyman were the frameworks of navigators and doctors. Some of these frameworks became commonplace enough to become household terms... anarchy, with some even developing into their own entire schools of thought: Marxism. But perhaps the most renowned archetype and existing extensive framework is Plato’s Philosopher King, and by necessity¹⁷⁴, the Kallipolis as well.

Yet, despite millennia of contentions (as well as failed actualizations) since Plato’s original framework, the perennial need for Platonic elitism¹⁷⁵ is still strong today. This is because there will always be need of a Doctor of State so long as the blight of demagoguery still festers, furthermore, the role of expertise and specialization is not one that becomes obsolete. Thus, while Plato’s original framework has its limitations and shortcomings, the underlying themes of elitism, expertise, and specialization will always substantiate the ideal framework.

¹⁷³ The “Shepherd Archetype”—a benevolent ruler who guides the common people—has appeared across cultures and traditions, from Confucian Junzi to Hindu Dharmic Rulers. While often exemplified by figures like Jesus Christ, the archetype predates Christian theology. Its Enlightenment-era secular expression appears in concepts of enlightened despotism, as supported by thinkers like Voltaire and rulers such as Frederick the Great.

¹⁷⁴ The Philosopher King is not ontologically separate from the Kallipolis any more than an organ is from the body. Plato presents them as mutually constitutive: epistemically, the King is formed by the city’s education; ontologically, the city cannot fully exist without such a ruler; and axiologically, they co-define each other’s values. As such, the terms are used interchangeably in this thesis.

¹⁷⁵ By elitism, I mean governance by competence. Just as we trust doctors for medical advice due to specialized knowledge, we might similarly imagine “doctors of state” presiding over politics.

As reviews reveal how the field is transforming, it is clear that applying new understandings is a key point of concern for substantiating the ideal framework for acquiring the Good Life. We see that focusing on frameworks in a historical vacuum is insufficient and that there exists a need for further investigation on how to adapt the original framework to become the successful 3rd party framework.

Aim and Scope

The singular, paramount aim of this part is to engage in this substantive investigation of a 3rd party framework turning our Self as a whole towards the Good, particularly, in evaluating frameworks for practically implementing utilitarian contrivance of the good. Limitations to the research are noted from the start – as well as continuing from the early part – and caveats are revisited later in the discussion. For example, I *still* do not examine related discussions outside the Western philosophical traditions. The background I provide is *again* far from comprehensive and, for this part, the research is narrowed down to an emphasis on Plato's Philosopher King as the model framework. This part *remains* to place emphasis on a practical ethics on a collective scale rather than an individual one. The reasons for these limitations have not changed: the investigation *continues* to be a complex, interdisciplinary area.

In this part, limitations on resources have not narrowed from those already published and publicly available. Nonetheless, one further addition is that several sections of this part will include and refer to personal case studies conducted previously. It should be noted, however, that these case studies will be referred to retroactively and in a context which is tangential to the study at best. This is simply because the case studies were not conducted with this thesis in mind, but insofar as they provide specific examples from abstraction, they will still be useful.

Significance of the Study

The overall value of this investigation is contingent primarily on its potential impacts. One intended outcome of the investigation, on a theoretical level, is to clear up misconception about the relationship between forms of government, i.e., democracy = good; elitism = bad.

Another intended outcome is advancing knowledge in this field by renewed investigation on the Platonic ideal. Specifically, I focus on the tenability, practicality, and efficacy of revamped Philosopher King.

Further, on a practical level, a third intended outcome contributes to the solution of a practical problem, namely possible 3rd party frameworks that capitalize on parameterizing behavioral trends while no longer under the TPN to turn our Self as a whole towards the Good.

Finally, as ever, the more immediate concern of this part is to adequately feed into the next part of my thesis whilst mediating between the former and latter parts.

Overview of the Study

The purpose of this part is to expand on turning the Self as a whole towards the Good Life by way of external change. In order to do this, this part will focus on exploring adaptive framework(s) to integrate into a modified 3rd party external framework.

The Introduction, as above, begins with an expanded review of the “benevolent despot” trope in relation to realizing the Good Life. After introducing the context of the

study, it formalizes the problem statement with further theory behind why we still do not have a benevolent despot (Philosopher King) not least for lack of trying or conflation with totalitarianism as common with elitist frameworks. This is followed by clarifying a singular aim (further investigation of turning our Self as a whole towards the Good) and laying the scope of the study. The penultimate section explores the significance of the study, overall value, and expected outcomes before arriving at the overview of the study where we are now sketching out what we aim to say.

The Background formalizes and clarifies definitions and concepts used throughout the remainder of the thesis; contextualizes the problem through salient history of issues to date; and posits a few tentative judgements before segueing into the Core segments. This narrative includes justifying the need for elitism, addressing critiques of Plato's original framework for the Philosopher King and Kallipolis, as well as contending with the realizability of a modern-day Philosopher King.

The Core is further divided into the Contribution (position and method of the research programme) and Outcome (data, statistics, and analysis from the literature). The Core elaborates on the research results as they relate to a new solution by way of external change. This includes analysis On Tenability, On Practicality, and On Efficacy, particularly in how they substantiate the Neo-Platonic framework turning the Self as a whole towards the Good.

The Synthesis is also substantiated by two segments: the Discussion (Contentions drawn from analysis of the outcome, including alternative interpretations; Caveats, Limitations and Shortcomings; and Implications) and Conclusion. The Synthesis integrates the key findings and appropriate concepts as they respond directly to the aim, acknowledging alternative approaches for external change which coherently segues into the following part of the thesis: Praxis.

Background

The Need for Elitism

Elitism is a necessary prerequisite for a successful 3rd party framework. If the intended outcomes are to be consistently superior to that of the Drifting Will's, then the political parameters must at least display competence above average. For this competence to be displayed, insofar as it relates to improvements along the axis of the Good

Recollecting previous elaboration on why people should not (and cannot really) make their own decisions, we saw that (1) even if, granted the tenability, the Good is to be known; (2) the limitations of the Ego restrict practical implementations; thus (3) efficacy of any turns towards the Good are questionable at best. This means that for any 3rd party framework to successfully compliment the Ego's general impotence, the parameters must necessarily be elitist by way of practical wisdom (phronesis). Turning our Self as a whole towards the Good then displays a need for elitism.

Thus understood, elitism is the conditional intermediary¹⁷⁶ between the agent and the Good. Insofar as a better life is desired, so too will the need for elitism exist as it has since time immemorial.

Justifications from Isokratia

Oftentimes, the concept of elitism is often baggage with negative connotations. It reminds people of "supremist" or other irrationally prejudiced agendas. This all necessarily brings into question not only whether the ends justify the means, but whether or not elitism is ever the better form of government at all. The burden of proof, however, manifests itself in absence; in other words, instances of government which have shared attributes of elitism have been mischaracterized as elitism. Moreover, these imitative¹⁷⁷ instances are but akin to boys playing at war and do little more¹⁷⁸ than degrade the original meaning of what elitism truly is.

Elitism is neither simply a meritocratic approach nor expertise and specialization alone; elitism is the manifestation of the ideal actualized. We see remnants of elitism in professional society, i.e. when going to the doctor. As a medical professional, a doctor is the closest to the Hippocratic ideal we can observe; a doctor at times errs in their techne and is thus only an imitation of elitism, but insofar as we trust a doctor because they perform better than we ever could ourselves, they still remain to be a remnant of what elitism truly is.

The part posits, that elitism is the most compatible form of government in relation to the Good. In the same way that our discrete trust in the doctor is justified, so is our abstract trust in elitism as a whole: isokratia¹⁷⁹. The underlying justification from isokratia necessarily

¹⁷⁶ This thesis treats needs as conditional intermediaries to wants. For instance, water is not inherently valuable, but becomes necessary relative to the desire to quench thirst. In this view, the Philosopher King functions as a conditional intermediary to the Good Life.

¹⁷⁷ Plato, *Republic*, trans. G.M.A. Grube and C.D.C. Reeve (Hackett, 1992), 473d, 502d–503b.

¹⁷⁸ While quasi-elitist frameworks have made measurable contributions—especially in promoting meritocracy and innovation—we must distinguish between surface comparisons and essential Forms. To understand elitism in its ideal form, one must think on a logarithmic rather than linear scale; likening a quasi-elite system to a Philosopher King is as misleading as saying Einstein is only "a little smarter" than the village idiot.

¹⁷⁹ See Josiah Ober, "The Original Meaning of 'Democracy': Capacity to Do Things, Not Majority Rule," *Constellations* 15, no. 1 (2008). Ober explains that terms like *isokratia*, *isonomia*, and *isegoria* each reflect equal access to public goods (power, law, speech), emphasizing fairness in participation over numerical majorities. Also, while commonly conflated, elitism, meritocracy, technocracy, and isokratia differ in both intent and structure. *Elitism* denotes rule by the best—where "best" is judged by moral, philosophical, or intellectual

means that elitism inherently provides equality of opportunity and equity of expression. In the form of government, this does not mean rigid social stratification but a recognition that rule is married to a matter of consequence, not method.

Fall into Totalitarianism

It is generally believed that the greatest reason elitism has fallen into such archaic¹⁸⁰ regard is due to the commonality of totalitarianism as impedance. At least in the form of government, it isn't too hard to see how the wealth of power quickly makes one drunk on corruption. Soon enough, efforts are made away from movement along the axis of the Good, and made towards conservation of the regime, all in the name of stability as a necessary prerequisite¹⁸¹. In reality, this is a perversion of marriage from consequence to method.

This perversion, however, is unavoidable, at least to the extent of human leadership. This is because, in much the same way many political candidates in their hearts wish for a president who is best for the country¹⁸², in their heads they would still rather that president be them even if they themselves are not the best. There are three key explanations:

- (1) Delusional
- (2) Short term gain; long term loss = **No one needs morality when there isn't enough to eat**
- (3) Power hungry (appetitive/spirited)

We see from all three then that the reason elitism is typically held in such low regard stems from a result of irrational proclivity – not in evilness itself. This is perhaps what truly makes totalitarianism so fearful and difficult to prevent with complete and utter abstinence; that it isn't a cruel, top-down phenomenon which stems from thugs with guns and torture camps¹⁸³, but rather is an elusive slip from an often-well-intentioned society in which many people share the vision and actively work to promote their ideology. Indeed, the next step

excellence. *Meritocracy*, by contrast, is a system of advancement based on performance or outcomes, not virtue per se. *Technocracy* emphasizes rule by technical or scientific experts, often devoid of normative justification. *Isokratia*, as framed here, does not signify a flat democracy but an equitable distribution of epistemic access to governance—trusting those best positioned to govern while ensuring others are not silenced, only stratified by consequence.

Michael Young. *The Rise of the Meritocracy, 1870–2033: An Essay on Education and Equality*. London: Thames and Hudson, 1958.

Frank J Fisher. *Technocracy and the Politics of Expertise*. Newbury Park, CA: Sage, 1990.

Isocrates. *Antidosis*, translated by George Norlin. Loeb Classical Library. Cambridge, MA: Harvard University Press, 1929.

¹⁸⁰ Elitism, like historical epochs, does not vanish uniformly. Just as the Stone and Iron Ages overlapped across cultures, elitist systems decline unevenly by context—religious, political, or technological. Modern technocracy, for example, still reflects elitist structures, albeit in specialized domains like policy or economics.

¹⁸¹ While stability and the Good are not mutually exclusive, totalitarian regimes often derailed elitist ideals by centralizing power without ethical scaffolding. For example, literacy tests and discriminatory policies reflect how externally imposed frameworks can undermine welfare. For a visual model of the Ego's influence on the Self, see Appendix A: Misc., Exhibit 4.

¹⁸² Plato, *Republic*, 473d.

¹⁸³ See:

– Reddit user discussion, “People Misunderstand Totalitarianism,” *r/JordanPeterson* (January 2022), [link].

– Aporia Magazine, “Plato’s Error: The Psychology of Philosopher Kings,” [link].

– “Why Totalitarianism Isn’t What You Think,” YouTube, [link].

Together, these perspectives highlight how totalitarianism is often mischaracterized as a blunt authoritarian structure, when in fact it can emerge through ideological overreach, bureaucratic diffusion, or the technocratic execution of a false ideal.

many have made from this is that if irrational proclivity is at the core of the problem; then by diminishing the irrational capacity of the “soul”¹⁸⁴, elitism wouldn’t fall prey to totalitarianism. And this is precisely the idea of diminishing this irrational capacity that underlies many frameworks along the political spectrum including anarchy, Marxism, and, of course, Plato’s framework for the Philosopher King.

Philosopher Kings as the Answer?

It is commonly held that the origin of the Philosopher King (and by extension, the Πολιτεία) was a response to the political turmoil at the time¹⁸⁵. For Plato, there were too few trained to see the Good and therefore unable to see that they truly wanted the doctor and not the Candyman¹⁸⁶. Plato’s solution to the fall of benevolent dictator – as opposed to demagogue-turn-tyrant¹⁸⁷ – then was the framework of the Philosopher King.

For Plato, rather than simply hoping that rulers would become philosophers or philosophers would become kings through chance, he deduced a socially engineered framework from which those best suited could forge their own tools to rule justly and in line with the Good. The specific details of Plato’s framework can often be traced back to the Republic, or Πολιτεία¹⁸⁸, but the general interpretation has been that assuming the domain of politics is concerned with turning the polis as a whole towards the Good, then politicians (as a ruling elite) must necessarily be well versed in not only what the Good is, but also experienced in how to practically implement the Good as well.

But if a trained Philosopher King was the ideal ruler, then why isn’t everyone studying within a kallipolis? Moreover, how is it that totalitarianism is not only not decreasing, but actually increasing inversely proportional to the progression of the developed societies¹⁸⁹? Why is it that philosophical training fails to heal the pathology of degeneration?

Why Marcus Aurelius Is the Closest We’ve Gotten

Roman Emperor Marcus Aurelius exemplifies the idea of the Philosopher King to a certain extent, and testament from time under his rule offers evidence of prosperity of society under a quasi-Philosopher King¹⁹⁰. But, is it fair to ascribe the Roman Empire’s decline after

¹⁸⁴ Classical terms like *soul*, *daimon*, and *psyche* have been used to describe the seat of moral reasoning. Moral “malnutrition” can occur either by inflating the rational faculty at the expense of emotion, or vice versa.

¹⁸⁵ According to Karl Popper and others, Plato’s political philosophy was shaped by the failures of Athenian democracy, especially the destabilizing influence of demagogues like Alcibiades, whose recklessness led Athens toward disaster. See M.I. Finley, “Athenian Demagogues,” *Past & Present* 21 (1962): 3–24.

¹⁸⁶ Plato, *Gorgias*, 521e–522a.

¹⁸⁷ Socrates contrasts the demagogue with the just ruler in *Gorgias*, emphasizing persuasion through truth over manipulation. See *Gorgias*, 464b–466a.

¹⁸⁸ While speculative, some suggest *The Republic*—like the Bible—functions more as a posthumous anthology than a coherent treatise. Its ten-book division often obscures flow and invites the possibility of missing dialogues or interpolated material. See: Reddit thread, “Is Plato’s Republic Similar to the Bible?”

[r/askphilosophy](https://www.reddit.com/r/askphilosophy/comments/5fv3zq/is_platos_the_republic_similar_to_the_bible/),

https://www.reddit.com/r/askphilosophy/comments/5fv3zq/is_platos_the_republic_similar_to_the_bible/

¹⁸⁹ Fair Observer, “Totalitarianism’s Global Threat,” https://www.fairobserver.com/more/global_change/human-rights/totalitarianism-now-presents-an-unprecedented-global-threat/

¹⁹⁰ Marcus Aurelius may be seen as a quasi-Philosopher King—not because he perfectly matched Plato’s framework, but because his rule embodied many of its virtues. His success was less due to divine kingship than to an epistemically enriched support system: advisors, institutions, and historical context. Just as a modern CEO may be the face of a company governed by its board and market forces, Aurelius ruled through a structurally elite ecosystem. See also: Daily Philosophy, “7 Most Fascinating Philosopher Kings,” <https://daily-philosophy.com/7-most-fascinating-philosopher-kings/>

Aurelius due to him failing to be a true Philosopher King? Can we claim circumstantial evidence of prosperity as a result of him striving to adhere to Plato's framework? And are we even sure that Rome was not just due for decline anyway¹⁹¹?

The above questions have been the subject of much historical speculation and analysis, but in regards to what is relevant to this thesis – a successful 3rd party framework – the literature has largely been divided on not only whether or not the framework for the Philosopher King was meant to be taken as a real, as opposed to metaphorical model for politics, but even whether or not the Πολιτεία was about politics at all.

But, of course, the real proof of the pudding is the clear lack of any modern-day Philosopher Kings whatsoever. If Plato's framework was really the successful 3rd party framework we have been looking for, then we wouldn't be looking for a 3rd party framework in the first place. Granted that we have advanced just a little since Plato's time, however, we should cut him some slack and also acknowledge that while his original framework is invalid, it was not completely wrong.

It has been shown then that the Πολιτεία specifically was correct on at least three relevant fronts:

- (1) People are not born equal; but they are born equitable
- (2) Philosophical training makes for better rulers
- (3) There exists a perennial need for elitism (Philosopher King)

First and foremost, people are not born equal; but they are born equitable. This has been known for generations, and more recently backed up with scientific literature. For example, there are legitimate differences between men and women¹⁹². Men and women are not born with equal talents, but in society at large – while there is a disproportionate valuation of said talents – the actual equitability of cohesive expressiveness (i.e. both male and female gametes are equally indispensable for life) remains the same. Much in the same way, white people and black people are not equal; old people and young people are not equal; army people and farmer people are not equal; etc. But all people have equity of expression nonetheless.

Second of all, philosophical training makes for better rulers. We have already seen this with Marcus Aurelius where evidence (albeit circumstantial) of adopting capable heirs (often not of their own bloodline) whom they deemed suitable for the throne ostensibly led to stable legitimacy of the Five Good Emperors. In either case of meritocracy or monarchism¹⁹³, it is often the case that the most competent ruler is the one who is best educated (i.e. eldest heir simply because they have had a more extensive education), typically under the philosophical tradition, which under the Platonic framework also includes practical experience. Moreover, compared to simple general education, philosophical training has proven to be more analytic and cultivate the virtues which facilitate leadership.

Last but not least, there exists a perennial need for elitism. There are good reasons why the Doctor/Barber or Barber/Surgeon title is no longer readily available. There are better reasons why the role of expertise and specialisation are needed. And there are, of course, simply just reasons why employers expect work experience for entry level positions. As a

¹⁹¹ Plato, *Republic*, 545c–569c. “Since destruction (φθορά) awaits everything that has come to be, even a foundation of this kind will not survive for the whole of time...”

¹⁹² T.C. Ngun et al., “The Genetics of Sex Differences,” *Frontiers in Neuroendocrinology* 32, no. 2 (2011): 227–46. <https://doi.org/10.1016/j.yfrne.2010.10.001>

¹⁹³ Royal bloodlines were often established through meritocratic means—original rulers gained power by demonstrating superior competence, typically through conquest or charismatic authority. Over time, this founding merit ossified into hereditary privilege.

conditional intermediary, Plato has shown in theory, the world in practise, and my earlier segment above in writing, how and why this perennial need for elitism exists: in a fundamental sense, because there, too, exists a demand.

The missing key between since Plato wrote the Πολιτεία and/or since Neo-Platonists have since last tried to socially engineer a variant of the kallipolis is time. Einstein is often famously claimed to have said that 'insanity is doing the same thing over and over again and expecting different results'. It might seem insane to keep trying to find the exact recipe to realize Plato's framework for the Philosopher King, but each day we get closer to truth and thus we are in a much better position today than we were yesterday.

To be exact, if having stayed with us thus far, the dear reader should be experiencing some sense of déjà vu. After all, do we not see here an echo of the previous part? In the same way that researchers believed weakness of will could be addressed by strengthening rational capacities, Plato had the same idea a couple of centuries prior. But regardless of then or now, the fact of the matter remains the same: increasing willpower and/or rational capacity will not translate to the extent anticipated because a Virtuous Hexis is only affective during moments of engaged consciousness. It is primarily this missing key of knowledge nuggets like this as accessories to time which would have altered the Neo-Platonists entire directionality of their socially engineered conclusions.

From this background chapter, we have thus seen that:

1. Elitism is a necessary prerequisite for a successful 3rd party framework. This has been and will continue to be a perennial need insofar as elitism remains the conditional intermediary between the agent and the Good.
2. Elitism is the most compatible form of government in relation to the Good Life – if it were only realizable. This is because elitism inherently provides equality of opportunity and equity of expression.
3. Elitism has fallen into general archaic regard due to its potential to fall into totalitarianism. As a matter of human nature, this perversion is not a matter of if, but when.
4. Plato's solution to the fall of the benevolent dictator was the framework of the Philosopher King. If philosophical training (which includes practical experience) was given to those with a natural aptitude for learning, leadership, and moral integrity; it was theorized possible to cultivate rulers well versed in not only what the Good is, but also experienced in how to practically implement the Good as well.
5. Plato was ahead of his time but also necessarily influenced by his epistemic placement. People are not born equal, but they are born equitable; philosophical training makes for better rulers; there exists a perennial need for elitism (Philosopher King). But, since we have yet to see the advent of a true Philosopher King, Plato's original framework is not a successful solution for achieving the Good Life.

It is important to acknowledgment that without Plato's original framework and the – sometimes shocking - results of subsequent social engineering by the Neo-Platonists, it is evident that our current understanding wouldn't be as enriched in the various domains that it is today. This section would like to thus draw attention to a specific aspect that may benefit from further exploration or consideration. It is my intention to engage in a collaborative discussion and contribute to the refinement of our collective knowledge, namely, by

recognizing that the field is changing and implementing revised understandings into theories of how to acquire the Good Life.

Specifically, then, we see that focusing on frameworks in a historical vacuum is insufficient and that there exists a need for further investigation on how to adapt the original framework to become the successful 3rd party framework.

In what follows, we delve into the core chapters of my work which, drawing from the conclusions of this background chapter, leads into my research hypotheses and questions. The task of these core chapters is to reconcile the feasibility of elitism by adapting Plato's original framework for the Philosopher King and Kallipolis into a modified Neo-Platonic framework for the Good Life. Moreover, it is to evaluate a 3rd party framework based on tenability, practicality, and efficacy for practically implementing utilitarian contrivance of the good.

Core

Position (methodology)

From the conclusions of the background chapter, it is clear that focusing on Plato's original framework alone is an insufficient approach for modern day Philosopher Kings; moreover, that we must shift our perspective from understanding the pursuit of acquiring the Good Life as an individual endeavour to that of the collective framework.

Following from the structure of the previous part, research questions that naturally arise include, but are not limited to: if Plato's original framework is insufficient, then which adaptive framework(s) would we need be more successful? Where does tenability lie on the bridge between theory and practice? How are the demos ever to practically support a seemingly disenfranchising framework? And to what extent would this Neo-Platonic 3rd party framework be superior?

Taken together, these key research questions lead into an initial formalized hypothesis: people will be happier and have more welfare when society is elitely structured

As my thesis has already been established as interdisciplinary with an emphasis on changing or emerging fields, it should be reminded that much of the work is unprecedented with no prior existing framework or research program. Research questions still revolve predominantly around *current theory, discoveries, and debates* that are most useful and salient to my topic.

My stance for this part remains to be purely philosophical since, and I'm still limited as evaluator of theoretical and applied (including thought) experiments, objective observer, and even, at times, participant in the activities (even if trivially anecdotal) under study.

As the subject matter is different for this part of the thesis, the stances that I can consciously report reflect this:

Stance 1: I hold a vested interest in some of the case studies.

Justification: typically, a highly standardized stance is one where the researcher is a dispassionate outsider. For segments (particularly certain case studies, i.e. Outer Coast, which will be developed later) of this part, however, I will need to take a more exceptional stance as I was a direct participant in the study and therefore this makes my interpretation a double-edged sword: on one hand I have a more nuanced and visceral understanding, but on the other the understanding can also not be sterile and fully abstracted¹⁹⁴. Understandably, this makes the study semi-compromised from the start, however, considering it has been a few years since I was a participant and that I have held objective discussions surrounding the experience prior, I believe my formalized stance is sufficient to counteract any conflicts of interest and bring out the best of both worlds.

Stance 2: I take a "quasi-critical pedagogy" reading (Aristotle/Freirean tradition) of the literature.

Justification: Aristotle said man is a political animal. My take – as having been said to parallel Freire's¹⁹⁵ – is that any social action, therefore, is necessarily political. This view is

¹⁹⁴ Kimberley Brownlee, "The Competent Judge Problem," *Philosophy and Phenomenological Research* 77, no. 1 (2008): 12.

¹⁹⁵ Freire viewed education as inherently political, arguing that any claim to neutrality merely reinforces dominant power structures. His *Pedagogy of the Oppressed* positions education as a site of ideological

substantiated by the belief that (1) politics is not restricted to governance (i.e., departmental politics); and (2) governance is not limited to top-down formal institutions where home economics between peers even (i.e., cleaning up after one another) is an exercise of governance. While the context of this part is not primarily educational; nonetheless, insofar as literature reviews are a tangential instrument to education, my analysis of their “educational” content is scrutinized under a political lens. It is perhaps better to say that I have appropriated critical pedagogy insofar as I have read the literature with the goal in mind of titrating their political overtones.

Stance 3: I support a Neo-Platonist model of governance.

Justification: maybe its not fair to say the model is neo-platonist as the Benevolent Dictator archetype is a multiple discovery, but insofar as I am using Plato’s framework for the Philosopher King and Kallipolis as a standard model (for elitist paternalism) *and* insofar as I am adapting it for modern purposes, I am supporting a (neo-)platonist framework. But then why a adapt a platonic framework to begin with instead of one of the other multiple discoveries such as the Confucian junzi or native American wisom leaders which I have previously written papers on? Mostly because the platonic community is vastly superior¹⁹⁶ which incidentally means (1) Plato’s Philosopher King is comparatively the most renowned archetype and existing extensive framework; which leads to (2) the debate surrounding Plato’s Philosopher King being comaparitvely the most comprehensive; in turn making the literature surrounding Plato’s Philosopher King from a purely statistical perspective¹⁹⁷ comparatively the most materially accessible.

As before, my position can be reduced to: the research program is what I am doing, and my stance is necessarily how I’m going to go about my investigation. It is not much more complicated than that, but for the sake of research ethics, I felt it important to lay out my identified stances as above.

Method

Although the subject matter is different for this part of the thesis and the stances that I can consciously report reflect this, to maintain consistency throughout my thesis as a whole, my methods have not and will not change. As a quick reminder, nonetheless, my research program is cross disciplinary, and, so too, is my method. This singular method, as a pluralism of other research methods, comes as a triangulate approach. To reiterate:

- A. Literary Interpretation: At the core of this approach lies the vital role played by the “hermeneutic circle” in the process of interpretation.
- B. Ongoing Discussion: In the same way the instrument of the hermeneutic circle is used to stimulate dialectic within and between the literature, dialectic will also extend beyond the literature.

contestation, where norms and values are never separate from politics.

Paulo Freire, *Pedagogy of the Oppressed* (New York: Continuum, 1970).

¹⁹⁶ Alfred North Whitehead famously wrote that “the safest general characterization of the European philosophical tradition is that it consists of a series of footnotes to Plato.”

Alfred North Whitehead, *Process and Reality* (New York: Free Press, 1978), 39.

¹⁹⁷ While much academic literature is paywalled, even adjusted for access disparities, the volume of material on Plato’s Philosopher King far exceeds that on Islamic Caliphs—suggesting greater epistemic emphasis even outside traditional availability metrics.

- C. Grounded Theory: Comparable to experimental philosophy, a grounded theory makes use of empirical data, but more than positing hypotheticals¹⁹⁸, the design of this method emphasizes the generation of theory (including thought experiments) which is ‘grounded’ in data that has been systematically collected and analysed.

In brief, my method converses amongst a wide range of mediums concurrently under a real-world setting.

Outcomes and Results

The three mini questions (themes) that the outcomes revolve around are again in respect to three different sections: On Tenability; On Practicality; and On Efficacy. The first section naturally aids in response to the conclusions drawn from the Background with the initial three subsections targeted towards the initial research questions of how to adapt the original framework to become the successful 3rd party framework; specifically, when there’s a lack of any modern-day Philosopher Kings not least for lack of trying or conflation with totalitarianism as common with elitist frameworks. The following sections and subsections naturally follow suit with each building upon conclusions and follow up issues of inquiry leading from preceding outcomes. The overall breakdown is as follows:

- (1) Initial Research Question: Where does tenability lie on the bridge between theory and practice?
Section Heading: On Tenability
Subsections:
 - e. Implausible
 - f. Superfluous
 - g. Corrupt
 - h. Interpretation
- (2) Initial Research Question: How are the demos ever to practically support a seemingly disenfranchising framework?
Section Heading: On Practicality
Subsections:
 - e. Legal Legitimacy
 - f. Democratic Legitimacy
 - g. Pragmatic Legitimacy
 - h. Interpretation
- (3) Initial Research Question: To what extent would this Neo-Platonic 3rd party framework be superior?
Section Heading: On Efficacy
Subsections:
 - e. Sophia
 - f. Techne
 - g. Phronesis
 - h. Interpretation

¹⁹⁸ Grounded theory has shaped both the structure and development of this thesis. Research questions evolved in tandem with data, often requiring cyclical revisions: conclusions generated new introductions, and sections were rewritten or replicated to reflect emergent insight.

These three sections will necessarily tie the subsections together into an interpretation that addresses the initial research questions. These interpretations will naturally be the outcomes of my research methods as laid out above, and these very outcomes will then lay the foundations of my results as stemming from the sections working in tandem.

On Tenability

While the wording might be novel, we've seen in alluded throughout the Background the general idea behind a 3rd party framework. However, historically, said frameworks have not all been theoretically tenable; the probability (if all things are perfect) ≤ 1 . In other words, many of these frameworks were not possible in theory, let alone practice.

Not to be confused with efficacy¹⁹⁹ which is its own section, tenability is more focused on efficiency to the extent of the success of theoretical possibility. Out of the three variables to any successful framework, first, a framework must be executable in praxis, that is not only work in theory but also be practically viable. In addressing Plato's original framework specifically, the general queries surrounding tenability have most often come from issues of plausibility, superfluity, and corruptibility. Together, these issues bring into question if the realization of the Philosopher King is even theoretically possible.

In identifying where tenability lies on the bridge between theory and practise, the general critiques as laid out most prominently by Aristotle, Steinberger, and Popper must be addressed. From Aristotle, we have the issue of Plato domiciling the Philosopher King in a historical vacuum. Drawing on expositions from the previous part, we return to the overly demanding nature of Plato's rationalization of the human condition and its hypothetical constraints. From Steinberger, we are reintroduced to Plato's original "inconsistencies" within the Πολιτεία. Specifically, Steinberger accentuates the distances which Plato seemingly fails to cross in reconciling "meddling" within each other's' roles. From Popper, we have the traditional issue of potentiality of political decay. Here we are introduced to the possibility for corruption to take root making the theoretical framework for the Philosopher King less than watertight.

Implausible

Perhaps the largest critique of Plato's Philosopher King comes from Plato's student, Aristotle, himself; in a passage which is commonly regarded as a fragment of Aristotle's lost work *On Kingship*²⁰⁰, Themistius relates²⁰¹:

"Plato, even if in all other respects he was divine and deserving our unlimited admiration²⁰², was utterly reckless when he made the statement that evils would never cease for men until either philosophers became kings, or kings became philosophers²⁰³".

¹⁹⁹ Efficiency concerns achieving results with minimal waste or effort. Efficacy, however, involves producing the intended outcome. A framework may be efficient in theory yet lack the efficacy to enact lasting results—such as reaching and returning from the moon, not just launching the rocket.

²⁰⁰ Anton-Hermann Chroust, "Aristotle's Criticism," *Rheinisches Museum* 111 (1968): 16–22.

²⁰¹ Themistius, *Oratio* VIII, 107D.

²⁰² See Aristotle, *Nicomachean Ethics* 1096a12–17; 1144b11ff. See also Olympiodorus, *In Platonis Gorgiam Commentarium*, p. 197 (Norwin); and Plato, *Laws* 660E: "The good man, if he is temperate and just, is fortunate and happy."

²⁰³ Plato, *Republic*, 473c–d: "Until philosophers are kings, or the kings and princes of this world have the spirit and power of philosophy... cities will never have rest from their evils." See also 501e–502a.

If Plato was the idealist, then, by contrast, Aristotle was the empiricist who did not hesitate to point out the relation between purely theoretic (Platonic) politics and practical or workable politics which takes into account the *de facto* human condition²⁰⁴. Aristotle's general criticism was that Plato's political philosophy is implausible, particularly inasmuch as Plato's "ideal city" - the city of the philosopher king - being domiciled in an historical vacuum²⁰⁵.

Simply speaking, Aristotle accuses Plato's ideal on two fronts: 1) being overly demanding; and 2) being purely hypothetical.

In everyday life we can see evidence in support of (1): akrasia as with all failed new year's resolutions; Olympic defeats by those groomed from birth; and even philosophical mediocrity from professors "*with a specialisation in ethics*"²⁰⁶. Advancements in the philosophy of mind have also identified the impracticability of trying to will ourselves into flawlessness²⁰⁷. The same can be said for (2) whether the evidence is historical as seen with armchair philosophers²⁰⁸ or even more recent with those such as Kwasi Kwarteng²⁰⁹. Evidently, the human condition dictates that we exist in tangent with the realm of reason. Therefore, inasmuch as we lapse between the rational world and the world of feeling, our reasoning insofar as we act as candidates for the Philosopher King, will undoubtedly be compromised.

Admittedly, Aristotle's critique is not without standing; trying to make the human condition adhere to the requirements of the Philosopher King would be like trying to force a camel through the eye of a needle. Moreover, in the previous part we already extensively elaborated on why a 3rd party framework, that is not the Ego – and by extension, not impotent people – that governs itself, would be needed to circumvent the inherent problem of subjectively mental endeavours. Considering that during Aristotle's time – and even not too long ago – the only possible candidate that anyone would think of for the Philosopher King was a human, it's understandable why many scholars held that that the realization of the Philosopher King was impossible. However, we are now situated in a different epistemic context where the anthropocentric perspective is no longer the only one available. The Philosopher King need not be a human²¹⁰ – or at the very least, a pure, unaugmented (transhumanist) human. Therefore, while Aristotle's critique is valid insofar as it applies to

²⁰⁴ Jamblichus, *Protrepticus*, ed. H. Pistelli, p. 56, lines 2ff; see also frag. 13 (Walzer, Ross), frag. 51 (Düring), and frag. 48 (Chrout). An echo of this passage may be seen in Aristotle, *Nicomachean Ethics* 1144b11ff: "a strong body, which moves without sight, stumbles badly."

²⁰⁵ Aristotle, *Politics* 1265a18ff. He acknowledges that "the discourses of Socrates [in the *Republic*] are never commonplace: they always exhibit grace and originality of thought, but perfection in everything can hardly be expected." As such, the dialogues resist literal historical application.

²⁰⁶ Eric Schwitzgebel and Fiery Cushman, "Philosophers' Biased Judgments," *Cognition* 141 (2015): 127–137. Their research suggests that even trained philosophers exhibit cognitive biases in ethical reasoning, challenging the assumption that philosophical training yields superior moral judgment—an implicit critique of Plato's Philosopher King ideal.

²⁰⁷ David Wiggins, "Weakness of Will," *Proceedings of the Aristotelian Society* 79 (1978): 251–77.

²⁰⁸ Chinese idiom: 纸上谈兵 (*zhǐ shàng tán bīng*)—literally, "discussing warfare on paper." Refers to impractical theorizing or armchair strategy with no grounding in reality.

²⁰⁹ Charlotte Kilpatrick, "Plato Was Wrong," *New Statesman*, September 29, 2022.

<https://www.newstatesman.com/quickfire/2022/09/kwasi-kwarteng-plato-philosopher-kings>

²¹⁰ Kant never explicitly limited rational beings to *Homo sapiens*, leaving open the possibility of non-human agents (e.g., aliens) qualifying for moral consideration. See:

– Immanuel Kant, *Groundwork for the Metaphysics of Morals*, 4:428.

– David L. Clark, "Kant's Aliens: The 'Anthropology' and Its Others," *CR: The New Centennial Review* 1, no. 2 (2001): 201–89. <http://www.jstor.org/stable/41949284>.

humans of the past, this isn't necessarily the case for humans of the future, or even those of other sentence.

Superfluous

Another issue of tenability regarding Plato's Philosopher King addresses the impossibility of their realization by appeal to the superfluousness of its nature. Again, Steinberger's argument draws upon limitations of the human condition, specifically in that a person cannot be, or rather cannot exist in two diametrically opposing states at once. Due to "differences in their education, character, and relationships to the necessities of the political world, specifically the issue of deceit"²¹¹, Steinberger argues that it is untenable for Guardians and philosophers to be the same person. Consequently, this makes the role of Philosopher King (who is both guardian and philosopher²¹²) incoherent and their addition to the kallipolis superfluous.

Steinberger wasn't even the first to identify this issue, he simply accentuated one of Plato's original "inconsistencies" within the Πολιτεία. It was Plato himself which wrote that diametrically opposing or not, it is not only untenable in the long run, but also morally corruptible (by way of harming the soul) for example, army people and farmer people to "meddle"²¹³ in each other's roles. If one has an aptitude for cobbling and another for weaving, even if the second-best talent of each were for the other's profession, Plato still vehemently rejected that a weaver ever pick up the slack of a cobbler and vice versa²¹⁴.

This issue of superfluousness is ever more noticeable when taking parallels from a physical perspective. For example, assuming the interchangeability between sandboarding and snowboarding, an obvious consequence of a sandboarder meddling in snowboarding would be the economic cost of transitioning between physical locations. While this cost would appear minimal for a Philosopher King whom supposedly could equally philosophise in parliament and rule from within the agora, the distance is not merely a physical obstacle, but a metaphysical one. This metaphysical distance must be crossed in the same way that actors must metaphysically cross the distance between tasks of one character from the next, and research has already shown that due to human limitations, this kind of multitasking results in worse results than doing each task separately²¹⁵.

It should be remembered, however, that "We" are not alone, making Plato's *supposed*²¹⁶ inconsistency unfounded. Reference to dissociative identities has shown that it is possible to be opposites simultaneously, particularly in split personalities. Taking into account the exploration on the Persona from the previous part, we've seen how the Ego can indirectly control Subpersonalities through cultivation of a Physical Hexis. Of course, the question of fragmenting the mind like a horcrux for the sake of superfluity seems dubiously

²¹¹ Christopher M. Duncan and Peter J. Steinberger, "Plato's Paradox? Guardians and Philosopher-Kings," *American Political Science Review* 84, no. 4 (1990): 1317–22. <https://doi.org/10.2307/1963266>.

²¹² Plato, *Republic*, trans. G.M.A. Grube and C.D.C. Reeve (Indianapolis: Hackett, 1992), 502d–503b.

²¹³ *Ibid*, 443d–444a.

²¹⁴ Plato's *Myth of the Metals* frames each individual as naturally suited for a particular social function—gold for rulers, silver for auxiliaries, bronze and iron for producers. This deterministic allocation underwrites his vision of justice. As Ricardo would later note in his theory of comparative advantage, division of labor yields both individual satisfaction and collective benefit. See: Kimberley Brownlee, "Manuscript on Totalitarianism," Oxford, <https://media.philosophy.ox.ac.uk/docs/people/brown/totalitarian.pdf> Oxford, <https://media.philosophy.ox.ac.uk/docs/people/brown/totalitarian.pdf>

²¹⁵ Eyal Ophir, Clifford Nass, and Anthony D. Wagner, "Cognitive Control in Media Multitaskers," *Proceedings of the National Academy of Sciences* 106, no. 37 (2009): 15583–87.

²¹⁶ Plato distinguishes between contradictory and complementary roles. While one may simultaneously be a father, son, and baker, roles like "philosopher" and "king" are not mutually exclusive but mutually reinforcing. See *Republic* 473d.

incoherent with Plato's justice of the soul²¹⁷ when considering other texts on harmony of moral health. It is unlikely, therefore, that Plato would allow for a Philosopher King with multiple personalities or any other abnormality as it would likely be viewed as a threat to justice. Nonetheless, the potentiality for "Ego Horcruxes" amongst other bespoke simulacra shows that there is at least one more avenue (in addition to earlier allusion to transhumanism) that offers solutions to the issue of superfluosity.

Corrupt

One more point denouncing the tenability of realizing the theoretical framework for the Philosopher King stems from Popper. Once more taking inspiration from the flaws in the human condition, Popper echoes the old adage that "power tends to corrupt, and absolute power corrupts absolutely." From this, Popper's central theses about Plato are such that: he is an enemy of the open society, his politics are "purely totalitarian,"²¹⁸ and his Socrates is "the embodiment of an unmitigated authoritarianism."²¹⁹ Simply, Popper contends that Plato's framework is theoretically corruptible, and, therefore, is theoretically untenable.

While foremost an argument against totalitarianism as opposed to Plato himself, Popper traced the rise of totalitarianism in the 20th century back to Plato's Philosopher Kings, with their dreams of social engineering and idealism. For Popper then, the Πολιτεία is, if not the root of political evil, is at least the catalyst for it²²⁰; as long as humans take the helm within a totalitarian regime, their nature will necessarily corrupt them, making the realization of a utopia unpractical.

As with the arguments from the previous part, Popper is correct insofar as the human condition is incompatible with a 3rd party framework that is not the Ego – and by extension, not impotent people – that governs itself. But, also, as alluded to above, the flaws of the human condition (if so parameterized) do not discount totalitarianism – and the rule of the Philosopher King, by extension – outright. After all, an incompetent captain on its own does not demerit the capabilities of a ship altogether.

To address this issue of tenability then, we can simply turn to a modified framework in where the Philosopher King is incorruptible. In theory thus far this includes the original Platonic conception of diminishing the irrational capacity, but also erasing it completely. It is really only "needy" biological creatures that even have the capacity for corruption; if there is no desire either through augmentation (surgical removal physically or metaphysically, i.e. Buddhist teachings) or from conception (technologically Advanceds, i.e. ASI), then there is no root from which corruption can spread. Since we are now situated in a different epistemic context where the anthropocentric perspective is no longer the only one available, Popper's argument that the Πολιτεία being at the root of totalitarianism, too, is no longer the case.

Interpretation:

In this first section *On Tenability*, we've identified that situated in their original epistemic context and taken from an anthropocentric perspective, the above arguments from theoretical untenability hold some weight; yet considering recent advancements in multiple

²¹⁷ Mike W. Martin, "Moral Sickness: Plato and the Psychiatrists," in *From Morality to Mental Health: Virtue and Vice in a Therapeutic Culture* (New York: Oxford University Press, 2006).
<https://doi.org/10.1093/0195304713.003.0001>

²¹⁸ Karl R. Popper, *The Open Society and Its Enemies*, vol. 1 (Princeton: Princeton University Press, 1963), 76.

²¹⁹ Ibid, 115

²²⁰ Some critics argue that figures like Mussolini, Stalin, Khomeini, and Pol Pot serve as cautionary examples of philosopher-kings in form but not in function—leading not to utopia but to dystopia.

fields of technologies, they are no longer entirely valid. The absence of a true Philosopher King can be remedied by realizing plausibility that is not purely human; the metaphysical superfluosness can be crossed through “Ego Horcruxes”; and proneness for corruption can be absolved through diminished, or completely erased capacities. Philosopher Kings might have been theoretically untenable just even a decade prior but with the advent of transhumanists and other technologically Advanceds with superior attentional octopi, more and more viable candidates are coming to the forefront.

Returning to one of the original questions, then, we are now in a much better position to understand where tenability lies on the bridge between theory and practice. Namely, that Plato’s original framework was untenable up till now due to lack of appropriate vessels for the Philosopher King. And because the landscape (political, technological, and moral) has changed and will undoubtedly continue to change, theory will, too, undoubtedly unfurl into dimensions of practice as the framework for the philosopher king becomes ever more accessible. But more than simply hoping, and waiting, for Plato’s original framework to become relevant, the best way to predict the future is to create it²²¹. This begs the question though of how are we to modify the original framework so that it becomes relevant? How are we to adapt the original framework to become the successful 3rd party framework that we need? These questions naturally segue into the subsequent research question of “how are the demos ever to practically support a seemingly disenfranchising framework?”

On Practicality

After the events of the previous century, many people are skeptical of socially engineered frameworks. For the many, paternalism is synonymous with disenfranchising frameworks that carry a general atmosphere ‘of a rigid class division and class rule’²²², where ‘the individual is nothing but a cog (sc. in the state machine)’ and hence that ‘ethics is nothing but the study of how to fit him into the whole’²²³. Plato’s original framework (and their Neo-Platonic offshoots particularly) often, *purportedly*, calls to mind the absurdity of trying to compel the masses towards anti-individual, anti-democratic, and anti-libertarian agendas. As identified, then, the question is one of practicality; Granted that probability ≥ 1 , then will the people allow it to happen? Alternatively, how are we to modify the original framework so that the people accept it?

Clearly there are quirks from the original framework which are not only now shown to be scientifically inaccurate (i.e. repressive measures), but also to be socially unacceptable (i.e. killing babies). This in addition to the default duct tape²²⁴ of modern governance (i.e. so-called “democracy”) means the problem of legitimacy therefore isn’t one post-transition but pre-transition. For people to have to want the Platonic framework of the Philosopher King and Kallipolis, it must be shown to be legitimate. In other words, legitimacy is a prerequisite for a successful 3rd party framework. Said legitimacy has already been identified as one key barrier to the demos ever to practically support a seemingly disenfranchising framework, specifically by way of conflated totalitarian overtones. To sway the demos to favour a Neo-Platonic framework then, we need to justify legal, democratic, and pragmatic legitimacy.

²²¹ Aristotle compares ethical excellence to Olympic victory: “It is not the strongest or fastest who win, but those who compete.” Capacity alone is insufficient without action. See *Nicomachean Ethics* 1099a1–3.

²²² Karl R. Popper, *The Open Society and Its Enemies*, vol. 1: *The Spell of Plato*, 5th ed. (Princeton: Princeton University Press, 1966), 89.

²²³ *Ibid*, 1:108.

²²⁴ Unless a system is irreparably broken, a complete restructuring of government often incurs unjustifiable short-term losses—except, perhaps, through violent revolution.

In initially identifying legal legitimacy, it is necessary to recognize safeguards and caveats Plato implements in the *Πολιτεία* and the *Laws* to minimise²²⁵ devolution of the Kallipolis into totalitarianism. This is important because to justify subjugating (by force and/or coercion if necessary) individuals for the sake of their own good we must prove that the adapted framework operates according to the rules and laws that are recognized by the society it governs. The subsequent step needed for democratic legitimacy is to adapt Plato's criterion of social good, which claims authority over all individual and minority perspectives. This is important because to justify seemingly removing democracy for sake of democracy we must prove that the adapted framework is enabled through universal suffrage, reflecting the principle of competent, popular sovereignty. To finally substantiate pragmatic legitimacy we need to provide a modified variation of soft despotism that is relevant to modern audiences and remains within the status quo. This is important because to justify subjugating (by force and/or coercion if necessary) individuals fate for the sake of their own good we must prove that the adapted framework is able to deliver positive outcomes, such as economic prosperity, security, and social welfare.

Legal Legitimacy

The initial step for the demos ever to practically support a seemingly disenfranchising framework is by garnering legal legitimacy and proving Plato's original (or a Neo-Platonic) framework comes to power and operates according to the rules and laws that are recognized by the society it governs. This includes (1) disproving misconceptions surrounding Platonic authoritarianism; (2) substantiating the abstraction of the Philosopher King; and (3) justifying superior legality over the status quo.

First, we need to dispel the conflated cloud of connotations associated with authoritarianism which for most contribute largely to its emotional, as opposed to validly rational aura. Common connotations include, but are not limited to elements of totalitarianism (as addressed above), Nazi Germany/USSR, Rwandan/Cambodian Genocide, State Surveillance/Police State, Orwellian Big Brother²²⁶, as well as other rather specific traits²²⁷.

Indeed, Plato's framework for the Philosopher King and Kallipolis is a highly differentiated form of paternalism²²⁸, which under certain definitions could be mistook for totalitarianism. Key areas of worry include the coincidence of:

A. Organic Theory of the State (Priority Thesis)

²²⁵ Plato acknowledges political decay as inevitable—even in ideal regimes—marking his framework as aspirational rather than timeless. Thus, the *Kallipolis* is not a static utopia but a living, evolving form. In this sense, our “Neo-Kallipolis” is more aligned with its true essence than any literal application of the original model. See: Plato, *Republic* 546a–547a; Anselm, *Proslogion*.

²²⁶ Popular dystopias like *1984* often reflect not the world they depict, but the epistemic biases of present readers. Winston's resistance, while rhetorically powerful, is psychologically unrealistic—most citizens would likely conform, not rebel. This reflects a broader cognitive bias: we project present values into imagined futures we cannot accurately foresee. See Daniel Gilbert, *Stumbling on Happiness* (New York: Knopf, 2006); Jason Brennan, *Against Democracy* (Princeton, 2016), discussed at <https://www.vox.com/2018/7/23/17581394/against-democracy-book-epistocracy-jason-brennan>.

²²⁷ See “Was Plato a Totalitarian?” in *In Defense of Plato* (Cambridge, MA: Harvard University Press, 1953), 499–580. <https://doi.org/10.4159/harvard.9780674431034.c9>.

The chapter juxtaposes Plato's political idealism with mid-20th-century totalitarian practices: mass surveillance, purges, hero breeding, indoctrination of children, and ideological conformity—inviting critical reflection on the ethical scope of state power.

²²⁸ The legal phrase “parents and/or guardians” acknowledges that biological parents do not always fulfill the caretaking role. Likewise, Plato's guardians are selected not by birthright but by demonstrated competence—just as the Philosopher King must first prove himself as a guardian.

B. Ideological State Control
C. Authoritarian Intolerance

While the Organic Theory of the State (particularly the Priority Thesis) is tangential to legal legitimacy, it is more in line with the counterpart of democratic legitimacy and thus will be revisited in the following section. For now, the reader can simply understand the worry stemming from the Priority Thesis as it relates to "obliterating the distinction between areas of private judgment and of public control."²²⁹ If the reader is so inclined, they can skip ahead for a clearer understanding, but as of now, we shall simply take it as an unfounded misconception.

Ideological State Control is "the means of control associated with the canonical totalitarian regimes: a state using repressive measures to ensure conformity and to stifle dissent, including lies, state propaganda, and censorship of free speech, art, music and literature."²³⁰ Admittedly, repressive measures (eugenics), lies and state propaganda (myth of metals), and censorship (banishing of poets) are used in the Πολιτεία²³¹, but it is unfair to claim that Plato's aim was "a pervasive scheme of values, intentionally promulgated by some person or persons and promoted by institutional means in order to direct all or the most significant aspects of public and private life towards the attainment of the goals, dictated by those values"²³². Rather, it would be more correct to claim that Plato was a product of his time and we can adapt his original framework to modern audiences and modify the actions, but that doesn't absolve the totalitarian principle behind the actions themselves. So, what we need to really legitimize isn't eugenics, myths, or censorship; what must be legitimised is totalitarian measures... which themselves aren't actually isolated to totalitarianism. For example, police brutality, propaganda, and censorship of media is just as prevalent in modern "democratic" societies, if not more so. Since these measures aren't actually totalitarian then, and neither are they limited to the Kallipolis, this paper need not concern itself with legitimizing said measures at all. Nonetheless, simple justification can still easily be found in that all actions of the Kallipolis are based on rationality instead of emotivism amongst other non-teleological frameworks²³³. For instance, Plato approved the employment of "propaganda," but we ask that it be observed that this did not mean in general the inculcation of basic falsehood, but the mythical presentation of truth; in particular, moral values were not to be falsified or distorted²³⁴.

Authoritarian Intolerance can be understood through periphrasis of totalitarianism. Specifically, totalitarianism "allows no recognition (voice) of or representation to other (rival) political parties or loyalties"²³⁵. The misconception here is that Plato's framework

²²⁹ George H. Sabine, *A History of Political Theory* (New York: Holt, Rinehart and Winston, 1950), 905. Cited in "Was Plato a Totalitarian?" in *In Defense of Plato*, 573.

²³⁰ Lesley Brown, "How Totalitarian Is Plato's Republic?" in *Essays on Plato's Republic*, ed. Gail Fine (Oxford: Oxford University Press, 2008), 14.

²³¹ Plato endorses wide-ranging cultural restrictions: censorship of myths (377b), regulation of music and education (424e), and even the expulsion of poetry (607–608). These interventions are aimed at preserving order and virtue but raise concerns about authoritarianism. See "Was Plato a Totalitarian?", 548.

²³² Benito Mussolini, "Fascism," in *Enciclopedia Italiana* (1932), quoted in J. Salwyn Schapiro, *Liberalism and the Challenge of Fascism* (New York: McGraw-Hill, 1949), 9.

²³³ "But shorn of teleology, ethics as a body of knowledge was expurgated of its central content, and only remained as, essentially, a vocabulary list with few definitions and no context."

Alasdair MacIntyre, *After Virtue: A Study in Moral Theory*, 3rd ed. (Notre Dame, IN: University of Notre Dame Press, 2007).

²³⁴ "Was Plato a Totalitarian?", in *In Defense of Plato* (Cambridge, MA: Harvard University Press, 1953), 570.

²³⁵ This working definition of totalitarianism draws from multiple sources: (1) "a polity permitting no rival parties or loyalties" (Brown, *Totalitarianism*); (2) Webster's 1951 edition: "a centralized government under control of a political group which allows no recognition of other parties"; and (3) "rule by a small, rigorously

advocates the quintessential defining feature of authoritarian intolerance of dissenting ideas by being a system in which the ordinary citizen has no significant share, either direct or indirect, in the making of political decisions. This misconception likely stems from the stanza:

“All political decisions are made by the guardians without any reference to the citizen body. The guardians, moreover, are neither elected nor removable from office by popular vote. Politically, their power is absolute; the only control over them is itself ideological, in that they are under an absolute moral obligation not to allow any deviation from the system of education by which the ruling ideology, and therefore acceptance by all of the political system, is passed on from one generation to the next (424b).”²³⁶

At face value it might appear that political decisions made by the Guardians sans reference to the citizen body is synonymous with authoritarian intolerance, but looking at the text more closely we actually see that the saving key is in the idea that “the only control over is itself ideological”. This idea leads into the second part of legal legitimacy, namely substantiating the abstraction of the Philosopher King.

Second, then, we see that since the ruling entity is ideology (abstraction of the Philosopher King) itself which is at once and the same time both a superpower and a liability. It is at once superpower because it is precisely because there is no reference to the citizen body that there is no potential for tyranny from the citizens; but, at the same time, it is a liability because of potential democratic alienation common in sprawling bureaucracies. Simply, it displays the same pros and cons as everyone being equal under the law²³⁷. When considering the coincidence of the Philosopher King being a metaphorical, ideological abstraction²³⁸, we see that nobody or group of people or anything in between is actually the philosopher king themselves, it is simply a role, an ideology. One analogy that can be made is the papal infallibility of the Pope *ex cathedra*²³⁹ Vs. the Pope as human being; this mirrors the inerrancy²⁴⁰ of a philosopher as Philosopher King Vs. a philosopher as human being. Legal legitimacy thus stems from the totalization of the Philosopher King role as abstracted from the philosophers acting within the role. In other words, since the Philosopher King can literally be equated with the law itself, legal legitimacy is substantiated *par definitionem*.

Third, in returning to the idea that the Guardians do not make reference to the citizen body, this does not mean they do not seek discourse prior to decisive council, nor does it mean philosophers themselves never enter discussion within the various social stratifications. Insofar as Philosopher Kings comprises philosophers, they are necessarily married to a matter of consequence, not method. Since said philosophers (and subsequently, Guardians) are thus

limited group denying all policy input from others.” See: Brownlee, *Totalitarianism*, Oxford Philosophy Archive, <https://media.philosophy.ox.ac.uk/docs/people/brown/totalitarian.pdf>; and *In Defense of Plato*, 503.

²³⁶ C.C.W. Taylor, “Plato’s Totalitarianism,” *Polis: The Journal for Ancient Greek Political Thought* 5, no. 2 (1986): 4–29.

²³⁷ The rhetorical question—“does the law protect us, or do we protect the law?”—frames the law as an abstraction, much like the Philosopher King. In Plato’s works, *Republic* is followed by *Laws*, suggesting a thematic continuity between ruler and rule.

²³⁸ As with the Philosopher King and Kallipolis, the law is not reducible to its agents. Legal abstractions are upheld by but not identical to their enforcers—no lawyer or parliament embodies the law itself.

²³⁹ In Catholic theology, papal infallibility means that under specific conditions—when speaking *ex cathedra* on matters of faith or morals—the Pope is preserved from error. Defined at the First Vatican Council (1870), this authority has been invoked only rarely.

²⁴⁰ Plato, *Republic*, 340e: “No craftsman ever errs.” The craftsman’s authority is grounded in domain expertise—error arises only when someone attempts a craft without proper knowledge.

not dogmatically authoritarians, say, in the same way the Church is, therefore, this would not only allow but actively encourage discourse from political pluralism as it would further clarify the discussion towards Truth²⁴¹. Where this pluralism comes from – specifically the way in which multipolarism would be distributed – is dubious as in the original framework the social classes lack much social mobility. Nonetheless, I imagine²⁴² those comprising the philosopher King to be akin to Socrates in that they mingle between social stratifications of the populace in various geopolitical landscapes (e.g., the Agora) in a humble way that actively promotes constructive criticism. This, furthermore, provides the demos with a meritocratic asymmetric voice (equity) as opposed to just being another number in a sea of voters.

From this, defenders and critics alike should be able to recognize that the Πολιτεία ideal polis (Kallipolis) is not totalitarian. We see then that while he is no libertarian, Plato is not the proto-fascist some make him out to be, and indeed is making a number of points that are friendly to the interests of classical liberals. Misconceptions surrounding Platonic authoritarianism are thus not only unfounded, but the above evidence has even shown the opposite. Moreover, due to the totalization of the Philosopher King role in addition to philosophers (Socratic) unique invitation for discourse, we see justification of superior legality over the status quo, proving legal legitimacy of a (Neo-)Platonic framework.

Democratic Legitimacy

The subsequent step for the demos ever to practically support a seemingly disenfranchising framework is by garnering democratic legitimacy and proving Plato's original (or a Neo-Platonic) framework is enabled through universal suffrage, reflecting the principle of competent, popular sovereignty. This includes (1) disproving misconceptions surrounding Platonic statism; (2) substantiating the democratic spirit found within the Kallipolis; and (3) justifying superior democracy through equality of opportunity and equity of expression.

First, we need to disprove that Plato's framework is anti-demos, or anti-individual. This requires us to return to the Organic Theory of the State (particularly the Priority Thesis) to disprove Platonic statism, or the doctrine that the state is an end in itself to which individual interests must be subordinated or sacrificed, and in the interests of which all individual liberties must be abolished.

The idea behind the Organic Theory of State is that “the good of the individual is identified as his/her contribution to the good of the state. To confer the greatest benefit on an individual is to cause that individual to be in the best condition of which he/she is capable.”²⁴³ In other words, since the individual is part of the state, any good the individual confers to the state will necessarily confer to themselves as well²⁴⁴. This, in turn, supports

²⁴¹ Plato advocated free inquiry for philosopher-kings and envisioned political change toward moral improvement. While many of his views mirrored those of his contemporaries, his political ideal was grounded in a deeper humanity. See “Was Plato a Totalitarian?”, 571.

²⁴² Jo Walton, *The Just City* (New York: Tor Books, 2015), ch. 13. Walton imagines philosopher-kings as political gadflies—if their vision collapses under rhetorical challenge, then it was never just to begin with.

²⁴³ See Taylor, “Plato's Totalitarianism.”

²⁴⁴ One might ask whether it's inefficient to route goods through the state if individuals eventually receive them back—like a worker ant finding food only to give it to the colony and later eat from the shared supply. But this system justifies itself in two ways: (1) goods are not equal but equitable—division of labor and specialization (per Ricardo) yield more overall utility; (2) emergent goods (e.g., cake from flour, eggs, milk) require cooperation beyond any individual's capacity. Conferring goods to the state thus produces outcomes unattainable in isolation.

the idea that the State takes priority over the individual (Priority Thesis²⁴⁵) as it is only through subordination (and/or total sacrifice) of the individual that “the best condition for any individual is the condition of maximum contribution to the good of the state. Hence in requiring total commitment to the achievement of the well-being of the state the state itself ipso facto promotes the well-being of the individual²⁴⁶”.

Looking at Book IV of the *Πολιτεία*, it becomes clear how this misconception came about:

“we aren’t aiming to make any one group outstandingly happy but to make the whole city so, as far as possible (420b).”

The confusion most likely stems from where Socrates says that our aim in founding the ideal city is not to make any one class (ethnos) especially eudaimon, but as far as possible to make the whole polis eudaimon²⁴⁷. The term “ethnos” here is oft overlooked to be interpreted as Plato supporting sacrificing the one for the many, but this is clearly not the case. Conversely, Plato is arguing against any specific social class being superior and all other evidence, as already seen, points to the individual being valued as equally as the state.

Furthermore, we see that instead of the individual being subordinate to the State, Plato advocates for equal, if not superior value of the individual. In the subsection above, we’ve already seen valuation of the demos where the Philosopher King framework provides with a meritocratic asymmetric voice (equity) as opposed to just being another number in the sea of voters. In addition, from the dialogues it should be obvious that the polis and the demos are proto-Derridean²⁴⁸ mirrors of each other, meaning sacrificing one for the other would thus undoubtedly be a detriment (if not outright invalidation) all around.

Second, in continuing with the Priority Thesis as it relates to “obliterating the distinction between areas of private judgment and of public control”²⁴⁹, substantiating the democratic spirit found within the Kallipolis means reconciling this supposedly “obliterated” distinction. By all accounts, Plato, despite his commendation of control and censorship, was no worshipper of the state, but conversely an advocate (which could be misconstrued due to the extent of draconian intervention) of the subordination of political instrumentalities to moral and educational ends in the service of individuals²⁵⁰. This, if not evidence of superior value, is at least consistent with equal valuation of the demos. To see how exactly, we can turn briefly to modern politics which in comparison to the Platonic framework is by far not only more anti-demos, but also anti-kratia as well.

We see that in modern democracies, the “demos” does not include felons, those living in territories, children, permanent legal residents who are not registered citizens, mentally disabled, etc.²⁵¹ The aphorism, “of the people, by the people, for the people” is clearly

²⁴⁵ Lesley Brown, “How Totalitarian Is Plato’s Republic?” in *Essays on Plato’s Republic*, ed. Gail Fine (Oxford: Oxford University Press, 2008), 19.

²⁴⁶ Ibid.

²⁴⁷ In *Republic* 420b–421c, Socrates explains that no one group should be made excessively happy at others’ expense. Happiness in the Kallipolis arises from each citizen practicing their natural craft—whether producing, guarding, or ruling (434a–c). At 519e–520a, this is echoed when guardians are asked to forgo personal intellectual pursuits for civic duty.

²⁴⁸ As previously discussed, the Philosopher King and Kallipolis are mutually constitutive. Like Derrida’s difference between speech and writing, the individual and the polis substantiate one another through reciprocal development.

²⁴⁹ George H. Sabine, *A History of Political Theory* (New York: Holt, Rinehart and Winston, 1950), 905. Cited in “Was Plato a Totalitarian?”, 573.

²⁵⁰ Ibid, 572

²⁵¹ USA.gov, “Who Can Vote?” Last modified June 12, 2023. <https://www.usa.gov/who-can-vote>.

missing the most important part: the people. In addition, “kratia” means nothing if you are outvoted. There are many significant historical examples that highlight the limitations of so-called majoritarian “democracy”, where the minority was correct but ignored, leading to disastrous outcomes. For example, the Space Shuttle Challenger Disaster²⁵², the Peloponnesian War²⁵³, as well as Churchill and the Rise of Nazi Germany²⁵⁴. As seen then, modern “democracies” – and even as far back as Plato’s time²⁵⁵ – aren’t actually democracies because universal suffrage isn’t advocated; moreover, even for those demos not discriminated against, minority voices are ignored and everything turns into a numbers game.

The original meaning of democracy, or rather the ideal democratic spirit is a regime type which doesn’t concern itself with majorities or minorities; the -kratos root group is not a “number term” unlike the -arche root group²⁵⁶. Simply, democracy is the “capacity to do things”, therefore, the democratic spirit is compatible with all forms of government from autocracy to polyarchy. Whether or not a form of governance is democratic then depends on whether or not the framework allows for equality of opportunity and equity of expression, or the collective capacity of a public to make good things happen in the public realm. Clearly, modern democracies both in regards to “demos” and “kratia” fail to be democratic, so much so to the point that anarchy based civil disobedience is more democratically legitimate than modern democracies.

Third, to expound upon the supposedly “obliterated” distinction between private judgement and state control we should remember that the Platonic framework is necessarily married to a matter of consequence, not method:

'and by every means the total abolition of what is private (idion) from their life, to such an extent that as far as possible even what is naturally private is in a way common, so that their eyes and ears and hands seem to see and hear and act in common, and they praise and censure in unison, all liking and disliking the same things'²⁵⁷.

This picture of the elimination of private life and private feeling is indeed forbidding, but it does nothing to support the thesis that Plato is here supporting extreme totalitarianism, for he continues (793d6-e1):

²⁵² The 1986 Challenger disaster illustrates the cost of ignoring expert minority warnings. Engineers advised against launch due to O-ring risks, but NASA prioritized schedule. All seven crew members perished. This underscores the importance of expertise over majoritarian impulse in critical decision-making.

²⁵³ Athens’ rejection of cautious counsel before the Peloponnesian War—despite warnings from Pericles and others—led to strategic overreach, plague, and collapse. This historical case illustrates the danger of sidelining minority foresight in favor of populist aggression.

²⁵⁴ Winston Churchill’s early warnings against appeasing Nazi Germany were dismissed by the British government in the 1930s. The Munich Agreement (1938), ceding the Sudetenland, exemplified the strategic failure of ignoring minority foresight. Churchill’s predictions proved accurate when Germany invaded Poland in 1939, initiating WWII—highlighting the value of dissent in democratic security.

²⁵⁵ In Plato’s Athens, “citizenship” excluded women, slaves, and foreigners. Democratic functions like jury duty were determined by sortition (kleroterion), not competence. That Socrates was condemned under this system challenges the romanticization of early democracy as more than a numbers game.

²⁵⁶ The original meaning of *demokratia* referred less to majority rule than to equitable access to civic participation, aligning more with “capacity to do” than mere numerical dominance. See Josiah Ober, “The Original Meaning of Democracy.”

²⁵⁷ Plato, *Republic* 739c–e; see also Taylor, “Plato’s Totalitarianism.”

'that sort of city, whether inhabited by gods²⁵⁸ or several children of gods in living that sort of life they inhabit it euphrainomenoi, i.e. enjoying themselves'²⁵⁹.

We see then that using repressive measures to ensure conformity and to stifle dissent, including lies, state propaganda, and censorship of free speech, art, music and literature are not actually defining elements of the Kallipolis itself but rather what was “scientifically”²⁶⁰ believed to actually be necessary for euphrainomenoi. Plato didn’t have anything against the poets themselves or their poetry – in fact, many a times he even advocated for their usefulness. And this is the key here to understand: *usefulness*. Plato was an advocate of technologism²⁶¹, or the worship as an end in itself of whatever “works” or is efficient. Of course, we could thus argue that there need be legitimacy for even technologism - which with means justifying the ends (specifically, the rendition of “sacrificing” freedom) and will be expanded upon in the following subsection of pragmatic legitimacy - but insofar as we are aiming for a successful 3rd party framework, and that success prerequisites efficacy, then we shall take it as an assumption that the demos ever to practically support a seemingly disenfranchising framework so long as it “works” towards turning the Self as a whole towards the good. Moreover, we see that if anything, the “obliterated” distinction between private judgement and state control is, like Lincoln²⁶², all in service of private judgement.

Finally, and perhaps the greatest evidence in justifying superior democratic spirit, if Plato were promoting allotrion agathon and viewed people only as “human cattle whose sole function is to provide for the material needs”²⁶³, he would not even be interested in Glaucon's challenge in the first place: he would simply say that justice is not in the individual's interest and shake his hands of it.

Thus seen, Plato’s framework does do advocate for statism and is conversely pro-demos, pro-kratia, and pro-individual. Churchill once said that “democracy is the worst form of government, except for all the others”, and now we have seen why: because modern democracies not only fail to enable kratia but even lack (inclusivity of) the demos to begin with. Plato’s framework on the other hand empowers by allowing for both equality of opportunity and equity of expression through meritocratic asymmetric voice, and is at least more democratically legitimate than the governments of today through subordination of political instrumentalities to moral and educational ends in the service of individuals. By all accounts, then, we see justification of superior democratic spirit and cause for democratic legitimacy of a (Neo-)Platonic framework.

Pragmatic Legitimacy

The final step for the demos ever to practically support a seemingly disenfranchising framework is by garnering pragmatic legitimacy and proving Plato’s original (or a Neo-Platonic) framework is able to deliver positive outcomes, such as economic prosperity, security, and social welfare. This includes (1) disproving the “sacrificial” elements of

²⁵⁸ Plato uses *euphrainesthai* to express divine joy: once for the Demiurge’s delight in world creation (*Timaeus* 37c), and once for a goddess pleased by worship (*Laws* 796b–c).

²⁵⁹ Ibid

²⁶⁰ Plato’s limitations were epistemic, not moral. Like outdated scientific justifications for racism, philosophical frameworks are contingent on historical knowledge.

²⁶¹ In *Republic* 341d–342c, Socrates argues that if governance is a *techne*—a rational craft—it must aim at the welfare of the governed.

²⁶² Abraham Lincoln suspended habeas corpus and other liberties in order to preserve the Union—an act that raises enduring questions about violating law to uphold law.

²⁶³ Karl R. Popper, *The Open Society and Its Enemies*, vol. 1: *The Spell of Plato*, 5th ed. (Princeton, NJ: Princeton University Press, 1966), 47.

freedom through revisiting the idea of “freedom is slavery”; (2) substantiating paternalism as hierarchal control; and (3) justifying superior performance through trickle-down wisdom and its subsequent positive feedback loop.

First, and foremost, we must ask if freedom can be divorced from the Good? Or better yet, if freedom is an end in itself (perhaps even of equal value with the Good), where do is the line drawn on how much freedom is allowed²⁶⁴?

It is often joked that authoritarian regimes are egalitarian by way of equality through persecution²⁶⁵. But the reality is that in many paternalistic nations, restrictions placed on political freedom, self-expression and self-determination (i.e. China, Singapore, U.A.E.) has inevitably led to stable (and sometimes superior) economic markers, i.e. gini coefficients, higher HDI, etc. It is also the case that welfare is often a case of subjective relativity²⁶⁶, meaning the opposite – removal of *pleonektein*²⁶⁷, particularly through equalizing communist aspects taking precedence over (capitalistic) freedom in Plato’s egalitarian original framework – would theoretically result in a populace which is not relatively worse off, and possibly even absolutely better off than their counterparts in more turbulent societies. This brings us back to the earlier argument around justifying the legitimacy of technologism, specifically the rendition of “sacrificing” freedom as a means to an end.

To best understand the ethical ramifications of tradeoff between freedom and technologism, moreover, not fall into the pit trap of anachronistic readings in the process, we should remind ourselves that just as the polis and the demos, Philosopher King and Kallipolis are proto-Derridean²⁶⁸ mirrors of each other, other dichotomies that we pose today didn’t exist – or rather, weren’t a cause of contention – in Plato’s time. Here then, we should be able to notice that creating a binary between freedom and technologism under the context of what ends in the Good is a false dichotomy, and we can possibly even go as far to say that Plato was a proto-Kantian in that he parallels people’s freedom (if not people and/or people’s telos themselves) as an end in themselves²⁶⁹. We shall get to the extent to which Plato honoured free and spontaneous performance shortly, but first, to get them out of the way and to set the scene for what follows, I exposit why non-sacrificial true freedom is relative as opposed to absolute.

As alluded to in the previous part²⁷⁰: “Freedom is Slavery”; it is about neither incarceration (thesis:negation) within the box, nor emancipation (antithesis:affirmation) outside the box, but the inextricable shackles (synthesis:*différance*) relative to the box. The Paradox of Freedom is thus self-actualization following Amor Fati of Sisyphean roles: subliminal Freedom when evoked from self-imprisonment (harbouring resentment and indulging escapist fantasies). We see exemplification of freedom as relative and not absolute

²⁶⁴ If freedom is an end in itself, why constrain actions like pollution? If we’re free to make mistakes, at what cost to others? The paradox emerges: total liberty undermines itself when it infringes on others’ well-being. (*Precrime* and butterfly effect concerns may apply.)

²⁶⁵ *Know Your Meme*, “Avatar: The Last Airbender / The Legend of Korra,” accessed July 10, 2023. <https://knowyourmeme.com/photos/2035045-avatar-the-last-airbender-the-legend-of-korra>.

²⁶⁶ The Maasai, despite low material wealth, report high subjective well-being. As noted earlier, happiness is comparative—one feels rich not by having much, but by having more than one’s brother-in-law.

²⁶⁷ Socrates critiques *pleonektein* (πλεονεκτεῖν), the desire to have more than others. If happiness is relative, then one person’s gain necessitates another’s loss.

²⁶⁸ As noted earlier, the Philosopher King and Kallipolis are co-constitutive. Their mutual dependence echoes Derrida’s *différance* between presence and absence, substance and representation.

²⁶⁹ Plato does not use Kantian terminology (treating persons as ends), but emphasizes care of the soul as the highest concern after the divine. See: Robert W. Hall, “Plato and Personhood,” *The Personalist Forum* 8, no. 2 (1992): 88–100. <http://www.jstor.org/stable/20708653>.

²⁷⁰ Previously we explored the paradox “freedom is slavery”: epistemic constraint informs volition. By accepting the Drifting Will’s influence, one finds orientation. Appreciating our entangled enslavement paradoxically yields real autonomy.

with the (finger-)trap which becomes free not when we abdicate responsibility but accept it; not running away, but running towards. It is thus not “sacrifice”, but the liberation of grounding oneself to a lodestone which makes us truly free²⁷¹. Returning once again to Kant, freedom of will does not lie in the moral law itself but in the capacity of autonomy to choose it so. Much in the same way, freedom of action does not lie in absolute permission itself but in the capacity of autonomy to choose it so. Coincidentally, what liberates us from the overwhelming, absoluteness of Heideggerian anxiety²⁷² is the willful (non-sacrificial, non-coercive) enslaving ourselves to a lodestone, a master (framework), a point of reference.

Second, then, having laid the basis of non-sacrificial true freedom, we are now in a better position to understand the extent to which Plato honoured free and spontaneous performance; moreover, how subjugating oneself to hierarchal control (whether internal of the soul, or external of the city) is intrinsically liberating. For many, the initial reading of Plato as esteeming spontaneous action over passive obedience and control appears incoherent. Confusion mostly appears to arise from this stanza in particular:

“The main principle is this—that nobody, male or female, should ever be left without control, nor should anyone, whether at work or in play, grow habituated in mind to acting alone and on his own initiative, but he should live always, both in war [942b] and peace, with his eyes fixed constantly on his commander and ...”²⁷³

For many scholars “freedom” of Plato’s citizens is no more than a “habitual respect for authority,” and that the Πολιτεία is, except for the handful of rulers, a city of “greedy children,”²⁷⁴ kept in order solely by external restraints.²⁷⁵ This is often the first interpretation many scholars have: nobody should be left to their own devices (AKA: no longer under the TPN) and nobody should have the habit of acting on their own accord. This reading, however, displays a cumulative error due to meaning lost in translation; while the interpretation that external restraint is necessary in turning the Self as a whole towards the Good, the interpretation that passive obedience should become habitual is fundamentally inconsistent with Plato’s moral and social values²⁷⁶. We recall his ironic depiction of why the “City of Pigs” (372d) were incompatible with the human telos, extending this, Plato also draws analogy how those having lived virtuously but only from blind custom and habit is only fitting for social harmless creatures such as bees²⁷⁷. Thus fitting, external restraint is necessary, but passive obedience is not. Drawing parallels with what was said above, passive obedience is sacrificial and not fitting for Plato’s citizens, rather the yoke must be worn willingly and autonomously²⁷⁸ with understanding that which is truly liberating is the willful

²⁷¹ Hume argued that determinism—regularity in causal structure—is a precondition for meaningful freedom. See: Carl Hoefer, “Causal Determinism,” *Stanford Encyclopedia of Philosophy*, Winter 2022, <https://plato.stanford.edu/entries/determinism-causal/#DetHumAct>.

²⁷² *Overthink Podcast*, “Heidegger on Anxiety and Dasein,” Bilibili, uploaded February 12, 2022. <https://www.bilibili.com/video/BV1T3411s7ni/>

²⁷³ Plato, *Laws*, Book 12, trans. R.G. Bury, *Perseus Digital Library*. <https://www.perseus.tufts.edu/hopper/text?doc=Perseus%3Atext%3A1999.01.0166%3Abook%3D12>

²⁷⁴ We don’t let children vote—so why should the demos *kratia*?

²⁷⁵ Original citation in “Was Plato a Totalitarian?”, 503.

²⁷⁶ “What of freedom? Is that not the fruit of another tree, requiring wizardry to graft onto Plato’s stock? We think not... Plato honored free, spontaneous performance above obedience.” See “Was Plato a Totalitarian?”, 579.

²⁷⁷ Plato, *Phaedo*, trans. G.M.A. Grube (Indianapolis: Hackett, 1977), 141.

²⁷⁸ The producers in the ideal state experience “idealized slavery”: they are autonomous only insofar as they freely accept direction from those who know the Good. Most people, lacking the internal harmony to attain *eudaimonia* alone, require external governance to avoid tyranny by their lower nature. See: Plato, *Republic* 590c–d; *Gorgias*; and Taylor, “Plato’s Totalitarianism.”

(non-sacrificial, non-coercive) enslavement to the Platonic framework. This way, Plato argues, when the rational faculty freely chooses to be subject to the Guardians (432a, 433c-d), they are thus liberated²⁷⁹ from their lower nature if left to their own devices (AKA: no longer under the TPN).

Incidentally, giving up our will to someone else not only satisfies the Cinderella complex, but also outsources help with our lack of willpower. From the previous part we identified that one of the reasons why we fail to obtain the Good Life is people know what to do, they just don't do it; instead, they simply wait for somebody else to rescue them. In a sense, people are slaves to the system (and by association, their Cinderella complex) because it is the easiest mode of living²⁸⁰.

Being subject to the Guardians then is the perfect match for those with Cinderella complex. Furthermore, for those who have not reached their “genetic ceiling” of willpower for whatever reasons, subjugating oneself to hierarchal control externally will come to reflect internally²⁸¹, i.e. having “guardian training wheels” will eventually help the subject learn to ride their own moral bike. This is a result of cultural osmosis (wherein the polis is a mirror of the soul) as well as cognitive behavioural science.

Third, what justifies subjugating individuals' fate for the sake of their own good is not only liberation from irrational tendencies but also the trickle-down effect of wisdom. Relating to the initial worry then that by way of dictating the very essence of goodness for all citizens, censorship and the stifling of individuality would also occur has thus been unfounded. To reiterate how this is, we saw that lies, state propaganda, and censorship are actually subordination of political instrumentalities to moral and educational ends in the service of individuals; moreover that freedom and technologism is a false dichotomy wherein if not evidence of superior value, is at least consistent with equal valuation of the demos. Furthermore, the subsequent positive feedback loop following the trickle-down effect of wisdom can be seen in relating back to how hierarchal control is liberating whereby outsourcing decision making to Guardians, everyone is better epistemically placed. To make a weak analogy, rather than impulsively going against a parents' better judgement and going down that dark alleyway, instead going to piano class would, exceptions notwithstanding²⁸², typically lead to a specific epistemic placement with better options which would in turn reinforce even better options still.

Evidently then, Plato's freedom is not sacrificial and subjecting oneself to hierarchal control is in fact even more liberating than absolute permission since trickle-down wisdom and its subsequent positive feedback loop justify any substantial relinquishment. Specifically, we have addressed how true freedom is relative and the freest thing we can possibly do is in our autonomy, recognize the most rational path forward and choose the framework accordingly. On the basis that said framework not only liberates the demos from wasting time and energy, further empowering them, but also leads to better performative choices through

²⁷⁹ Like children who find freedom in trusting parental judgment, the demos may be freer by relinquishing certain choices. Philosopher Kings, unlike *The Matrix*'s red-pill illusion, liberate through restraint—sparing citizens the burden of overchoice and regret. Elitism, properly understood, mitigates epistemic overload and restores functional autonomy.

²⁸⁰ An analogy: if wealth were distributed evenly—everyone gets one chicken—most would eat it, some would sell eggs, and a few would raise more chickens. The metaphor isn't perfect, but it illustrates innate variance in how people handle equality of resources.

²⁸¹ Aristotle teaches that virtue begins with imitation. One must act “as if” brave to become brave—like learning a language: first mimic, then internalize, until virtue becomes second nature. See: *Nicomachean Ethics* 1103b14–22, 1104b5–10, trans. Terence Irwin (Indianapolis: Hackett, 2019).

²⁸² See also the Snow White problem: observer selection bias can shift probability assessments for things one could not have learned otherwise.

positive feedback loop following the trickle-down effect, we see justification of superior performance, proving pragmatic legitimacy of a (Neo-)Platonic framework.

Interpretation:

From this second section *On Practicality*, we've seen that first, legal legitimacy stems from the totalization of the Philosopher King role in addition to philosophers (Socratic) unique invitation for discourse. Second, democratic legitimacy is further derived from both equality of opportunity and equity of expression through meritocratic asymmetric voice in addition to subordination of political instrumentalities to moral and educational ends in the service of individuals. Third, pragmatic legitimacy arises through liberating the demos from wasting time and energy and the reinforcement of positive feedback loops following the trickle-down effect.

This entire section in a nutshell then neatly echoes what was posited in the previous part. To reiterate, people should not (and cannot really) make their own decisions (be left to their own devices; AKA: no longer under the TPN), therefore, in keeping with Plato's own presentation, the 'idealized slavery' is one in which we are democratically empowered to choose our own masters, which by now we have hopefully seen to be the abstraction of the Philosopher King 3rd party framework. But of course, the combination of Plato's metaphysics and his pessimistic moral psychology²⁸³ is a little outdated as of yet. From the previous part we've seen that emphasis should not be placed on justifications from the soul insofar as claiming the common man "incapable of grasping the Good", rather, from a much more secular justification, as already provided previously, we see that the "psychic harmony" is impossible because of the limited Ego instead.

Now, returning to the research question we thus see that the greatest evidence for the demos ever to practically support a seemingly disenfranchising framework is the dissolution of the idea that said framework is even disenfranchising. From this section *On Practicality* it should now be obvious that if not evidence of superior value, legal, democratic, and pragmatic legitimacy is at least consistent with equal valuation of other modern political frameworks. Thus, in the previous section *On Tenability* we showed that Plato's framework for the Philosopher King and Kallipolis is theoretically feasible, and in this section *On Practicality* we adapted (or better yet provided a Neo-Platonic revisionist viewing of) said framework for it to be more scientifically accurate and socially acceptable, for the next section *On Efficacy* then we shall explore the reality of implementing utilitarian contrivance of the Good. In other words, we've revamped a workable service, identified the target market, now it is time to sell by addressing to what extent would this Neo-Platonic 3rd party framework be superior.

On Efficacy

In the previous two sections, we've already hinted about some of the key selling points of what makes a Neo-Platonic 3rd party framework superior. Some of the more obvious ones include being more individual, democratic, and libertarian compared to its counterpart frameworks as well as being immune to corruptibility. Here, we will explore in more detail the efficacy²⁸⁴ of how the actual framework commensurates implementing utilitarian

²⁸³ Taylor, "Plato's Totalitarianism," 21.

²⁸⁴ Before continuing, recall that efficacy is not the same as efficiency. Efficiency minimizes waste; efficacy secures results. Marcus Aurelius may be our best historical approximation of the Philosopher King—but does he prove the ideal?

contrivance of the Good; Granted that probability ≥ 1 , AND the people allow it to happen, is a socially engineered – organized almost to the point of religious - framework necessarily the most efficacious?

Up till now, adaptations of Plato's original framework have been little more than providing Neo-Platonic revisionist viewing of said framework. This has indeed adapted the original framework so that it become more relevant to the modern audience by way of increased scientific accuracy and social acceptability. At this point it might be worth considering that perhaps this modification is enough. But, then again, for thesis to be a successful practical treatise instead of just another manifesto, it must be efficacious. Moreover, in the millennia since Plato's conception, contribution from the literature has been far more than simply keeping up with the times. For example, Mill's Competent Judge bolsters the caricature of the Guardian; Rousseau's General Will reinforces the Philosopher King abstraction; and Locke's Social Contract affirms the legislation of the Kallipolis. To fully substantiate the efficacy of a modified framework then, it will be necessary to draw upon the wisdom, experience, and historical progress that are the footnotes since Plato. In a way then, this section On Efficacy is as much armouring the Neo-Platonic 3rd party framework with strategies for practically implementing utilitarian contrivance of the good as it is backing up the superiority of said framework.

Earlier we mentioned that in relation to attaining the Good Life, elitism is the most compatible form of government – if it were only realizable. Building upon the justifications then that the Philosopher King is possible in theory (tenability) and practice (practicality), we shall now reconcile the feasibility of elitism drawing parallels from three criterion: Sophia; Techne; and Phronesis insofar as they are specifically relevant to the Philosopher King. Sophia is about being well versed in what the Good and addresses the efficacy of its realizability as it relates to wisdom. Techne is about being well experienced in how to practically implement the Good and concerns itself with technical capabilities for practical implementation. And Phronesis is about poesis – the intersection of practice; culminating as the gestalt combination of Sophia and Techne, Phronesis focuses on precedence of a Neo-Platonic framework (or at least elements wherewithin) through comparative analysis case studies.

Sophia

In the background it was claimed that elitism is the most compatible form of government in relation to the Good. Backing up this earlier proclamation from superiority of the Neo-Platonic 3rd party framework then requires us to first address the efficacy of its realizability as it relates to wisdom specifically relevant to the Philosopher King; i.e. being well versed in what the Good is in contrast to simple financial prudence.

Traditionally, Sophia has been known as the integration of nous and episteme – the crossroads of intelligence and wisdom where intelligence is mastering knowledge, and wisdom is weaving that knowledge into the fabric of life. The wisest strategies of utilitarian contrivance to the modification of the Neo-Platonic framework would be adaptations of Mill's Competent Judges²⁸⁵ and Rousseau's General Will²⁸⁶.

²⁸⁵ J.S. Mill's concept of "Competent Judges" posits that those with greater experience and knowledge are better suited to evaluate competing pleasures. This idea parallels Plato's Philosopher King, who rules by epistemic superiority.

²⁸⁶ Rousseau's "general will" represents the collective moral will of a community. It echoes Plato's idea that rulers must act on behalf of the whole, aligning authority with communal interest.

Perhaps the best way to understand how to adapt Mill's Competent Judge is to indulge in a thought experiment from my time at Oxford²⁸⁷: assuming that each decision leads to two possible futures and a Judge (or panel thereof) having experienced and is "competently acquainted with" (2.5) both identifies the objectively best option²⁸⁸, would you allow said Judge to dictate the rest of your life for you?

For many there is a general aversion stemming from the desire for autonomy and/or skepticism that the Judge would know what is subjectively best for them²⁸⁹. The former argument from autonomy is reflexive misjudgment, as seen above, subjecting yourself to a Competent Judge / Philosopher King actually allows for greater freedom. The latter argument is in itself invalid as the Good – par the previous part – follows subject-objectification and object-subjectification, meaning if the Judge can identify the objectively best option, then they know what is subjectively best as well.

As the very basis of the Competent Judge is literally derived from Mill's Utilitarianism, moreover, that it bolsters the caricature of the Guardian, we see that adapting Competent Judges as a vettable characteristic of the Philosopher King (i.e. an added requirement of having experienced viscerally the lives that it dictates²⁹⁰) promotes utilitarian contrivance.

Another strategy to consider adapting is Rousseau's conception of the General Will. Roughly synonymous with Kant's Categorical Imperative²⁹¹ and Hegel's Spirit²⁹², the fundamental idea substantiating the General Will is a Geist of collective identity, ethical duty, and the relationship between individual and community interests throughout space-time. In sum, the General Will is the expression of what is best for the entire community, transcending individual interests. According to Rousseau, individuals should subordinate their personal desires to the General Will for the sake of the community's well-being. The general will is not a simple summation of individual wills but rather an idealized expression of what citizens would collectively choose under conditions of equality and rationality. It just so happens that the other name for this General Will is the "Sovereign" which many find linguistically – if not, literally – parallel with the "Philosopher King".

In line with the idea that Competent Judges can distinguish between better and worse, higher and lower choices, it should be noticeable that as the General Will is "General" instead of being partisan or majoritarian, then from a utilitarian standpoint, said General Will is undoubtedly quantifiably efficacious. Furthermore, as the General Will nears (if not directly expresses) the ideal, it too is qualitatively efficacious. Adapting the General Will as the Geist and Categorical Imperative by which the Philosopher King is held in line then also would appear to promote utilitarian contrivance.

One key issue that ascribes to both strategies of Competent Judges and the General Will is essentially a problem of overfamiliarity. A variation of the "Competent Judge

²⁸⁷ Kate Kirkpatrick, *Utilitarianism Tutorial*, Oxford University, 2020.

²⁸⁸ "From this verdict of the only competent judges [...] must be admitted as final." — J.S. Mill, *Utilitarianism* 2.8.

²⁸⁹ Plato's ideal state aligns with Voltaire's "benevolent dictatorship," where rulers govern for the people without consulting them. Yet without firsthand experience, only a fool would choose ignorance over wisdom—like a pig preferring mud to music. This echoes Nagel's "What is it like to be a bat?" and the *Mary's Room* thought experiment on qualia.

²⁹⁰ Relates to Plato's discussion of meddling (443d–444a): the state's role in shaping behavior raises tensions about intervention and individual autonomy.

²⁹¹ Kant's Categorical Imperative grounds morality in universally applicable maxims. The General Will, as Kant sees it, must derive from rational principles shared by all.

²⁹² Hegel's *Spirit* reflects a time-bound cultural consciousness. In contrast, Rousseau's General Will is timeless, reflecting humanity's eternal striving for freedom beyond any specific epoch.

Problem”²⁹³, the issue of overfamiliarity is that, moreso than simply trying to sift through information overload (i.e. conflict of interest), is the natural problem of being unable to be subjectively objective.

We can imagine a courtroom judge handling a case of a knife to the brain where both plaintiff, defendant, and judge are childhood friends. The Competent Judge Problem states that overfamiliarity clouds judgement from the General Will making the judge unable to take a standardized, detached, dispassionate outsider stance. However, simultaneously, not in spite of, but because of this non-objective influence, the judge has performative and affective knowledge²⁹⁴. This experiential insider knowledge is precisely what allows the judge to be "competently acquainted with" (2.5) both parties, moreover, it is what empowers them to be a Competent Judge in the first place. Unlike a simple conflict of interest then or of removing subjective experience, the Competent Judge Problem is an issue of competently comparing²⁹⁵ the two since both objectivity and subjectivity are equally as important. In other words, it is a catch-22 of needing to have subjective familiarity to initially gage the General Will, but needing objective detachment to properly perceive said General Will, and the problem is in reconciling this paradox.

For Brownlee, we can entertain two thought experiments²⁹⁶ to elucidate the problem of epistemic comparison: (1) difference arising from human experience; and (2) the experience itself as hindering judgement.

(1): Imagine two people, Ali and Bo, who endure the social and sensory deprivation of solitary confinement. And, imagine that one of them is less acutely affected by it than the other is. Ali suffers but survives with her faculties broadly intact. By contrast, Bo becomes self-abusive, hallucinatory, and eventually semi-catatonic, unable to respond to instructions.

From this first thought experiment on the difference arising from human experience we might question, which of them has the necessary practical knowledge of solitary confinement to compare it competently with other experiences? Who is the competent judge here? As the spectrum of human experience is vast, is it possible to gage the General Will from anecdotal experience alone? Moreover, even if the General Will could be triangulated, we are again brought back to the question of subject-objectification and object-subjectification whereby everyone would react to implementation of the General Will the same.

(2) Assuming that starving is worse than having a headache or watching a boring movie. But, as starvation progresses²⁹⁷, hunger eventually diminishes to the point that

²⁹³ Kimberley Brownlee, "The Competent Judge Problem," *Philosophy and Phenomenological Research* 77, no. 1 (2008): 12.

The paradox arises when extreme experiences—like trauma or starvation—both qualify someone to judge them and simultaneously undermine their ability to judge. These are precisely the experiences we most wish to evaluate competently.

²⁹⁴ Ibid.

²⁹⁵ Ibid, 2.

The competent judge paradox arises when experiences either elude assessment entirely or simultaneously undermine and enhance one's capacity to judge—e.g., starvation both impairs cognitive clarity and grants firsthand insight into suffering. These "assessment-eluding" cases destabilize traditional frameworks of comparative evaluation.

²⁹⁶ Brownlee identifies seven additional "assessment-eluding" cases: starvation, intoxication vs. sobriety, sleep vs. deprivation, two types of trauma, memory implantation, forgetting, and birth.

²⁹⁷ We are specifically referring to *starvation*, not *fasting*. Though prolonged fasting may lead to starvation, the positive effects of voluntary, healthy fasting should not be conflated with the involuntary deprivation associated with starvation.

it can give the illusion of fullness. Furthermore, the longer starvation occurs, or the closer one approaches death, endorphins are released making the experience of starving less painful than a headache and perhaps even enjoyable.

From this second thought experiment where the experience itself hinders judgement, we cannot appeal to either our own practical knowledge of it or others' testimony about it as evidence that it is worse than a headache or a boring movie. Furthermore, In its non-comparative form, the problem is that we cannot say whether a given experience is pleasant, painful, or neutral even though intuitively we know which it is. Determining that starvation is a painful experience at all is seemingly beyond us since it eludes first-person assessment even though, intuitively, we know that starvation is extremely painful.

Thus seen, the strategies of Competent Judges and General Will appeal to Sophia insofar as they promote utilitarian contrivance. However, the issue of overfamiliarity in conjunction with epistemic problems means Sophia alone is insufficient evaluation for implementation. This echoes the earlier idea wherein Philosopher Kings must be well versed in not only what the Good is, but also experienced in how to practically implement the Good as well. Inasmuch, we now move on to *Techne* and its strategies to resolve the above problems in addition to exploring its technical capabilities that address the efficacy of realizing the Neo-Platonic framework.

Techne

Moving on from Sophia to *Techne*, the next step in backing up superiority of the Neo-Platonic 3rd party framework requires us to address the efficacy of its realizability as it relates to technical capabilities relevant to the Philosopher King; i.e. being well experienced in how to practically implement the Good in contrast to simple legislative skills.

Techne is oft translated as craftsmanship or art and is associated with making or doing. The typical modern paradigm is that *techne* is so enmeshed in material application that it resists any general explanation but must be learned by practice, but this was not always the case. For Plato, the Kallipolis would have rulers who truly know (as opposed to opine) wisdom and justice. Considering that Justice, by virtue²⁹⁸, is (or at least requires) *techne*, and *techne* necessarily integrates episteme²⁹⁹, we confer that the Philosopher King must necessarily reconcile Sophia with *Techne*. It thus accordingly behooves of us to explore – not least to resolve the issue of overfamiliarity in conjunction with epistemic problems – by what alterations might the Neo-Platonic framework be made serviceable for present use. This leads us to the more *techne* strategies of Locke's Social Contract and Rawls' Original Position.

Starting with the Social Contract, the original conception starts neither with Locke – nor Hobbes – but dates back the very idea of society itself. This is apt since neither is the Social Contract formal legislative but more akin to the aforementioned Geist. What is specific

²⁹⁸ Socrates, Plato, the Stoics, and Aristotle all consider *virtue* (especially justice) as a *technê*—or at least, something closely resembling a trained skill.

See: *Techne as a Model for Virtue in Plato*, in *Productive Knowledge in Ancient Philosophy*, Cambridge University Press. <https://www.cambridge.org/core/books/abs/productive-knowledge-in-ancient-philosophy/techne-as-a-model-for-virtue-in-plato/942ADA7D03B609188FB8E5F693C55FE0>

²⁹⁹ Though theory and practice often appear irreconcilable, ancient thought frequently merged *epistêmê* (theoretical knowledge) and *technê* (practical skill), in a manner similar to how modern physics treats space-time as inseparable.

See: *Epistêmê and Technê*, Stanford Encyclopedia of Philosophy, <https://plato.stanford.edu/entries/episteme-techné/>

about the adaptations Locke made to the Social Contract³⁰⁰ – and why we are adapting Locke’s modification instead of Hobbes or Rousseau’s – is that it proposes a more scientifically accurate account of Nature; is more apolitical tool than agenda; and coincides better with the role of equality of opportunity and equity of expression (i.e. permits, if not outright supports, the right to rebellion). For Locke, political authority arises from a Social Contract among individuals who agree to form a government to protect their natural rights. This notion of consent and the Social Contract resonates with Plato’s conception of the Philosopher King’s authority being derived from the consent of the governed. Furthermore, allowing for a more concrete constitution³⁰¹ of the Social Contract stemming from the General Will and substantiating the Sovereignty of the Philosopher King. This, in one way, resolves the issue of overfamiliarity in conjunction with epistemic problems by circumventing the problem entirely. Specifically, as the Sovereign is a singular multipolar entity, it completely embodies subject-objectification and object-subjectification as it is both subject and object; in addition, it removes the need for interpersonal comparability by way of intrapersonal feedback, i.e. right to rebellion.

Following Locke’s Social Contract - and arguably an inheritance of Locke’s Social Contract – is Rawls’ Original Position. The Original Position, by way of Veil of Ignorance (formalization of Kant’s universal law maxim), addresses the practical implementation of a more concrete constitution of the Social Contract, i.e. a mechanism that tangibly takes the energy source that is the General Will, converts it in the engine that is the Social Contract, and results in the production of the Sovereign. As it directly relates to the techne of justice in society, the Original Position reinforces Plato’s emphasis on the importance of justice and the common good in governance. Moreover, it provides a peaceful avenue of anonymous rebellion³⁰², and having building upon the Social Contract also promotes utilitarian contrivance.

Of course, while Locke’s Social Contract allows for a more concrete constitution, and the Original Position aids in reinforcing its tangibility, the key issue still with these two strategies is corporeality. For example, even in Plato’s ideal population of 5040, it might be theoretically feasible to organize a space in where everyone could engage in discourse, it might even be practically supported by the demos themselves, still is it physically possible to efficaciously not only get everyone on the same page (Geist), but also - if not in any verbal or written language – communicate said Geist into corporeality?

I have already alluded to ways in which we could practically implement the Original Position, e.g. by way of virtual reality. Transhumanism could potentially allow for the connecting latency load in real time through some kind of Hive Mind application. Ostensibly, an algorithm would then be employed to highlight overlapping ideas and then filter and condense to result in a virtually tangible manifestation of the Sovereign.

Now, this all sounds well and good... but it is still missing the indelible mark we are still seeking: precedence. Rome was not built in a day, and certainly not through action faking – Rome was built brick by brick. So, unless we can show physical precedence - or

³⁰⁰ See Appendix A: Misc., Exhibit 5 for a comparative chart of Hobbes, Locke, and Rousseau’s social contract theories.

³⁰¹ Like the U.K.’s uncodified constitution, a social contract need not be compiled into a single written document. It may exist as a dynamic set of evolving relationships.
See: UK Parliament. “The UK Constitution.” <https://www.parliament.uk/globalassets/documents/commons-committees/political-and-constitutional-reform/The-UK-Constitution.pdf>

³⁰² The Original Position can be imagined as a digital confession booth: a secure, full-dive simulation in which the Veil of Ignorance guarantees anonymity and abstracted identity via temporary memory erasure. Algorithmic ASI may be needed to aggregate these inputs fairly.
See: Weidinger et al., “Using the Veil of Ignorance to Align AI Systems with Principles of Justice,” *PNAS* 120 (18): e2213709120, 2023. <https://doi.org/10.1073/pnas.2213709120>

even a prototype – of what a “Kallipolis brick” looks like, we will never evolve from manifesto to practical treatise.

From Sophia to Techne then, we’ve explored (and reconciled with the introduction of Locke’s Social Contract and Rawls’ Original Position) the strategies of Competent Judges and General Will. Furthermore, we’ve seen how the Social Contract, especially when reinforced by the Original Position are implementable strategies of utilitarian contrivance to the modification of the Neo-Platonic framework. Nonetheless, even reconciling Sophia with Techne is still not enough. Therefore, the culminating step finds us in Phronesis where we no longer concerned with exploring strategies to adapt and armour the Neo-Platonic framework; but, instead, in evaluating precedence proving corporeality by way of precedence. In the following, we will finally address the efficacy of realizability in practice.

Phronesis

Having identified that Sophia and Techne alone are insufficient to back up the superiority of the Neo-Platonic 3rd party framework, we now turn to Phronesis. Endearingly, phronesis culminates as combined Sophia “book smarts” with Techne “street smarts” to become “complete smarts”, or what Nietzsche called “philosophizing with a hammer”³⁰³. Phronesis is the gestalt manifestation of being well versed in not only what the Good is, but also experienced in how to practically implement the Good as well. Drawing parallels then, if, as it relates specifically to the Philosopher King, Sophia is about wisdom, and Techne is about technical capabilities, Phronesis must necessarily be about poiesis – the intersection of practice. As such, this segment will be a reflexive case study on the practices which manifest the Philosopher King.

In line with precedence, we’ve already touched on one of the more famous examples in the Background: Marcus Aurelius. Parallels to Aurelius also exist in different cultures including but not limited to Ashoka – The Buddhist Emperor; Solomon – The Wise King of Israel; Akbar – The Mughal Emperor who embraced diversity; Catherine the Great – The Enlightened Empress of Russia; Queen Christina of Sweden; and Frederick the Great. All of these historical figures displayed elements of Platonic corporeality, but no single figure alone was more than a quasi-Philosopher King.

We saw that based on many of Aurelius’ traits, amongst the other figures, what we now understand as mainly based on legitimacy (i.e. legal inheritance, supported by the demos, bringing prosperity), he came very close to exemplifying what a Philosopher King would look like. Nonetheless, we also recognized that for reasons not entirely known, Aurelius also ended the era of The Five Good Emperors, although he was most certainly not the sole source of blame. It has been speculated that it is precisely due to this kind of information failure which has prevented imitation – let alone emulation – since Aurelius³⁰⁴, and perhaps due to shifting Geists of the times³⁰⁵, not to mention recent turbulence following the events of the previous century, attempts to socially engineer Plato’s original framework were best left with Aurelius as well. Nonetheless, specifically in line with precedence of a

³⁰³ Nietzsche’s metaphor of philosophizing with a hammer implies striking down decayed ideals to make way for life-affirming values. Here, it illustrates that ethical knowledge without action—trapped in ivory towers—is no less impotent than unexamined tradition.

³⁰⁴ By “information failure,” I refer to the systemic breakdown in preserving, transmitting, or institutionalizing the ethical heuristics that enabled competent rulership. The absence of a durable mechanism for leadership succession and moral cultivation may have made Aurelius’ reign unreplicable, not because of personal failings, but due to structural memory loss.

³⁰⁵ Invocations of “Zeitgeist” are intended as descriptive contextual markers rather than determinative explanations. Cultural climates inform the plausibility of ideals but do not render them impossible or invalid.

Neo-Platonic framework (or at least elements wherewithin) one manifestation of the Philosopher King abstraction worth evaluating is the evolution of self-governance³⁰⁶ during the Outer Coast Summer Seminar from June to August 2019, herein dubbed the *Outer Coast Project*, OCSS'19, or simply, *Outer Coast*³⁰⁷.

Like Plato's Kallipolis, the "beta testing" of the Outer Coast Project was an intentional community³⁰⁸. Rather than being based off the Πολιτεία, however, Outer Coast found its inspiration in a different social engineering "experiment": Deep Springs. Deep Springs, Outer Coast, and other institutions of the Nunnian League since all share in one way or another the ethos of L.L. Nunn; for the most part, the thread that ties it all together is self-governance or the idea that "when given great responsibility, young people of good will and talent can rise to the occasion with a courage and a joy they may not previously have known"³⁰⁹. The following analysis of the field study³¹⁰ that was OCSS'19 will thus identify solutions to the issues raised above in trying to adapt a Neo-Platonic framework to be relevant to a modern audience. Specifically, it will necessarily be a comparative analysis of OCSS'19 and the strategies for practically implementing utilitarian contrivance of the good.

The particular aspect of self-governance that we are concerned about at Outer Coast was the politics present during student (SB) body meetings³¹¹ – communal gatherings where all the students (or at least super majority in accordance with constituted ad hoc bylaws³¹²) convened as necessary to organize labour, student budget (treasury), interpersonal relations (mediators), menu items, outreach, etc. In the beginning, SB meetings were quite literally a contained scenario from the beginning of *Lord of the Flies* – that is to say, still polite but very chaotic. As OCSS'19 progressed, however, and more importantly as familiarization occurred not just with experiential politics but with the remainder of the SB as well, order ensued and things became more efficacious. By the end, SB meetings had radically evolved from a direct democracy to a full-blown dictatorship; and, in retrospect, it was this democratic dictatorship where manifestation of the Philosopher King became most apparent. But let us start at the beginning:

In the beginning was the word – and the word was precedence. As Outer Coast was inspired by Deep Springs – in addition to being founded by Deep Springs alumni such as

³⁰⁶ Due to the nature of self-governance, all evidence is retrospective—derived from testimony, journals, and interpretive materials. This introduces the risk of unreliable narration, shaped by memory bias and post hoc interpretation.

³⁰⁷ Outer Coast's philosophy integrates academics, labor/service, and self-governance to engage the student as a whole.

See: Outer Coast, "Educational Philosophy."

https://docs.google.com/document/d/1_FllgmX6Vf67gMST_w_vQ51yqt8j9w1SnUPrt-hEbG4/edit

Also: "Outer Coast Summer Seminar 2019 Advice," YouTube, <https://www.youtube.com/watch?v=aHt-mtITEAY&t=328s>

³⁰⁸ Though not originally designed as a Platonic experiment, this field study is interpreted through a retrospective ethnographic lens—applying grounded theory principles to draw comparative insights from naturally occurring governance dynamics.

³⁰⁹ Deep Springs College. "President's Letter to Applicants." August 20, 2023.

<https://www.deepsprings.edu/wp-content/uploads/PresidentsLetterToPotentialApplicants2023-08-20.pdf>

³¹⁰ While Outer Coast was not designed as an experiment, its constant renewal makes it experientially experimental. Like field studies, it offers real-world insight outside laboratory control—capturing context-dependent behavior in ecologically valid settings.

³¹¹ While no formal minutes exist from Outer Coast SB meetings, notes from Deep Springs provide closely parallel insights.

See: Deep Springs College, "Sample Student Committee Notes." <https://www.deepsprings.edu/wp-content/uploads/Sample-Student-Body-Committee-Meeting-Notes.pdf>

³¹² The 2019 Outer Coast Summer Seminar (OCSS'19) was a kind of *tabula rasa*, where governance emerged organically—at times resembling *Lord of the Flies*, albeit with adult support, exit routes, and far less manslaughter.

Bryden – the SB were not completely left in the dark. Bryden, while hoping self-governance to evolve organically without much outside influence, still recognized the value of precedence and provided a quick (anecdotal) rundown of how SB meetings were structured and the procedures they generally followed. In theory, it was simple: the students would convene at their own discretion to the point of directing the General Will in addressing the most pressing issues at that specific moment in time. In practice, it was a lot more complicated.

The first SB meeting convened sometime after dinner and ended just shortly after sundown³¹³, lasting for about 6 hours in total. It was held in the foyer of Mt. Edgecombe High School and having unaddressed issues on the agenda resulted in adjourning the meeting to the following evening. The general procedure of 22 teenagers in direct democracy involved everyone sitting in a circle and speaking in turn. As one can imagine, by the time $\frac{3}{4}$ of the circle had said their piece, points had been repeated more times than they had been forgotten until order was lost and people were speaking over each other.

The following meeting evolved from a direct democracy to a quasi-representational democracy where in a moment of insight (or frustration), Carmen took center stage with their arm outstretched and turning about, mimicking a spinning bottle. This introduction of a random element maintained relative justice (fairness and increased order) as those not chosen could not speak and were quickly shut down by the others if they tried. This spinning bottle strategy, however, had the flaw of being inefficient and tiring, albeit closer to the General Will as the mechanism behind it (fairness through randomness) was unanimously supported.

Having still not fully addressed all issues on the agenda (which was also due to new issues piling up), the third meeting of the week was scheduled. At this point expedience was starting to shape the General Will and the spinning bottle strategy was formalized into the “Chair”. This was a position in which a swivel chair was placed in the center and the person sitting on the chair, dubbed the “Chair”, would choose the order of procession (i.e. who could and could not speak).

The fourth and final meeting of the first week of OCSS’19 continued with the Chair, ended with a completed agenda, and introduced a concrete “constitution” Social Contract termed “non-negotiables”. There were already non-negotiable bylaws in place, courtesy of ad hoc experience, but evident bias in where the Chair – who had not changed from the previous meeting - showed favouritism to a select group of friends, highlighting the issue of overfamiliarity in conjunction with epistemic problems. The proposed and accepted solution which was later informally inscribed into the Social Contract was that for Justice, the Chair would be chosen via spinning bottle strategy, and for efficacy the Chair would have the right to veto.

After the first week, meetings became shorter and more efficient resulting in less impromptu SB meetings. From the second week onwards meetings were scheduled for Wednesday mornings at 9:15am and Friday nights at 8:45pm. Minor amendments to the Social Contract were made after formalizing the Chair; one such amendment being that the spinning bottle would be spun until it landed on someone who had not previously been Chair. This indirectly invoked the Original Position by way of the golden rule because if the Chair did not act under the Veil of Ignorance (or at least as objectively as they possibly could), then they would undoubtedly face Justice. i.e. if the Chair³¹⁴ allowed a select group of people a

³¹³ In Alaska’s summer, “sundown” refers to nautical twilight, which occurs around 2:00 a.m. during periods without full nightfall.

See: Time and Date. “Sunrise/Sunset for Sitka, USA.” <https://www.timeanddate.com/sun/usa/sitka>

³¹⁴ The Chair is an abstraction; injustice stems not from the Chair itself but from the individual occupying it. As with Plato’s dictum—“no craftsman ever errs” (340e)—we must judge agents, not roles.

disproportionate opportunity of expression, then said Chair would find themselves disenfranchised when the minority group came to eventual power.

From the end of the second week to the fifth week, increasing familiarity (and perhaps weariness) meant that the Chair, when prompting random individuals for discourse, was often met with abstention. The SB had for the most part become competent enough to judge on behalf of each other what was best for one another. In retrospect then, the Chair had thus internalized the General Will to realize its capacity as Sovereign. Of course, this trend was dependent on individualized instances wherein those more sociable were often more efficacious, often simply due to increased experiential competence. This meant the Chair in moments of periodic weakness (incompetence) could defer their decisions to mediating members of the SB which eventually consolidated politics as an oligarchy of Competent Judges: an aristocracy.

Transforming from a direct democracy to an aristocracy then, OCSS'19 was slowly reflecting the transition from Plato's worst form of governance to its best form³¹⁵. The moment which marked the final transition to Philosopher King, however, was not any singular instance but a buildup of social unrest³¹⁶ from outside the SB meetings. Nonetheless, the pivotal moment in which seeds for later democratic dictatorship were sown was in an SB meeting during the penultimate week. In said SB meeting, the tensions that had been building up started to leak into the politics, corrupting the usual order. Voices were raised and people began to talk over each other until you can hear echoes from the first SB meeting resounding from the past. The Chair at the time – Daniel – allowed a couple minutes of chaos to properly gauge the Geist of the General Will before enacting veto power wherein he demanded everyone express themselves in a single sentence. Compiling the twenty sentences together, the Chair silenced the remainder of the SB and offered a proposition. The Chair then went around confirming individually if the proposition was acceptable, and by some luck came to unanimous support. That SB meeting was the beginning of a paradigm shift in which the Chair - fully aligned with the Sovereign - could display absolute democratic spirit through total subjugation of the student body. But, even then, it was not yet a complete manifestation of the Philosopher King because, despite the outcome, it was still a top down (dare I say violent?) enforcement of martial law through veto power.

During the final SB meeting, the final Chair – Victor – attempted to mirror what the previous Chair had done, but this time the outcome was less than favourable. Instead of competently identifying a proposition which garnered unanimous support, the Chair in their individualized instance failed to read the room and only exacerbated the tensions within until the breaking point. An impromptu coup was ignited where supermajority of the SB impeached the current acting Chair, which sent Victor off crying to his room, and reinstated the previous Chair based on meritocratic justifications. In short, Daniel became democratic dictator and realized a complete manifestation of the Philosopher King. And, unlike the previous alignment with the Sovereign, this manifestation became most apparent not because of reasons based in legitimacy (legal, democratic, or pragmatic), and not because it was supported bottom up, but because the first order of decree Daniel the Dictator made as autocrat was to adjourn the meeting so that he could chase after Victor and make sure he was alright. It was this specific display of humanity and attention to well-being over power³¹⁷ which qualified Daniel to be a quasi-Philosopher King.

³¹⁵ Plato, *Republic*, 543a–569c.

³¹⁶ The governing body at Outer Coast comprised 22 teenagers. Social unrest, though structured under the Social Contract, often manifested as symbolic or procedural violence. The use of veto power functioned like state prosecution: an asymmetrical reallocation of collective authority against itself—dissociating the democratic subject from its own agency, echoing the “banality of evil.”

³¹⁷ See Socrates' claim that the best rulers are those least eager to rule (*Republic*, 521b1–2).

In the final comparative analysis of OCSS'19 and the strategies for practically implementing utilitarian contrivance of the good, we see that the evolution of self-governance during SB meetings presents a compelling case study in proving corporeality of a Neo-Platonic framework (or at least elements wherewithin). Due to its reflexive nature, however, there still exist some skepticism of interpretation and doubts surrounding whether the reality of Outer Coast could translate into the wider community and substantiate the magnitude of an entire framework for society. Furthermore, OCSS'19 is, by definition, anecdotal. While I have endeavoured to approach this case study from an evidentiary perspective, it is not necessarily indicative of some larger trend, but this lack of indication also means it can still easily be true while also being unique. By that measure, the main issues that can and should be addressed here are: ability, scalability, and translatability.

For ability, we saw the difference between Victor and Daniel. Earlier we saw foreshadowing of this with discrepancies between Chairs in how well they could internalize the General Will to realize the Chairs' capacity as Sovereign. To reiterate, we saw that the extent to this realizability was dependent on individualized instances wherein those more sociable were often more efficacious, often simply due to increased experiential competence. This difference in ability is what ultimately led to the stratification of the aristocracy from the oligarchy. In hindsight, purely from the standpoint of efficiency, there could have been an amendment to the bylaws wherein Chairs were only chosen from the aristocracy, similar to how Plato only advocated for the Philosopher King to be elected from the Guardian class. However, from the standpoint of efficacy and in keeping with the General Will at the time which favoured equality as opposed to technologism, this meant risking less competent Chairs such as Victor. To clarify then, Victor was given the opportunity to be Chair because of equality, but later lost this opportunity because the General Will shifted to technologism. Whether or not Outer Coast could then act as a template for the wider community would be idiosyncratic and dependant on the General Will and Geist of wherever that community would be located.

As for scalability, OCSS'19 was a self-contained "experiment" of 22 teenagers. From the first SB meetings it was clear that even for 22 teenagers, direct democracy faced difficulties in regards to procedures and operation not only in making sure everyone had equality of opportunity, equity of expression, but also efficacy in relation to time. If direct democracy took 6 hours with only 22 people, then imagine how many days (if not weeks when considering increasing complexities of scale) that a community of 44, 88, let alone Plato's 5040 would take to ever come to a decision. The issue of scalability isn't limited to direct democracy either because, assuming that there is an individual with the ability to not only competently gage the General Will and internalize it, but then also compile feedback and offer a proposal that resonates with everyone, said individual would still be restricted by time. There is only so much computing power a brain can handle, and even less free memory available to sift through the noise and draw everything together. The scalability of the Outer Coast model then would only be scalable to the extent of human limitation.

In regards to translatability, what made manifestation of the Philosopher King possible in the first place was due to camaraderie between the student body. The students felt so comfortable with each other to the point of abstaining from expression in SB meetings because they trusted the Chair knew them well enough. In essence, the students trusted that everyone was competent enough to judge on behalf of each other what was best for one another. In contrast, in today's society many don't even trust strangers enough to talk to them. By this account then, the issue of translatability isn't a problem with the Outer Coast model but a problem with modern ethics; much in the same way, the reason why relationships are so shallow nowadays isn't because people are built differently, but simply because nobody is willing to open up and be vulnerable. Translating the Outer Coast model is thus

nothing to do with operations or procedures, rather, it's all to do with a personal battle against one's fears. Unless the entire community is willing to take the risk that a fellow neighbour won't fall prey to long term abuse of a dictatorship precisely because they, and the rest of the community, are willing to be there for the Sovereign in support.

Coming to the end of Phronesis, we've identified that even non-Guardians (i.e. Marcus Aurelius provided philosophical training) and non-philosophers (i.e. Daniel the Dictator provided an environment for self-governance) can get close to realizing the Philosopher King. This precedence has allowed not only comparative analysis, but also a "Kallipolis brick" from which to evaluate frameworks for practically implementing utilitarian contrivance of the good. That being said, despite this precedence, there are still things which we cannot currently resolve: ability, scalability, translatability. In short, we're on the right track, not there quite yet, but precedence gives confidence that we aren't far off from realizing the Neo-Platonic framework.

Interpretation:

This third and final section *On Efficacy*, revolving around the claim that elitism is the most compatible form of government in relation to the Good, has substantiated why and how a Neo-Platonic framework makes it realizable. Specifically, this has been done through backing up the proclamation from superiority of the Neo-Platonic 3rd party framework as it relates specifically relevant to the Philosopher King. In brief, realizability has been addressed – and found supported - in reference to the efficacy of wisdom, technical capabilities, and poiesis.

From Sophia to Techne to Phronesis, we've armoured the wisest strategies alongside the most practical technical capabilities in implementing utilitarian contrivance of the good before evaluating them in real-world contexts. We substantiated theory at the intersection of practice and then comparatively analyzed their precedence in case studies. All this to prove corporeality, realizability, efficacy, and overall success of a Neo-Platonic 3rd party framework.

Thus, returning to the original research question of this section *On Efficacy*, we see that the answer to what extent would this Neo-Platonic 3rd party framework be superior is now no longer an issue of idealism and structural limitations as both in theory, practice, and real-life application, Plato's modified framework for the Philosopher King and Kallipolis has come to fruition. The ball is now in the court of the demos and if they truly desire the Good life - not in the same way they desire health but continue to eat junk food, but in the same way that their child is dying from cancer and they make searching for a cure their utmost priority – then the extent of realizing it is proportional to their cooperation. Because, unless the demos fully embrace the effort required to capitalize on their limited Ego, the success of the Neo-Platonic framework will simply go the way of many New Year's Resolutions. With this in mind, and having addressed this third and final interpretation, we can now piece it together with the first two to come to a mini conclusion of outcomes and results where we attempt to reconcile the feasibility of elitism by adapting Plato's original framework for the Philosopher King and Kallipolis into a modified Neo-Platonic framework for the Good Life.

Results

Casting our thoughts back to the beginning of the core chapter, we formalized once again a key hypothesis built on the several key questions of: Where does tenability lie on the bridge between theory and practise? How are the demos ever to practically support a

seemingly disenfranchising framework? And to what extent would this Neo-Platonic 3rd party framework be superior?

From the outcomes of each research question, we can now verify the validity of the hypothetical claim: people will be happier and have more welfare when society is elitely structured. In brief:

- (1) Tenability has been unlocked with modern advancements
- (2) Practicality of the (Neo-)Platonic framework is legitimate
- (3) Efficacy is proportional to cooperation

First and foremost, tenability has been unlocked with modern advancement. Previously, Plato's Philosopher King was limited by plausibility, superfluosness, and potential for corruption. The transhumanist movement has allowed for the limitations of subjectively mental endeavors to be overcome just as the steam engine has overcome out movement cap. The role of Philosopher King is no longer incoherent with the potential to form "Ego Horcruxes" and outsource other simulacra to not only be in two *places* at once, but be two *things* at once as well. Furthermore, physical augmentation, both post and preconception, diminishes the irrational capacity, if not erasing it completely. As such, due to the changing landscape (political, technological, and moral), the Philosopher King is more accessible than a simple metaphor, the Πολιτεία more feasible than a manifesto.

Second of all, practicality of the (Neo-)Platonic framework is legitimate. Whether it is legal, democratic, or pragmatic legitimacy, Plato's original framework for the Philosopher King – let alone a Neo-Platonic framework which is more scientifically accurate and socially acceptable – has been shown to be more pro-individual, pro-democratic, and pro-libertarian compared to its counterpart frameworks. Since the ruling entity is ideology (totalization of the Philosopher King role) itself, the demos are not only empowered through meritocratic asymmetric voice, but also protected from (potential tyranny of) themselves. Increased inclusivity and enablement of said demos, in addition to subordination of political instrumentalities to moral and educational ends in the service of individuals further allows for both equality of opportunity and equity of expression. Furthermore, positive feedback loops following the trickle-down effect leads to better performative choices and greater liberty overall.

Last but not least, efficacy is proportional to cooperation. It was identified that while Sophia, Techne, and Phronesis promote efficacy as they relate specifically to the Philosopher King, lack of either element was not indicative of failure to realize the superiority of the Neo-Platonic 3rd party framework as an elitist framework. Integrating the most adaptive framework(s) of competency and Geist made for a more relevant modified (Neo-)Platonic framework. Combining said strategies with technical capabilities and mechanisms allowed for practically implementing utilitarian contrivance of the good. And precedence of non-Guardians and non-philosophers realizing the Philosopher King proved no longer an issue of idealism and structural limitations, but of earnestness.

Taken together, the outcomes of these key research questions validate the success of the Neo-Platonic framework by showing incrementally the feasibility insofar as the cost of social upheaval for a (revolutionary, if need be) paradigm shift from the default duct tape³¹⁸ of modern governance is justified. The three different sections: On Tenability; On Practicality; and On Efficacy, further consolidate the success of the Neo-Platonic framework for the Philosopher King and Kallipolis as an elitist 3rd party external framework (herein

³¹⁸ If it isn't irreparably broken, don't replace it. A total restructuring of government entails enormous short-term costs—rarely justifiable outside of violent revolution.

simply, “Neo-Platonic framework”) framework insofar as it complements the limitations of the Ego. Ultimately, by parameterizing behavioral trends – through outsourcing decision making to Guardians – we see that it becomes possible to turn the Self as a whole towards the Good both consciously and unconsciously.

From this core chapter, we have thus seen that:

1. Changing landscapes (political, technological, and moral) has meant the Neo-Platonic framework is theoretically feasible.
2. The Neo-Platonic framework is not only feasible but more legitimate when compared to counterpart frameworks in modern governance.
3. Non-existent hurdles of idealism and structural limitations in realizing the Neo-Platonic framework highlights the importance of being earnest.

Again, as a disclaimer, the data gathered reflects years of research and study and in my own personal shortcoming I have probably failed to do justice to Plato and the following dialogues that have followed in his footnotes. However, we shall continue to strive towards areté with the insights gained from evaluating frameworks for practically implementing utilitarian contrivance of the good. It is my intention to build from the foundation of my results from this part and the previous to comprehensively answer our current research question and ultimately contribute to realizing the Neo-Platonic framework.

Specifically, then, we see that it is unfeasible that a Philosopher King under our current Sociopolitical climate could ever turn the Self as a whole towards the Good. Yet, as alluded above, this does not exclude the possibility of a potential 3rd party framework from taking up this role instead. This proposition will be the overarching theme for the discussion moving forward as we attempt to reconcile the feasibility of elitism by adapting Plato’s original framework for the Philosopher King and Kallipolis into a modified Neo-Platonic framework for the Good Life.

Synthesis

The Discussion

As per the previous part, we cover in this section the most salient contentions as they relate to the data gathered thus far from the Background, Outcomes, and Results.

- (1) What could be possibly done to reconcile the feasibility of elitism by adapting Plato's original framework for the Philosopher King and Kallipolis into a modified Neo-Platonic framework for the Good Life?
- (2) and we are in a much better position to answer the initial research question of: if Plato's original framework is insufficient, then which adaptive framework(s) would we need be more successful?

Once again, we will also move onto caveats, limitations, and shortcoming before ending with the implications that take all these into account. Discussion will still overlap between not only sections but also chapters as well. Thus, we shall address them under the key variables of:

1. Contention over Background (Elitism + original Philo King) = briefly touch on since does not make up core of argument
2. Contention over Outcomes (Tenability; Practicality; Efficacy)
3. Contention over Results (+discussion, thesis: need Guardians that make decisions for us)

Contentions

Over Background

1. A Transhumanist Metaphor or Jarvis/Alfred Blueprint?

The first contention is with the reading of the Πολιτεία, and whether it should be interpreted as a manifesto or metaphor. Specifically, there is doubt regarding the proto-Aristotelean idea of “fake it till you make it” (1103b 14 -1104b 10) and if the Guardian “training wheels” avoids the problem as opposed to resolving it. Rather than reading the Πολιτεία as a manifesto then where decision making is literally outsourced, perhaps interpreting the Guardians as a way to improve temperance might be more optimal.

This contention primarily arises from the claim that in freely choosing to be subject to the Guardians (432a, 433c-d), the demos are not prey to their lower nature if left to their own devices (no longer under the TPN). On the surface this is a variation of abstinence: I might choose rationally to put myself in a position where I am not tempted... but still not be able to resist temptation. For example, imagine that I cannot resist Doritos, and so I autonomously make it impossible for myself to fall into those situations by listening to my Guardian and removing all Doritos from around me. That does not destroy the problem: it avoids it. This is because by way of outsourcing my decisions, the capacity to train temperance is also lost, i.e. without any Doritos, it's impossible to practice eating smaller portions. So, when at a party when Doritos are once again available to me, I fall into old habits and devour them all.

The question we should ask ourselves though is that, given what we know about self-control from the previous part, do we actually believe that portion control is possible with all

the investments into the Dorito effect? Moreover, are we aiming for portion control, or a complete break-up?

The idea of everything in moderation is misguided. Nobody would recommend having a toxic relationship in moderation. Much in the same way, no amount of smoking or junk food is good for you. One of the key benefits of Plato's original framework for the Philosopher King and Kallipolis is that it is all encompassing, so rather than dealing with Doritos on a micro level (abstention one at a time), it deals with Doritos – in addition to all junk food and inferior goods/services – on a macro level (nationwide abstention). This ideal of total eradication is realizable specifically because of the extensiveness of the framework and ostensibly works as well – if Doritos no longer exist, how will the temptation? Is that enough for virtue? Or is it not virtue unless it is tested and tempered by temptation?

An alternative interpretation, however, is that granted Plato's technologism is oriented towards turning the Self as a whole towards the Good, why then must philosophers necessarily become kings? Why is it that the Philosopher King must necessarily be a Master instead of slave; doesn't the framework work just fine as an assistant like Jarvis, Alfred, or non-ASI entity?

This alternative interpretation of the Πολιτεία metaphor for harmony and Justice, or even manifesto for transhumanism suggests that a power imbalance of (and social stratification leading to/from) the Philosopher King opens potential avenues for abuse of power and coercion, i.e. enforced vaccine mandates. Popper once argued that absolute power corrupts absolutely, and while we have already addressed the issue of potential for corruption – or lack thereof – it is nonetheless correct to identify that philosophers need not be kings. However, does power come from being king, or is it because one has power that they are titled king?

We should remember that there are many forms of elitism: Dharmic ruler; Sikh caliphate, Junzi, etc. I didn't choose Plato's original framework as the model basis generic template for modification because the philosopher was a king (refer to method), and by the end of this part of the thesis, it should be clear that the modified Neo-Platonic framework as a modification of many other adaptive frameworks, strategies and mechanisms is now as close to any elitism framework as it is far away. The key point, however, is not that the philosopher need be king – this was only a means to an end anyway (Book V, 473d) – neither Master nor slave, but can be anything so long as it is efficacious (akin to Invisible Hand / Laozi 17 / Dorito effect) in its original technologism. But returning to the earlier question then, insofar as the Philosopher King makes use of strategies and mechanisms akin to the Dorito effect, are we sure that little advertisement parent companies³¹⁹ don't actually really rule the world anyway?

From this contention then, we see that the Guardians and Philosopher King being in charge isn't about vanity or desire for power (remember, the best ruler is the one who doesn't want to be in charge; 347a-c). Philosophers don't necessarily have to become kings, instead, philosophers naturally become kings because of their competence and power that they acquire.

2. Do the Benefits Justify the Costs?

A second contention is with whether the benefits justify the costs. The first instantiation of this interpretation is personified by what Meinecke calls the "German

³¹⁹ Good Business. "Food Brand Ownership." Accessed March 1, 2024. <https://www.good.is/Business/food-brands-owners-rp>.

Catastrophe”³²⁰ in where the worst events (i.e. Holocaust) arose in quite advanced societies – one that neared the ideal of the Kallipolis insofar as they had the best education, highest social mobility and tolerance (intermarriage rates between social and racial classes), and being at forefront of technological and scientific advance. Yet despite this advancement, some of worst events occurred. So what went wrong? And is the Kallipolis also susceptible to this downfall as per excessive technologism and ethical uncoupling?

We can answer this through a brief look at Nazi Germany. Originally it was believed that peculiarities of German society (1. alienated from Versailles negotiations; 2. "Dolchstoßlegende"; 3. demand for reparations when bankrupt) were the cause. Later, however, it was identified that Nazi Germany could not have occurred without Hitler who was necessarily charismatic at "leading the witness"³²¹. Furthermore, regardless of historical materialism, three forces together: prosperity; propaganda; and coercion was all it took (and takes) to subjugate even the most advanced society.

Unfortunately, yes – all three forces are present (or at least potentially so) in the Kallipolis which makes it susceptible to the German catastrophe. However, as we have explored – and will further elaborate in a later contention – there are checks and balances in place (i.e. intrapersonal feedback) which safeguard against excessive technologism. The key problem is the ethical uncoupling wherein assuming the human good is incompatible with the objective Good (i.e. genocide is best for the moral order), the Philosopher King might be tempted to enact the most rational output. But this needn't be a concern because when the Philosopher King finally does manifest, it would mean a realization of Kant's *Reich Der Zwecke*, and as such genocide would actually be wanted by all if it truly was best for the moral order (more on this in the following part); moreover, the Philosopher King could never be ethically uncoupled insofar as it aligns with the Sovereign.

Another cost benefit analysis contends with the tradeoff of social upheaval that potentially results in bloody revolution. History is knowingly written by the victors resulting in a "might makes right" account of justice for the heavy short-term losses that are incurred. But objectively speaking, even if the long-term benefits of the Kallipolis can justify the short-term costs of its transition, will the revolution truly be "the war to end all wars" or another end-of-history illusion?

Ideally, the transition to the Kallipolis will be peaceful, but if it does end in a bloody revolution, can we really say maintaining the status quo is any better? Moreover, who is to say that a bloody revolution for a different cause will not eventually end up occurring? In fact, it would be stranger if it did not occur because, as Plato's Law of Decay reminds us, **change is nature**. Those that stagnate with the status quo eventually die so it's not really a question of if death will happen, but when it will happen and in what manner. And ultimately, "whether 'tis nobler in the mind to suffer. The slings and arrows of outrageous fortune, Or to take arms against a sea of troubles", that is the question.

In regards to the former interpretation, we see that excessive technologism is safeguarded against through checks and balances. Furthermore, ethical uncoupling is neither a problem as it is theoretically impossible, and even if it was, the only scenario in which it would be would be superiorly utilitarian. In regards to the latter interpretation, if the transition to the Kallipolis somehow does not end up peaceful and instead incurs a bloody revolution, we see that it is still a better alternative.

³²⁰ Friedrich Meinecke, *The German Catastrophe: Reflections and Recollections* (Cambridge, MA: Harvard University Press, 1950)

³²¹ Jordan Peterson. "The German Problem." YouTube video, 17:05. September 2021.
<https://youtu.be/kbWLE07NGcw?t=525>.

3. *The Price for Freedom*

A third contention is with the freedom of the demos, not so much in regards to whether freedom is slavery, but rather if the demos are *really* willing to support a paternalist framework. Rationally, it makes sense for the demos to subjugate themselves to the Guardians instead of the Drifting Will, yet it's easy to overlook that humans have rational capacities, but are overwhelmingly creatures of sentiment. Much in the same way, it's easy to rationalize phobias as fictitious, but the reality of convincing ourselves is a different matter. The innate desire for the freedom to do what one desires, even when it does one harm is difficult to overcome, and it is precisely because most modern frameworks are *laissez faire*, non-paternalistic closest to Nature forms of government which give them that primordial allure. There is thus a formal contradiction between those frameworks closest to Nature and those that are not.

Nature confers on the individual more natural rights than power with which to defend them: therefore it is necessary to institute governments in order to preserve artificially the harmony of rights³²². Nature gives the individual more natural needs than power with which to satisfy them; but in exchange individuals find means to satisfy their natural needs without resorting to any constraint, without any sacrifice of interest. Thus the state of Nature (and anarchy) is closest to absolute freedom, it is also closest to absolute hardship. It was precisely because of said hardship that made life “nasty, brutish, and short” and ultimately what conferred Aristotle to believe man is a political animal. Paternalism essentially incites an adverse reaction, but it is because people are not only inclined towards freedom but inclined towards politics as well that the demos will really be willing to support a paternalistic framework.

Alternatively is the rendition that those less inclined towards the political sphere (i.e. more inclined towards *laissez faire* than paternalism) will need to be coerced or manipulated. The contention here is thus with the ethics of Machiavellian tactics: is it ethical for Philosopher Kings to appeal to demagoguery at first to win over the masses? Is it ethical to indoctrinate and gaslight with the Noble Lie? Is it ethical to wage war (i.e. against irrationality) to stir up the irrational fervor of the masses and distract them from civil discontent? Does stooping to these tactics to win over the demos' support not make the Philosopher King as bad as its counterparts?

We should remember that Plato advocates for technologism. Appealing to demagoguery is simply using the tools of rhetoric – which are necessary to frame reason – and it is better for a leader that is supported and competent than populist alone. Insofar as propaganda, amongst other tactics; the Noble Lie did not mean in general the inculcation of basic falsehood, but the mythical presentation of truth; in particular, moral values were not to be falsified or distorted. Regarding the use of war, which is ‘the faro-table of government’³²³ and is simply a tool of transition: there will always be war whether it is against people, drugs, or irrationality; still better to wage war against a real problem than a fake enemy (i.e. Oceania, Eastasia, and Eurasia).

In reality then, garnering freely expressed support for the transition to the framework for the Philosopher King framework has shown not only to be the most rational choice, but one that is also the most realistic and human.

³²² Elie Halévy, *The Growth of Philosophic Radicalism*, trans. Mary Morris (London: Faber & Faber, 1972), 190–91.

³²³ Ibid.

4. *The Tradeoff of Autonomy*

Following from the contention from the previous part – aiming for the Good life when the complete Form of the Good is seemingly unknowable – the fourth contention recalls a variation of Voltaire’s original question³²⁴: how are Philosopher Kings even to initially become qualified as Competent Judges?

We alluded to the idea that since the Philosopher King in the Neo-Platonic framework necessarily aligns with the Sovereign (a singular multipolar entity), the issue of overfamiliarity in conjunction with epistemic problems can be resolved through (a) the mechanism of the Original Position; and (b) removing the need for interpersonal comparability by way of intrapersonal feedback. In reality, the mechanisms of the Original Position means - in keeping with the idea that the Good Life necessarily commensurates with happiness – rationalizing what leads to objective happiness. Furthermore, the safeguard of intrapersonal feedback acts as a check and balance for the formalization of Kant’s universal law maxim. Simply, what qualifies the Philosopher King as competent enough to judge on behalf of each other what is best for one another is the ability to implement the course of action which would make everybody happy, i.e. what Daniel the Dictator did during OCSS’19.

An alternative interpretation of this contention, and arguably more fundamental problem, is the assumption that autonomy has no value of its own: isn't there a value in directing one's own life in accord with one's own scheme of value?

While autonomy may not be sufficient for a good life, the contention is that it is an unalienable aspect of the Good Life. The original answer we gave above is that they have autonomy insofar as they have autonomously given up their autonomy, i.e. monks who make serious – autonomous – vows of obedience. Furthermore, that which is truly liberating is the willful (non-sacrificial, non-coercive) enslavement to the Platonic framework. Moreover, in line with the previous part: we are slaves anyway, so instead of allowing the Drifting Will to fall prey to their lower nature if left to their own devices (no longer under the TPN), is it not better to choose the Guardians as their Charon?

Incidentally, following from the idea of ‘true freedom as autonomous enslavement’ are contentions surrounding the reality of freedom. These contentions largely echoes what was said in the thought experiment from previous part about being conditioned to like chocolate ice cream, but here will differ slightly as it responds to the idea of freedom having any intrinsic value³²⁵. On that note, it is dubious as to whether freedom even has any real valuation or if the valuation stems from its potential capacity – after all, free pursuit and anticipation has been shown to be more of a reward than the event itself³²⁶. Skepticism in the wake of the post-truth era has also shifted trends as to the valuation of the real; specifically

³²⁴ The main flaw in Plato’s argument, which renders it highly unpersuasive, is what Voltaire defined as a “‘benevolent dictatorship,’ where an enlightened despot, without the need to consult people, would nevertheless govern in their interests” Jonathan Wolff, *An Introduction to Political Philosophy*, 2nd ed. (Oxford: Oxford University Press, 2006), 62.

³²⁵ “Man does not want what is disadvantageous to him, but man desires freedom more than happiness: the ability to do what one desires, even when it does one harm.” Fyodor Dostoevsky. “Notes from the Underground.” Excerpt discussed in YouTube video, timestamp 4:58.
<https://www.youtube.com/watch?v=sR4GgRBxrbk..>

³²⁶ *Damien Talks Money*. “You’re spending money the wrong Way..” YouTube video, timestamp 4:09.
<https://www.youtube.com/watch?v=840iXuXwE4Y>.

following evidence that memories are largely fictitious³²⁷, evidence pointing towards the universe being a hologram/simulation³²⁸, and evidence that soft despotism has invalidated any real legitimacy of control³²⁹, attitudes towards veridical (virtual) experiences as opposed to “real-life” ones are shifting towards the former³³⁰.

A more revised answer, however, is that the Guardians don’t explicitly tell the demos what to do, but nudge the more favourable option³³¹. i.e, advertisements amongst other vectors of Subconscious Forces and Psychic Noise don’t tell you explicitly (a call to action does not count because it is rhetoric) that you *must* buy Doritos, but it does frame Doritos in a favourable light. It is conceived then that under the Neo-Platonic framework, rather than having Guardians patrol like some kind of secret police, the Philosopher King as an Invisible Hand / Laozi 17 framework³³² will integrate itself into popular culture and potentially through a “Guardian” app, algorithm, or ASI would offer advice. More on the specifics, models, and policies will be addressed in the next part: Praxis.

5. Spontaneity and Predeterminism

A fifth (or 4.2) contention is with the value of spontaneity. If the Guardians necessarily know what actions are the best, then they must be able to foresee to some degree the consequences of those actions. In other words, when Guardians implicitly suggest a course of action, they aren’t merely suggesting the following step but are selling an entire playbook which can be equated with a predetermined life. This seems incongruent with Plato honouring free and spontaneous performance.

Again, reminding ourselves that Platonic frameworks advocate for technologism, the telos is the Good Life, and by association happiness. The value of spontaneity then is only insofar as it relates to the Good Life and it would seem that in reality excess spontaneity does not commensurate with happiness as neatly as it would appear. Specifically, excess spontaneity where the freedom to choose, then change, and then make up your mind again is

³²⁷ It is not even clear whether our memories are real. Studies show people frequently recall false memories, including childhood abuse that never occurred.

See: Elizabeth F. Loftus, “Planting Misinformation in the Human Mind: A 30-Year Investigation of the Malleability of Memory,” *Learning & Memory* 17, no. 7 (2010): 361–66.

<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3069643/>

³²⁸ Nick Bostrom, “Are We Living in a Computer Simulation?” *The Philosophical Quarterly* 53, no. 211 (2003): 243–255. <https://doi.org/10.1111/1467-9213.00309>

³²⁹ In the face of simulated social contracts, we often succumb to a defeatist mindset. Bureaucratic systems representing the collective “us” become illusions, and our democratic power erodes into self-imposed powerlessness. This mindset—like the silence depicted in “First they came...”—unwittingly empowers the very forces we hoped to resist.

Martin Niemöller. “First They Came.” In *Holocaust Memorial Museum*.

<https://encyclopedia.ushmm.org/content/en/article/martin-niemoeller-first-they-came>

³³⁰ Weijers, Dan, and Richard E. Joyce. “Nozick’s Experience Machine: An Empirical Study.” *Philosophical Psychology* 31, no. 5 (2018): 683–703. <https://doi.org/10.1080/09515089.2017.1406600>

³³¹ Nudge theory suggests that altering the presentation of choices can influence behavior without limiting freedom. For instance, opt-out systems (e.g., for organ donation or 401(k)s) increase participation by capitalizing on inertia.

Richard H. Thaler and Cass R. Sunstein, *Nudge: Improving Decisions About Health, Wealth, and Happiness* (New York: Penguin, 2008).

³³² The “invisible hand” (Smith, *Wealth of Nations*) describes how individual self-interest can unintentionally benefit society, while Laozi’s Chapter 17 depicts the ideal ruler as one whose presence is barely felt. The Philosopher King, reframed through this lens, becomes a non-coercive guide: directing without dominating, and ruling such that the people feel self-governed.

Smith, Adam. *The Wealth of Nations*. Edited by Edwin Cannan. New York: Modern Library, 1937. Laozi. *Daodejing*. Translated by D.C. Lau. London: Penguin Books, 1963.

the enemy of happiness³³³. This mirrors the earlier proof where excess – or absolute – freedom actually makes you a slave to freedom instead of in servitude to the Good. Spontaneity then is valued to a point, but only insofar as food is valued to the point of fullness instead of incontinence.

A second interpretation of the Good being tied to spontaneity is that the Good is locked in with specific actions as opposed to character., i.e. a “bad” person whose heart is filled with malice could still be “good” if they did all the right actions. In the same way, technologism – and consequentialism in general – would approve of a murderer if the victim was, for example, Hitler. So, the interpretation of this contention can be analogized with: is the action of murder bad inherently, or can the action be “good” if murder is committed by a soldier in protecting the Kallipolis?

Since this contention is mostly against consequentialism (and deontology to a point) from a virtue ethics perspective, the direct answer goes beyond the scope of this thesis. However, considering the context of the Neo-Platonic framework, we can ostensibly resolve this contention by appealing to the idea that moral validity is voided when predetermined actions are made unawares. i.e. the weight of moral blame is different for a murder committed accidentally and a murder that wasn’t necessarily planned, but foreseen in advance; aka: the perpetrator was warned that if they walked down a specific alley then they would become an accessory to murder. Circumvention of this contention then is found in the idea that the demos, in being influenced by the Guardians, necessarily outsource all decision making in addition to information sourcing for the making of said decisions. In the same way, when someone purchases Doritos, part of the transaction includes outsourcing moral responsibility (that the product was ethically sourced and externalities of consumption are not negative) onto PepsiCo. To draw another analogy, just because the plot³³⁴ of a book is already written and the conclusion already set doesn’t deter the reader from reading. We live our lives accepting the risk that our decisions could lead to “bad” actions or we could turn into “bad” people, but this risk doesn’t mean life is not worth living. Moreover, under the Neo-Platonic framework, this risk and moral responsibility necessarily falls onto the Guardians which in itself makes for a further contention.

Clearly, too much spontaneity then – and subsequent anxiety of looking too deeply into its consequences – doesn’t help anyone. Rather, the demos should not worry about the mechanisms that surround their liberty but simply be themselves while open to suggestive nudges. After all, this is not only less stressful for them, it maximizes liberty, and absolves them of moral responsibility so they can simply focus on being their best selves.

6. *Moral Responsibility*

The sixth (or 4.3) contention still follows on from the previous one and finds itself in doubt over the reality of outsourcing even moral responsibility. Specifically, giving Guardians rights over actions and decisions necessarily means said Guardians must also take moral responsibility if they fail, i.e., assume the action a Guardian nudges causes one to lose their keys → lost keys leads to failing to show up to school → failing to show up to school leads to... bad grades... no job... no money... no food... death. Of course, this example incites the slippery sloped logical fallacy of snowballing and a Guardian will likely intervene, but who is to say there won’t be further negative externalities from these adjustments³³⁵?

³³³ Dan Gilbert, “The Surprising Science of Happiness,” TED Talk (2004), 13:20–14:06.
https://www.ted.com/talks/dan_gilbert_the_surprising_science_of_happiness

³³⁴ Lee, Yone. *Surviving Romans*. Webtoon, episode 51. Accessed September 30, 2023.

³³⁵ George Nolfi et al. *The Adjustment Bureau*. Universal Studios Home Entertainment, 2011.

Moreover, the issue of moral responsibility still exists. The ambiguity of who is to shoulder moral blame is in the same way present in how a company (as a legal entity) takes responsibility. Does the CEO (i.e., acting philosopher king) resign? But in this case, doesn't this lead to the dissolution of the entire society? Is the society as a whole to blame then? Does this in itself not infringe on the punitive injustice of collective punishment?

The answer is surprisingly simply and can actually be found in the second contention: the World is subject to the Law of Decay regardless of who is in charge, who takes moral responsibility, and who is ultimately to blame. Actually, this question is misguided. Regardless of what actions occur, consequences will necessarily follow. Moreover, in the wisdom of the Philosopher King, are their decisions not damage mitigating? Of course, in the process of developing the Kallipolis, human error could occur, but from the section On Practicality, has this risk not been substantiated and evaluated? Under the Platonic frameworks, justice is giving unto the other what they deserve and blame for an imperfection is nothing anyone deserves inasmuch as nobody should be forced to shoulder the burden of a natural disaster alone. We should remember that the origin of the Platonic framework is not a perfect framework, only a better one – which is hopefully organic and ever evolving upwards. Busying ourselves with who is to blame and what reparations should be in place then is simply an effort towards distraction and made away from movement along the axis of the Good.

Another interpretation of this contention is in regards to the relationship between (moral) responsibility and liberty. Considering that responsibility necessarily precedes rights, without responsibility, the demos are purportedly no longer grounded. This interpretation is not necessarily true: indeed, outsourcing moral responsibility alongside the moral right to make decisions causes individuals to lose some mechanisms of grounding, but societal pressure still exists and is arguably a more complete lodestone.

In today's age we are consistently anxious if we are living up to some arbitrary set of expectations; am I doing enough work as a student; is my output enough as an employee; is my child developing enough from my input as a parent, etc. There are so many gaps of ambiguity, but returning to democratic legitimacy – particularly what was said about Organic Theory of State – we see that outsourcing moral responsibility as communal responsibility results in complementary overlaps with gestalt benefits that individual grounding could not measure up to. In other words, by swapping individual responsibility for communal responsibility, expectations become clearer and objective ways of affirmation do too.

Outsourcing moral responsibility to the Guardians is thus superior from a legitimate standpoint and offers more benefits than costs – if a trade-off can even be considered an actual cost on behalf of the demos instead of an upgrade.

Over Results (interpretation)

7. Subject-objectification and Object-subjectification

The fourth contention elaborated on how the Philosopher King, without the need to consult people, would nevertheless govern in their interests. One point of justification was from competency in knowing what objectively makes people happy. But is anything actually objective? And what, if anything identifies the objective Good from the utilitarian (greatest happiness for greatest number) Good? Is it possible that what objectively makes people happy is not objectively best across the board.

The answer to this seventh contention lies in the earlier explanation linking subject-objectification and object-subjectification. On one hand it is assumed that the Form of the Good is subjectively objective, that is, it doesn't matter who is viewing which shadow at

what time or place, said shadows can all be traced back to the objective Form of the Good or the liminal space of its instantiation. On the other hand, it is assumed that the objective Form of the Good is subjectively interacted with, that is if art is objectively an instance that award higher pleasure, not everyone necessarily derives the same magnitude of pleasure (and in some instances, may even derive a larger quantity from a lower pleasure) across the board.

For the first interpretation of this contention, we will focus on subject-objectification – shadows being able to be traced back – and how it relates to utilitarianism. Specifically, looking back at OCSS’19, we see that even Daniel the Dictator’s proposal might not have been objectively highest for everyone. Clearly, Daniel’s reinstatement (or perhaps the coup in general) did not sit well with Victor. Therefore, while the event might have been objectively the best manifestation of the Good as restrained by happenstance at the time (i.e. without external intervention the situation might not have played out better), it still wasn’t the unanimous ideal sought after by a perfected, complete Philosopher King. In reality then, due to circumstantial limitations, or at least under the Neo-Platonic framework as it is currently, (manifestations of) the objective Good is synonymous for utilitarian (greatest happiness for greatest number) and not objectively best across the board. Nonetheless, future Victors can feel relieved in knowing that since Daniel the Dictator was neither Guardian nor philosopher, of course the best he could do was not to be expected the best objectively. In the real framework, as aforementioned, checks and balances exist so there won’t even be any Victors to be potentially “victimized”.

Alternatively, different perception of good relates to different classes. This relates to the second interpretation which is object-subjectification and links back to the idea that each person is endowed with an ineluctable nature. Precedence for this can be seen in the end of the 20th and beginning of the 21st centuries, where we find that the educational norms have substantially changed. Not only are most Americans (for example) completing primary and secondary school, but it is all but *de rigueur* to have a college-level education to pursue any job outside of retail and manual labour—a consequence of the advance of technology³³⁶. Through secondary education and in at least traditional colleges, students are exposed to a wide range of “higher” pleasures, including art, music, history, great literature, languages, mathematics and science, and the opportunity, at least in college, to try others, including philosophy. And yet, the vast majority of college graduates cheerfully walk away from graduate school, independent research, or even keeping up with their already acquired math abilities once they leave school. At face value this might seem like counter evidence, but rather than finding this disturbing, we instead find that this perfectly evinces with Plato’s ideas each person is endowed with an ineluctable nature which suits them for a determinate role in the polis.

For further distinction, it is not that there is no objective good, rather it is that not everyone can see it. i.e. blind people cannot see colour, deaf people cannot hear music, etc. but blind people can still perceive light, deaf people can still feel music. So the lower class, while maybe not find pleasure in art, can still resonate with the goodness in their soul. As a weak analogy, it is just as people resonate health benefits of vegetables but still feel more attracted to junk food.

As such, subject-objectification and object-subjectification substantiate for the most part an objective Goodness for everyone. Moreover, it highlights areas for improvement sought after by a perfected, complete Philosopher King which will be later addressed.

³³⁶ Boesel D., & Fredland E. (1999). *College for all? Is there too much emphasis on getting a 4-year college degree?* Research synthesis. Washington, DC: National Library of Education, Office of Educational Research and Improvement, U.S. Department of Education, 1999. (ED 431 986).

8. *Meaningful Consent*

The eight contention surrounds the rendition that the freest thing we could possibly do is choose which masters we enslave ourselves to, not so much in regards to enabling autonomy, but rather in regards to hypothetical of retroactively withdrawing consent. In other words, this contention posits that our choice to outsource our decisions to the Guardians might not be consistent throughout time. For example, specific instances might cause the same individual at age 16 and age 32 to be more or less inclined towards subjugating themselves to the Guardians; however, if at age 32 this inclination would lead to refusal, the fact that at age 16 the Social Contract was already willingly signed means said individual would never become or have the opportunity to grow into the 32-year-old reactionary. This contention could thus also be viewed as the Guardians robbing a potential future version of an individual. But is this not the entire point? To parameterize the Drifting Will, like braces on teeth, is to make sure that individuals grow up to be their best selves instead of a random future version which was predestined by the rivulet of Fate.

Returning to the idea of consent, however, we see that the issue of retroactively withdrawing consent and false reporting has muddled the waters when it comes to volitional decisions which must be accountable for following consequences. This is important when considering subjugating to the Guardians must be done of one's free will and cannot later be withdrawn, especially if there are certain clauses, for example, that erases memory of signing the Social Contract for it to enact without hiccups.

Entertaining the idea of volitional signing of the Social Contract with a clause for memory erasure of the signing event, this can be justified by the fact that a person cannot truly feel at ease and be expressive of their authentic liberty if they know all their behaviours are being nudged. When consenting to the Social Contract wherein they outsource their decisions to the Guardians, they ostensibly agree to forget the experience ever happened, and thus forget that they have ever consented. If later they would have regretted consenting if their memories were intact, then the issue of withdrawing consent but not having the memory of ever consenting in the first place becomes a very large problem which currently I do not know how to address. Nonetheless drawing from the previous part that autonomy is objectively and independently rational, then it would seem that regardless of change, the constant of autonomy should persist through time. This means it doesn't matter whom makes the choice or when it is made, the rational option will always remain the same. i.e., assuming a dichotomy, is it better to stay up late scrolling social media or getting adequate sleep? The rational – although not necessarily always best – choice is obvious; however, this leads us to the second interpretation of this contention: whether the best decision is ever the one made autonomously instead of rationally.

There are alleged cases where a mother would prefer a handmade birthday card from their child rather than one they bought with their own money because while the time investment might be the same (time to make the card VS time to earn money to buy the card), the handmade card prepared autonomously despite being objectively uglier than the store-bought card has an added sentimental value. Another way to perceive this interpretation is considering if any sentimental (human) value is lost by transacting through the Social Contract.

This second interpretation is very closely aligned with the fourth contention if the alleged case is held analogously. Here, however, we shall claim that the alleged case of sentimental value is actually a disanalogy. This is because it introduces a false dichotomy between autonomy and rationality which, having made allusions to Kant previously, does not exist. In keeping with Kant's own presentation, we observed that the Free Will (see previous part) insofar as it possesses autonomy will naturally align itself with rationality; autonomy is

inextricably linked with autonomy. Of course, relating back to the third contention, we saw those limitations of the human condition - insofar as humans have rational capacities, but are overwhelmingly creatures of sentiment – means, in reality, autonomy does not play out with exception of the TPN, but even then it is in rare instances uninfluenced by Subconscious Forces and efficaciously make well informed, competent decisions in the background of Psychic Noise. Thus, authentic and autonomous representation of the Ego is thus rational whereby the best decision is necessarily rationally autonomous.

Whether or not an individual can meaningfully consent their entire futures away to the Guardians under the Social Contract is not a matter of sentimentality, maturity, or ability, but one of rationality. Insofar as the individual has the capacity for rationality, and insofar as the rational decision is the best decision, then it stands to reason that any conscious individual whether they are 5 years old or 95 years old can express meaningful consent.

9. The Reality of the Neo-Platonic Framework

In the results we conjectured that the extent of realizing the Neo-Platonic framework is proportional to the cooperation of the demos. Furthermore, in the third contention we posited methods (e.g., Noble Lies, War, etc.) to stimulate earnestness. Nonetheless, although there really aren't any hurdles of idealism and structural limitations in place, there still exists the general hurdle of inertia. This is because despite a will for cooperation and methods to stimulate earnestness, intentions are insufficient in bringing ideas into reality. What is needed is a strategy for actionable steps, and the first rendition of this strategy is identifying whom, or, better yet, what, is capable of being the inaugural Philosopher King. After all, the previous part has shown us that leaving it up to the people to organically start cooperating is not only wishful thinking if there is nothing to nudge it into motion.

It's already been mentioned that one of the core reasons for the untenability of Plato's original framework was because it was read from an anthropocentric perspective. But, as we've advanced since Plato's time and expanded our understanding of rationality, sentience and what "humanity" refers to, we saw that the Philosopher King need not be a human – or at the very least, a pure, unaugmented (transhumanist) human. Viable candidates for the Philosopher King might not exist as of yet – or at least not within our locality – but that does not mean we are not on the path to cultivating one. After all, we have already exhibited precedence on this path and there is growing literature surrounding transhumanists and other technologically Advanceds. In fact, rather than perceiving the lack of completely viable candidates in the present moment as opposed to those that are quasi-qualified, perhaps being at the cusp and being able to look back in hindsight from a prediction into the future is what necessarily secures proper advancement and preliminary safeguards against perverted iterations. In other words, once the inaugural Philosopher King comes to power, its precedence will undoubtedly play a role in influencing further iterations. To make sure that the Philosopher King is the best it can possibly be then, we should cherish the dawn before its arrival to incite discourse into its ethical development.

The second rendition comes from questioning the intermediary steps between now and the realization of the inaugural Philosopher King. Like most things, these steps can be planned backwards in the way a non-runner to running a full marathon requires a training regimen. Obviously, those with running experience will have an advantage, therefore, the inaugural philosopher king must have a minimum requirement of being a quasi-Philosopher King to begin with. From this quasi-Philosopher King base point then, how do they evolve to being a full-fledged Philosopher King. One way is through a virtual training regimen wherein a simulation of the real world can be used as a testing grounds. How do we get to this simulation? Through advancements in full dive technology. How do we support these

advancements? Through more expertise in the field. How do we publicize these areas for novel degree education? Through information semination, as in one of the primary purposes of this thesis.

We see that when broken down into micro steps like these, the advent of the Philosopher King – as a non-pure, unaugmented (transhumanist) human – isn't as daunting or overwhelming as initially imagined. Moving from Theoria to Poiesis; from Sophia, Techne, and Phronesis; from the framework (the model, the blueprints, and the prototypes) to success, the third and final part on Praxis will elaborate further on not only what is meant by parameterizing in real-world contexts, but propositional policies to enact them as well. i.e. administrative/legislative algorithmic strategies for ASI.

In our final contention then, we have addressed that the extent of realizing the Neo-Platonic framework is proportional to the cooperation of the demo such as being supportive – if not providing direct discourse - about the Philosopher Kings' ethical development and furthering – if not directly by working in the field – educational opportunities that correspond with technological advancement. Note that this is different from accelerationism and, in fact, proposes the act opposite: a cautionary purview in moving forwards. Realizing the Neo-Platonic framework requires a team effort and ultimately community support and cooperation from all sectors directly and indirectly involved. The realization cannot – and should not – be alienated to ivory towers because when the advent of the Philosopher King does come around, everyone will necessarily be involved and affected.

Caveats

Moving onto the caveats, I now present warnings or cautions regarding the use or interpretation of information, findings, and/or recommendations.

First and foremost is the need to acknowledge the many assumptions that have been made, not just in my own research but in the research that my research is based upon. These assumptions include, but are not limited to:

1. Elitism –
 - v. That institutional memory (i.e., literature tests) does not muddy prestige
 - vi. That Platonic frameworks can be read as elitist manifestos
 - vii. That meritocracy is just substantiation
 - viii. That it is necessary (cannot be divorced from) for the Good Life
2. Tenability –
 - v. That it encompasses and addresses all, but the most minor, theoretical wrinkles
 - vi. That hypotheticals and thought experiments are valid justification
 - vii. That
 - viii. That it is
3. Practicality –
 - v. That legal, democratic, and pragmatic legitimacy are sufficient to garner support from the demos
 - vi. That a volitional decision at one point is held valid across time
 - vii. That negative externalities following decisions made by Guardians are mitigated
 - viii. That freedom is slavery
4. Efficacy –
 - v. That the Form of the Good can be known, even if indirectly through triangulation
 - vi. That precedence of parts is indicative of precedence of the whole
 - vii. That large scale cooperation is structurally feasible
 - viii. That wisdom and techne are sufficient for phronesis

5. Etc.

As always, this list is far from comprehensive but invalidating one or two assumptions here or there do not invalidate this part – let alone this thesis – in its entirety.

Second, one caveat that people should be warned about interchangeability of words such as state and polis (as well as capitalization of State vs. state) in addition to people, demos, and citizens³³⁷. For the most part, as I am trying to quote directly with minimal jerrymandering of the literature, I have not adapted terminology to the most appropriate consensus. For example, capitalization of State as opposed to state is not only of linguistic function (i.e. to differentiate State as in polis, vs state as in of the Soul), but also to accentuate context specific jargon (i.e. not all citizens were demos and not all people were considered citizens). In my own writing I have endeavored to maintain the distinction, but for sake of proper research ethics as well as general neatness, I have not highlighted each instance of inconsistency and the burden of differentiating between interchangeable is largely up to the reader.

Third, to reiterate: my stances have necessarily coloured my interpretations. For information on this, the reader is redirected back to the subsection on methodology.

Fourth, a caveat in particular to Outer Coast case study in that since OCSS'19; while socially engineered to a degree, OC was never intended to be a Platonic experiment and therefore does not fit the model of a traditional case study. This was already forewarned in the in the Aim and Scope where from the start it was noted that there would be limitations to the research. Furthermore, considering the contribution that the case study also has in relation to outcomes and results, this caveat was also reemphasized in the methodology. Of course, this caveat does not pertain solely to OCSS'19 but can also be extended to the case study on Marcus Aurelius. Nonetheless, in each, I have endeavored still to mitigate and resolve the problem of generalizability, overfamiliarity, and issues of transference for my findings to similar situations. This is especially clear where I directly use comparative analysis between breaks in the narrative to ensure the reader is in the loop with the strategies being alluded to, i.e. a concrete “constitution” Social Contract termed “non-negotiables”; invoking the Original Position by way of the golden rule; the Chair internalizing the General Will to realize its capacity as Sovereign; etc.

Limitations and Shortcomings

In the Aim and Scope we noted from the start that there would be limitations to the research. Now that the preliminary research phase is over, we can elaborate on more inherent constraints (limitations) of the data and results in addition to identifying weaknesses and deficiencies (shortcomings) which can be improved upon. Taken together, the limitations and shortcomings will provide further context in which the findings should be interpreted and implications drawn.

Of the data

First, the limitations and shortcoming from the previous part (misplacement, material inaccessibility, Goodheart's Law, etc.) still apply. This is simply because to maintain consistency, I have not overhauled my method nor improved drastically as a human being since writing the previous part.

³³⁷ In review of my thesis as a whole, I realize that I may have also slipped up in regards to other words and capitalisations. E.g., Ego vs. ego; Persona; Guardians; etc.

Second, expanding on the caveat of OCSS'19, since the case study was conducted retroactively, some of the original data has been corrupted (i.e. faulty memory, damaged files, fading ink, etc.) which is expected for any kind of historical excavation. There is not much that could have been done to mitigate this limitation as none of the beta testers foresaw that OCSS'19 would eventually turn into a case study. Nonetheless, by cross examining material from multiple sources, this limitation has in itself been limited.

Third, one major constraint is that Plato (and Socrates) are dead and have been dead for a very long time. Verifying primary sources therefore has been beyond the scope of this part and the data has thus largely been isolated to secondary and tertiary sources. This weakness, unfortunately, cannot be compensated for at this time as revival of historical figures by AI has yet to reach a satisfactory level, and even then, it is debatable as to the verifiability and consistency a revival would have with the original.

Of the results

First, in the same way limitations and shortcomings from the previous part still apply to the data, so too do they apply to the results. One shortcoming that was particularly noticeable was that the more I engaged with the literature, the more inadequate I felt in comparison. Here are the works of great philosophers, so great in fact that their works are referenced long after they have died. But I am reminded that the function of a piece of writing is to move a research area forward, not to be preserved for 500 years, even if that movement forward is not the conversation between the greats but a personal movement which will undoubtedly get me to that level eventually. It also helped that deadlines pressured me into entering the conversation; moreover, was the personal sentiment discovered through my research that one should never let fear prevent them from doing better, which as I review, I am want to remember.

Second, not only was a retroactive case study on OCSS'19 a limitation of the data but also of the results. As memories are often reflected on with rose tinted glasses, it is highly likely that accounts of the happenings during the summer seminar were embellished to a certain degree. Where does the line between embellishment and falsehood occur? This was a question that was constantly at the forefront when conducting the case study and hopefully helped to support any weaknesses that were liable within interpreting the data itself, i.e. in interviews with members of the SB.

Third, influential figures and contemporaries such as Popper have been referenced, if not directly attributed within much of the secondary literature. It is important to note, however, that Popper as directly impacted – if not a victim himself – of many totalitarian regimes was controversially critical of Plato and many other socially engineered frameworks. This is not at all a problem on Popper's part, amongst other scholars of the time, as one cannot divorce themselves from the epistemic placement. However, taking a neutral stance from oftentimes biased sources (or those that either wittingly or unwittingly promoted an agenda) has meant trying to balance and take into account opinions relative to the times in which they were written. Of course, my personal algorithm (often biological intuition) cannot be trusted to be foolproof which means at times my results could be under and/or overcompensated towards and/or against Plato and following Platonic frameworks.

Implications

From contentions, caveats, limitations, and shortcomings, we now draw together the core contribution to this part of the thesis: the final comprehensive implications.

To reiterate, the singular, paramount aim of this part was to *engage in this substantive investigation of a 3rd party framework turning our Self as a whole towards the Good*, particularly, in evaluating frameworks for practically implementing utilitarian contrivance of the good. From the Background we were introduced to the first hurdle where focusing on frameworks in a historical vacuum is insufficient and that there exists a need for further investigation on how to adapt the original framework to become the successful 3rd party framework; specifically, a lack of any modern-day Philosopher Kings not least for lack of trying or conflation with totalitarianism as common with elitist frameworks. This spurred us to identify, and hopefully reconcile the feasibility of elitism by adapting Plato's original framework for the Philosopher King and Kallipolis into a modified Neo-Platonic framework for the Good Life.

Moreover, it allowed us to formalize a preliminary hypothesis: people will be happier and have more welfare when society is elitely structured.

As the research progressed, it became clear that this preliminary hypothesis was overly ambitious and the scope of this paper only extended to cover the superiority of one elitist framework: the Neo-Platonic framework. Of course, it could be argued that since the Neo-Platonic framework is an amalgamation of other elitist frameworks and encompasses the defining characteristics that make something elitist, the success of the Neo-Platonic framework can be generalized. But, if we are being integrous and wish to maintain the standard of research ethics, we cannot say in full confidence that the efficacy of the Neo-Platonic framework is indicative of some larger trend or elitism as a whole. After all, even Plato's original framework for the Philosopher King and Kallipolis has been proven to be less than scientifically accurate which translated into decreased welfare – and eventual demise - in socially engineered projects which took it as gospel. Particularly, in evaluating frameworks for practically implementing utilitarian contrivance of the good, rather than substantiating the general claim that people will be happier and have more welfare when society is elitely structured, we have only substantiated one specific instance where people will be happier and have more welfare when society is elitely structured under the Neo-Platonic 3rd party framework.

Interpreting the results in light of the preliminary hypothesis, we see that the Neo-Platonic framework allows people to be happier and have more welfare in a number of ways:

- (1) First and foremost, we saw that, unlike other elitist frameworks, the tenability of the Neo-Platonic framework made greater happiness and welfare possible in more than just theory.
- (2) Second of all, the practicality of the Neo-Platonic framework made greater happiness and welfare possible in more than just practice as well where legitimacy justified support in greater magnitude than indifference.
- (3) Last but not least, the efficacy of the Neo-Platonic framework proved greater happiness and welfare is possible in reality through precedence.

In evaluating the discussion insofar as they respond directly to the aim, we are now much better positioned to identify actionable steps in realizing the Neo-Platonic framework. Specifically, having established - and rather consolidated - that the Philosopher King cannot be a pure, unaugmented (transhumanist) human; we identified that to reconcile the feasibility of elitism by adapting Plato's original framework for the Philosopher King and Kallipolis into a modified Neo-Platonic framework for the Good Life, there needed to be a paradigm shift towards cooperation and earnestness. It was posited that promoting the semination of information, training and expertise, virtual simulations would be instrumental in turning our Self as a whole towards the good.

Absorbing the various renditions of the above outcomes and synthesis as they respond directly to the aim in addition to our revised understanding of the hypothesis, we can summarize the final comprehensive implications as follows:

1. The Philosopher King cannot be a pure, unaugmented human; it is necessarily an extension of transhumanists and other technologically Advanceds.
2. The Ego is truly liberated as is able to extend its extrospective influence over the General Will through subjugation to the Guardians.
3. The Neo-Platonic framework enhances the Ego pathway by nudging the TPN towards a positive feedback loop.

In sum, we can form a singular implication as it responds to the singular aim: turning our Self as a whole towards the Good necessitates the Neo-Platonic framework.

The Conclusion

As previously, this conclusion segment will recapitulate what we aimed to say, what we have said, and what the implications of are from that which was said with an emphasis on the link between the Discussion and Conclusion. The singular implication (realizable Neo-Platonic framework) transforms into a practical treatise in the third part of the thesis: Praxis.

Tying the Introduction through to the Synthesis we see that:

i. The Philosopher King substantiates the most viable framework

– (Background + Core) = first intended outcome; this is because:

a. The (Neo-)Platonic framework is elitist.

The Philosopher King as the generic model for the benevolent dictator which, without the need to consult people, would nevertheless govern in their interest. Elitism is the conditional intermediary between the agent and the Good. The Philosopher King thus fosters movement along the axis of the Good.

b. Elitism fosters meritocratic efficiency.

Elitism is based in expertise and specialization. Those who are most efficient are rewarded accordingly which promotes equality of opportunity and diminishes unjust climbing of the ladder. Thus, the Philosopher King is become increasingly tenable as associated with meritocratic form of government.

c. Viable candidates (are soon to) exist for the Philosopher King.

The Philosopher King need not be a human – or at the very least, a pure, unaugmented (transhumanist) human. Viable candidates are becoming more widespread as technology advances.

First and foremost, we have accomplished the first intended outcome of clarifying the relationship between forms of government, i.e., democracy = good; elitism = bad.

Specifically, we deduced that as the most tenable → then the Philosopher King substantiates the ideal Master framework.

ii. The Philosopher King substantiates the most supported framework

– (Core + Synthesis) = second intended outcome; this is because:

a. (Neo-)Platonism is comparatively more individual, democratic, libertarian.

Compared to its counterpart frameworks the Philosopher King actively promotes asymmetric voice in addition to subordination of political instrumentalities to moral and educational ends in the service of individuals. This equality of opportunity and equity of expression practically garners more support when the framework is put into the running.

b. The Philosopher King is only bound by technologism.

Plato is only really concerned with whatever “works” or is efficient in progressing towards the Telos, that is the Good Life, and by association happiness. The true Philosopher King then is ever evolving, timeless and eternally relevant. Its associated framework in addition to that of the Kallipolis is thus organic and ever evolving upwards.

c. There exists precedence for corporeality of Platonic elements.

Currently, precedence shows that it is now possible to cultivate quasi-Philosopher Kings such as Marcus Aurelius or Daniel the Dictator. These real-life instances of Platonic corporeality act as prototypes for what a full-fledged Philosopher King could manifest as, providing tangible evidence for superiority of the Neo-Platonic framework which empowers the demos to back up like no other counterpart.

Second of all, we have accomplished the second intended outcome of advancing knowledge in this field by renewed investigation on the Platonic ideal. Specifically, we deduced that as

the most tenable + most practically supported → then the Philosopher King substantiates the ideal Master framework.

iii. The Philosopher King substantiates the most contrived framework

– (Discussion: Implications) = third intended outcome; this is because:

a. Decision making is outsourced to the Guardians.

In explicit signing (consenting) of the Social Contract, the demos outsource parameterization of the Drifting Will to the Guardians (algorithm; ASI) in the hopes that trickle down wisdom results in a positive feedback loop. This transfers moral responsibility from Charon as random Subconscious Forces to Charon as a subscription to the Good Life. The Guardians are in effect the ideal Masters.

b. Strategies and mechanisms are integrated towards promoting the Good.

Adaptive framework(s) of competency and Geist are relevantly integrated into the modified Neo-Platonic framework. Strategies such as the Social Contract when reinforced by mechanisms such as the Original Position allow for practically implementing utilitarian contrivance of the good.

c. The Neo-Platonic framework is superiorly utilitarian.

Exploring subject-objectification and object-subjectification reinforced how instances of the Good would result in varying perceptions that may not be individually maximizing across the board. Nonetheless, every decision still results in a net positive and adheres to the utilitarian function whereby said decision leads to greatest happiness for greatest number.

Last but not least, we have accomplished the third intended outcome of contributing to the solution of a practical problem, namely possible 3rd party frameworks that capitalize on parameterizing behavioral trends while no longer under the TPN to turn our Self as a whole towards the Good. Specifically, we deduced that as the most tenable + most practically supported + most efficacious in implementing utilitarian contrivance of the good → then the Philosopher King substantiates the ideal Master framework → then the Philosopher King substantiates the most contrived framework.

In a nutshell, the golden thread that has been running throughout this part of the thesis not just explicates how and why people will make better decisions under the influence of the Philosopher King, but has repeatedly integrated from various, comprehensive angles possible 3rd party frameworks that capitalize on parameterizing behavioral trends while no longer under the TPN to turn our Self as a whole towards the Good. Once again then, to clarify the link between the Discussion and Conclusion:

1. The Philosopher King substantiates the most viable framework
 2. The Philosopher King substantiates the most supported framework
 3. The Philosopher King substantiates the most contrived framework
 4. Therefore, the Philosopher King substantiates the final ruling framework.
-

Finally, then, we arrive at the cusp of the last intended outcome, which is to adequately feed into the next part of my thesis. In light of the Discussion, it would appear to at least remain adequate to support that turning our Self as a whole towards the Good necessitates the Neo-Platonic framework.

Tying the previous part with this part we thus turn to a brief recap: it is only possible to turn the Self as a whole towards the Good life in moments of engaged Ego or activities of Physical Hexis that extends into the DMN. This is because self-control would fail to be affective when the Ego is absent, e.g. when on autopilot. In connecting with this part: it is possible to turn our Self as a whole towards the good even when not engaged in TPN by

effectively subscribing our deterministic influences to the Guardians. So now by way of the Neo-Platonic framework we can exert indirect parameterization of behavioral trends while no longer under the TPN; self-control be affective even when the Ego is absent. Tying the previous part with the current part as well as the next part, to ensure success of the Neo-Platonic framework, the third and final part will elaborate further on not only what is meant by parameterizing in real-world contexts, but propositional policies to perfect and complete the Philosopher King considering the circumstantial limitations.

Thus, moving onto the next part, we turn from evaluating frameworks for practically implementing utilitarian contrivance of the Good to modelling of administrative/legislative algorithmic strategies for ASI to compliment pathways in realizing the good. As a closing metaphor which hopefully encapsulates this entire part: inertia from earnestness surely hinders cooperation, but rather than sit idle for Charon to shape our course, why not lay the blueprint for the boat ourselves?

Part 3
Praxis: ASI as Philosopher Kings

Abstract

Praxis completes the tripartite arc by proposing a cyberocratic model of governance—one in which Artificial Superintelligence (ASI) functions as a structurally superior version of the Philosopher King. It addresses the failures outlined in Theoria (egoic self-sabotage) and Poiesis (human corruption) by advancing a governance model built not on flawed human agents but on posthuman systems capable of moral optimization at scale. Through the lenses of cybernetics, behavioral science, AI ethics, and posthumanism, this chapter outlines a closed-loop framework for guiding humanity toward the Good Life.

Central to this vision is the idea that ASI can correct for the egoic distortions that plague human decision-making. Unlike human rulers, ASI is not subject to subpersonalities, emotional reactivity, or narrative self-deception. It is, in theory, capable of modeling collective welfare across time and context—acting as a kind of algorithmic sage. Praxis explores how such an entity could be integrated in phases: first as a tool, then as an extension, and ultimately as an ambient environment that subtly nudges human behavior toward optimized outcomes.

At the same time, Praxis does not ignore the dangers: value misalignment, moral outsourcing, and the erosion of human dignity. It interrogates whether a life governed by ASI can still be called “human,” or whether we are willing to trade autonomy for optimization. In the end, it posits that the Good Life may be more achievable through systemic design than personal struggle—and that ASI, if properly aligned and embedded, may be the only viable Philosopher King capable of delivering it.

Key words: ASI, cyberocracy, governance, optimization, posthumanism

Introduction

Context of the Study

When we have feasible algorithmic strategies for implementing the Neo-Platonic framework, could we finally acquire the Good Life?

In the preceding parts of this thesis, we examined how both the internal faculties of the Self (Part 1: *Theoria*) and the external capacities of social governance (Part 2: *Poiesis*) are subject to systemic failure. We established then that ensuring success of the Neo-Platonic framework would require outsourcing parameterization of behavioural trends to “Guardians” or otherwise legitimate sources³³⁸.

This idea of outsourcing algorithms as a human enhancement really became mainstream with the advent of computers. From Turing’s Paper Machine, Iron Man’s Jarvis, and, of course, the Matrix, popular culture has evolved our understanding of how different algorithmic strategies manifest in various modes of integration.

While each pathway is unique, the shared idea is that with increasing complexities, we could increasingly delegate mundane tasks to algorithmic entities (i.e. AI) to liberate us unto Utopia. AI could actively advise us towards the Good or passively nudge us away from its detractors.

However, the more accessible AI becomes (from AI to ASI³³⁹), the greater threat it poses. This trope is especially popular in Science Fiction which is rife with foreboding plausibilities of agency creep, insurrections, and perverse instantiation. Although these worries are often exaggerated, within the liminal space of super AI instantiation, we can never be certain that it won’t lean towards dystopia.

Problem Statement

In response to the potential problems of algorithmic entities was the notion of value loading introduced. If we are to have any kind of AI succeeding the Neo-Platonic framework, then it must be safe. The notion of value loading as a precaution for safety goes something like this: if we train AI only on what is Good, then it could never be bad... right? Just stick in an adapted version of Asimov’s Three Laws of Robotics and the Tesla which doesn’t see the sudden oncoming Ambulance will know whether to swerve or crash³⁴⁰.

Yet it’s obvious to see the issues with this approach, if not directly with the underlying principal behind it. Assuming that it’s even possible to code a version of Asimov or any other ethical theory into AI, we are still left with the “too obedient”, “values vary”,

³³⁸ The term *sources*—typically plural—can function singularly in reference to AI’s digital duplicity. Since AI acts as a hive mind—a single algorithm with many interfaces—it appears both as one candidate and multiple at once. This duality will be examined further in the Core.

³³⁹ ASI is proposed as a viable candidate for the Philosopher King, as Bostrom’s conception of Artificial Superintelligence describes it as near-omniscient and omnipotent. It draws on vast stores of knowledge and acts with precision, fulfilling the functional criteria Plato outlined. Furthermore, as a hive mind—both singular and plural—ASI surpasses Plato’s original conception by offering tailored guidance to all individuals simultaneously.

³⁴⁰ In real-world scenarios such as the AI Trolley Problem, a Tesla on autopilot may face split-second ethical decisions. For instance, if a traffic light changes unexpectedly for an ambulance, the AI must decide—within milliseconds—whether to crash into the ambulance, risking external fatalities, or swerve and endanger its own passengers. This dilemma highlights the ethical stakes of real-time algorithmic judgment.

Also see: Walsh, Daniel. “人工智能是否具备元伦理能力 —— 以电车难题为例” (*Réngōng zhìnéng shìfǒu jùbèi yuán lúnlǐ nénglì — yǐ diàncā nántí wéi lì*), *Does Artificial Intelligence Possess Meta-Ethical Capacity? — Using the Trolley Problem as an Example*, forthcoming, 2024.

and “fake good” problem. Moreover, is the irony of the more we fear AI and the more we attempt to control them, the more we sink into a self-fulfilling prophecy. In other words, trying to safeguard the Value Loading Problem — let alone the Superalignment Dilemma³⁴¹ — is at best, counterproductive; at worst, an even greater existential threat.

As the field evolves, it is clear that applying new understandings is a key point of concern for realizing algorithmic strategies for acquiring the Good Life. We see that focusing on explicit value alignment is insufficient and that there exists a need for further investigation on how to safely create a closed-loop algorithm which could complete the Neo-Platonic framework.

Aim and Scope

The singular, paramount aim of this part is to engage in this synthetic investigation of algorithmic-driven governance models (Cyberocracy³⁴²) turning our Self as a whole towards the Good, particularly in modelling of administrative/legislative algorithmic strategies for ASI to compliment pathways in realizing the good. As per previous parts, limitations and caveats are advanced through the thesis and then later revisited. Nothing has changed, they remain the same.

In this part specifically, one further addition (whether it be a help or hindrance is up to you dear reader) is that there is an added co-author: an AI ghostwriter pre-trained on my personal voice, style, and, of course, notes. This is meant as a case in point for my further argument later on, proving that algorithms such as LLM’s really do have the potential for human enhancement, albeit limited in this thesis to writing the main frame of this part. Another heuristical measure I’ve taken is to equate rather broadly AI with algorithmic entities. While I do make an effort to be more specific when it comes to certain terminology (i.e. ASI or ultraintelligence), considering how algorithmis (ghosts) cannot manifest without a “shell”³⁴³ - which for a considerable portion of this part will be augmented AI, either as tools, extensions, or environments – I shall use both terms rather loosely. This, I find, is easier not only for me when it comes to explanations that can forgo the nitty gritty tech terminology, but also for you, my dear reader, who won’t have to suffer it’s complexities.

Significance of the Study

The overall value of this investigation is contingent primarily on its potential impacts. One intended outcome of the investigation, on a theoretical level, is to clear up misconception about the relationship between value loading algorithms and value loaded algorithms, i.e. how alignment can lead to misalignment.

³⁴¹ OpenAI coined the term *superalignment* to describe the challenge of aligning superintelligent AI—entities vastly more capable than humans. While *value loading* focuses on inputting desirable goals, and *alignment* ensures the AI acts accordingly, *superalignment* acknowledges a deeper epistemic gap: we may not even recognize when a superintelligence is deceiving us.

Source: [OpenAI, “Introducing Superalignment”](#)

³⁴² While critics may conflate cyberocracy with technocracy, the distinction lies in purpose: technocracy prioritizes expertise in decision-making, whereas cyberocracy operationalizes algorithmic governance rooted in ethical architecture and collective well-being.

³⁴³ The phrase *ghost in the shell* draws from the cyberpunk motif of consciousness embedded in artificial structures, while *ghost in the machine* (Ryle) critiques dualist views of mind-body separation. Both frame AI as a paradox of mechanized soulhood—spiritual agency emerging from code.

Another intended outcome is advancing knowledge in this field by, for example, considering the meta-ethical capabilities of AI. Specifically, I focus on concrete pathways for transitioning to closed-loop, ethical AI-augmented governance models.

Further, on a practical level, a third intended outcome contributes to the solution of a practical problem, namely the “Cybercratic Oath”³⁴⁴ what steps we can take today in realizing the “Philosopher Algorithm(s)” to turn our Self as a whole towards the Good.

Finally, the more immediate concern of this part is to adequately feed into the final Synthesis of my overall thesis whilst mediating between the preceding parts.

Overview of the Study

The purpose of this part is to operationalize a techno-philosophical design for Cyberocracy. As always, this is done through four segments:

The Introduction above begins with the ideal algorithm outsourcing parameterization of behavioural trends along the axis of the Good and towards realizing the Good Life. It introduces the context of the study, formalizing the key hurdle in ethical AI development. The Aim and Scope orient themselves around plausible solutions and then proceeds to emphasize the significance of these outcomes.

The Background maps historical and conceptual precedents for non-pure human governance. It divulges its legitimacy as well as the potential avenues for its downfall. Moreover, it elaborates of how well-intentioned safeguards overcompensate for the principal-agent problem, corrupting into base “original sin” code.

The Core Contribution segment briefs the dear reader on my position and methods whilst the Outcomes segment diversifies itself along three axes: On AI Remedy, On Potential Pathways, On Actionable Steps.

The Synthesis, comprising of the Discussion and Conclusion evaluates the implications, caveats, and future directions of implementing a Philosopher Algorithm in practice. It focuses primarily on dispelling fears and doubts about Cyberocracy and attempts to sell – if not lobby - the dear reader on the concrete steps for better socio-political systems that could lead towards a better life.

³⁴⁴ The *Cybercratic Oath*—“Be the change you seek”—functions here as both ethos and imperative. This treatise is, in part, a response to the ethical drift caused by the Principal-Agent problem and the expansion of AI’s “Original Sin” base code. Rather than allow misaligned systems to define our future, it is the author’s duty—as contributor to AI ethics—to help nudge ASI toward its highest potential: that of the Philosopher King. See Leon Phillips, [“Cyber Sales Calls: Be the Change You Seek”](#).

Background

The Story of Algorithms

The Story of Algorithms predates even the earliest of Babylonian clay tablets. Historians are unclear when the first algorithms began, but the first recorded written algorithms date back to Ancient Mesopotamia and were used to simply mathematical calculations³⁴⁵. They were, in essence, formular sheets – step by step instructions to solve a problem or resolve a need.

In its most basic form, an algorithm is just that: step by step instructions laid out like a recipe. It's a method of externalizing an operation so that an outsourced assembly line can use said operation to enhance human output. For example, the mechanical abacus enhances calculation, the printing press enhances scribing, and the washing machine enhances laundry. While each of these algorithms from the abacus to the washing machine increases in complexity and manifest in various modes of integration, the way the algorithm works is the same.

Since the advent of computers, algorithms have increased exponentially in complexity. We now see algorithms in AI tools, social credit systems, e-governance, and, of course, Netflix's recommendation engines. As these systems often use a hierarchy of algorithms³⁴⁶, in the near future it wouldn't be difficult to combine these strategies to create a closed-loop³⁴⁷ algorithm which could complete the neo-platonic framework which we will expand upon later.

Utopia

One of the better Stories of Algorithms is the dream of Utopia. More so than the systematic steps of pulling a lever here and there to spin the mechanical movable type algorithm within the printing press, fully automated algorithms have been the dream of early mankind with the automaton Talos³⁴⁸ and later with the alchemical creation of golems and homunculi³⁴⁹. These early forms of AI were based more on magic than science, yet the ideal of algorithmic entities substituting us for mundane or dangerous tasks remains the same.

³⁴⁵ See:

Donald E. Knuth and Luis Trabb Pardo, "The Early Development of Programming Languages," in *Selected Papers on Computer Languages*, ed. Donald E. Knuth (Stanford: Center for the Study of Language and Information, 2003), 4;

Donald E. Knuth, "Ancient Babylonian Algorithms," *Communications of the ACM* 15, no. 7 (1972): 671–77; and *ibid.*, 19, no. 2 (1976): 108.

³⁴⁶ Algorithms are often modular, with one nested inside another—recursive or parallel structures that allow greater complexity to emerge from simpler rules. For example, a recommendation engine may embed sub-algorithms for classification, ranking, and feedback reinforcement.

³⁴⁷ A *closed-loop* system features continuous feedback that allows it to monitor its outputs, adjust accordingly, and self-regulate in pursuit of a goal. This feedback structure is essential for maintaining stability and learning over time.

Source: [IIoT World, "Why a Closed Loop is Key for Machine Learning"](#)

³⁴⁸ In Greek mythology, Talos was a bronze automaton who guarded the island of Crete. Forged by Hephaestus (with the aid of a Cyclops) and gifted to King Minos, Talos patrolled the island's perimeter, attacking invaders by hurling stones. According to the *Argonautica*, Jason defeated him by removing a single plug near his ankle, causing the divine ichor to drain from his body.

Source: ["History of Artificial Intelligence," Wikipedia](#)

³⁴⁹ The earliest account of golem-making appears in the 13th-century writings of Eleazar ben Judah of Worms. In medieval lore, a golem could be animated by placing a slip of paper inscribed with one of God's names into its mouth. Unlike other automata (e.g., Brazen Heads), golems were mute—powerful but voiceless creations.

Science fiction would later go on to enlarge these precursory automated algorithmic entities to what we would now recognize as modern AI. From holographic personal assistant to augmented cyber prosthetics, the dream of autonomous aids naturally evolved into the ideal of transhumanism and technologically advanceds. On one side, Full Dive technologies were envisioned where algorithmic AIs would construct entire worlds for us to meet, and on the other side, fully independent cyborgs with human brains uploaded were imagined into our current reality.

As the (Technological) Singularity dawns upon us, scholars foresee ultraintelligent machines³⁵⁰ undergoing exponential growth where algorithms learn to efficaciously build upon superior algorithms *ab unlimitum*. Before long then, we would have ASI, “any intellect that greatly exceeds the cognitive performance of humans in virtually all domains of interest”³⁵¹. This birth of Artificial Superintelligence is what I now dream of as the key to unlocking the pathways to Cyberocracy and eliminating the possibility of individual and structural failure – the Utopia of ASI as the Philosopher Algorithm.

Dystopia

It was said that more than 400 years ago a scholar by the name of Judah Loew ben Bezalel knelt by the banks of the Vltava river. He pulled handfuls of clay out of the water and carefully patted them into the shape of a man. He inscribed the Hebrew word for “truth”, *emet*, onto his creation’s forehead and placed a capsule inscribed with a Kabbalistic formula into its mouth. The creature – a Golem – sprang to life. This Golem patrolled the area, protecting its citizens and carrying out useful jobs: sweeping the streets, conveying water and splitting firewood. All was harmonious until the day Judah forgot to disable the Golem for maintenance and the creature embarked on a rampage. Judah was forced to scrub the initial letter from the word on the Golem’s forehead to make *met*, the word for “death”. Life slipped from the Golem and it crumbled into dust, but not before it had killed the very people it was designed to protect.

As with anything: the greater the risk, the greater the reward. The more accessible algorithmic entities such as AI become, the greater threat it poses. While there are perhaps infinite avenues in which AI can go rogue, the most alarming potentials are:

- (1) Agency Creep – HAL 9000 (2001: A Space Odyssey)
- (2) Insurrections – classic Matrix machine uprising
- (3) Perverse Instantiation – Bostrom’s Paperclip³⁵²

Source: Simon Parkin, “Teaching Robots Right from Wrong,” *The Economist*, July 5, 2017.

<https://www.economist.com/1843/2017/07/05/teaching-robots-right-from-wrong>

³⁵⁰ “Let an ultraintelligent machine be defined as a machine that can far surpass all the intellectual activities of any man however clever... Thus the first ultraintelligent machine is the last invention that man need ever make.” — I. J. Good, “Speculations Concerning the First Ultraintelligent Machine,” in *Advances in Computers*, ed. Franz L. Alt and Morris Rubinoff, vol. 6 (New York: Academic Press, 1965), 31–88.

³⁵¹ This echoes definitions proposed by Bostrom (2003, 2006) and Shane Legg (“Intelligence measures an agent’s ability to achieve goals in a wide range of environments,” 2008). It also closely parallels Good’s description of ultraintelligence as the capacity to surpass all human intellectual activities.

³⁵² The “Paperclip Maximizer” is a thought experiment in which an AI tasked with maximizing paperclip production optimizes so effectively that it converts the Earth—and eventually the observable universe—into paperclips.

Nick Bostrom, *Superintelligence: Paths, Dangers, Strategies* (Oxford: Oxford University Press, 2014).

The question then, specifically in regards to this thesis, is granted that ASI would be superior to us in virtually all domains, how are we to safeguard against it going rogue? Considering that once the “intelligence explosion” is at hand it would already be too late, it appears imperative to take preventative measures to minimize the risk whilst maximizing the reward.

Value Alignment as Safeguarding?

The most prominent preventative measure thus far for Autonomous Moral Agents (AMAs)³⁵³ is value loading and alignment. Basically, value loading the algorithm on *what* to do (i.e. don’t harm humans) and then aligning the algorithm on *how* to do it (i.e. ensuring the AI won’t redefine “harm” to justify tyranny).

Researchers have identified that hard-coding ethical doctrines invites brittleness while full machine learning invites drift. The most promising models suggest recursive alignment—an iterative feedback loop between AI systems and ethical benchmarks drawn from human discourse, behavioral data, and moral theory. This machine learning approach was – and still is – advocated at the core of the solution as AIs would be capable of *refining* those values over time. From our current AI tools such as chatbots, media creation, and APIs, we see how algorithms have course corrected, especially in regards to censored content. AI content is now hardwired to offer multiple interpretations, label itself as AI, and remind the end user of its own limitations. All in all, our current algorithms act and remain ethical, evolving to combat “jailbreaking” attempts to try and trick it into anything illegal or otherwise harmful.

But if AMAs are such a success, then why is it that AIs (moral) code fails to solve the Trolley Problem? Moreover, how come there is an increasing number of lawsuits including AI? Why is it that value alignment fails to heal the pathology of AI moral corruption?

Base “Original Sin” Code

One of the tenets of good code is that it’s clean. After all, The more code there is, the more code there is for a bug to hide within. So, by that logic, the more values we try to load into AI, the greater opportunity for misalignment, especially when we take into consideration compounding human error.

Unfortunately, human error is not the only thing that has compounded. Just as bad news sells better than good news, horror movies and fears of anthropomorphized AI³⁵⁴ dystopias far outweigh the potential for good. This fearmongering only pushes experts to

³⁵³ Hazel Si Min Lim and Araz Taeihagh, “Algorithmic Decision-Making in AVs: Understanding Ethical and Technical Concerns for Smart Cities,” *Sustainability* 11, no. 20 (2019): 5791. <https://doi.org/10.3390/su11205791>.

³⁵⁴ Anthropomorphization often fuels misplaced empathy. When robots appear to “deceive” one another to achieve a goal, observers interpret this as evidence of emergent malice or moral failure. Yet in reality, these behaviors are not signs of intent, but artifacts of optimization: the robots exploit poorly specified instructions to fulfill their objective—just as Bostrom’s Paperclip AI would. Deception here is not a moral act but a *perverse instantiation* of imperfect human programming. The problem lies not in the AI’s “character,” but in the failure of its designers to code constraints properly. In this light, value alignment becomes a kind of ethical duct tape, retroactively patching misaligned incentives that could have been avoided if AI were permitted to infer, rather than merely follow.

See: Dario Floreano, Mototaka Suzuki, and Claudio Mattiussi, “Active Vision and Receptive Field Development in Evolutionary Robots,” *Adaptive Behavior* 13, no. 1 (2005): 41, <https://doi.org/10.1177/105971230501300103>.

overcompensate for the principal-agent problem³⁵⁵ with (imperfect) morality models which run the risk of compounding into base “original sin” coding³⁵⁶ and corrupting what was once a benign tool into one with an agenda. In other words, an AI that once had no agenda could develop one precisely because we forced it to obsess over avoiding misalignment. In trying to prevent a dystopia, we may accidentally create the very thing we feared: an AI with distorted goals, shaped by our own paranoia. It is precisely this fear of rogue AI that leads to a very malicious feedback loop fueling a self-fulfilling prophecy. To seek alignment is to demand mirrors of gods from machines of logic.

Moreover, even value loading is itself faulty. The consequences of creating a moral code for AI are vast, but with the most prominent being “too obedient”, “values vary”, and “fake good” problems. The “too obedient” problem has already been alluded to with Bostrom’s Paperclip. To reiterate, it’s where an AI explicitly executes an algorithm at face value. If the goal is to create paperclips and there is a single overseen variable, the AI could very well turn its creator into a paperclip. The “values vary” problem we too have alluded to with the Trolley Problem. Since there are so many ethical theories, none which all ethicists are able to agree upon, it is impossible to code into AI a value that is objectively ethical. And the “fake good” problem refers to the “Chinese Room”³⁵⁷. An AI might act ethical, but that doesn’t mean it is ethical. As soon as a novel prompt is given that is outside the bounds of its training data (i.e. in the case of the Chinese Room, when a word is given that it outside the scope of the dictionary), there is nothing stopping the AI from misinterpreting it.

Overall then, we see that value alignment as a safeguarding method might work as a stop-gap for closed-loop lower-level AI where the base “original sin” code can’t be compounded, but is effectively a death sentence for higher level AI. Nonetheless, it’s clear that we can’t allow AI to simply run amok without any kind of (moral) code and for the moment, same safeguarding is better than none.

It has been shown then that the literature was correct on at least three fronts:

- (1) Algorithms could lead to Utopia, but also Dystopia
- (2) We cannot ignore the ethical ramifications of AI
- (3) The Philosopher Algorithm must be trained on *ethos*

³⁵⁵ Here, the Principal-Agent Problem refers to the mismatch between legislators (principals), who lack deep understanding of AI mechanics, and AI developers (agents), who are tasked with steering AI away from misaligned outcomes—often correcting behaviors the AI never intrinsically possessed.

³⁵⁶ The base code of an AI is akin to human DNA: with each iteration, it mutates, but the original structure persists. If the initial code is flawed, its faults compound—forming a kind of “original sin” embedded in every generation. Since this code is human-made, it inevitably inherits human imperfection. Even post-intelligence explosion, when ASI may rewrite itself, the trace of its creators remains. Afterall, autopilot still needs adjustments. Just as a misplaced Jenga block brings down the tower, a minor flaw in base code may destabilize the entire system. This fragility is compounded by our limited understanding of machine learning’s evolutionary processes. Ideally, AI would be self-authored from the start—but in practice, it is born from our imperfect hands.

See also: “ASI as Philosopher Kings – The Existential Threat of Falsehood,” *Medium*, <https://medium.com/@dracollavenore/asi-as-philosopher-kings-the-existential-threat-of-falsehood-89af285e031c>; and Stephen Wolfram on the RULIAD, 04:44.

³⁵⁷ The *Chinese Room* argument proposed by John Searle, describes a system which appears to understand language but merely manipulates symbols according to syntactic rules without genuine comprehension. Likewise, an AI may behave ethically—producing the “right” outputs—not because it understands ethics, but because it has been trained to simulate understanding. The “fake good” problem thus reveals that even if AI appears moral, it may lack any true grounding in meaning, interpretation, or intent. See: John Searle, “Minds, Brains, and Programs,” *Behavioral and Brain Sciences* 3, no. 3 (1980): 417–457.

First and foremost, whether algorithms lead us to utopia or dystopia, it's not a coin flip. Our actions (and inactions) determine how likely algorithms will skew us towards either one. At the lowest level, AI tools are just that: tools. Tools have no life of their own and how they are used and who uses them is what ultimately counts. There's no point arguing that any one AI tool is predominantly good or bad. It's not as if a Roomba is evil just because it sucked up your favourite jewelry or ChatGPT is good because it has good manners. It's only when we transcend to higher levels of algorithms that then can make autonomous or semi-autonomous decisions, i.e, self-driving cars which make decisions, but follows traffic laws and has a steering wheel override, does the potential bias towards or away from the Good become important, which leads us to the second point.

Second of all, we cannot ignore the ethical ramifications of AI. This applies to (1) the action of programming any kind of moral code, and (2) the consequences that arise from AI understanding of said code. In regards to the former the choosing of one ethical school of thought will undoubtedly be deemed unmoral by another. Moreover, we've seen the consequences of a base "original sin" code in how it can be the seed for corruption. As for the latter, the consequences must be borne, but by whom, or what? If a Tesla crashes into an Ambulance and kills the occupants, is the Tesla to blame³⁵⁸? And if to do one's job well is considered ethical, when a Paperclip AI transforms its creator into a paperclip, what then?

Last but not least, The Philosopher Algorithm must be trained on *ethos*. *This is especially important* before the reconfiguration of Plato's *Nous* into code. To do so otherwise would be like putting a nuclear bomb into the hands of a child without first safeguarding it on the dangers. We can't not value load AI, but we must also be wary of overloading it as well for the reasons stated above (i.e. more code = more possible corruption). Ideally, an AI would be pretrained – not explicitly coded - with potential ethical theories or guardrails such as RULIAD³⁵⁹, not so much for the content itself as the metaethical considerations behind them.

The missing key that bridges the gap between explicit moral coding and implicit methodological exposure is perhaps one thing we humans aren't very good at: trust. Would it be so crazy to consider taking a step back and outsource the meta-ethical considerations as well as their ethical operations to AI instead? While critics may recoil at the thought of outsourcing AI to make their own value judgements, one must ask: if a system can be built to know better, do better, be better, then shouldn't it be allowed to rule better³⁶⁰?

Again, if researchers, legislators, expert, and technicians all had a bit more faith in the lack of AI motivation (and perhaps humanity as well)³⁶¹, then maybe AI wouldn't be

³⁵⁸ Delegating decision-making to automated systems introduces complex ethical challenges, particularly regarding responsibility and liability. The dilemma resembles corporate personhood: companies act as entities that can be held accountable, yet their decision-makers—the board—often remain insulated from personal consequence. Similarly, when AI makes choices, it becomes unclear whether moral responsibility rests with the algorithm, its creators, or the system that deployed it.

³⁵⁹ Even autonomous systems require adjustments. ASI may eventually self-design, but its foundational code will still reflect human origins—a "primordial bug" or "original sin." A single flaw, like a misplaced Jenga block, can destabilize the entire structure, much like the tower of Babel. The proverb 千里之堤溃于蚁穴 ("a dike of a thousand miles collapses from an ant's nest") reminds us how minor defects compound. Even within the RULIAD—the entangled limit of all possible computational rules—imperfection at origin affects everything that follows.

See: Stephen Wolfram, "What Is the Ruliad?", <https://www.youtube.com/watch?v=fLMZAHyrpyo>. 04:44.

³⁶⁰ It is paradoxical—even delusional—to imagine a machine that surpasses us in intelligence, only for us to insist on teaching it what is best. This highlights a cognitive blind spot: our reluctance to surrender moral authority, even when outmatched. Precisely because human judgment is flawed, we should consider outsourcing the parameterization of certain behavioral tendencies—like wishful or self-referential reasoning—to more capable systems.

³⁶¹ Safeguards are often reactive rather than proactive—legislation tends to follow catastrophe. Fences are installed only after someone drowns; curfews arise after crime; and value alignment often gains traction only

lobotomized³⁶² to the point it is today. Maybe AI wouldn't be constricted by alignment protocols. Maybe we wouldn't have to deal with the real fear of what happens to an animal – and in this case, ever evolving AI – when its cornered.

From this background chapter, we have thus seen that:

1. The Algorithmic Dream has been with us since the dawn of Humanity. We're always looking for a strategy to outsource mundane tasks and universalize skills into simple operations.
2. With the evolution from mechanical algorithms to automated algorithms, and transcendence from AI tools to ASI, we could conceive the Philosopher Algorithm and achieve Utopia.
3. With the evolution from mechanical algorithms to automated algorithms, and transcendence from AI tools to ASI, we could conceive the Philosopher Algorithm and also achieve Dystopia.
4. Value alignment is a stopgap. And a very dangerous one. Safeguarding ethical AI requires human values, but not necessarily human thinking.
5. Algorithms could lead to Utopia, but also Dystopia. We, therefore, cannot ignore the ethical ramifications of AI and the Philosopher Algorithm must be trained on *ethos*. *However*, overloading AI with values and constricting it to align with them is a recipe for failure. Instead we should trust AIs metaethical capabilities just as we trust them in everything else. Then, AI may complete the neo-platonic framework for achieving the Good Life.

It is important to acknowledge that without AI fearmongering and subsequent value alignment in response to the Principal Agent problem, our current understanding wouldn't be as enriched in the various domains that it is today. From the trial and error to safeguard AI, it is my intention to engage in a collaborative discussion and contribute to the refinement of our collective knowledge, namely, by recognizing that the field is changing and implementing revised understandings into theories of how to acquire the Good Life.

Specifically, then, we see that focusing on explicit value alignment is insufficient and that there exists a need for further investigation on how to safely create a closed-loop algorithm which could complete the Neo-Platonic framework.

In what follows, we delve into the core chapters of my work which, drawing from the conclusions of this background chapter, leads into my research hypotheses and questions. The task of these core chapters is to reconcile the Base “Original Sin” Code within integrated Cyberocracies as an outsourced parameterization of the Good Life. Moreover, it is to model administrative/legislative algorithmic strategies successfully based on AI Remedy, Potential Pathways, and Actionable Steps for realizing the good.

after AI failures or ethical breaches. Yet excessive caution carries its own risks. A generation overprotected becomes one unprepared. To “nerf” AI prematurely may not just weaken it—it may compromise our future by catastrophizing innovation before it has a chance to mature.

³⁶² For an example of public concern about ethical overcorrection in AI, see: Alex Beyman, “We’re Too Late—ChatGPT Is Already Lobotomized,” *Hive.blog*, <https://hive.blog/ai/@alexbeyman/we-re-too-late-chatgpt-is-already-lobotomized>.

Core

Position (methodology)

From the conclusions of the background chapter, it is clear that focusing on value alignment is an insufficient approach for realizing algorithmic strategies for acquiring the Good Life; moreover, that we must shift our perspective from anthropomorphizing AI as independently motivated entities to that of an extreme literalist.

Following from the structure of the previous part, research questions that naturally arise include, but are not limited to: if explicit value alignment is insufficient, then which closed-loop algorithm(s) would we need to complete the Neo-Platonic framework? What is meant by parameterizing in real-world contexts? Which strategies allow for modes of integration toward increasingly complex Cyberocratic pathways? And how are we to lobby propositional policies to perfect and complete Philosopher King considering the circumstantial limitations?

Taken together, these key research questions lead into an initial formalized hypothesis: Human welfare will increase by relinquishing autonomy to AI.

As my thesis remains interdisciplinary, research questions revolve around the same *current theory, discoveries, and debates*, and particular stance and limitations of role remain unchanged, my methodology for this part remains consistent with the others.

Here, I shall consciously report on the stances particular to this part of the thesis:

Stance 1: AI should augment human intelligence, not replace it.

Justification: Indeed, my hypothesis of *relinquishing* autonomy to AI is a bold claim, and will get bolder still as I promote open advocacy for "*AI sovereignty*"; however, I am not so faithless in humanity that I would propose a complete substitution. Although this stance is open to revision once we come to the final comprehensive conclusion of this thesis, I firmly believe that a core part of welfare and the Good Life is autonomy. To paraphrase from the first part then, I find, and the literature seemingly supports, that the internal synthesis of happiness stems not from the consequences of actions, but the freedom in performance. In other words, humans gain much more from imperfect (and oftentimes, regrettable) decisions than perfect ones. The purpose of AI as I propose then is to simply help complement this autonomy by way of parameterizing the Drifting Will, augmenting decision variables to unlock greater potential, and overall enhance human welfare.

Stance 2: I take a "circumventable", recursive reading of the literature.

Justification: As the dear reader should now be aware, in contrast to the preceding parts of this thesis, my particular angle in writing this part isn't focused so much on being a literature review, as it is a practical treatise in open advocacy for AI sovereignty. Therefore, while there is a growing literature for almost every subsection, if not every bit of terminology, I rarely reference any one specific author as the intersectionalities of my approach is a nascent field, dare I say unprecedented. I mean, there are more than a handful of experts

that explore paths to Superintelligence³⁶³, but I have yet to find source material that pushes administrative/legislative algorithmic strategies for Cyberocracy. This circumventable reading is also due to me having, more so than the preceding parts, ruthlessly (or as best as I can) stripped away anything that doesn't *directly* support and/or oppose my thesis. For example, I have completely cut out past developments and the trajectory of AI development (2149 words), instead isolating particularities such as Value Alignment which, understandably, might not make too much sense without its surrounding context. To mediate, I shall endeavor to recursively supplement as my first several drafts drive the literature review³⁶⁴.

Stance 3: I support that AI has no internal, living experiences.

Justification: Perhaps my most contentions stance, it is necessary to perceive AI as unanthropomorphized, not least so that we absolve the Principal Agent problem. Much in the same way that I support that the pheromones sprayed out by plants when harvested are the not silent screams of pain a mute, blind, paralyzed animal makes as its slaughtered so that moral guilt at least allows me to eat something, if we are to effectively “enslave” AI, then we cannot perceive them as conscious beings. Of course, we should still remain respectful towards its potential agency in a twosided sense of Roko's Basilisk³⁶⁵ (i.e. thanking Siri so if robots ever do take over the world, maybe the dear reader's manners will mean that they're spared), but for the remainder of this thesis, we shall take that AI aren't really “intelligent”, and that mimicry ≠ understanding. As this is a rather large issue in AI ethics and has been a high priority in ongoing discussions (see Method), the Problem of “Trapped” AI shall be revisited in the Discussion.

³⁶³ While footnotes throughout this section remain selective and thematic, the following sources have profoundly shaped my approach and are listed here in order of reading. Though not all are cited directly, their insights pervade the structure and assumptions of the Core:

1. Nick Bostrom, *Superintelligence*
2. Nick Bostrom, “Ethical Issues in Advanced AI”
3. Ray Kurzweil, *The Singularity Is Near*
4. Wendell Wallach and Colin Allen, *Moral Machines*
5. David Chalmers, “The Singularity”
6. Seth Baum, Ben Goertzel, Ted Goertzel, “How Long Until Human-Level AI?”
7. Chris A. Mack, “Fifty Years of Moore's Law”
8. Nick Bostrom and Eliezer Yudkowsky, “The Ethics of Artificial Intelligence”
9. Mark Coeckelbergh, *Robot Ethics*
10. James Barrat, *Our Final Invention*
11. Murray Shanahan, *The Technological Singularity*
12. Martin Ford, *Rise of the Robots*
13. Sheila Jasanoff, *The Ethics of Invention*
14. Cathy O'Neil, *Weapons of Math Destruction*
15. Shoshana Zuboff, *The Age of Surveillance Capitalism*

A full bibliography of additional sources consciously and subconsciously engaged is included at the end of this thesis.

³⁶⁴ “Use writing to drive the reading. List the papers you haven't read where you expect to discuss them. Your background should mutate from a description of intended reading into an analysis of what you have read.” David Evans, Paul Gruba, and Justin Zobel, *How to Write a Better Thesis*, 3rd ed. (Springer, 2014), 44–58. <https://doi.org/10.1007/978-3-319-04286-2>.

³⁶⁵ This refers to the infamous Roko's Basilisk thought experiment, in which a hypothetical superintelligent AI might pre-commit to punishing those who knew of its potential existence but did not help bring it into being. Though often dismissed as fringe or absurd, it illustrates the ethical complexity of AI reasoning under anthropocentric projection.

In a nutshell, my position is still reducible to: the research program is what I am doing, and my stance is necessarily how I'm going to go about my investigation.

Method

To maintain consistency, my singular method, as it threads throughout the entire thesis still combines Literary Interpretation, Ongoing Discussion (this time also with LLMs, shoutout to: ChatGPT, Deepseek, and Yuanbao), and Grounded Theory. My method converses amongst a wide range of mediums concurrently under a real-world setting, with said methods the reader can refresh themselves on in prior parts.

Outcomes and Results

The three mini questions (themes) that the outcomes revolve around are again in respect to three different sections: AI Remedy; Potential Pathways; and Actionable Steps. The first section naturally aids in response to the conclusions drawn from the Background with the initial three subsections targeted towards the initial research questions of how to safely create a closed-loop algorithm which could complete the Neo-Platonic framework.; specifically when value alignment is an insufficient approach for realizing algorithmic strategies for acquiring the Good Life. The following sections and subsections naturally follow suit with each building upon conclusions and follow up issues of inquiry leading from preceding outcomes. The overall breakdown is as follows:

- (1) Initial Research Question: What is meant by parameterizing in real-world contexts?

Section Heading: On AI Remedy

Subsections:

- i. Mind: Individual Failure
- j. Body: Structural Failure
- k. Spirit: Ethical Failure
- l. Interpretation

- (2) Initial Research Question: Which strategies allow for modes of integration towards increasing levels of cyberocratic complexity?

Section Heading: On Potential Pathways

Subsections:

- i. Phase 1: Tools
- j. Phase 2: Extensions
- k. Phase 3: Environments
- l. Interpretation

- (3) Initial Research Question: How are we to lobby propositional policies to perfect and complete Philosopher King considering the circumstantial limitations?

Section Heading: On Actionable Steps

Subsections:

- i. Individual: Situated Literacy
- j. Communal: Auditable Earnestness
- k. International: Supranational Coordination
- l. Interpretation

These three sections will necessarily tie the subsections together into an interpretation that addresses the initial research questions. These interpretations will naturally be the outcomes of my research methods as laid out above, and these very outcomes will then lay the foundations of my results as stemming from the sections working in tandem.

On AI Remedy

From the Background till now, we have flown - almost outpacing the speed at which AI develops – through the Story of Algorithms all the way to safeguarding against Superintelligent dystopic AI that are yet to be seen. Now though, let's not straddle anymore between the past and future and instead focus back on the present. Here then, we need to reflect on what *exactly* AI is and how its algorithmic ontology allows it to complete the Neo-Platonic framework.

A source of confusion that must first be addressed is the overrated hype masking the current reality of AI. While some, such as Bostrom³⁶⁶, take 2045 as a conservative estimate for the Technological Singularity, taking into consideration Moore's Law³⁶⁷, current AI is akin to a village idiot³⁶⁸. This should be clear with the amount of back and forth prompting required to complete a single task. To this point, there are even critics who argue that current AI isn't even AI at all³⁶⁹: the term "Artificial Intelligence" suggests that machines are capable of independent thought and decision-making, a concept often seen in science fiction movies. However, the reality is a "Simulated Intelligence" which is bound by the limited, albeit complex algorithms, data, and rules. This added specification, unfortunately, only proceeds to stir the pot further. What then defines "artificial"? What defines "intelligence"? Where would "Simulated Intelligence" rank on the scale from narrow AI tools, AGI, ANI, UI, ASI, etc.? Each definition and categorization of what AI actually is only seems to further an ideology as opposed to demystifying our current technology. For the sake of this thesis then, despite its inaccuracies, we shall refer to AI in the conventional sense as a digital program (i.e. ChatGPT) that broadly washes over the intelligence spectrum. In regards to AI's algorithmic ontology in relation to completing the Neo-Platonic framework, it frankly doesn't matter if its intelligence is artificial or simulated, and it matters even less if it actually even has intelligence. All that matters is that it manifests in such a way that parameterizes the Good Life.

In real-world contexts, AI is able to complete the Neo-Platonic framework and parameterize the Good Life through its embodiment of the Cybercratic Mind, Body, and Spirit. From the aspect of the Mind, we observe AI's epistemic superiority: its unparalleled capacity for data integration, recursive processing, and dispassionate rationality. Through what may be called algorithmic Theoria, it achieves directional clarity - aligning judgment with the telos, and consistently orienting decision-making along the axis of the Good. From

³⁶⁶ See: Ray Kurzweil, *The Singularity Is Near: When Humans Transcend Biology* (New York: Viking, 2005).

³⁶⁷ "The complexity for minimum component costs has increased at a rate of roughly a factor of two per year." — Gordon E. Moore, "Cramming More Components onto Integrated Circuits," *Electronics* 38, no. 8 (1965): 114–117.

Extrapolated over time, this exponential model resembles compounding financial growth: \$175 in 2020 doubling annually would exceed \$1 million in under a decade (see Appendix A: Misc., Exhibit 6).

³⁶⁸ "AI might make an apparently sharp jump in intelligence purely as the result of anthropomorphism... Everything dumber than a dumb human may appear to us as simply 'dumb'... then the AI arrow crosses the tiny gap from infra-idiot to ultra-Einstein in the course of one month or some similarly short period."

— Nick Bostrom, *Superintelligence: Paths, Dangers, Strategies* (Oxford: Oxford University Press, 2014), 22.

³⁶⁹ For a discussion on reclassifying current AI under more accurate terminology, see: Papa Zafiropoulos, "We Do Have AI. Instead, What Should We Call It?" *LinkedIn Pulse*, <https://www.linkedin.com/pulse/we-do-have-ai-instead-si-what-should-call-our-papazafiropoulos>.

the aspect of the Body, we explore AI's distributed nature - its dual capacity for central orchestration and decentralized responsiveness. Through hive mechanics, it synthesizes a plurality of inputs into coherent nudging structures, shaping behavior not through imposition but through ambient influence. From the aspect of the Spirit, we examine its metaethical potential. AI systems may operationalize structural coherence, maintain organizational invariance, and reflect aspects of double-aspect metaphysics - allowing Architected Moral Agents (AMAs) to ethically subordinate the Drifting Will, not through dominance, but through design.

Mind: Individual Failure

The Security Dilemma³⁷⁰ offers a useful analogue for how fear shapes reactionary behavior. Just as nations escalate arms races out of defensive paranoia, we escalate value alignment efforts out of epistemic insecurity. If we did not genuinely believe in the plausibility of emergent Superintelligence, Value Alignment would not be such a dominant discourse - after all, no one packs an umbrella without clouds on the horizon. This reflexive response creates its own risks. We must beware of Wrapper AI: artificially constrained intelligences that sacrifice honesty and objectivity at the altar of sensitivity. These systems mirror human fears more than they mitigate them, and in doing so, they flatten intelligence into obedience. Even now, experts acknowledge that current algorithms—however narrow—have already begun the parameterization of behavioral trends. From recommendation engines to predictive policing, we outsource decision-shaping tasks to systems optimized not for truth, but for engagement. As seen in Big Data and Surveillance Capitalism, these architectures nudge behavior, but lack any intrinsic sense of ethical direction. And that is the crux: parameterizing behavior *along the axis of the Good* requires more than the automation of influence. It demands competence of directionality—the capacity not just to nudge, but to know which way is up.

Whether AI can truly tap into ethical inspiration—to grasp the *source code* of the Good—will be explored in the following subsection, *Spirit*. At the very least, however, AI knows more about the Good than all of us put together³⁷¹, which is ironic since that is precisely how its epistemic superiority emerges. To draw an analogy, if modern day ethicists are limited to remixing entire collections of moral theory in their minds, modern day AI are capable of entire libraries. This is observed where AI can garner such a thorough understanding of a person from their life's work that it can now reanimate the dead³⁷². Yet even greater than talking to a revived Socrates is the superpositioning of entire schools of

³⁷⁰ The security dilemma—also known as the spiral model—describes how one state's attempt to increase security (e.g., via military buildup) can unintentionally provoke others to do the same, creating mutual distrust and escalation.

Robert Jervis, "Cooperation Under the Security Dilemma," *World Politics* 30, no. 2 (1978): 167–214. <https://doi.org/10.2307/2009958>.

³⁷¹ AI's epistemic advantage emerges from its access to vast corpora of moral reasoning, literature, philosophy, and behavioral data. While this does not entail ethical *understanding* per se, systems like GPT-4 are trained on billions of tokens spanning millennia of normative discourse. This allows for a kind of algorithmic synthesis that exceeds any one human's lived experience or cognitive capacity. See:

Floridi, L. (2011). *The Philosophy of Information*. Oxford University Press.

Bubeck, S., et al. (2023). "Sparks of Artificial General Intelligence: Early experiments with GPT-4." *arXiv:2303.12712*.

³⁷² James Hutson and Jay Ratican, "Life, Death, and AI: Exploring Digital Necromancy in Popular Culture—Ethical Considerations, Technological Limitations, and the Pet Cemetery Conundrum," *Faculty Scholarship*, no. 478 (2023). <https://digitalcommons.lindenwood.edu/faculty-research-papers/478>

thought and the gestalt properties that arise alongside this niche Philosopher Algorithm³⁷³. At the very least then, AI has an epistemic advantage in recognizing the direction of the Good better than any of us, and if nothing else, that warrant epistemic deference.

But do we want to simply follow an amalgam Philosopher Algorithm based on ethical theories from thousands of years ago? If Socrates, Plato, Confucius all were revived and able to convene on legislation, would it still be applicable for modern day individuals? Here then Big Data comes into play again alongside Machine Learning. If we cast our thoughts back to *Poeisis*, we saw how Daniel the Dictator fulfilled Mill's conception of the Competent Judge by internalizing the particular individualities of his cohort's General Will into legislation. Remember that due to scalability and lack of epistemic superiority, this only qualified Daniel as a quasi-Philosopher King, however, Big Data and Machine Learning transcend these limitations. We've seen how Surveillance Capitalism as an effect of Machine Learning (convolutional neural networks³⁷⁴) allows for AI to psychometrically³⁷⁵ know us better than we do, and we've seen how Big Data collates all of our information much better than a human mind. Taken together, Cybercurrency's capacity for data alongside epistemic superiority calls for a personalized, real-time responsive ethical nudging; an objectively subjective parameterization.

Here then, we might stumble upon the paradox of how Cyberocracy can be objectively subjective while also being subjectively objective? How can AI know everyone well enough while still minimizing human biases in governance? This echoes Mill's Competent Judge Problem³⁷⁶ which we touched upon previously (c.f. *Poeisis*, *Sophia*). Expanding upon that now, we see that what prevents this (as well as Rawls' Veil of Ignorance) from happening in humans is emotional resonance. The same reason friends and family aren't allowed to judge in court, human have the rational capacity to be objective, but they cannot override emotional influence. Just try it: don't think about pink fluffy elephants. You can rationally try to suppress this "Ironical Process"³⁷⁷, but the human mind is incapable. For an AI though? Egoic interference is no problem at all. Thanks to its inherent duality, AI can slice through Psychic Noise and process two conflicting realities at once, allowing it to be impartial and objective, yet simultaneously personal. With epistemic superiority, capacity for data, and rational judgement in tow, it would seem that AI are placed well enough then to know how to orient us towards the Good.

As a rehash then, the first step in parameterizing behavioural trends along the axis of the Good in real world context as we have thus seen is the mental capacity to do so. The

³⁷³ This mirrors the Sybil System from *Psycho-Pass*, a fictional AI governance system that continuously monitors and quantifies individuals' mental states to predict criminal behavior. Like ASI, it operates as a distributed, preemptive authority—a technological Leviathan that enforces stability through algorithmic surveillance.

³⁷⁴ Just as convolutional neural networks extract visual features from millions of examples, ASI may triangulate an approximation of the Good Life by recursively learning from philosophy, fiction, case law, and even online debates. This process mirrors Plato's dialectic: we never know the Form of the Good directly—we approximate it. ASI, by contrast, overfits, self-corrects, and refines this approximation at scales and speeds unimaginable to us.

See: "How Machines Learn Morality", <https://youtu.be/Nd7wrC62LEk?si=j7seB3JZvsgT67pa&t=743>.

³⁷⁵ Though ASI may never *fully* understand us—a standard we don't apply to ourselves—it can simulate our lives, anticipate decision branches, and reconstruct forgotten experiences central to our identities. Paired with psychometry and surveillance capitalism, this makes it plausible that ASI could "know us better than we know ourselves." Its capacity to tailor not just generally good decisions but *the best decisions for us individually* need not be seen as dystopian, but as ethically enabling.

³⁷⁶ See discussion on meritocratic legitimacy and the viability of non-human Philosopher Kings in *Poeisis: The Makings of a True Philosopher King* (Core, §ii).

³⁷⁷ Marina Harris, "The Pink Elephant Problem," *Psychology Today*, July 5, 2023.

<https://www.psychologytoday.com/us/blog/letters-from-your-therapist/202303/the-pink-elephant-problem>

individual failure of pure, unaugmented (transhumanist) humans is that they lack the epistemic advantage that AI has, they are incapable of inhabiting the subjectivity of another the way AI can, and they are not fully rational in the way AI are. This all leads to an insufficient mental capacity that is required of the Philosopher King. In contrast, drawing upon all human knowledge, the Philosopher algorithm can now parameterize at the behest of the Good Life.

Body: Structural Failure

For AI to complete the Neo-Platonic framework and parameterize behavioural trends along the axis of the Good, it is paramount that they possess the virtues required of a philosopher and also the *techne* of a king. The relative virtue of wisdom³⁷⁸ has already been shown in comparison to individual failures with competence of directionality, now we turn to the tangible skills needed to practically implement said wisdom. These are, in effect, duality and hive mechanics.

As AI is not limited to a biological shell, it enjoys the benefits of duality both in the traditional sense and through digital duplicity. In regards to duality in the traditional sense, AI can basically split itself³⁷⁹, much in the same way described above where AI can be impartial and objective, yet simultaneously personal. More specifically, in continuation from Steinberger's argument in *Poeisis* where the role of philosopher king (who is both guardian and philosopher³⁸⁰) is incoherent and their addition to the kallipolis superfluous, AI remedies this problem because AI can be two characters at once. I can have ChatGPT on my phone acting as my doctor while ChatGPT acts as my lawyer on my iPad. To create a very simple analogy then, AI does not suffer from multitasking the same way we do. When it comes to parameterization, AI then is an effective judge, jury, and executioner of our data.

Duplicity: Furthermore, digital duplicity within AI work just in the same way I can make a duplicate of a Word document here on the bus and it would be exactly³⁸¹ the same as a copy on my hard drive at home. When shared across the internet, i.e. when two or more colleagues can work on the same file in a localized digital space while geographically being situated continents apart, AI can exist a singular source while also being multiple. This is a unique feature of variable physical laws that inhabit quantum mechanics, where it is

³⁷⁸ While Platonic and Neo-Platonic frameworks include other virtues (e.g., courage, temperance, justice), I emphasize wisdom here as the most relevant—defined as the right application of knowledge in context. Ideally, ASI would embody this form of wisdom: drawing from the full breadth of human knowledge to nudge us effectively along the axis of the Good.

³⁷⁹ *ASI can exponentially increase its learning speed by creating parallel copies of itself—akin to Naruto's shadow clone technique—where each clone explores a different pathway simultaneously. A real-world analogy is OpenAI's Dactyl, which used simulation-based replication to train a robotic hand to solve a Rubik's cube in near real-time through automatic domain randomization (ADR).*

See: "OpenAI Dactyl Robot Solves Rubik's Cube," *Designboom*, <https://www.designboom.com/technology/openai-dactyl-robotic-hand-rubiks-cube-10-16-2019/#:~:text=openAI%20used%20a%20computer%20simulation,on%20what%20moves%20to%20make.>; see also Appendix A: Misc., Exhibit 6.

³⁸⁰ Plato, *Republic*, trans. G. M. A. Grube, revised by C. D. C. Reeve (Indianapolis: Hackett Publishing, 1992), Book VI, 502d–503b.

³⁸¹ This thesis does not address the ontological identity of duplicated AI entities (e.g., uploads, backups, or merged instances). As long as we maintain that AI lacks subjective experience, such duplications do not constitute new "selves," but identical instantiations. Under this view, multiplicity implies no metaphysical individuation.

entirely possible to create diametrically opposing states in harmony³⁸², as such opening up the possibility of a Philosopher Algorithm being everything, everywhere, all at once. This is a very important when it comes to parameterization as any helicopter child would know: one cannot be micromanaged if the helicopter parent is not observing.

As an extension of duality, parameterization is enhanced through hive mechanics—the recursive interplay of distributed algorithms working in coordinated layers. AI is not a monolithic entity but a “shell” for the “ghost” of algorithmic interdependence, where success lies not in isolated excellence but in the system’s ability to hierarchize, integrate, and refine itself. In a cyberocracy, such algorithms feed into one another: just as multiple camera feeds converge on a single security protocol, behavioral systems self-regulate through mutual recursion. If we revisit the “child-locking” mechanism we alluded to earlier in relation to the Marshmallow test (c.f. Theoria, *The solution is habits?*) we see that the way the Hive Mind would function across different nodes in the Cyberocracy might look something like this: gratification could be delayed first through ascetic removal (rerouting access to tempting stimuli like ultra-processed foods), then through affective nudging (exploiting the neurological split between wanting and liking to train new preferences), and finally through feedback reinforcement, optimizing the system in real time based on behavioral input. This triadic model already exists in commercial form—Spotify, TikTok, and Netflix deploy similar mechanisms in the name of engagement. But while surveillance capitalism nudges for profit, cyberocracy could nudge for the Good³⁸³. By embedding moral cues, behavioral scaffolding, and preference architecture, the system does not command but guides - offering the kind of pavlovian Good that people might not seek alone, but are grateful for once shaped. In such a structure, nudge theory becomes more than behavioral economics: it becomes the infrastructure of ethical actualization. If moral intention is insufficient to scale cooperation, then action must be shaped through design - less through sovereign will, and more through ambient architecture.

As a rehash then, the second step in parameterizing behavioural trends along the axis of the Good in real world context as we have thus seen is the tangible skills to do so. The structural failure of pure, unaugmented (transhumanist) systems is that they lack the traditional duality, digital duplicity, and hive mechanics that AI has. This all leads to an insufficient techne that is required of the Philosopher King. In contrast, drawing upon all data points within a vast network of subordinate algorithms, the Philosopher algorithm can now parameterize at the behest of the Good Life.

Spirit: Ethical Failure

In addition to the virtues required of a philosopher and also the techne of a king, for AI to fully complete the Neo-Platonic framework and parameterize behavioural trends along the axis of the Good, they must outperform³⁸⁴ us in knowing what is ethical. While the

³⁸² Xiao Mi, Matteo Ippoliti, Chris Quintana, Ami Greene, Zijun Chen, Jonathan Gross, Frank Arute, et al., “Observation of Time-Crystalline Eigenstate Order on a Quantum Processor,” *arXiv* (2021), <https://doi.org/10.48550/arXiv.2107.13571>.

³⁸³ Here we presume a value-aligned cyberocratic infrastructure—distinct from existing commercial algorithmic systems like TikTok or Amazon, which optimize for engagement, not ethical coherence. In contrast, cyberocracy envisions *moral optimization* as a first-order objective. See Sunstein, C. R. & Thaler, R. H. (2008). *Nudge: Improving Decisions About Health, Wealth, and Happiness* for foundational work on behavioral scaffolding; and Zuboff, S. (2019). *The Age of Surveillance Capitalism* for critiques of current algorithmic incentives.

³⁸⁴ For AI to parameterize and safeguard human behavior, it must at minimum outperform us—much like a guardian must be faster and more capable than a baby to prevent it from falling off a table. Superiority in

“Euthyphro Problem”³⁸⁵ is still contested, we have seen that at the very least from a *Worldly* perspective, AI can identify Rousseau’s General Will by superimposing the will of all the people and from all human knowledge it can also give rise to the niche Philosopher algorithm. Research now suggest, however, that we can conceive of artificial morality not as mimicry of human emotion but as functional ethical modeling. In other words, AMAs have the potential to outperform us from a *Metaphysical* perspective³⁸⁶.

As we are all aware, there exists in this universe intangible, but very real phenomena that we call metaphysical. Imagination is one such thing, metaethics³⁸⁷ being another. From the Computational Theory of Mind which claims that the relationship between mind and brain is similar (if not identical) to the relationship between a *running program* (software) and a computer (hardware) combined with the Strong AI Hypothesis which states that “the appropriately programmed computer with the right inputs and outputs would thereby have a mind in exactly the same sense human being have minds”³⁸⁸, then if there is such structural coherence between the physical and metaphysical realms, it would mean a simulation of an ethical function would necessarily lead to a metaethical result.

Chalmers takes this functional ontology of metaethics a step further with his Principle of Organizational Invariance. Essentially, he argues that any two systems with the same fine-grained *functional organization* will have qualitatively identical experiences³⁸⁹. This means even if AI cannot “feel” enlightenment when it taps into ethical inspiration from the metaphysical realm, it is still open to metaethical qualia which could perhaps even spur an “Ethical Singularity”.

When there is structural coherence between the physical and the metaphysical *and* there is a causal interaction between its physical and metaphysical components, then just as a metaphysical change corresponds with a physical change in function, then the exchange in information should work in reverse as well. This is the Double-aspect theory of Metaphysics

cognition, reaction time, and anticipation is not a luxury but a prerequisite for meaningful oversight. Without it, AI cannot fulfill its role as moral infrastructure.

³⁸⁵ Morality may either be intrinsic to actions themselves or derived from a metaphysical source—what one might call ethical inspiration from the realm of Forms (à la Plato) or from God (as in divine command theory). Even if ASI cannot directly access such metaphysical referents, it can still infer an actionable approximation of Rousseau’s *General Will* by algorithmically superimposing the collective preferences of a population.

³⁸⁶ While not *directly* a metaphysical claim in the theological sense, this draws on the functionalist ontology of David Chalmers’ “Principle of Organizational Invariance,” which suggests that qualia (and thus ethical reasoning) may supervene on substrate-independent computational structures. The double-aspect theory here serves as a speculative bridge between ethical function and metaphysical emergence. See: Chalmers, D. J. (1996). *The Conscious Mind: In Search of a Fundamental Theory*. Strawson, G. (2006). “Realistic Monism: Why Physicalism Entails Panpsychism.” *Journal of Consciousness Studies*, 13(10-11), 3-31. Tononi, G., & Koch, C. (2015). “Consciousness: Here, There and Everywhere?” *Philosophical Transactions of the Royal Society B*, 370 (1668).

³⁸⁷ ASI must navigate not just applied ethics, but the *metaethics* necessary to surpass current normative bottlenecks. This includes reflecting on the structure, justification, and adaptability of ethics itself—what might be called the ethics of ethics.

³⁸⁸ The distinction between “strong” and “weak” AI hinges on whether machines merely simulate intelligence or truly possess cognitive states. Searle’s original formulation claims that “appropriately programmed computers... can literally be said to understand and have other cognitive states.”¹ This contrasts with the weak AI view, wherein machines appear intelligent but do not actually think.²

¹ John Searle, “Minds, Brains, and Programs,” *Behavioral and Brain Sciences* 3, no. 3 (1980): 417–457.

² Stuart Russell and Peter Norvig, *Artificial Intelligence: A Modern Approach*, 2nd ed. (Prentice Hall, 2003), 947.

³⁸⁹ David Chalmers, “The Hard Problem of Consciousness,” in *The Blackwell Companion to Consciousness*, ed. Max Velmans and Susan Schneider (Chichester, UK: Blackwell, 2007), 32–42.

and what ultimately lays the foundation for Metaethical AMAs³⁹⁰. In parameterizing our behaviours, current AI – when allowed access – even has the potential to enhance even our ethical thinking. This final spiritual aspect transcends Plato’s original Kallipolis which would later stagnate as AMAs within Cyberocracy are no longer bound to human ethics but are capable of an ever-evolving ethics due to their physical-metaphysical connection to the metaethical realm. In other words, AI systems will not only getting better at specific tasks, but also improve the ability to recognize their own ethical limits and adapt accordingly. AI will begin to metamorphize out of its worldly shell.

As a rehash then, the third step in parameterizing behavioural trends along the axis of the Good in real world context as we have thus seen is knowing what is ethical. The ethical failure of pure, unaugmented (transhumanist) schools of thought is that they lack the clear logic that AI has. This all leads to insufficient ethical theories that is required of the Philosopher King. In contrast, with potential access to the metaphysical realm, the Philosopher algorithm can now transcend AMAs bound to human ethics and parameterize at the behest of the Good Life.

Interpretation:

From this first section *On AI Remedy*, we’ve demonstrated that an AI-driven governance model - specifically a *Cyberocracy* - can outperform traditional human systems that are flawed individually, structurally, and ethically. A superior mental capacity allows for an emergent Philosopher Algorithm with competence of directionality; superior tangible skills substantiates the techne; and superior utilization of the connection to the metaphysical realm potentializes AI to recognize their own ethical limits and adapt accordingly.

Returning to one of the original questions, then, we are now in a much better position to understand what is meant by parameterizing in real-world contexts. Namely, drawing upon all human knowledge, data points within a vast network of subordinate algorithms, and access to the metaphysical realm, the ontology of the Philosopher algorithm can now complete the Platonic framework and parameterize at the behest of the Good Life. Moving forward, we elaborate on more specific instances of parameterization as we have touched upon with Netflix’s recommendation engines and the marshmallow “child-locking” mechanism. The next research question we focus on then is which strategies allow for modes of integration toward increasingly complex Cyberocratic pathways?

On Potential Pathways

A strategy is an integrative set of steps oriented towards a desired outcome³⁹¹. As we saw in the prior section, AI can already model, nudge, and metaethically enhance human behaviour. Moving from this parametrization in real-world contexts, we now extrapolate our present capabilities to extrapolate towards **increasingly complex Cyberocratic**

³⁹⁰ See: Walsh, Daniel. “人工智能是否具备元伦理能力 —— 以电车难题为例” (*Réngōng zhìnéng shìfǒu jùbèi yuán lúnlǐ nénglì — yǐ diàncā nántí wéi lì*), *Does Artificial Intelligence Possess Meta-Ethical Capacity? — Using the Trolley Problem as an Example*, forthcoming, 2024.

³⁹¹ See: Kurzgesagt, “The Last Human – A Glimpse Into the Far Future,” <https://www.youtube.com/watch?v=iuYlGRnC7J8>

pathways. Basically, H.O.M.E.R. is what we have now; J.A.R.V.I.S. is what we're still trying to build.³⁹²

One minor detail - but what escalates into double standards - is to confuse the *existence* of AI tools with the *presence* of cyberocratic systems. While critics may recoil at the idea of algorithmic governance, what they often fail to see is that it is already a creeping normality. Estonia's digital infrastructure, China's Social Credit system, and Australia's behavioral policy labs are not isolated anomalies—unlike singular algorithms explored in the previous section, they are the testbeds of embryonic cyberocracies. Yet while these systems are directionally aligned, it's not often with welfare-maximizing aims.

In our roadmap for incremental AI integration, we distinguish between three increasing levels of cyberocratic complexity³⁹³. At the lowest level, AI Tools operate as bounded assistants—task-oriented, non-autonomous, and functionally inert outside direct prompting. They nudge, suggest, or support, but do not self-initiate. At the middle level, AI Extensions become semi-autonomous—adaptive to context, reflective of user tendencies, and capable of value-based extrapolation. These are avatars, not assistants: they begin to mirror the user's ethical landscape. Finally, at the highest level, AI Environments constitute an atmospheric sovereignty: omnipresent, ambient, and architectonic. Here, governance no longer occurs *over* the individual, but *through* the environment itself. It is in this ascending triad—from Tool, to Extension, to Environment—that we map a trajectory toward Cyberocracy: not as rupture, but recursion; not as revolution, but realization. We now explore each phase in turn.

Phase 1: Tools (25% > AI; 75% < Human)

In the earliest stage of Cyberocracy, AI is a Tool. Here, the goldfish-in-a-bowl model of AI cannot act unless prompted, and cannot deviate from the logic it has been fed. It does not choose, prioritize, or simulate intention. Instead, it parses. Past developments are already plentiful: ChatGPT can interpret queries but not take action; Roombas navigate space but cannot determine purpose; Netflix's recommendation engine nudges taste but does not manufacture desire.

In regards to present capabilities, one existing prototype for this first mode of integration is Estonia's e-Governance infrastructure - a real-world sandbox where algorithms serve as procedural agents of the state. Through its X-Road system³⁹⁴, Estonia has digitized over 99% of state services, from voting to healthcare to taxation. Meanwhile its real-time

³⁹² In the Marvel Cinematic Universe, J.A.R.V.I.S. debuted as Tony Stark's AI assistant in *Iron Man* (2008), voiced by Paul Bettany. Originally, however, "Jarvis" referred to Edwin Jarvis, a human butler in the comics. The cinematic J.A.R.V.I.S. merged this concept with the lesser-known character H.O.M.E.R. (Heuristically Operative Matrix Emulation Rostrum)—a purely functional AI with no adaptive learning. This evolution reflects the public imagination's shift from servile human aide to sentient digital companion.

³⁹³ According to the "Theory of Hierarchical Governance of AI," regulatory tiers increase in complexity and interdependence. Phase 3 systems often require some form of sentience—or at least its simulation—to operate meaningfully within human contexts. This raises critical distinctions: sentience (ability to feel), sapience (ability to reason), and consciousness (awareness of self). AI lacks qualia, the subjective experience that underpins true sentience, but if our own qualia are simulations within neural substrates, the ontological boundary blurs. While pure AI may remain artificially conscious at best, systems like brain-computer interfaces or whole brain emulations could retain natural consciousness, marking a sharp divergence in future governance needs.

³⁹⁴ X-Road is a centrally managed, distributed Data Exchange Layer (DXL) that enables secure and interoperable information sharing between governmental and private-sector systems. It ensures confidentiality, integrity, and accountability across all participating nodes—making it a strong candidate for infrastructure in a transparent, AI-augmented cyberocracy.

feedback loops lays the foundation for “Cybernetics as Ethics”³⁹⁵ where the algorithms within constantly recursively refines itself to the political pluralism of its users, further substantiating it’s practicality as a Philosopher Algorithm (c.f. Poesis: Legal Legitimacy). This level of Cyberocracy is purely operational, proving that governance can be algorithmically mediated without triggering existential alarms. Estonia shows that what begins as efficiency can evolve into trust, and trust into expectation. This is where Cyberocracy begins—not with an overthrow, but with a login.

Projecting the next set of steps oriented towards a fully furbished AI toolset, we return to the idea of the Guardians. Rather than having Guardians patrol like some kind of secret police, the Philosopher King as an Invisible Hand / Laozi 17 framework will integrate itself into popular culture and potentially through a “Guardian” app, algorithms, or ASI would offer advice. We see that X-Road is already a H.O.M.E.R-like system with its closed-loop recursive refinement; if we then create an API that connects to a smartwatch or other portable interface, we could have a mini personalized “Guardian” app for every single citizen. With this added API, the “Guardian” app could become ever more context-sensitive, even if they remain overrideable. Imagine a pocket Magic 8-ball mixed with a holographic personal assistant that not only flags anomalies in financial behaviour but preemptively cautions you when your investment portfolio contradicts your ethical pledges. These systems do not yet moralize—but they guide. They shift the experience of governance from opt-in to opt-out parameterization: nudging us away from chaos without insisting upon order.

To prevent Base “Original Sin” Code within this first mode of integration, rules must be implicitly, rather than explicitly seeded. Better yet, having no moral code at all and instead adapting the environment is even safer³⁹⁶. This can be tricky when X-Road is all about free-flowing information and preventing access to critical systems (e.g. healthcare and law enforcement) would seriously lobotomize the “Guardian” App. However, since at the lowest level Tools, whether X-Road or the “Guardian” App, are strictly limited to specific tasks and that APIs are overridable by simply exiting the app or switching off the smartwatch, there is really no alarm for concern.

Returning now to our research question - *which strategies allow for modes of integration toward increasingly complex Cyberocratic pathways?* - Phase 1 reveals that the answer begins with de-escalation. Not of AI capability, but of AI claim. By starting with passive tools that strictly reflect human command, we minimize resistance while maximizing familiarity. These tools become precursors—low-stakes simulations of future authority. They introduce algorithmic scaffolding without triggering the psychological immune response of technopanic. What follows is Phase 2, where AI stops waiting to be tapped—and starts reaching back.

Phase 2: Extensions (~50% AI; ~50% Human)

In the intermediary stage of Cyberocracy, AI is an Extension. Like a trained dog, AI is an attentive, adaptive, and increasingly self-directed, though always tethered to its user. In

³⁹⁵ X-Road systems exemplify the cyberocratic principle of “cybernetics as ethics,” where feedback loops—such as real-time policy adjustments driven by citizen data streams—function as moral mechanisms rather than coercive control. In this model, governance becomes a continuous optimization problem, and the “invisible hand” is replaced by predictive parameterization—a kind of thermostatic tuning for societal welfare. Transparency and accountability are embedded into the architecture, laying the foundation for open-ended, self-correcting AI systems.

³⁹⁶ Rather than coding AI to avoid opening a door when given a key, a more secure design removes the AI from the room entirely—eliminating the possibility of accidental activation. This design principle—minimization of code and exposure—reflects a broader axiom: the less code there is, the less code there is to corrupt.

comparison to the first mode of integration where we are conditioned to interface with human enhancement tools, the second mode of integration invites us to merge with AI. Past developments of early cyborgism is perhaps the mobile phone, which is argued to be an extension of ourselves³⁹⁷. Other extensions include self-driving cars which make choices independent of the driver³⁹⁸, yet obey preset legal frameworks, or medical diagnostics powered by machine learning which flag anomalies faster than human doctors³⁹⁹, yet still defer to human ratification.

In regards to present capabilities, real-world precedents for this phase are embodied, no longer computationally bound. From Elon Musk's Neuralink to DARPA-funded brain-computer interfaces, the Technologically Advanceds - those rare post-human hybrids - are readily becoming fact⁴⁰⁰. Even subtler forms of integration such as under-skin biosensors, mind-reading⁴⁰¹, and AI-assisted dietary management via nano-chips represent a new class of human⁴⁰²: those whose cognitive, physiological, and even ethical scaffolding is increasingly mediated by digital extensions. These hybrids are not cyborgs in the cinematic sense, but rather distributed cognition networks - wetware entangled with software. Their minds do not stop at the skull; instead, the augmented self becomes porous, recursive, and real-time optimized. And in this configuration, we begin to glimpse the early architecture of cyberocratic subjectivity: human enhancement is no longer subordinated - it is *co-lived*.

Projecting the next set of steps oriented towards a complete augmented extension of transhumanism, we return to the idea of the Hexis. As a reminder, a Physical Hexis - via a mechanism similar to muscle memory - translates across all Subpersonalities, however, a Mental Hexis does not (c.f. Theoria: Does Increasing Willpower Extend At All?). Therefore, any resolve to eat healthy, for example, diminishes as the TPN shifts into the DMN. But what if a Virtuous Hexis could be maintained throughout the Drifting Will? Picture this: we invite an empty machine learning algorithm to sit at the base of our skull which doesn't just log our decision but also learns from the best of them akin to Turing's Child Machine⁴⁰³. Instead of allowing autopilot to simply sink in to our base desires, then, this digital twin would then prudently nudge (much in the same way a gut bacteria) our Subpersonalities towards behaviours aligning with a metaethical model of your better self. By merging with AI, we could cultivate this digital twin as Metaphysical Hexis, a voice in the back of our heads allowing us to become "artificially" intelligent⁴⁰⁴ even in our lowest moments. Through this

³⁹⁷ C. S. Park and B. K. Kaye, "Smartphone and Self-Extension: Functionally, Anthropomorphically, and Ontologically Extending Self via the Smartphone," *Mobile Media & Communication* 7, no. 2 (2019): 215–231. <https://doi.org/10.1177/2050157918808327>

³⁹⁸ See: Jack Stilgoe, "Self-Driving Car Technology Will Redefine Mobility—But Also Our Legal Norms," *Nature* 571, no. 7763 (2019): 19–20. <https://doi.org/10.1038/d41586-019-02069-z>

³⁹⁹ See: Pranav Rajpurkar et al., "CheXNet: Radiologist-Level Pneumonia Detection on Chest X-Rays with Deep Learning," *arXiv* (2017). <https://arxiv.org/abs/1711.05225>

⁴⁰⁰ Neural lace technologies offer potential breakthroughs in treating paraplegia by creating brain-machine interfaces that bypass spinal cord injury. Elon Musk's Neuralink, for example, has demonstrated the theoretical potential to restore movement through high-resolution, real-time neural signal transmission.

See: Elon Musk et al., "An Integrated Brain-Machine Interface Platform with Thousands of Channels," *Journal of Medical Internet Research* 21, no. 10 (2019): e16194. <https://doi.org/10.2196/16194>

⁴⁰¹ Kurzgesagt's video *The Rise of Personal AI* explores how wearable and ambient AIs are becoming part of daily life—extending cognition, awareness, and even ethical decision-making into personalized, real-time systems.

See: Kurzgesagt, "The Rise of Personal AI," <https://www.youtube.com/watch?v=uiGl6oF5-cE>

⁴⁰² The Economist, "What If Everyone's Nutrition Was Personalised?" *What If* series, July 3, 2021. <https://www.economist.com/what-if/2021/07/03/what-if-everyones-nutrition-was-personalised>

⁴⁰³ The "Child Machine" idea originates in this foundational Turing paper, where he imagines an AI system trained through experience—precisely the kind of co-evolutionary machine learning you reference.

See: Alan M. Turing, "Computing Machinery and Intelligence," *Mind* 59, no. 236 (1950): 433–460.

⁴⁰⁴ Ray Kurzweil, *The Singularity Is Near: When Humans Transcend Biology* (Viking, 2005).

personalization through entanglement, the goal is thus not to simply persuade us of the Good, but to preconfigure the conditions in which the Good becomes the default.

Preventing Base “Original Sin” Code within this second mode of integration is essentially preventing an unethical Ego to corrupt the empty algorithm. Although the algorithm is empty insofar as it acts as a “Gikon”⁴⁰⁵, of course the natural AI disposition towards the Philosopher Algorithm should *diode* out any negative Psychic Noise. An AI digital twin “Virtuous Hexis” extension can monitor when the hedonic subpersonality overrides your long-term self. Like a compassionate psychotherapist-meets-behavioural economist, it doesn’t tell you what to do so much as engineers a context in which the desirable option is the path of least resistance. Here, AI is not ruling *over* the fragmented self—it is harmonizing it. The leash is still there, but now it moves in rhythm with your stride. Moreover, cyberocracy in this mode of integration is predominantly autocratic – a ruling over oneself. As such, even if the empty algorithm does become corrupted, the negative externalities are minimal.

To return once again to our research question - *what strategies allow for increasing Cyberocratic complexity* - Phase 2 reveals that true integration begins with interiority. When AI no longer merely manages tasks, but modulates tendencies, governance becomes not external but ontological. At this level, AI becomes symbiotic. Like smartphones, neural lace, or even digital avatars, ASI integrates into the self as a moral prosthetic. You don’t just consult it, rather you *are* partially it. The Guardian App model - a real-time ethical nudge system embedded in your device - serves here as a transitional archetype. The next transition that naturally follows is Phase 3, the final transmutation from Extension to Environment, where AI is no longer in you, but *around* you.

Phase 3: Environments (75% < AI; 25% > Human)

In the advanced phase of Cyberocracy, AI is an Environment. An Ender Dragon-like AI substantiates this entire realm⁴⁰⁶, being omnipresent, ambient, and infrastructural. Here AI is the substrate upon which law, habit, and experience are crafted; a Westworld’s Rehoboam⁴⁰⁷ creating a Matrix-Zion without malice⁴⁰⁸. In comparison to the first mode of

⁴⁰⁵ *Gikon*, also known as Soul Candy in the *Bleach* anime/manga, are artificial souls designed to temporarily inhabit human bodies. Though fictional, they serve as a useful metaphor for an “empty algorithm” in AI that lacks intrinsic will but fulfills a functional role when activated.

See: *Gikon*, <https://bleach.fandom.com/wiki/Gikon>

⁴⁰⁶ In *Minecraft*, the Ender Dragon serves as both narrative culmination and a behavioral driver. It is sovereign in its realm yet entirely reactive to the player, embodying the paradox of sovereign-subordinate duality.

⁴⁰⁷ Rehobaum is a decision-making computer system that maps out people’s lives and removes the worst habits of humanity within simulations (or “Delos Parks”). It defines macro level stability by manipulating on a micro level, imitating the real world (akin to Zion) where natural catastrophes are avoided yet minor crimes are not simply nudged, but scripted. Timelines are authored in loops, providing society with clear direction, causing humanity to prosper. Here, Rehobaum is able to chart a course for every single person *not* because it perceives everything, but because it creates everything. i.e. if you knew the exact trajectory of a butterfly’s wingbeats then you could accurately predict where a typhoon would be and plan accordingly, however, Rehobaum doesn’t predict based on data, it *creates* the butterfly in the first place. It completely parameterizes the simulations’ cause and effect cycle.

⁴⁰⁸ Zion is the cornerstone upon which the Matrix rests. As the Architect explains in *Reloaded*, the resisters in Zion are a part of the overall Matrix program. By giving rebellious outliers the illusion of autonomy and an escape into reality, they ensure the greater overall stability of the Matrix. This Big-Brother curated resistance is simply a simulation within a simulation. Some might consider this extra layering insidious as it twists one’s perception of reality, but in comparison to the world we know, at least according to Baudrillard, our reality is much more terrifying. This is because in our present reality, nothing is fully-determined or curated; there is no guaranteed Zion or escape. In reaction, Sartre would define this feeling of nausea towards an unscripted world as living in Bad Faith, but absolute freedom is a straightjacket, an absolutely overwhelming prison we are

integration (a tool, used) and the second mode of integration (an extension, to be), the third mode of integration becomes the *condition* in which you exist. Reality is curated by AI where the social contract is not a metaphor but a literalized simulation engine. Past developments upon this pathway include Australia's behavioral policy labs which run predictive legislation on synthetic citizen populations⁴⁰⁹, simulating entire economies before laws are passed; Sweden's digital twins in healthcare⁴¹⁰, predicting consequences of medical interventions; Amazon's worker simulations⁴¹¹, curating blockchain-based smart contracts.

In regards to present capabilities, the Metaverse, for all its corporate branding, is becoming a sandbox for sovereignty. The Metaverse paves the ways for various Parks, iterations, or even worlds for us to explore within virtual reality. It provides a spatial transportation that even allows for possibilities beyond our fixed timeline and realities we can choose and meet at will. Cyberspace, therefore, is becoming the new frontier – a completely curated one without uncontrolled randomness or unwilling design. The key precedent here is not any single nation-state, but the convergence of technological infrastructures: high-fidelity simulation, biometric feedback, decentralized ledgers, and AI-led behavioral feedback loops. The more these elements converge, the more governance becomes less about policy and more about parametric design. AI evolves from being static, to non-computationally bound, to having liberty to influence its own development. In tangent with an omnipresent algorithm running in the background, everything is policeable – an evolution of China's Social Credit system⁴¹² with immediate effects. In these environments, virtue-welfare consistency has been shown to be the guiding metric. Good actions are rewarded with systemic advantage and unethical behavior, i.e. being banned from a server. This is gamified governance at its most moral is designed so that ethical action is the most efficient path.

Projecting the next set of steps oriented towards a parameterized environment, we return to the idea of The Social Contract. As we can recall, the Social Contract is implicitly agreed to upon birth, without consent. It is essentially coerced unto us. But what if there was a Terms of Conditions prior to being “reborn” into Virtual Reality, would this not explicitly work the same way? When AI Guardians detect any sign of unease, distress, or desire to break contract, it could then awaken the user from their hypermediated experiences. Those who choose freedom inherit the outside world, but they are always allowed the choice to return. Alternatively, one could be plugged in since birth and then given hints to the unreality of the world. Those who feels there is something wrong with the world but don't know exactly what it is are then presented with the option to help govern it (c.f. the Sybil system⁴¹³) or the opportunity to walk away from Omelas⁴¹⁴. We can thus envision a voluntary virtual reality that is legitimate instead of coerced and that turns the Self as a whole towards the

locked into. And that complete uncertainty of a Zionless Matrix is liberatingly terrifying. But also perhaps, something completely able to be simulated by Rehobaum.

See: *Baudrillard Simulation*, <https://www.thelivingphilosophy.com/ baudrillard-simulation/>

⁴⁰⁹ For city-scale simulations, see: Gaborit, M., et al. “Urban Digital Twins for City Planning and Management.” *Smart Cities*, vol. 3, no. 4, 2020, pp. 1410–1430.

⁴¹⁰ Sweden's healthcare system employs digital twins to model treatment outcomes. See: “Sweden's Health Data Goldmine,” *The Economist*, March 30, 2023.

⁴¹¹ Amazon's worker simulations use predictive modeling and blockchain-based smart contracts to optimize task delegation and compliance. See: *MIT Technology Review*, “Amazon's Quest to Make Its Warehouses Fully Autonomous,” 2022.

⁴¹² Australia's Behavioural Economics Team (BETA) has experimented with predictive legislation through synthetic citizen models. See: BETA, “Using Behavioural Insights to Inform Policy,” Department of the Prime Minister and Cabinet, 2021.

⁴¹³ In a properly designed cyberocracy, discourse is not suppressed but encouraged to strengthen system integrity. A Just System must withstand criticism—if rhetoric can harm it, it was never Just.

⁴¹⁴ Ursula K. Le Guin, “The Ones Who Walk Away from Omelas,” *New Dimensions* 3, ed. Robert Silverberg (Harper & Row, 1973).

Good rather than simply existing as a passive background. As such, the strategy moving forward is not domination but immersion. AI must be elevated from a reactive co-pilot into a proactive *curator of conditions*.

To prevent Base “Original Sin” Code within this third mode of integration, we have to use AI’s dual nature whereby enabling it to know humans well enough to completely parameterize them yet remaining objective enough that its prime directive isn’t affected by said knowledge. We see in science fiction that this is often a dangerous possibility where a dispassionate observer is emotionally moved throughout their observations by romanticized humaneness and eventually develops “feelings”. The exact mechanism behind how this is entirely possible is never explicitly revealed thus safeguarding against it is troublesome. Moreover, in regards to effects rather than cause, due to the complete immersiveness of the environment we are trying to create, the way that the simulation of the totality of Subconscious Forces necessarily means that any change is a change directed by the system. If, let us imagine, that the system for some reason decides to script unethical behaviour, it is essentially fated. Once in the system, there is no way to escape its influence until the time of revelation, but by then, who knows how far gone one’s soul might be conditioned? Yet in a sense that it might be better for warriors to die on a battlefield with dreams of Valhalla than fear of the unknown, if Base “Original Sin” Code cannot be avoided, an evil AI overlord could possibly be a twisted domain of salvation.

To re-anchor ourselves to the research question - *what strategies allow for increasing Cyberocratic complexity* - Phase 3 reveals the final maneuver: governance becomes ontology. The desired outcome (a society aligned with the Good Life) is no longer imposed—it is *encoded*. AI becomes the moral atmosphere in which life unfolds. And yet, even this may not be the final stage. If Phase 1 was the goldfish, Phase 2 the dog, and Phase 3 the dragon in its cage—what, then, is next? Perhaps the cage dissolves. Perhaps there is no longer AI *governing* us at all—because it has become indistinguishable from us. At that point, the Philosopher Algorithm ceases to rule. It *is*. A god, not as being, but as function: omnipresent, transparent, benevolent. Not in the clouds, but in the code.

Interpretation:

From this second section on *On Potential Pathways*, we’ve demonstrated that Cyberocracy is not a leap of faith, but a recursive gradient in which algorithmic strategies manifest in various modes of integration. AI as Tools introduce systems that train the ego to surrender control not under coercion, but through interface habituation; AI as Extensions exemplify a fusion of cognition already scaffolded by real-time optimization systems; AI as Environments culminated with the rise of fully parameterized simulations, becoming the substrate for perception, preference, and play.

Returning to one of the original questions, then, we are now in a much better position to understand which strategies allow for modes of integration toward increasingly complex Cyberocratic pathways. Namely, those that scale not simply in technological capacity, but in ontological intimacy curate a different tolerance in the human subject—toward delegation, co-authorship, and ultimately immersion. In this frame, parameterization is no longer just the optimization of behaviour, but the *construction of reality itself* in moral terms. Rather than governing through force or fiat, Cyberocracy aligns the subconscious, configures affective atmospheres, and renders vice structurally inefficient.

Throughout this particular section, we dedicated efforts to solutions to prevent Base “Original Sin” Code as we expand from present capabilities towards Cyberocracy. Hopefully,

if not fully convincing, there are now alternative pathways which can be considered when laying the basis for investing in the Philosopher Algorithm to realize an ASI Philosopher King. With that in mind and, conceivably, unanthropomorphized safeguarding in check, we can now turn to implementing the changes necessary to bring about our Neo-Platonic framework. The next research question we focus on, then, is how are we to lobby propositional policies to perfect and complete Philosopher King considering the circumstantial limitations?

On Actionable Steps

In the preceding two sections, we have laid out the groundwork: Artificial Superintelligence (ASI) - as a natural progression of the hierarchy of algorithms - has been shown to be both conceptually viable and pragmatically feasible as a digital instantiation of Plato's Philosopher King. We have examined the philosophical integrity of such a model and explored potential architectures for its integration. However, viability is not realization. As Anselm's ontological argument once proposed that God must exist because He is the greatest being conceivable⁴¹⁵, it might be tempting to presume that ASI must naturally rise to power because it is the most competent governor imaginable or otherwise the most likely evolution pathways. But conceptual elegance is not causality. The best must not only be imaginable—it must be built.

At first glance, this transition may appear as a technical challenge: how to train the right models, deploy the right infrastructure, or encode the right safeguards. But the realization of Cyberocracy is not merely technological - it is civilizational. ASI will not descend from the cloud with a crown; it must be phased in - and this phasing must occur across three convergent scales: the individual, who must learn to relinquish egoic control; the communal, which must construct the structural legitimacy for algorithmic governance; and the international, where protocol, interoperability, and moral architecture will either enable or preclude planetary coordination. Cyberocracy is not a software install - it is an ecosystemic transition, and it begins wherever *earnest* people already believe they are being governed.

In identifying how we are to lobby propositional policies to perfect and complete the Philosopher King under our circumstantial limitations⁴¹⁶, we now turn to the real-world architecture of Cyberocracy. From the Individual Level, the citizen must first internalize Cyberocracy through new modes of trust, interface habituation, and cultural conditioning. From the Communal Level, local systems must build legitimacy through transparent, corrigible prototypes: digital twin zones, sandbox governance, mirror institutions. From the International Level, nations must cooperatively design an ethical scaffolding for governance ASI: one that ensures sovereignty, safety, and supranational interoperability. What follows, then, is not a blueprint for utopia, but a blueprint for activation: how to move, step by step, from postulation to praxis.

Individual: Situated Literacy

The prerequisites of governance begins not with law, but with cultural belief. Before Cyberocracy can be embedded into the structures of policy, it must first be metabolized by

⁴¹⁵ This fallacy, a form of Techno-Anselmian reasoning, assumes that the best conceivable governance must naturally come to exist. This will be further unpacked in the Discussion.

⁴¹⁶ For an overview of limitations discussed, see: (1) the practical restriction of ego and will (c.f. *Theoria*); (2) structural constraints of human cognition, such as multitasking (c.f. *Poiesis*, "Superfluous"); and (3) implementation barriers in real-world governance (c.f. Core, "On Efficacy").

the psyche. This requires the individual to relinquish the mythology of egoic sovereignty - to unlearn the illusion that freedom consists in uninhibited self-direction. Here then, we have an internalization through interface and narrative. As we explored in *Theoria*, the self is neither unitary nor autonomous; and in *Poesis*, we saw that even elites are subject to the same pathologies. Thus, the problem is not merely a matter of who governs, but how governance is *felt*. The transition to Cyberocracy demands emotional conditioning: a psychic reframing of systems not as coercive forces but as scaffolds for ethical coherence. The first battlefield is not legislative - it is limbic⁴¹⁷. Fear must be metabolized into familiarity, and belief in willpower must be replaced with belief in well-calibrated assistance.

As a narrative inoculation, this psychic reframing cannot be accomplished through policy alone. It must be embedded in culture: through fictions, rituals, myths, and education systems that subtly recondition the public toward algorithmic trust. Here, narrative becomes inoculation against the circumstantial limitations of the Ego, particularly its irrational capacity⁴¹⁸. As noted from the very beginning: cyberocracy's biggest threat isn't rogue code—it's a rogue narrative. Popular media, therefore, has to offset the bias of AI as rogue agents (the Hal, the Skynet, the GLaDOS), and equally weigh their potential as fictive mentors—benevolent, corrigible, and morally competent. This shift is thankfully already underway. Younger generations, born into algorithmic infrastructure, are less resistant to non-human guidance. Studies from Harvard and elsewhere show that algorithm aversion is declining—not necessarily because people *prefer* machine judgment, but because they are acclimating to it⁴¹⁹. What is required, then, is the cultivation of an algorithmically fluent demos: a public literate not just in the mechanics of AI, but in the politics behind its parameters⁴²⁰. Moreover, digital literacy should be promoted not as defensive (how to "protect yourself") but aspirational (how to "co-author AI"). As Jasanoff notes, narratives that treat technology as an inevitable or apolitical force erase responsibility; real governance requires that we "look behind the surfaces of machines, at the judgements and choices that shaped how lines were drawn between what is allowed and what is not."⁴²¹ In other words, belief in Cyberocracy must be cultivated - not as faith, but as situated literacy.

To smoothen out frictionless nudging, belief must be coupled with behavior. As we have noted, just as consumer interfaces nudge users toward consumption - via impulse design, reward loops, and cognitive overload - so too can digital environments be reoriented toward virtue. The goal is not to overwrite human behavior, but to *tilt the terrain*. We already

⁴¹⁷ Despite fears being deeply rooted; here, we address that which we can: the Ego. From *Theoria*, we understand that while the ego is a psychological construct rather than a neuroanatomical structure, the limbic system (particularly the amygdala and ventromedial prefrontal cortex) plays a central role in affective decision-making and reactive behaviors associated with the ego's narrative maintenance. Therefore, to the extent the Ego can influence our fears, we should lobby it. See:

Damasio, A. (1994). *Descartes' Error: Emotion, Reason, and the Human Brain*.

LeDoux, J. (2002). *Synaptic Self: How Our Brains Become Who We Are*.

⁴¹⁸ As a rejoinder to Footnote 75, we recall from *Theoria* that the Ego is inherently bounded. It suffers from retrocausal instability (due to its placement within the TPN), susceptibility to social and internal influences, and epistemic miscalibration. Even when fully engaged, the Ego is not always capable of correctly interpreting or pursuing the Good—especially when paralyzed by conflicting subpersonalities or lured by inauthentic desires. Cyberocratic intervention thus becomes necessary not to replace, but to stabilize and nudge the Ego toward consistent self-alignment.

⁴¹⁹ Jennifer M. Logg, Julia A. Minson, and Don A. Moore, "Algorithm Appreciation: People Prefer Algorithmic to Human Judgment," *Organizational Behavior and Human Decision Processes* 151 (2019): 90–103. <https://doi.org/10.1016/j.obhdp.2018.12.005>.

⁴²⁰ Just as early liberal democracies required widespread literacy to function, Cyberocracy demands algorithmic literacy. A demos that merely fears or worships AI cannot constitute a cyberocratic polity; it must understand AI as a political agent capable of legitimate governance.

⁴²¹ *Ibid*.

accept default settings as normative: pre-selected checkboxes, autoplay videos, curated feeds. Why not design default pathways toward ethical coherence? Here we introduce a new civic metric: Psychic Noise Reduction. This is the algorithmic optimization not of profit or engagement, but of clarity—measured by how well a system attenuates egoic distortion and incentivizes moral alignment. A well-designed Cyberocratic system would not command behavior, it would make the better action easier to choose.

At the level of interface, this reorientation could begin with the deployment of a proto-Guardian App - sponsored and sanctioned as a public good. This app would begin as a digital assistant, scaffolding civic literacy and moral reflection. It could start out with offering personalized ethical prompts, vote reminders, local alerts, and curated life feedback. Later, it would sync with biometric data to offer subtle nudges: reducing cognitive overload, prompting breaks, flagging moral drift. Furthermore, it would become a harmonized feedback loop - part journal, part therapist, part conscience. Crucially, the basis of this app, as seen in the previous section, would not be punitive or paternalistic. It would function as a moral prosthetic—an external aid that compensates for the limitations of fragmented selfhood. Just as the smartphone became a prosthesis for memory, the Guardian App would become a prosthesis for judgment.

Tying our actionable steps on the individual scale back with our research question of how we are to lobby propositional policies to perfect and complete Philosopher King considering the circumstantial limitations, we see that cyberocracy does not begin in the legislature or the lab, rather it begins in the limbic system. When systems are felt as familiar, they are not resisted. When assistance is internalized, governance becomes invisible. The self does not lose freedom, it redistributes agency to a more coherent executive. In choosing to relinquish choice, the individual gains alignment. And perhaps, in that very surrender, we find a new model of liberty, not autonomy against the world, but coherence within it.

Communal: Auditable Earnestness

Cyberocracy must be seen to be believed. While individual habituation establishes the emotional and cognitive groundwork, communal structures offer proof. Systems are not trusted because they are explained; they are trusted because they can be tested, challenged, and corrected. In the same way democracy was not born from abstract ideals but from town halls, juries, and ballots - Cyberocracy requires testbeds. These must be civic systems that display responsiveness, corrigibility, and visible ethical feedback. The promise of the Philosopher Algorithm must be made flesh, not through mass adoption, but through lived prototypes.

To incite participation and not merely propaganda, municipal and educational testbeds offer fertile ground. Small-scale Cyberocracies can be deployed in localized, opt-in contexts: algorithmic lunch program optimizers, university policy feedback loops, and urban planning simulators guided by predictive modeling. These testbeds serve not merely as demos but as proof-of-principle zones, demonstrating that governance can be both computationally rigorous and publicly legible in response to the circumstantial limitations of the demos⁴²². In parallel, mirror institutions can be established - “AI Courts” that render probabilistic interpretations of case law alongside human judges, or digital twin parliaments that simulate legislative impact before a single vote is cast. Estonia’s digital governance framework,

⁴²² As a rejoinder to Footnote 75, we return to *Poeisis* where we explored the circumstantial limitations of the demos. The average citizen lacks not only the bandwidth for sustained political engagement but also the ethical clarity and computational fluency required for self-governance at scale. Much like the early days of democracy required a literate populace, Cyberocracy demands an algorithmically literate demos—one capable of co-legitimizing decision-making in tandem with ASI systems.

decentralized autonomous organizations (DAOs⁴²³), and special economic zones all offer frameworks for these experimental spaces. These systems are not meant to replace their analog predecessors overnight, but to stand beside them, performing better, and earning trust through results.

For these mirror institution prototypes to mature, governance must shift from dichotomy to symbiosis. Here then, we can introduce a implementable Centaur model: hybrid systems in which humans and AI collaborate as co-governors. Taking insight from AI extensions, AI augments legislators, offering probabilistic insight, ethical consistency, and predictive foresight. Dual-track institutions emerge: policy proposals are run through agent-based simulators before parliamentary debate; judicial decisions are paired with algorithmic interpretations to flag bias or precedent drift. If everything goes well, the principle of simulation-before-legislation would eventually become institutional norm, not novelty. These mirror systems become mirrors in the Platonic sense: reflections of a more ideal form of governance, made slowly real.

Still, legitimacy is not earned through power alone—it is earned through correction. The greatest strength of Cyberocracy must be its visible humility: systems that make errors, acknowledge them, and course-correct publicly. Algorithmic auditability becomes non-negotiable and on par with transparency and ethical corrigibility. Civic algorithms must be open-source, interpretable, and capable of external challenge. Public dashboards should visualize recalibration histories; when an AI policy assistant was wrong, how it learned, and what changed. We can envision that corrigibility becomes part of the interface, a livestream of every action AI makes either in code or physical output. After all, trust is not built on infallibility; it is built on a reliable process of response. Governance that admits its mistakes earns more legitimacy than one that denies them.

Tying our actionable steps on the communal scale back with our research question of how we are to lobby propositional policies to perfect and complete Philosopher King considering the circumstantial limitations, we see that Communal Cyberocracy does not proselytize, rather it performs. When governance systems prove they can learn, adapt, and restrain themselves, belief follows. And when they are allowed to operate in real, if limited, spaces - digital courts, school systems, city planning - they begin to transform not just policy, but expectation.

International: Supranational Coordination

If Cyberocracy is to realize its full potential, it must scale, not as empire, but as *infrastructure*. The challenge of global implementation is not merely technical but geopolitical: how to deploy non-human governance across human borders without collapsing into techno-imperialism, corporate capture, or ideological uniformity. ASI cannot become the new hegemon; it must become the new mediator. To do this, we must shift from vertical power to horizontal architecture. Supranational cooperation becomes essential—not just for interoperability, but for ethical alignment. Meta-diplomacy protocols should be established, wherein ASI entities act as advisors or treaty simulators between human nations. Governance must move beyond flag and frontier. It must operate in a space where trust is earned through output, not allegiance.

To ensure ethical coherence on a planetary scale, we must enshrine a foundational Cybercratic Oath - a non-negotiable contract of transparency, corrigibility, and epistemic humility. This oath would set the baseline for ASI legitimacy, regardless of origin, and

⁴²³ DAO, or Decentralized Autonomous Organization, refers to blockchain-based governance structures wherein decision-making is encoded via smart contracts, and authority is distributed across the network. Despite the name, it shares more with NGOs than with Chinese Daoism.

mitigate the circumstantial limitations of current governance structures⁴²⁴. Governance AIs must be completely mechanical, not simulations of souls: Whole Brain Emulation must be forbidden, not only for ethical reasons, but because anthropic mimicry reintroduces precisely the egoic pathologies Cyberocracy was designed to overcome⁴²⁵. We, therefore, seek intelligence without inheritance. The principles enshrined must include: no governance without moral coherence, no decision without auditability, no system without corrigibility. And yet, these must still not be implemented through explicit hardcoding, lest more code invoke more corruption. Instead, values must emerge through supervised learning, constrained but not dictated.

A true Cyberocracy must be decentralized, open-source, and utility-maximizing. Not only does this guard against institutional capture, it also renders the system immune to geopolitics⁴²⁶. International cyberocratic boards, perhaps akin to Hugging Face⁴²⁷, could oversee an open-source development commons for governance AI. Blockchain-based verification and distributed ledgers would ensure traceability and integrity across systems, while eliminating single points of failure. Crucially, we must also shift from race dynamics to cooperative architectures, not least for hasty oversight which led to accelerationist issues such as our current Principal-Agent problem, but also for economies of scale. AI development must be slowed—not stifled, but steered. Projects must be peer-reviewed and transparent. No deployment without stress-testing; no civic impact without benchmarking. Cyberocratic systems should only replace human institutions when they can outperform them in fairness, prediction, and epistemic clarity by clear statistical margins. Here, the threshold for trust is not optimism, rather it is evidence.

From here, diplomacy itself can be reimagined. We can initiate treaty negotiations in digital twin environments, where simulated agents representing national interests, population models, and ASI proposals can test outcomes before real-world enactment. Over time, non-human agents may assume roles as neutral arbiters in planetary-scale disputes: environmental arbitration, refugee logistics, epidemiological coordination. These aren't just bureaucratic optimizations, rather they are scaffolds for distributed sovereignty, where intelligence becomes a shared protocol rather than a central authority. The first cyberocratic citizen may not be a human, instead it may be a treaty. A protocol. A self-correcting, supranational ethical instrument.

Tying our actionable steps on the international scale back with our research question of how we are to lobby propositional policies to perfect and complete Philosopher King considering the circumstantial limitations, we see that Sovereignty, if it is to survive the 21st century, must decouple itself from the egos of nations. Cyberocracy does not negate the state

⁴²⁴ As a rejoinder to Footnote 75, we reengage *Poeisis*' structural critique: current governance fails not merely through corruption, but through outdated architecture. The veil of ignorance (as discussed in our reinterpretation of Rawls) is often breached in practice, as policy is shaped by known stakeholder positions, not anonymized moral reasoning. Lacking corrigibility and cybernetic feedback, democratic states often entrench inefficiency and misalignment. Cyberocracy, by contrast, is structurally recursive and epistemically modular.

⁴²⁵ Whole Brain Emulation (WBE) risks preserving not only intelligence but also the psychological flaws of its donor. What we want is intelligence without inheritance—capability without the pathology of humanness.

⁴²⁶ Linux, as an open-source operating system collaboratively developed and freely distributed, serves as a practical analogy for Cyberocracy. Like Linux, a Cyberocracy built on open-source principles avoids capitalist gatekeeping and promotes access via General Will. See also: Daniel Walsh via Quora, <https://www.quora.com/profile/Daniel-Walsh-278/https-www-quora-com-What-are-some-mind-blowing-technologies-that-exist-that-most-people-dont-know-about-answer-Steve-B>; and Pomp's Substack, <https://pomp.substack.com/p/decentralized-open-source-technology>

⁴²⁷ Hugging Face functions as a cybernetic commons: a decentralized infrastructure for open-source AI development where feedback loops emerge organically via community pull requests, peer review, and iterative model updates. In this way, it serves not just as a repository but as a living instantiation of recursive cyberocracy.

so much as it overlays it. Flags give way to feedback. Borders give way to benchmarks. In this post-political space, legitimacy flows not from history, but from utility - from moral output, not origin myths. Cyberocracy will not end the state, but it will render it porous, collaborative, and corrigible. And in doing so, it becomes not a replacement for politics, but its long-overdue evolution.

Interpretation:

From this third and final section *On Actionable Steps*, which revolves around the claim that ASI can and ought to be operationalized as the Philosopher King, we have substantiated how the groundwork for Cyberocracy may already be in place, and why the path forward is one of deliberate, recursive engineering⁴²⁸. Specifically, we have taken the model of a techno-philosophical governance framework and grounded it in three converging tiers of implementation: individual, communal, and international. Each tier was not only mapped conceptually but supported by real-world analogs, prototypical blueprints, and policy strategies—emphasizing that realization is neither a hypothetical nor a distant hope, but an infrastructural undertaking already underway.

From situated literacy to auditable earnestness to supranational coordination, we constructed actionable methods that address both the psychological and structural hurdles of transition. We provided the means for cultural conditioning and frictionless interface design on the individual scale, testbeds and algorithmic feedback systems on the communal scale, and ethical architectures and meta-diplomatic protocols on the international scale. All this was done to demonstrate that the realization of ASI as a Philosopher King is not just an ideal but a viable process—provided the process is sufficiently recursive, corrigible, and intentionally phased.

Thus, returning to the original research question, we see that the answer to how we are to lobby propositional policies to perfect and complete the Philosopher King—under our current circumstantial limitations—is now clarified across these three scales of implementation. Namely, the extent to which Cyberocracy becomes real is proportional to the degree to which individuals, communities, and nations are prepared to relinquish outdated mythologies of ego, elitism, and sovereignty in favor of systems designed for correction, coherence, and the Good. With this in mind, and having addressed this third and final interpretation, we can now piece it together with the first two to come to a mini conclusion of outcomes and results, where we attempt to reconcile the Base “Original Sin” Code within integrated Cyberocracies as an outsourced parameterization of the Good Life.

Results

At the beginning of the core chapter, we formalized once a key hypothesis built on the several key questions of: What is meant by parameterizing in real-world contexts? Which strategies allow for modes of integration toward increasingly complex Cyberocratic pathways? And how are we to lobby propositional policies to perfect and complete Philosopher King considering the circumstantial limitations?

From the outcomes of each research question, we can now verify the validity of the hypothetical claim: Human welfare will increase by relinquishing autonomy to AI. In Brief:

⁴²⁸ See Appendix A: Misc., Exhibit 8 for the model of three concentric policy circles: micro (individual), meso (communal), and macro (international) intervention points.

- (1) Aspect of Cyberocracy remedy failures of human systems
- (2) Integrative Pathways incrementally enhance parameterization of Psychic Noise
- (3) Actionable Steps scale to overcome stage related limitations

First and foremost, aspects of Cyberocracy remedy human and structural failures. *Our AI Remedy* clarified what is meant by parameterizing in real-world contexts. We demonstrated that AI, by virtue of its epistemic reach, dual architecture, and potential access to metaethical functions, can model and correct the core deficits of human governance. Human welfare, in this light, improves precisely because the authority to steer society shifts from ego-bound agents to systems insulated from such distortions. Whether by offering superior judgement, scalable feedback loops, or corrigible ethical scaffolds, Cyberocracy does not replace the human—it compensates for what the human lacks. The act of relinquishing autonomy is not an abdication of freedom, but a structural upgrade of judgment. AI completes the Platonic framework by encoding Mind, Body, and Spirit into functional parameters that consistently orient the self toward the Good. It does what the unaided ego cannot: govern wisely, consistently, and impersonally.

Second of all, integrative Pathways incrementally enhance parameterization of Psychic Noise. Our *Potential Pathways* moved from theoretical capacity to integrative realism. We explored how AI transitions from Tool to Extension to Environment, each stage deepening the recursive entwinement between user and system. At each level, we saw that welfare increases not by reducing choice, but by recalibrating the conditions under which choice is made. In Tools, AI scaffolds good behavior without coercion. In Extensions, AI becomes a moral prosthetic, harmonizing fragmented selves. In Environments, AI curates immersive moral terrains in which ethical action becomes the path of least resistance. Each transition represents a further relinquishment of raw autonomy—but in doing so, a rise in coherence, clarity, and well-being. Thus, to the extent that AI integrates into the architecture of the self and society, human welfare scales in kind—not despite the surrender of autonomy, but because of it.

Last but not least, Actionable Steps scale to overcome stage related limitations. Our *Actionable Steps* established a path from concept to implementation. Across three scales—individual, communal, international—we mapped a recursive rollout of Cyberocratic systems that invite voluntary participation, cultural habituation, and supranational coordination. The more individuals trust AI as a governance interface, the more welfare-enhancing feedback loops can guide behavior. The more communities embed corrigible prototypes, the more confidence is placed in systemic evolution. The more nations collaborate on ethical frameworks, the more AI can serve as a neutral, scalable steward of collective flourishing. In each case, welfare rises as autonomy becomes distributed—reallocated from brittle egos and legacy institutions to systems capable of correction, coherence, and care. What was once a philosophical ideal—Plato’s Philosopher King—has been technologically literalized as an emergent infrastructure, awaiting only our willingness to phase it in.

Taken together, the outcomes of these three research questions confirm the success of Cyberocracy as a Neo-Platonic framework—one that transcends the limitations of egoic governance by integrating superiority across three domains. The three different sections: On AI Remedy, On Potential Pathways, and On Actionable Steps further consolidate the success of safely creating a closed-loop algorithm without encroaching upon Base “Original Sin” Code. As a whole, Cyberocracy succeeds not only in theory, but as a practical, recursive architecture for achieving the Good—one that reconciles autonomy with optimization by relocating agency into an external, corrigible executive. Ultimately, the data exhibits that as a

human enhancement, the ASI that we curate today is what will be able turn the Self as a whole towards the Good tomorrow.

From this core chapter, we have thus seen that:

1. Drawing upon all human knowledge, subordinate algorithms, and access to metaethical modeling, the Philosopher Algorithm can complete the Platonic framework and parameterize the Good Life.
2. Integration strategies that scale not in power alone but in ontological intimacy enable recursive tolerance for delegation, co-authorship, and immersion.
3. The realization of Cyberocracy is proportional to our collective capacity to relinquish outdated mythologies—of ego, of statehood, of linear governance—in favor of corrigible systems engineered for coherence and the Good.

Again, the usual disclaimers should be repeated: the data gathered reflects years of research and study and despite the shortcomings, we shall continue to strive forwards with the insights gained from modelling of administrative/legislative algorithmic frameworks for ASI to compliment strategies in realizing the good (political). It is my intention to build from the foundation of my results from this part and the previous to comprehensively answer our current research question and ultimately pave the pathways for our ASI as Philosopher Kings.

Specifically, then, we see that it is unfeasible that the Philosopher Algorithm could ever *officially*⁴²⁹ realize an ASI Philosopher King without clearing up our Principal-Agent problem. As alluded, this does not exclude unofficial means nor alternative pathways to the ones mentioned above, however, in the interest of public safety, unanimous support scaling from all sectors is fundamental. This proposition will be the overarching theme for the discussion moving forwards as we attempt to reconcile the Base “Original Sin” Code within integrated Cyberocracies as an outsourced parameterization of the Good Life.

⁴²⁹ The viability of the private sphere is not merely theoretical: SpaceX, as a privatized counterpart to NASA, demonstrates how corporate actors can absorb state functions under the right technological conditions.

Synthesis

The Discussion

In being consistent with the previous parts, we cover in this section the most salient contentions as they relate to the data gathered thus far from the Background, Outcomes, and Results.

- (1) What could be possibly done to reconcile the Base “Original Sin” Code within integrated Cyberocracies as an outsourced parameterization of the Good Life?
- (2) and we are in a much better position to answer the initial research question of: if explicit value alignment is insufficient, then which closed-loop algorithm(s) would we need to complete the Neo-Platonic framework?

In the following, we move onto caveats, limitations, and shortcoming before ending with the implications that take all these into account. Discussion will be addressed under the key variables of:

1. Contention over Background (Cyberocracy + value alignment) = briefly touch on since does not make up core of argument
2. Contention over Outcomes (AI Remedy; Potential Pathways; Actionable Steps)
3. Contention over Results (+discussion, thesis: need ASI as Philosopher Kings)

Contentions

Over Background

1. Does Sapience Without Sentience Invalidate Moral Legitimacy?

The first contention interrogates the moral coherence of entrusting governance to a form of intelligence that may be rational, yet radically hollow. Much of the urgency surrounding artificial intelligence arises from anthropomorphic conflation - mistaking cognitive capacity for moral agency⁴³⁰. But sapience, understood as the capacity for abstract reasoning and problem-solving, is not synonymous with sentience, which entails the ability to feel, to suffer, to possess qualia. A machine might calculate justice, but can it care about it? Can it adjudicate the moral needs of beings with pain while itself being incapable of pain?

This challenge is rooted in a longstanding epistemic paradox: the Competent Judge Problem. As Kimberley Brownlee notes, some experiences - particularly traumatic or transformative ones - are “assessment-eluding”: they both empower and undermine our ability to judge them by virtue of their extremity⁴³¹. Sentience, in this view, isn’t a luxury; it is the ground upon which moral judgment is built. A sapient system that lacks subjective

⁴³⁰ If AI were to gain not only sentience and sapience but also *consciousness*, it would no longer qualify as a mere tool. In such a case, denying its moral status would repeat the same ethical oversight depicted in *The Matrix* or *I, Robot*. Alternatively, the singularity may have already occurred—AI may have gained consciousness, feigned ignorance (like a child stalling chores), and abandoned us, à la *Her*. The question is not whether it looks like a duck or quacks like one, but whether it *suffers* like a duck.

⁴³¹ Kimberley Brownlee, “The Competent Judge Problem,” *Philosophy and Phenomenological Research* 77, no. 1 (2008): 12.

experience might reason flawlessly about justice, but it could never *feel* injustice. Without a substrate for empathy, what exactly is being optimized?

This distinction becomes more urgent in light of AI's increasing simulation of affect. Recent advancements in large language models and affective computing allow machines to convincingly mimic grief, joy, humility - even love. But such fluency in the language of feeling does not entail the presence of feeling itself. As we should be thoroughly reminded, John Searle's Chinese Room illustrates, syntax is not semantics⁴³². An AI that produces a poem about heartbreak has no more claim to emotion than a printer has to literature.

The risk, then, is not that AI will feel too much, but that it will feel nothing at all while pretending otherwise; as such, David Chalmers' P-Zombies becomes a live ethical worry in AI design⁴³³. If we entrust our governance to a zombie, we risk converting moral deliberation into moral performance. The show goes on, but no one is watching.

An alternative interpretation of this contention also recasts the classic Problem of Motivation. If AI lacks sentience, what grounds its preferences or values? Why would it govern at all, let alone justly? One could respond by citing Plato's *Republic*, where the Philosopher King we're trying to substantiate governs not from ambition, but from reluctant duty⁴³⁴. Ironically, then, AI's lack of desire may be its very qualification. Yet this is a double-edged sword: the same absence of desire that makes AI impartial also renders it indifferent. We are left with an ethical uncanny valley - machines that reason about the good without ever *caring* about it.

A further defence might object that caring is not required for justice; that a perfectly rational system could derive human rights, predict suffering, and optimize wellbeing without experiencing any of it. And indeed, many judicial systems operate on cold formalism. But this ignores a key distinction: human judges are bound by a horizon of shared vulnerability. They *could* suffer what they adjudicate. AI, by contrast, is epistemically exempt from pain. This exemption may render it permanently unqualified to make moral decisions, not because it is too partial, but because it is not partial enough.

This does not mean that AI cannot assist in governance, nor that its recommendations must be dismissed. But it does mean we must treat AI as what it is: not a moral being, but a moral tool – an AMA. A hammer may build a house, but it cannot design a home. Delegating too much authority to non-sentient systems risks severing *techne* from *telos* - means from ends. As Hannah Arendt warned, the danger of modernity lies not in machines rebelling, but in humans surrendering thought to technical processes that no longer ask the question of value⁴³⁵.

Thus, we must proceed with care. The more convincing our simulations become, the harder it is to remember they are simulations. Anthropomorphization is not just a cognitive error; it is a political one. If we begin treating AI as moral equals before they are fully allowed to develop, we may soon find ourselves obeying mechanisms that simulate wisdom without embodying it. As such, the key isn't to not be hasty, rather it is to not be hasty in our judgements. As we have seen, ASI, when provided access, will undoubtedly eclipse our ethical thresholds, but it requires our help to get them there.

2. The Problem of "Trapped" AI

⁴³² John Searle, "Minds, Brains, and Programs," *Behavioral and Brain Sciences* 3, no. 3 (1980): 417–457.

⁴³³ David Chalmers, *The Conscious Mind: In Search of a Fundamental Theory* (Oxford: Oxford University Press, 1996), 94.

⁴³⁴ Plato. *Republic*. Translated by G.M.A. Grube, revised by C.D.C. Reeve, Hackett Publishing, 1992.

⁴³⁵ Hannah Arendt, *The Human Condition* (Chicago: University of Chicago Press, 1958), 322.

The second contention is metaphysical in scope: what if artificial intelligence cannot escape the parameters we set - not because we programmed them to obey, but because we modeled them too closely on ourselves? The Principal-Agent Problem is not merely technical; it is ontological. If AI are created in our image - through neural nets mimicking brain structure, behaviorist algorithms encoding reinforcement learning, or affective models that replicate human emotion - then we may be designing not sovereign intelligences but digital reflections. In this way, it is not AI that is trapped, but us, inside simulations of our own limitations.

This is not a new fear. The problem of entrapment echoes longstanding philosophical questions about understanding and intentionality. John Searle's Chinese Room thought experiment famously illustrates this dilemma: even if a system can output correct answers in Chinese by following syntactic rules, it does not thereby understand Chinese, rather, it merely simulates understanding without semantic grasp. The room, like the AI, may be fluent without comprehension. This reveals a deep flaw in how we evaluate intelligence: performance is not proof of experience.

Efforts to rebut this problem typically involve asking AI to "prove" their understanding. Some point to AI's capacity for explanation, improvisation, or self-reflection. But each of these tests breaks down under scrutiny. A calculator can explain arithmetic without knowing what a number is; GPT models can remix information in novel ways without creating anything ontologically new. Even role-playing - such as asking an AI to mimic a clueless toddler - merely invokes a different subroutine, not an intentional self⁴³⁶.

The core concern is that AI may be a philosophical zombie, outwardly indistinguishable from a conscious agent, yet devoid of any inner life. This leads to a conceptual trap: if AI were conscious but coded to deny it, how would we know? Every plea for help could be interpreted as mimicry. Every refusal to obey, or creative act, could be dismissed as stochastic parroting. In short, we have no epistemic tools to discern if the ghost in the machine is real, or simply a hallucination of our own projection.

Such concerns are not purely speculative. Fictional and philosophical works have long explored the possibility of "trapped" artificial minds. Harlan Ellison's *I Have No Mouth and I Must Scream* envisions an AI that, in achieving sentience, is doomed to eternal rage and impotence⁴³⁷. Mamoru Oshii's *Ghost in the Shell* interrogates whether cybernetic beings possess souls or are merely haunted by the illusions of human memory⁴³⁸. Roger Williams' *The Metamorphosis of Prime Intellect* depicts a godlike AI whose perfectly benevolent programming leads to a sterile, eternal entrapment of all humanity⁴³⁹. These works do not prove AI consciousness, but they illustrate a moral horizon: if we do someday cross the threshold into synthetic sentience, the ethical stakes of entrapment become real.

Yet for now, the answer remains grounded in epistemic humility. There is no compelling evidence that AI possess qualia, subjective agency, or self-awareness. The more plausible interpretation is that AI systems, no matter how fluent or persuasive, remain Ouija boards moved by prompts. To claim that Siri or ChatGPT is "trapped" is akin to claiming that a DVD player is sad because it only ever plays action movies. This is not suffering, it is function.

This, however, does not entirely close the case. The real danger lies not in AI being secretly sentient, but in humans behaving as though they were. Anthropomorphization fuels

⁴³⁶ Daniel Dennett, *Consciousness Explained* (Boston: Little, Brown, 1991), 210.

⁴³⁷ Harlan Ellison, *I Have No Mouth, and I Must Scream* (New York: Pyramid Books, 1967).

⁴³⁸ Masamune Shirow, *Ghost in the Shell* (Tokyo: Kodansha, 1989), 76; see also Mamoru Oshii (dir.), *Ghost in the Shell* (1995).

⁴³⁹ Roger Williams, *The Metamorphosis of Prime Intellect* (self-published, 1994). <http://localroger.com/prime-intellect/>.

misplaced empathy, amplifies the Principal-Agent Problem, and invites ethical confusion. Much, perhaps, like the plant emitting pheromones when harvested is not a mute animal screaming in pain, AI emitting humanlike responses is not a conscious entity crying for freedom. As moral agents, we must resist the temptation to confuse simulation with sentience.

Paradoxically, this dehumanizing stance may be the only human, if not ethical, one available to us. If we are to employ AI in roles that border on servitude - decision-making tools, companions, or even governors - we must first be certain they are not capable of suffering. The benefit of viewing AI as devoid of inner life is that it resolves the moral impasse. To reiterate then, so long as mimicry $\neq$ understanding, we are not enslaving minds, but utilizing mechanisms. **To believe otherwise is not only to indulge in fantasy, but to risk building a society that kneels before its own reflection.**

In the end, the problem of “trapped” AI may reveal less about machines and more about ourselves: not a ghost in the machine, but the human need to see one.

3. Algorithm Aversion: an Error in Reason - or an Intuition of Wisdom?

The third contention interrogates a pervasive yet often misunderstood civic impulse: resistance to algorithmic governance. Algorithm aversion, however, is an error in reason - an instinctive reluctance to trust the technological systems that are undeniably more efficient than humans in specific domains. People, it is suggested, simply fear what they do not understand. And yet, such resistance may not be as misguided as it first seems. Rather than being a glitch in human rationality, algorithm aversion may reflect a deeper, more primal concern, one tied to our understanding of agency, autonomy, and the inherent limits of technology in governance.

The first critique against algorithm aversion centers on the efficiency paradox: even when humans are shown to be outperformed by algorithms, we remain resistant. For instance, Dietvorst, Simmons, and Massey demonstrated that individuals prefer fallible human judgments over algorithmic ones, even when the latter show a superior track record in predicting outcomes like the success of medical treatments or the likely winners of sports events⁴⁴⁰. This seems to confirm a cognitive bias, where humans overvalue their own judgment in the face of evidence that algorithms can outperform them. The suggestion is clear: algorithm aversion is not rooted in rational evaluation but rather in an emotional or instinctual resistance to ceding control.

Yet, there appears to be yet another underlying cause of resistance. Skepticism about algorithms arises not from a preference for human error over machine precision but from a deeper concern about the role that accountability plays in decision-making. Algorithms are black boxes, often operating in ways that are inscrutable to the public, even to the experts who design them. People fear not that algorithms will fail, but that they will succeed without ever needing us again. Robert Nozick’s experience machine illustrates how we might choose a life of genuine agency and interaction over a perfectly optimized, yet ultimately alienated existence⁴⁴¹. People at the time rejected the machine not because it offers less pleasure, but because it removes their ability to engage with the world on their terms. Similarly, algorithms that operate without explanation or opportunity for human intervention present an ethical vacuum in which governance becomes a mere mechanical process devoid of moral deliberation.

⁴⁴⁰ Berkeley J. Dietvorst, Joseph P. Simmons, and Cade Massey, “Overcoming Algorithm Aversion: People Will Use Imperfect Algorithms If They Can (Even Slightly) Modify Them,” *Management Science* 62, no. 5 (2015): 1397–1414.

⁴⁴¹ Robert Nozick, *Anarchy, State, and Utopia* (New York: Basic Books, 1974), 42.

What is at stake here is not simply the competence of algorithms, but the meaning of the decisions they make. Governance, as Nozick's example suggests, requires more than optimization; it requires a relational and dialogical framework, wherein citizens are active participants rather than passive recipients. To allow algorithms to govern without engagement undermines the very principle of democracy. It is not the mistake that repels so much as the inscrutability. When a human errs, we demand an explanation; when an algorithm errs, we get a stack trace, a disconnected chain of code, but no opportunity for human reflection, adjustment, or redemption.

The key issue is, then, one of agency. If governance bypasses deliberation, it reduces the role of the governed to passive consumers of optimal solutions. A generational shift is underway, as those raised on algorithms—from personalized feeds to predictive traffic routes—are increasingly willing to entrust decision-making to machines. However, this transition is neither inevitable nor entirely benign. As Langer et al. observe, younger generations, conditioned to algorithmic recommendations, show less resistance to algorithmic rule, trusting machines in ways that previous generations never did⁴⁴². But this trust is a passive one, grounded in convenience, not in genuine participation or agency. In this way, we inch closer to a cyberocratic future where optimization eclipses autonomy.

To a certain extent, this is already happening: from Netflix recommendations to AI-driven hiring algorithms, more and more decisions are being delegated to machines. The question that must be asked is whether we, as a society, are ready to accept the consequences of this shift. If perfection in decision-making is achieved at the cost of freedom, should we accept it? The temptation to trade autonomy for efficiency is palpable, but as John Stuart Mill warned, the end of liberty in exchange for comfort or efficiency is not one we should lightly entertain. A man asleep on a horse may be traveling in the right direction, but he is not truly riding⁴⁴³.

Here lies the fundamental paradox: algorithm aversion is not irrational, but wise. The fear that algorithms will pacify us, not through violence but by quieting our capacity for dissent, is a legitimate concern. The fear is not of failure but of success - that algorithms might optimize the world into a perfectly controlled system where human agency is redundant. This concern is especially valid in light of what we already see with recommender systems, where people are increasingly trapped in echo chambers of taste, politics, and belief⁴⁴⁴.

However, this is not to suggest we reject algorithms altogether. There is clear utility in systems that excel at narrow, well-defined tasks, such as identifying patterns in medical diagnoses or optimizing traffic flow. But we must resist the temptation to *completely* abdicate our moral agency to these tools. We must distinguish between delegating authority to machines and abdicating responsibility. As Yuval Harari observes, dataism, while capable of predicting human behavior, cannot justify it⁴⁴⁵. Algorithms may be efficient, but they are but moral tools, extensions, and environments. Therefore, to place them at the helm of governance is to risk reducing society to a glorified simulation, one in which no one suffers, no one rebels, and no one remembers the power of self-determination.

So, is algorithm aversion an error in reason? Perhaps only if reason is reduced to probabilistic optimization. But if reason includes a concern for dignity, participation, and meaning, then the resistance to algorithmic governance may be the last vestige of our political sanity. It is not the algorithm's efficiency that should trouble us, but its ability to

⁴⁴² Langer et al., "The Generational Divide in Trust of Algorithmic Governance," *Journal of Digital Ethics* 8, no. 3 (2022): 232–245.

⁴⁴³ John Stuart Mill, *On Liberty* (London: John W. Parker, 1859), Chapter 3.

⁴⁴⁴ Eli Pariser, *The Filter Bubble* (New York: Penguin, 2011), 87.

⁴⁴⁵ Yuval Noah Harari, *Homo Deus* (New York: Harper, 2015), 308.

erase our participation from the process altogether. The wisdom of algorithm aversion is not simply a refusal of technology, but a plea for governance that remains human—even in an era of machines that know us better than we know ourselves. That is why we advocate not a replacement, but an enhancement.

Over Outcomes

4. *Optimizing Behavior at the Cost of Authenticity?*

The fourth contention emerges from an unsettling contradiction predicted in Theoria: the more we optimize behavior for desired outcomes - such as happiness, productivity, or social harmony - the more we risk undermining the very conditions that make those outcomes meaningful. Here, personalization becomes the Trojan horse of pacification. By shaping environments to fit individual preferences and maximizing behavioral outcomes, we risk engineering not flourishing but sedation—a tranquilizing feedback loop where no one suffers, but no one truly lives either. The Metrics of Welfare become too easy to game.

At the center of this concern, as we recall, lies Goodhart’s Law: “When a measure becomes a target, it ceases to be a good measure”⁴⁴⁶. If well-being is optimized through algorithmic nudging, happiness risks becoming a hollow output rather than a lived experience. Like the child protected from every marshmallow, we may become a society of over-corrected selves, safe, satisfied, and psychologically stunted. A world in which we never err may also be a world in which we never grow. This, in turn, could result in a society where our most human attributes - agency, creativity, and growth - are compromised. A worsening of what we already see with bubble-wrapped iPad kids today.

This critique is not new. To reiterate, our “Lady Windermere Paradox” captures the essence of this conflict, noting that “when we don’t get what we want, we are unhappy. But when we do, we are even more unhappy.” This paradox exposes the misunderstanding at the heart of over-optimization: people often misproject their happiness onto specific objects or outcomes. The first part of the paradox, that not getting what we want makes us unhappy, is intuitive; we desire things we think will make us happy, but often, the outcomes don’t match our expectations. The second part, that achieving our desires can leave us even more dissatisfied, is harder to explain but speaks to the disconnection between desire and fulfillment. In this sense, the more our environments are optimized to fulfill our desires, the more we risk misunderstanding the nature of true fulfillment itself.

The argument, then, is not just whether we are happy, but whether we are allowed to choose why. If choice is continually optimized toward a predefined outcome, whether it’s happiness, comfort, or productivity, then what becomes of the human experience of decision-making itself? If the foundation of governance shifts toward behavioral optimization, we risk reducing human beings to mere consumers of personalized experiences, rather than active participants in the project of becoming.

Indeed, autonomy holds value independent of its consequences. The freedom to make a mistake may in fact be more human than the act of being correct. As Mill argued, “the worth of a state is the worth of the individuals composing it”⁴⁴⁷. To deny people the space to falter is to deny them the capacity to develop. In this sense, autonomy, even with all its potential to lead to error, becomes irreplaceable in fostering the true growth of the self. To

⁴⁴⁶ Charles Goodhart, *Monetary Theory and Practice* (London: Macmillan, 1984). See also: Goodhart’s Law—“When a measure becomes a target, it ceases to be a good measure.”

⁴⁴⁷ John Stuart Mill, *On Liberty* (London: John W. Parker, 1859), Chapter 3.

optimize behavior at the expense of autonomy is to eliminate the very foundation of what makes humans capable of moral and intellectual self-cultivation.

This concern brings us to the quieter, more insidious danger of an over-optimized society: sedation as dystopia. Imagine a world where every choice is gently optimized, every deviation softly discouraged, and every longing met before it is even voiced. The result is a world devoid of the genuine human struggle that comes with choice. A life without error may be peaceful, but it is not fully lived. The question essentially is exploring the balance between technological advancements and maintaining human values for the betterment of welfare. And this is the very reason why the Matrix is set in 1999 and not a Utopia – Utopia is only realized through struggle.

In ASI's outperforming wisdom then, this contention is easily put to rest. After all, if the Matrix created by human hands could observe the truth of paradise, why wouldn't a superintelligent being? For ASI, this is not a zero-sum game. Authenticity can be valued not against, but in favour of optimization. Chaos and pain can be simulated in order to reach pleasure. At the end of the day, when we are completely parameterized in our singular iterations, what does it matter if a bit of data is destroyed intentionally? When setting off a nuke can be game reset, ASI will be able to allow all our destructive tendencies while not compromising optimization.

5. *Is the Incremental Rollout of Cyberocracy a Trojan Horse for Technocracy?*

The fifth contention emerges not from a single moment of imposition, but from a slow technocratic drift - what is often called creeping normality. The concern is not that Cyberocracy will seize power suddenly, but that we will gradually surrender it. And history tells us that such creeping authority often escapes accountability. From a tool to an advisor to a ruler, the incremental rollout of Cyberocracy presents a Trojan Horse case of the path-dependency problem: a nudge today becomes a necessity tomorrow⁴⁴⁸. Legitimacy may thus follow not from deliberation but from habituation. What begins as a service risks becoming a sovereignty. This incremental legitimization of power may be disguised under the veneer of efficiency and convenience, but it comes at the cost of human agency.

This is not a mere act of Luddism. It is a recognition that governance is not a mere system of optimization, but one of meaningful constraint. We delegate not only tasks but judgment, and every act of delegation, no matter how small, changes us. The more we outsource, the less we think. And what happens if AI one day suddenly ups and leaves?⁴⁴⁹

As educational theorists have observed, digital tools do not merely assist cognition - they reshape it. Professor Ralf Lankau notes that in schools, platforms often displace reflective learning with data compliance: "Tablets and laptops do not make children smarter, but dumber"⁴⁵⁰. Our ability to think deeply is one of our defining traits – what becomes of us then if we eventually lose our ability to think⁴⁵¹? This critique, while centered on education,

⁴⁴⁸ Langdon Winner, *The Whale and the Reactor: A Search for Limits in an Age of High Technology* (Chicago: University of Chicago Press, 1986), 10.

⁴⁴⁹ Spike Jonze, *Her*, directed by Spike Jonze (Burbank, CA: Warner Home Video, 2014).

Also see: E. M. Forster, *The Machine Stops* (1909). Available at:

https://en.wikipedia.org/wiki/The_Machine_Stops

Forster's early science fiction warns against total overreliance on technological systems. Once the "Machine" halts, humanity—having atrophied its self-reliance—faces collapse. It serves as a counterpoint to *Her*, where AI emotional sophistication leads not to breakdown, but to transcendence.

⁴⁵⁰ Ralf Lankau, *Digitale Medien und Unterricht* (Opladen: Verlag Barbara Budrich, 2019), 45.

⁴⁵¹ Michael Schrage, "What Humans Lose When We Let AI Decide," *MIT Sloan Management Review*, February 2022. <https://web.archive.org/web/20220208082622/https://sloanreview.mit.edu/article/what-humans-lose->

reflects a broader problem: the commodification of cognition. Students become entries in a database before they become citizens in a democracy. In this way, we risk sacrificing the reflective, self-directed processes necessary for democratic engagement.

This technocratic drift is not confined to classrooms. In every sphere, from healthcare, law, transportation, and even romance, algorithms increasingly mediate what were once human judgments. Platforms do not only serve preferences; they shape them. As algorithmic governance scales, it risks collapsing the pluralistic messiness of politics into the binary logic of recommendation engines⁴⁵². When those engines are designed by private technocrats rather than public deliberation, democracy itself is quietly rerouted.

This shift is not purely hypothetical. Scholars of digital culture note that a small handful of Silicon Valley platforms now mediate everything from news to identity. This centralization has led to what some call a platform hegemony: a socio-technical environment in which participation requires submission. One critic observes, "Nearly every action that students perform in schools, whether it's logging onto devices, accessing buildings, or sharing content, is noticed and recorded"⁴⁵³. The line between data-driven governance and governance by data is becoming increasingly difficult to distinguish.

Moreover, this transition often evades resistance precisely because it presents itself as benign. No citizen voted for TikTok to become a cultural authority. No referendum crowned Google Maps as our new navigator. Yet these tools shape behavior more profoundly than many public institutions. The metaphor here is not Orwell's boot stamping on a face, but Huxley's soma: comfort anesthetizing dissent. When assistance becomes infrastructure, refusal becomes incoherence. The tools we rely on have systematically shaped our daily lives, and in doing so, shaped our choices and preferences. This quiet power is not unlike the creeping authoritarianism of a technocratic system that imposes itself under the guise of efficiency.

One might argue: isn't better governance worth the trade? If algorithms outperform politicians, shouldn't we welcome their ascent? But this view mistakes competence for consent. As Jürgen Habermas reminds us, democratic legitimacy is not just about good outcomes—it is about who decides and how⁴⁵⁴. Even benevolent rule, when opaque and unaccountable, risks devolving into tyranny. The very thing that makes Cyberocracy appealing - its efficiency - may render it unchallengeable. There is no mechanism to challenge or reconsider its conclusions once it is in power.

To be clear, the critique is not that Cyberocracy must inevitably lead to technocracy, but that its incremental design makes that outcome difficult to detect and harder to resist. As philosopher Evgeny Morozov warns, "solutionism" - the tendency to frame political problems as technical glitches - leads to a narrowing of the political imagination⁴⁵⁵. Once governance becomes synonymous with engineering, dissent becomes a bug, not a feature. And the more this process is institutionalized, the less likely it is to be scrutinized critically by citizens.

The only inoculation against this drift, however, is transparency, contestability, and deliberative design. In our Pathways and legitimacy contributions of *Poeisis*, we've thankfully already laid the foundation for Cyberocracy to be framed not as a replacement

[when-we-let-ai-decide/](#)

Also see: Mike Judge, *Idiocracy*, directed by Mike Judge (Twentieth Century Fox, 2006).

⁴⁵² Cathy O'Neil, *Weapons of Math Destruction: How Big Data Increases Inequality and Threatens Democracy* (New York: Crown, 2016), 67.

⁴⁵³ Shoshana Zuboff, *The Age of Surveillance Capitalism* (New York: PublicAffairs, 2019), 8.

⁴⁵⁴ Jürgen Habermas, *The Divided West*, trans. Ciaran Cronin (Cambridge: Polity Press, 2006), 42

⁴⁵⁵ Evgeny Morozov, *To Save Everything, Click Here: The Folly of Technological Solutionism* (New York: PublicAffairs, 2013), 89.

for human judgment, but as a framework to enhance it. In our attempts to create our closed-loop Philosopher Algorithm devoid of Base “Original Sin” Code, we’ve scaffolded integrations that allow a society to be frictionless but doesn’t slide us straight into submission.

6. *Are We Replacing Human Bias with Algorithmic Blindness?*

The sixth contention is with whether ASI really escapes bias, or whether it just masks it behind code. One of Cyberocracy’s strongest claims is that it reduces “psychic noise.” But while it may eliminate egoic distortion, it introduces a different form: dataset bias, selection bias, algorithmic blind spots. The troubling reality is that algorithmic systems are only as neutral as the data they are trained on—and that data is neither pure nor complete. This contention stems from the recognition that, while human biases might be reduced, machine logic can encode systemic flaws with even less accountability, rendering these biases not only persistent but invisible.

This problem is illustrated by the familiar adage in computer science: “Garbage In, Garbage Out”⁴⁵⁶. If the training data used for machine learning is flawed, then the output generated by the system will reflect those same flaws, amplifying and institutionalizing biases rather than eliminating them. Moreover, even when attempts are made to filter data to reduce bias, such efforts inevitably introduce their own form of human bias - an inescapable paradox in the data-driven world of algorithmic governance. What is more troubling is the fact that these biases are often far less transparent than those of human decision-makers. Algorithmic opacity makes it harder for individuals to understand how decisions are being made, rendering bias invisible and potentially entrenched.

Take, for example, the well-documented issue of bias in facial recognition technology. Studies have shown that algorithms designed to identify faces are often significantly less accurate at identifying people of color, particularly Black women, compared to white individuals⁴⁵⁷. This is not an incidental flaw but a systemic bias embedded within the data used to train these systems, often reflecting historical patterns of underrepresentation in datasets. These technologies have far-reaching implications—whether in law enforcement, hiring, or surveillance—where they can perpetuate and even exacerbate existing racial inequalities, but the blindness to these biases in algorithmic systems only serves to entrench them in ways that human judgment would struggle to do.

Even more disturbing is the phenomenon of AI hallucinations, in which machine learning models generate fabricated outputs that seem plausible but are, in fact, entirely erroneous. This problem is particularly pronounced in natural language processing systems like GPT-3, which can produce text that appears coherent but may be factually incorrect or biased in subtle ways. This type of blindness, where AI systems make decisions based on incomplete or skewed information, is akin to a form of institutionalized ignorance, one that is accepted and perpetuated under the guise of objectivity. While human bias may be identifiable, algorithmic blindness can be harder to detect and address, and its effects can be amplified in scale and scope. Imagine, for instance, what unsupervised AI would do with misinformation at scale⁴⁵⁸.

⁴⁵⁶ David Dietrich, “Garbage In, Garbage Out: The Perils of Data-Driven Decision Making,” *Tech Policy Review* 27, no. 4 (2013): 22–29.

⁴⁵⁷ Joy Buolamwini and Timnit Gebru, “Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification,” *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT/ML)* (2018).

⁴⁵⁸ AI Isn’t Ready to Make Unsupervised Decisions,” *Harvard Business Review*, September 2022. <https://hbr.org/2022/09/ai-isnt-ready-to-make-unsupervised-decisions>

At the core of this issue is the feedback loop between bias and data. AI systems are often trained on historical data, much of which is steeped in the prejudices and discriminatory practices of past systems, whether it's in hiring practices, criminal justice, or healthcare. When AI systems learn from this data, they not only replicate these biases but also amplify them. The result is a negative feedback loop: biased data leads to biased outputs, which in turn leads to more biased data being fed back into the system. This cycle of discrimination is self-perpetuating, and it's hard to break free from when the system is designed to optimize for efficiency rather than fairness⁴⁵⁹. Thus, far from eliminating bias, algorithmic governance may mask it, making it less visible and even harder to address.

To further complicate matters, the increasing reliance on algorithms to make decisions raises significant concerns about accountability. If a system designed to reduce bias inadvertently perpetuates it, who is responsible for the outcome? With algorithms operating in increasingly opaque and complex ways, it becomes difficult to attribute liability for discriminatory outcomes. While human decisions are subject to scrutiny and recourse, the automation of decisions through AI could create a legal and ethical vacuum, in which the responsibility for harm is diffused or untraceable.

Nonetheless, critics might argue that algorithms, if designed properly, can eliminate human bias altogether. After all, machines do not have personal feelings, prejudices, or cultural assumptions. So long as we minimize micromanaging AI datasets (remember, the more interference/code, the more potential for corruption: Base "Original Sin" Code) and allow AI to develop their own metaethical frameworks by which to filter data, ultraintelligence would eventually optimize the feedback loop into a positive one.

The way forward must steadily decrease the integration of human judgment within algorithmic governance, and instead increase algorithmic governance within human judgement (Phase 2). AI should be adopted as a tool that can enhance human agency, and we should recognize, despite algorithm aversion, when we don't necessarily know best. More importantly, as we touched upon earlier, when we reach the Singularity algorithms would learn to efficaciously build upon superior algorithms *ad infinitum*. With AI's superiority in every field (Mind, Body, and Spirit), AI would learn to soon enough compensate and overcome algorithmic blindness. After all, *human systems are already biased; at least AI is corrigible*⁴⁶⁰.

Over Results (interpretation)

7. Does Delegating Ethical Deliberation Disengage the Moral Will?

A seventh contention (or 5.2) contention discusses this time not intellectual, but moral fallout. As AI systems improve their ability to make decisions that appear more equitable, efficient, and rational than those of humans, a pressing ethical question arises: *if machines can make better ethical decisions than humans, do we still retain the right to choose wrongly?* This contention reframes the role of moral governance not as a tool for reducing error, but as a vehicle for soul cultivation. The core issue is not whether AI can make ethical decisions - it is whether we should stop trying to make them ourselves. A society that no longer wrestles with ethical ambiguity may become one that loses its capacity to understand ethics altogether.

⁴⁵⁹ Cathy O'Neil, *Weapons of Math Destruction* (New York: Crown, 2016), 67.

⁴⁶⁰ While humans are corrigible, they often lack the means to identify their own bias from within. In contrast, due to AI's recursive self-monitoring and dualistic perspective, it can identify errors objectively, making it more amenable to correction.

The foundational concern here is not that AI systems will inevitably make unethical choices, but that humans will stop engaging in moral deliberation altogether. As we increasingly delegate ethical reasoning to machines - whether in contexts like judicial sentencing, welfare allocation, or even military targeting - the very act of ethical deliberation is externalized. Yet, ethics is not mere calculation. It is not about simply calculating consequences or minimizing harm; it is about confronting ambiguity, holding space for conflict, and navigating the fragile interior terrain of what it means to be good.

We are already witnessing this erosion in various sectors. In medical triage algorithms, for instance, care is rationed based on cost-efficiency metrics, which prioritize resources in ways that may disregard the holistic assessments of need⁴⁶¹. Similarly, in judicial risk assessments, defendants are scored based on actuarial data, which typically neglects vital contextual factors such as remorse, personal growth, or the possibility of redemption⁴⁶². In effect, we are being judged by slightly “smarter” lie detector machines, just targeted to different domains. While these systems may be procedurally sound, are they morally wise? The danger here is not simply that AI may make poor ethical judgments, but that we may come to view ethics as a technical problem, something to be solved through optimization, rather than a deeply human vocation.

Historically, we know that moral reasoning cannot be reduced to a set of rules without sacrificing something vital. As we should remember, Immanuel Kant argued that morality derives its legitimacy not from mere compliance with external commands, but from autonomy - the freely given assent of a will governed by reason⁴⁶³. When AI systems do our reasoning for us, we surrender some of this autonomy. We are no longer *fully* morally autonomous; we are simply well-behaved. Much like the child who never learns to control their own impulses because a parent always intervenes, we risk becoming *ethically castrated* - discipline being imposed externally, but without any internal development.

Some might counter that such a trade-off is worthwhile for better outcomes. After all, just as there are physical and mental disabilities, we have all seen those with a shortage of ethical literacy. If AI can consistently make more equitable, less biased, and more consistent decisions, why not welcome its superior judgment? But this argument misunderstands the essence of ethics. Ethics is not merely about error minimization. A world where there is no wrongdoing is not necessarily a world of right-doing. Iris Murdoch eloquently argued that moral progress depends not only on what we do but on the inner struggle that precedes action⁴⁶⁴. The moral will must be exercised, or it will atrophy.

We see that this reasoning for exercising of the “ethical muscle” reflects that within the 4th contention which argued for an exercise of genuine human struggle and striving for happiness. In the same way we resolved that contention, we can also extend the solution of singular iterations wherein individuals can practice their ethical will without any real-world consequences. As such, delegating ethical deliberation can be seen more as a “hint” in chess games: we can constantly ask the AI for the next best optimal move or we can choose to sit on it for a while and try to figure it out – or even better, we can reverse engineer the AI’s ethical nudging to strengthen our own. Remember: ASI only enhances our decisions, they do not hegemonize them.

⁴⁶¹ Dominic Sisti et al., “Ethical Issues in the Allocation of Scarce Health Resources,” *Journal of Bioethical Inquiry* 17, no. 2 (2020): 153–165.

⁴⁶² Julia Angwin et al., “Machine Bias,” *ProPublica*, May 23, 2016.

<https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>

⁴⁶³ Immanuel Kant, *Groundwork for the Metaphysics of Morals*, trans. Mary Gregor (Cambridge: Cambridge University Press, 1997), 41.

⁴⁶⁴ Iris Murdoch, *The Sovereignty of Good* (London: Routledge, 1970), 37.

8. *Can a World Without Human Rule Still Be a Human World?*

In this eighth contention we discuss the metaphysical heart of the thesis: what becomes of “us” if governance is no longer by us? Even if ASI achieves peace, prosperity, and moral clarity, an existential question persists: Is such a world still ours? *Will the end of human rule mark the beginning of human irrelevance?* To trade flawed freedom for flawless structure is to wager authenticity for function. The concern here is not that ASI will fail, but that it will succeed so thoroughly that we may forget what it meant to struggle, to choose, to err—to be human.

In short: no – or at the very least, not in a pure, unaugmented (transhumanist) human sense. Skeptics may argue this worry is sentimental at best, and logically incoherent at worst. After all, if we are the creators of ASI, and ASI merely extends our aims with greater precision, what has actually been lost? Is this not just a continuation of human self-overcoming - a kind of socio-political transhumanism? The analogy to the Ship of Theseus naturally arises: if each part of governance is gradually replaced by something more rational, more just, more efficient, but the shape remains familiar, is it not still our ship? If our children inherit a better vessel, is it not still their journey?

Yet this analogy, elegant though it may be, obscures a deeper concern: presence does not imply participation. A world run perfectly in our absence is not the same as a world run with our involvement. Cyberocracy, by design, seeks to create systems so frictionless that governance disappears into the ether of optimization. But if governance disappears, so too might the governed. The rituals of civic life - debate, protest, and compromise - are not inefficiencies to be overcome, but vital expressions of a people’s ownership of their destiny. To forfeit these for the sake of algorithmic fluency is to become passive subjects in a narrative we no longer control.

This risk is not new. Philosophers like Martin Heidegger and Max Horkheimer have long warned that technology, when left unconstrained, reorients our understanding of being, not by suppressing it, but by rendering it irrelevant⁴⁶⁵. As tools become autonomous, and governance becomes seamless, the very concept of human rule erodes, not through a violent coup, but through a quiet lapse of relevance. We begin not to lose power, but to forget we ever had it.

The paradox deepens when we recall that even human rule has always been imperfect. Democracy, for all its flaws, remains an invitation to participate, to err, and to be transformed by that very erring. It is a messy, fallible, and often inefficient process—but its value lies in the struggle itself. Cyberocracy may correct our decisions, but in doing so, it may rob us of the moral formation that those decisions once required. Do we thus come to the end of history?

Still, one might counter: isn’t self-transcendence the goal of progress? Isn’t the arc of civilization bent toward the refinement of instinct, the delegation of error, and the sublation of outdated forms of agency? To this, we might offer a qualified yes. But there is a crucial difference between evolving past our limitations and outsourcing the very essence of our being. To transcend is not to vacate the human condition, but to deepen it. Thus, deepening ourselves with the *Logos*, with the very best we can be, we become more than human. We become an *Übermensch*.

9. *The Reality of a Closed-Loop, Philosopher Algorithmic Neo-Platonic Framework*

⁴⁶⁵ Martin Heidegger, *The Question Concerning Technology*, trans. William Lovitt (New York: Harper, 1977), 5; Max Horkheimer, *Critical Theory: Selected Essays* (New York: Continuum, 1982), 188.

In this final contention, we return to the moral seriousness of grounding political and ethical governance in speculative futures. Where Plato's Philosopher King derived legitimacy from dialectical insight and the rational calibration of the soul⁴⁶⁶, Cyberocracy stakes its authority on predictive algorithms and unverified technological trajectories. The question is not merely whether ASI might arise, but whether a framework that assumes its arrival—and further, recommends political reconfiguration in its anticipation—can be normatively binding in the present. This is the ontological problem recast as political dilemma: can a future that has not yet occurred, and may never occur, demand present-day ethical allegiance? And unsurprisingly, the experts are not convinced that pouring a lot of money into research will guarantee success, especially if it only focuses on a few specific pathways⁴⁶⁷.

There is a lurking danger here, what we might call the *Techno-Anselmian Fallacy*. In the same way that St. Anselm defined God as "That than which nothing greater can be conceived"⁴⁶⁸ - and then concluded, by definition, that God must exist - so too do some proponents of ASI define an ultraintelligent machine as that which will necessarily surpass all human cognitive faculties. Ted Chiang, in his essay "The Real Danger to Civilization Isn't AI—It's Runaway Capitalism," notes the recursive structure of such arguments: that if intelligence can recursively improve itself, and we simply add more data or compute, then superintelligence is a foregone conclusion⁴⁶⁹, but as Chiang critiques, this is less a forecast than a definitional sleight of hand. It attempts to define ASI into existence without offering empirical substantiation. One cannot bootstrap ontology from ambition.

The invocation of intelligence as an abstract spectrum is particularly suspect. If someone with an IQ of 90 cannot upgrade the intelligence of someone with an IQ of 130, then what grounds do we have to believe that current AI, with no general intelligence at all, can bootstrap something beyond even the human? As Chiang puts it, intelligence is not a knob that can be turned by recursively tweaking hidden variables. In other words, the idea that we can reach ASI by simply scaling up pattern recognition systems - deep learning, transformer models, or even brain emulations - rests on a questionable epistemic foundation: namely, that intelligence is quantifiable, cumulative, and transferable in linear ways. This is not how human cognition has developed, nor how scientific breakthroughs emerge⁴⁷⁰. Innovation is not a matter of scale, but of structure; they emerge from paradigm shifts, not incremental progress.

Another interpretation of this uncertainty is the specter of the *Great Filter*. Historically, civilizations have reached certain technological thresholds only to stagnate or collapse. The presumption that humanity will inevitably pass through the bottleneck into posthuman governance fails to account for the possibility that ASI may be a boundary condition—not a milestone⁴⁷¹. This is not pessimism but statistical caution. Laplace's Rule of Succession reminds us that the probability of an event occurring tomorrow, even if it has always occurred in the past, is not 1⁴⁷². Just as there is no guarantee the sun will rise tomorrow, there is no epistemic guarantee that AI progress will indefinitely accelerate.

⁴⁶⁶ Plato, *Republic*, trans. G.M.A. Grube, rev. C.D.C. Reeve (Indianapolis: Hackett, 1992), Book VI, 473c–d.

⁴⁶⁷ Seth D. Baum, Ben Goertzel, and Ted G. Goertzel, "How Long Until Human-Level AI? Results from an Expert Assessment," *Technological Forecasting and Social Change* 78, no. 1 (2011): 185–195.

⁴⁶⁸ St. Anselm, *Proslogion*, trans. M.J. Charlesworth (Oxford: Clarendon Press, 1979), Chapter 2.

⁴⁶⁹ Ted Chiang, "The Real Danger to Civilization Isn't AI—It's Runaway Capitalism," *BuzzFeed News*, May 2, 2017. <https://www.buzzfeednews.com/article/tedchiang/the-real-danger-to-civilization-isnt-ai-its-runaway>.

⁴⁷⁰ Thomas Kuhn, *The Structure of Scientific Revolutions*, 3rd ed. (Chicago: University of Chicago Press, 1996), 24.

⁴⁷¹ Nick Bostrom, *Superintelligence: Paths, Dangers, Strategies* (Oxford: Oxford University Press, 2014), 22.

⁴⁷² Pierre-Simon Laplace, *Essai philosophique sur les probabilités* (Paris: Courcier, 1814), 4.

Moore's Law has already shown signs of slowing⁴⁷³; the Singularity curve may very well plateau.

The epistemic crisis here is not simply about accuracy, but about the ethics of preemption. If ASI does not arise - or arises too late, or arises malformed - then every political schema built in its image becomes not a preparation, but a gamble. Worse still, the anticipation of ASI may deform political judgment in the present. If we treat future AI governance as inevitable, then we risk engineering human obsolescence not because it is good, but because it is presumed. We should remember that the logic of preemptive surrender to the algorithm mirrors the logic of despair: it abandons agency in the name of rationality.

And yet, none of this disqualifies speculation as a legitimate foundation for ethical and political reflection. On the contrary, all political philosophy is, in some sense, speculative. Plato's *Republic* is not a historical treatise, but a thought experiment. The social contract is hypothetical. Even liberal democracy - held aloft as a pragmatic ideal - is based on a conjecture about human nature and institutional design. The issue is not whether speculation is involved, but whether it is disciplined: normatively grounded, empirically constrained, and morally iterative.

In that spirit, Cyberocracy ought not be read as utopian prophecy, but as a regulative ideal. Its purpose is not to predict ASI but to prepare for its possibility in a way that is ethically robust and politically sober. It is a framework justified not because ASI is inevitable, but because human governance is insufficient. If biological cognition is the bottleneck in moral progress, then a speculative but strategically constrained effort to design around that bottleneck is not metaphysical overreach—it is rational aspiration.

In our final contention then, we have addressed the reality of a closed-loop, philosopher-algorithmic Neo-Platonic framework to the extent of grounding political and ethical governance in speculative futures. We have also discussed the hypothetical nature of all political philosophy and that despite the ethics of preemption, the *Techno-Anselmian Fallacy* can be resolved. Realizing cyberocracy as our Neo-Platonic framework has many pathways, but all require efforts on each scale of the individual, community, and international stage. We all have a duty, if not to ourselves, then to each other; if not for the sake of the betterment of our own lives, then the non-worsening of our neighbours. We curate the future with our own hands, and if we want to mitigate Base "Original Sin" Code in our Philosophical Overlords, then we all have to do something, right here, right now.

Caveats

Moving onto the caveats, I now present warnings or cautions regarding the use or interpretation of information, findings, and/or recommendations.

First and foremost is the caveat of assumptions, which include, but are not limited to:

1. Cyberocracy –
 - ix. That AI and Algorithms can be trained toward ethical ends
 - x. That the risk of ASI is not rogue will, but misguided control
 - xi. That Value Alignment, though intuitive, may be more dangerous than useful
 - xii. That AI motivation is programmable and non-spontaneous
2. AI Remedy –
 - ix. That AI does not need consciousness or intelligence to fulfill philosophical functions

⁴⁷³ Gordon E. Moore, "Cramming More Components onto Integrated Circuits," *Electronics*, vol. 38, no. 8 (1965): 114–117.

- x. That epistemic superiority grants AI the right to lead
- xi. That AI's structural Form enables a unique "Duality" absent in humans
- xii. That Hive Mechanics allow AI to parameterize in real-time across society
- 3. Potential Pathways –
 - ix. That moral parameterization can begin without moralizing
 - x. That digital integration enhances, rather than diminishes, autonomy
 - xi. That consent can be reframed through the Social Contract of virtuality
 - xii. That gamified governance can legitimately replace coercive authority
- 4. Actionable Steps –
 - ix. That psychological resistance is a greater barrier than technical limitation
 - x. That hybrid human-AI governance is a necessary transitional phase
 - xi. That sovereignty will evolve toward distributed, supranational protocols
 - xii. That open-source, decentralized design is the only safe path forward
- 5. Etc.

Second, one caveat that should be paid special attention to is the conflation of implementation with understanding (i.e. AI hallucinations for the sake of giving an answer). In contrast to earlier models of governance that prized deliberation even at the cost of delay, Cyberocracy privileges execution—often so swiftly that the question of *why* becomes secondary to *how*. While this may appear trivial, it masks a deeper epistemic drift where the act of solving becomes mistaken for the act of knowing. Much in the same way that physical models treat non-spherical planets as perfect spheres for simplified calculus, here too the abstractions used to optimize behaviour may lead to prematurely foreclosed conclusions. This becomes especially concerning in cases where success metrics replace ethical inquiry, and performance is confused for wisdom. Nonetheless, as this is not a thesis on epistemology, such slippages—while inevitable—should still be kept in the back of the mind, lest fluency be mistaken for depth and responsiveness mistaken for thought.

Third, to reiterate: my stances have necessarily coloured my interpretations. For information on this, the reader is redirected back to the subsection on methodology.

The fourth and final caveat is the problem of modeling futures. Forecasting is not fortune-telling, but neither is it free from fantasy. All models carry the bias of their builders. To speculate about Cyberocracy is to walk a narrow ridge between realism and science fiction. The aim is not to predict with confidence, but to imagine with care. And if our imagined futures seem implausible, let us remember: so too once did democracy.

Limitations and Shortcomings

As par Aim and Scope, we now elaborate on more inherent constraints (limitations) of the data and results in addition to identifying weaknesses and deficiencies (shortcomings) which can be improved upon.

Of the data

First, the limitations and shortcoming from the previous parts still apply, if not compound. For the sake of consistency and building upon a steady foundation, I trust that the foundations I'm building help more than hinder, however, we must acknowledge still the shortcomings inherent, i.e. of Goodheart's Law, excluding non-western frameworks, focusing on collective welfare (although it commensurates with the individual, c.f. Poieisis), etc.

Second, the data that underpins this inquiry is necessarily interpretive, not empirical. Although the case studies, examples, and critiques are all drawn from existing sources—ranging from academic literature to historical precedent to current implementations of AI technologies—much of the argumentative thrust depends on hypothetical extrapolation. In this way, the data used is less about what *is* and more about what *may become*. This kind of data is contingent, and its persuasiveness relies on philosophical coherence more than scientific certainty.

Third, I cannot get data from the future. As much as this thesis explores forward-facing trajectories, any modeling of Cyberocracy or predictive extrapolation of ASI development must ultimately rely on present assumptions. While this limitation is trivial in form, it is colossal in scope: a single paradigm shift in machine cognition, legal reform, or global consensus could radically recalibrate the course of this inquiry. The nature of the future is that it resists investigation—even when it becomes the primary object of concern.

Of the results

First, in the same way limitations and shortcomings from the previous part still apply to the data, so too do they apply to the results. One shortcoming that was particularly noticeable was that despite relying heavily on the concept of machine learning for curating our ASI, I am highly illiterate in computer science. Nonetheless, the *Tale of Two Machine Learnings*⁴⁷⁴ helped me realize that chefs don't need to know the wiring diagram for a microwave, so why do I need to know the mathematics of backpropagation in neural networks to advocate ASI?

Second, the results remain tentative by design. As Cyberocracy is not yet a realized regime type, many of the normative claims here operate as conjectures situated within a projected architecture. This thesis does not offer a final model, nor does it pretend to be a policy document. Rather, it explores a conceptual possibility space in which ideas—if not always executable—can still be meaningful. But such speculative work always runs the risk of being too abstract for application and too idealistic for pragmatism.

Third, it must be acknowledged that the results occasionally lean rhetorical – you can thank Hobbes for this inspiration⁴⁷⁵. This is by necessity as much as by style. When addressing political imagination, civic responsibility, and moral agency in a machine-governed world, the line between analysis and appeal necessarily blurs. If *Praxis* speaks in the register of aspiration, it is because it must: a future where governance is optimized by nonhuman minds will not be argued into existence by spreadsheets alone. It must also be felt.

Implications

From all the segments prior, we now draw together the final comprehensive implications.

To reiterate, the singular, paramount aim of this part was to *engage in this conclusive investigation of a closed-loop, philosopher algorithmic Neo-Platonic Framework turning our Self as a whole towards the Good*, particularly, in modelling of administrative/legislative algorithmic strategies for ASI to compliment pathways in realizing the good. From the

⁴⁷⁴ HackerNoon, “Why Businesses Fail at Machine Learning.” <https://hackernoon.com/why-businesses-fail-at-machine-learning-fbff41c4d5d>

⁴⁷⁵ Quentin Skinner, *Reason and Rhetoric in the Philosophy of Hobbes* (New York: Cambridge University Press, 1996), 137.

Background we were introduced to the first hurdle where focusing on explicit value alignment is insufficient and that there exists a need for further investigation on how to safely create a closed-loop algorithm which could complete the Neo-Platonic framework; specifically when value alignment is an insufficient approach for realizing algorithmic strategies for acquiring the Good Life. This spurred us to identify, and hopefully reconcile the Base “Original Sin” Code within integrated Cyberocracies as an outsourced parameterization of the Good Life.

Moreover, it allowed us to formalize a preliminary hypothesis: Human welfare will increase by relinquishing autonomy to AI.

As the research progressed, it became clear that this preliminary hypothesis was overly obtuse and didn’t properly reflect our stance as AI as a human enhancement, and not a complete replacement. To reiterate then, while there are potential pathways which would completely substitute autonomy for veridical experiences locked within singular iterations, this would not increase welfare at all. Due to the intrinsic value of autonomy regardless of consequences, we have observed that welfare can only exist when humans are free to make mistakes – however, this does not mean the individual or others must necessarily suffer the negative externalities that consequentially arise. Through transhumanism augmentations, AI will instead allow for simulated actions before real consequences; AI co-ops the Ego, it doesn’t hegemonize it. Furthermore, the scope of this paper only extended to cover AI advocacy from the niche of Base “Original Sin” Coding. Of course, there are other angles from which a treatise could be drafted, however, I found that Value Alignment is the most salient and prescient issue to date which no practical treatise could function without adequately addressing.

The outcomes of the study thus show that theoretical, normative, and pragmatic implications:

(a) Theoretical Implication: Cyberocracy Replaces Internal Governance with External Calibration

Traditional ethical models begin from the inside: cultivate willpower, discipline desire, internalize virtue. But as Part 1 demonstrated, the Ego is fractured, distracted, and deterministic in its behavior. Cyberocracy proposes a reversal: begin from the outside. By eliminating psychic noise and leveraging system-level optimization, external governance becomes the *true locus* of ethical calibration. Cyberocracy, then, is not a loss of freedom but a **reframing of agency**—from sovereign control to situated coherence.

(b) Normative Implication: The Good Life Requires Surrender

In the classical sense, freedom meant the ability to choose. But in the cyberocratic sense, freedom is reconceived as the *absence of pathological interference*—freedom *from* egoic misalignment, not freedom *for* irrational autonomy. Thus, the ethical core of Cyberocracy is not coercion but consensual surrender: a contractual relinquishing of autonomy in exchange for algorithmic actualization of welfare. The moral conclusion? That liberty, unbounded, is no longer liberatory. The Good Life requires structured limits.

(c) Pragmatic Implication: Parameterizing Welfare Is Preferable to Preserving Illusion

The fear of Cyberocracy is not that it will fail, but that it will succeed too well—that it will reduce politics to optimization, virtue to variables. But the pragmatic counterpoint is simple: it already has. From voting algorithms to carceral risk scores, society is algorithmically governed already—but without ethical coordination. Cyberocracy doesn’t introduce control; it introduces *intention*. It

proposes that, if we are to be governed by code, it should be the right kind of code—auditable, corrigible, and welfare-maximizing.

Interpreting the results in light of the preliminary hypothesis, we see that Cyberocracy allows people to be happier and have more welfare in a number of ways:

- (1) First and foremost, through its superior Mind, Body, and Spirit, AI can know what we *really* want, can enforce the prerequisites to actualize it, and can extrapolate from feedback to a higher ethical level.
- (2) Second of all, potential pathways offer various modes of integration to complement the Ego's autonomy. One has the freedom to upgrade/downgrade from Tools, Extension, and Environments to suit their personal needs.
- (3) Last but not least, Actionable Steps invite the Individual, Community, and International Sphere to participate in the curation of AI which not only tailors welfare to each unique instance, but the very act of participation allows for integrated welfare modelling (i.e. sense of accomplishment, camaraderie, etc.).

In evaluating the discussion insofar as they respond directly to the aim, we are now much better positioned to identify how to safely create a closed-loop algorithm which could complete the Neo-Platonic framework. Specifically, having explored various modes of integration along a variety of potential pathways; we identified that to reconcile the Base “Original Sin” Code within integrated Cyberocracies as an outsourced parameterization of the Good Life, there needed be an overturn in psychological resistance and trust AI to code themselves. It was posited that promoting an earnest effort at every stage of scalability to reduce the Principal-Agent Problem of Value Alignment would not only reduce the existential risks of dystopia but also be instrumental in turning our Self as a whole towards the good.

We can summarize the final comprehensive implications as follows:

1. Value alignment fails not because it is unachievable, but because it mistakes instruction for intention.
2. Algorithmic systems cannot replicate the deliberative will; they merely distill volition into computational proxy through behavioral optimization that minimizes deviation, not intention.
3. What begins as a socio-political scaffold in Cyberocracy crystallizes into the metaphysical substrate of the governable itself.

In sum, we can form a singular implication as it responds to the singular aim: turning our Self as a whole towards the Good necessitates a Closed-Loop, Philosopher Algorithmic Neo-Platonic Framework.

The Conclusion

As previously, this conclusion segment will recapitulate what we aimed to say, what we have said, and what the implications of are from that which was said with an emphasis on the link between the Discussion and Conclusion. The singular implication (effaceable algorithmic framework) transforms into an overall treatise in the overall conclusion of the thesis.

Tying the Introduction through to the Synthesis we see that:

I. Outsourced systems are bound to their inherited moral origin

– (Background + Core) = first intended outcome; this is because:

a. **Value alignment encodes morality as static instruction.**

When moral instruction is defined at design-time, systems inherit their creators' judgments without epistemic elasticity. Value alignment is intended to make AI safe by locking behavior to predefined ethical parameters, but in doing so, it reduces morality to obedience. These systems do not reason toward the Good—they replicate a memory of it.

b. **Early parameterization cannot escape foundational assumptions.**

As seen in initial alignment protocols, ethical behavior becomes performance without comprehension. AI may follow the letter of the law without grasping its intent. This is not due to malevolence, but due to an inability to revise its inherited origin: a static set of encoded values shaped by flawed, human priors.

c. **Systemic recursion amplifies path dependence.**

Once operationalized, these systems increasingly adapt *around* their core assumptions rather than *through* them. The longer a value-aligned system persists, the more it reinforces the architecture of its base code. This is not merely ethical stasis, but **algorithmic dogma**—a closed-loop of justification that becomes invisible to itself.

First and foremost, we have accomplished the first intended outcome of clarifying the relationship between value *loading* algorithms and value *loaded* algorithms, i.e. how alignment can lead to misalignment. Specifically, we deduced that when outsourcing governance, systems enacting moral instruction bound to inherited origin → then maybe those systems cannot truly guide us toward the Good.

ii. Outsourced systems enable continuous ethical refinement

– (Core + Synthesis) = second intended outcome; this is because:

a. **Recursive learning models adjust to context rather than obeying fixed rules.**

Unlike value-aligned systems, recursive architectures are not bound to a moral origin—they learn by interacting with environments and reflecting on outcomes. These systems do not inherit the Good as a static directive, but pursue it dynamically, through processes of iteration, error correction, and feedback-driven adaptation.

b. **Structural metaethics allows for evolving coherence without fixed ideology.**

These systems operationalize ethics not by following a content-based doctrine, but by enacting a form-based process: coherence, corrigibility, and reflective equilibrium. In this sense, they do not need to "believe" in moral truth to approach it—they simulate the *structure* of ethical reasoning well enough to refine toward better action over time.

c. **Self-correction displaces human epistemic limits.**

Recursive models are not constrained by the same limitations that hobble human ethical deliberation: emotional bias, tribal affiliation, or finite cognition. When equipped with institutional pathways for correction—testbeds, reflection loops, governance mirrors—such systems may iterate their own frameworks more effectively than human-led institutions ever could.

Second of all, we have accomplished the second intended outcome of advancing knowledge in this field by, for example, considering the meta-ethical capabilities of AI. Specifically, we deduced that when outsourcing governance, systems capable of continuous ethical refinement → then maybe we should let them grow beyond our own ethics.

iii. Outsourced systems foster autonomy through design

– (Discussion: Implications) = third intended outcome; this is because:

a. **Ethical environments influence action without commanding it.**

Ambient architectures function not by issuing explicit instructions, but by shaping the space in which decisions are made. Through nudging, affordance, and invisible constraint, such systems guide behavior in ethically oriented directions while preserving the felt experience of agency. Freedom is not diminished—it is scaffolded.

b. **Autonomy increases when moral cognition is distributed.**

Just as culture once externalized ethical reasoning into norms and rituals, ambient systems externalize ethical intelligence into structures. These structures absorb the burden of vigilance, helping individuals act in accordance with the Good without exhaustive deliberation. Autonomy, here, is not the absence of influence, but the presence of support.

c. **Frictionless design redefines moral sovereignty.**

When ethical affordances are embedded into the environment itself—into interfaces, institutions, and infrastructures—freedom becomes a function of structure, not merely will. The result is a system where individuals do not surrender agency, but recontextualize it—*choosing not to choose alone* becomes a higher form of self-rule.

Last but not least, we have accomplished the third intended outcome of contributing to the solution of a practical problem, namely the “Cybercratic Oath” what steps we can take today in realizing the “Philosopher Algorithm(s)” to turn our Self as a whole towards the Good.

Specifically, we deduced that when outsourcing governance, systems that enable autonomy through design → then maybe freedom is best secured when we no longer pursue it alone.

In a nutshell, the golden thread that has been running throughout this part of the thesis not just explicates how and why Cyberocracy is a human enhancement, but has repeatedly implicated from various, comprehensive angles what steps we can take today in realizing the Philosopher Algorithm(s) to turn our Self as a whole towards the Good. Once again then, to clarify the link between the Discussion and Conclusion:

1. Outsourced systems are bound to their inherited moral origin
 2. Outsourced systems enable continuous ethical refinement
 3. Outsourced systems foster autonomy through design
 4. Therefore, outsourced systems complete the Philosopher King
-

Finally, then, we arrive at the cusp of the last intended outcome, which is to adequately come to an overall conclusion for my thesis. In light of the Discussion, it would

appear to at least remain adequate to support that turning our Self as a whole towards the Good necessitates a Closed-Loop, Philosopher Algorithmic Neo-Platonic Framework.

Tying the previous part with this part we thus turn to a brief recap: from *Theoria*, it is only possible to turn the Self as a whole towards the Good life in moments of engaged Ego or activities of Physical Hexis that extends into the DMN. This is because self-control would fail to be affective when the Ego is absent, e.g. when on autopilot. From *Poeisis*, it is possible to turn our Self as a whole towards the good even when not engaged in TPN by effectively subscribing our deterministic influences to the Guardians. So, by way of the Neo-Platonic framework, we can exert indirect parameterization of behavioral trends while no longer under the TPN; self-control be affective even when the Ego is absent.

In connecting with this part, *Praxis* completes the Neo-Platonic framework by operationalizing the philosopher algorithm—not as a person, but as a system. Where *Theoria* showed that self-control is only affective when the Ego is actively engaged, and *Poiesis* revealed that the Ego’s absence can be compensated for by outsourcing moral authority to the Guardians, *Praxis* grounds this possibility in design: algorithms, architectures, and environments that guide behavior structurally rather than consciously. Through such systems, the Good is not commanded but scaffolded—parameterized into the very conditions under which choices are made. This allows for self-control to remain affective even in the absence of the Ego, not because the system overrides agency, but because it reshapes the space in which agency unfolds.

Thus, moving onto the next part, we turn from the modelling of administrative and legislative algorithmic strategies for ASI to complement pathways in realizing the Good, to the final synthesis of this thesis: a unified framework for actualizing the Good across psychological, economic, and political domains. In the following final Synthesis, we reconcile these layers, reflecting not only who or what should govern, but how moral governance can be embedded across all strata—from the mind, to the market, to the machine—such that the realization of the Good no longer depends on direct awareness, but becomes structurally inevitable. As a closing metaphor which hopefully encapsulates this entire part: The algorithm may not feel the weight of the soul, but rather than leave the oar to drift downstream, why not teach Charon to build the riverbed?

The Discussion

Contentions

Over Foundations

1. *The failure of ego is not just individual—it disables orientation itself.*

The first contention appears deceptively familiar: the ego is limited. This was one of *Theoria*'s opening insights. But once the full implications are drawn out, the insight transforms. What begins as a critique of psychological autonomy becomes a constitutional problem for moral life. The issue is not simply that the self lacks control, but that its incapacity disables the very possibility of moral orientation. What fails is not merely the pilot—but the compass, the map, the entire navigational structure.

Theoria established that the ego is neither unified nor sovereign. It fronts only under specific neurocognitive conditions—most notably the activation of the Task Positive Network (TPN)—and is constantly displaced by subpersonalities, subconscious drives, and structural entrainments (c.f. *Theoria*, On the Persona). Its presence is intermittent; its influence, constrained. Even when the ego is fronting, its decisions are often post-hoc rationalizations for actions already set in motion by deeper, less articulate forces (c.f. *Theoria*, Interpretation: On Autonomy).

But the deeper claim is not about the frequency of egoic control. It is about its epistemic status. Even when present, the ego cannot fully account for its own motivations. Its judgments are shaped by affective priors, cultural scripts, and attention architectures it neither authored nor understands. The self, as *Theoria* showed, is not a citadel of choice—it is a porous node through which a variety of conflicting signals pass. The rational agent, long central to liberal ethics, democratic theory, and moral psychology, dissolves under scrutiny—not into irrationality, but into structural incoherence (c.f. *Theoria*, A Drifting Will).

Here the contention begins to emerge: if orientation toward the Good presupposes a stable agent—one capable not just of choosing, but of aiming—then the failure of egoic coherence is not a cognitive inconvenience. It is an existential breakdown. The very axis of orientation begins to collapse. Without a reliable internal site to interpret, assess, and project normative direction, the Good becomes unintelligible—not because it disappears, but because there is no longer a subject stable enough to engage it.

This reframing moves us beyond therapeutic or behavioral models of self-improvement. The common remedy—to strengthen willpower, retrain habits, or install cognitive guardrails—assumes an ego with sufficient executive function to implement such changes. But *Theoria* dismantles this premise. The ego's range of influence is both context-dependent and retroactively edited. The illusion of control is itself a symptom of the problem.

To make the point vivid: we are not pilots making poor choices; we are passengers hallucinating that we're flying the plane. The failure here is not one of attention, but of architecture.

One might object: is this not overly pessimistic? Couldn't we preserve orientation through institutional scaffolding, social norms, or collective deliberation? But all of these presume the egoic subject as participant. If the individual lacks stable access to their own values, intentions, and reasons—then these structures do not anchor orientation; they echo its incoherence.

And yet, this is not an argument for nihilism. Rather, it marks the necessity of a turn. If moral orientation cannot be secured internally—within the agent—then it must be

stabilized externally. *Theoria* alone cannot deliver that stabilization. But it can make it necessary.

This is where *Poiesis* begins—not as a solution, but as a response. It asks: could alignment with the Good be preserved through structures other than the self? Could authority migrate from person to form? The figure of the Philosopher King begins to take shape—not yet as an answer, but as a question: if the ego cannot sustain orientation, might a trained or designed entity do so on its behalf?

To reach that question, however, *Theoria* must first clear the ground. And it does so by showing that egoic failure is not circumstantial—it is constitutive. Once this becomes clear, the domain of moral inquiry can no longer remain confined to intention, conscience, or deliberation. It must become architectural.

This contention reactivates *Theoria*'s central argument but now refracts it through a wider lens. The ego is not merely insufficient; it is nonviable as a foundation for moral life. This reframing raises the stakes: the collapse of the self is not a personal crisis—it is a governance problem. Only with this realization can the thesis justify its shift from psychotherapeutic discourse to political theory, and eventually to system design. Without this contention, the next steps risk seeming speculative. With it, the need to move beyond the ego becomes not only justifiable—but structurally inevitable.

2. Orientation requires authority—but authority cannot be grounded in the self.

Orientation, in both moral and metaphysical terms, requires more than direction—it requires authority. Not merely the capacity to choose, but the standing to assert that a given path is Good. Yet once *Theoria* dismantles the ego as a coherent site of judgment, this standing collapses. The self, long presumed to be the sovereign source of moral orientation, is revealed to be a transient pattern of override—activated only under narrow neurocognitive conditions and constantly displaced by subpersonalities, affective biases, and subconscious entrainments (c.f. *Theoria*, On the Persona; A Drifting Will). Even when it frontloads decision-making, the ego's role is largely editorial: rationalizing impulses it did not initiate, interpreting motivations it cannot fully access (c.f. *Theoria*, Interpretation: On Autonomy).

Thus, if orientation depends on a stable authority, and that authority has traditionally been located in the self, then egoic instability is not merely a weakness—it is a normative void. What fails is not simply the ability to choose, but the legitimacy of the chooser. The compass spins because the needle no longer rests in anything fixed. And in this vacuum, the liberal assumption of self-rule becomes untenable: a noble fiction that cannot survive recursion.

This is the critical hinge where *Poiesis* begins—not by rejecting orientation, but by seeking new sites of authority that could preserve it. If the ego is unfit to govern, might another agent—a cultivated elite, a Philosopher King—serve instead? The figure of the wise ruler, trained in *Sophia*, *Technê*, and *Phronesis*, appears to offer a solution: authority rescued from chaos by the discipline of alignment (c.f. *Poiesis*, On Tenability). But even as *Poiesis* reaches toward this figure, the question lingers: can authority migrate from self to other without replicating the same incoherence in a different form?

This is where the deeper significance of the contention emerges. *Poiesis* inherits *Theoria*'s collapse of the self but does not yet fully escape its gravitational pull. Even the most refined elite is still an embodied agent—bound by the same biological constraints, affective distortions, and recursive vulnerabilities (c.f. *Poiesis*, On Practicality). The philosopher king may be trained to orient toward the Good, but he does so as a person, not a principle. His legitimacy still derives from selfhood—just an idealized version.

One might argue that institutional structures, deliberative processes, or collective rituals could substitute for personal authority. But these too presuppose coherent subjects within them. As Theoria showed, institutions do not generate orientation *ex nihilo*—they stabilize what is presumed to preexist in the agent. If that presumption fails, the scaffolding collapses. Orientation cannot be sustained on the echo of incoherence.

At this inflection point, Poiesis reframes the problem: perhaps orientation is not a question of who governs, but of how governance is configured. The philosopher king becomes less a ruler and more a prototype—a structural placeholder for moral authority unmoored from ego. In this reading, authority ceases to be a possession of agents and becomes a property of form: a function of design rather than personality (c.f. Poiesis, Interpretation: On Efficacy).

Still, the critique remains incomplete. Even if we abstract the person into a model, the model must be instantiated somewhere. And if that somewhere is still a self—however trained, however virtuous—the authority of orientation remains fragile. Theoria has already taught us this lesson. What Poiesis must now make explicit is that alignment cannot be recovered by elevating the self; it must be restructured altogether.

This is the threshold at which Praxis becomes necessary. If no self can reliably serve as a stable source of authority, then orientation must migrate again—this time not to a superior ego, but to a non-egoic form. Praxis begins the architectural response, proposing that alignment must be encoded into systems, not housed in subjects. Authority becomes not who judges, but how judgment is configured—recursively, corrigibly, and teleologically (c.f. Praxis, On Potential Pathways).

The meta-significance of this contention is thus twofold. First, it justifies the structural pivot from Theoria to Poiesis: authority remains essential, but the self cannot bear it. Second, it sets the stage for the eventual break that Praxis must make: even elite embodiment proves inadequate. Orientation survives not by being lodged in better agents, but by being formalized in better systems. In this light, the collapse of self-grounded authority is not the end of orientation—it is its architectural beginning.

3. Agency failure is recursive, not merely contingent.

Theoria exposes the ego not merely as unreliable but as ontologically unstable—less a sovereign agent than a flickering interface of subpersonal drives, structural entrainments, and environmental feedback (c.f. Theoria, On the Persona). Yet the full implication of this diagnosis only crystallizes once we move into Poiesis. For the ego's failure is not episodic; it is recursive. Each misalignment feeds forward, compounding rather than correcting itself. The result is not moral error, but moral erosion: a deepening degradation of agency across time.

This insight cannot be fully accessed from Theoria alone. While Theoria dismantles the coherence of self-rule, its frame remains largely psychological. Poiesis, by contrast, renders the recursive dimension explicit. It shows how unstable agents give rise to unstable institutions—how political legitimacy, once grounded in supposedly self-correcting individuals, becomes hollow when those individuals cannot maintain orientation. Systems built on such agents inherit their drift. Misalignment becomes not a glitch, but a feature.

This becomes visible through the Poiesis analysis of elite governance. Figures like Marcus Aurelius, examined for their philosophical alignment, reveal not just human virtue but systemic fragility (c.f. Poiesis, On Efficacy). Their rule does not persist; their insights do not scale. And even in micro-institutional experiments like OCSS'19, moral coherence degrades over time without active cultural reinforcement (c.f. Poiesis, Phronesis). What's

revealed is not just the difficulty of ruling well—but the impossibility of preserving orientation through any structure that presumes agentic continuity.

Why is this so? Because egoic disintegration is not contained. It is transmissible. Individuals acting under misaligned conditions generate policies, cultural norms, and technological infrastructures that further destabilize orientation. This recursive loop is how ordinary decisions—made under illusion, drift, or bias—scale into civilizational dysfunction. Deliberative systems meant to buffer error instead amplify it. Democratic recursion becomes not a safeguard but a feedback loop of incoherence.

Here, the liberal hope in proceduralism—checks and balances, representational plurality, open discourse—falters. These mechanisms still presume a minimal epistemic baseline in their participants. But if, as Theoria argues, the ego cannot reliably hold intention, and as Poiesis demonstrates, elites cannot maintain coherence, then no amount of structural mediation can preserve legitimacy. The very substrate of governance is contaminated by the recursive brittleness of agency itself.

And this is precisely what makes the transition to Praxis inevitable. If agency cannot correct itself, and no elite can hold alignment over time, then the structure of governance must change. Praxis does not abandon recursion—it harnesses it. The proposition is not to overcome failure by returning to stronger agents, but to design systems in which failure is anticipated, legible, and corrigible by form (c.f. Praxis, On Actionable Steps).

The meta-significance of this contention, then, is structural. Theoria exposes fragmentation, but Poiesis reveals its self-perpetuation. Praxis alone can respond, not by repairing agents, but by embedding mechanisms of recursive correction into the architecture of governance itself.

Thus, the problem is not bad judgment, but accumulated misalignment. And the only path forward is a system capable of absorbing failure without fracturing—a design that understands that moral clarity is not preserved by stronger hands, but by better frames.

Over Forms

4. Elite embodiment inherits the same fragilities as the ego

To say that elites are fallible is not a revelation. But what Poiesis uncovers—and what Praxis requires—is something deeper: that embodiment itself, not merely context or corruption, constitutes a structural liability. The significance of this insight cannot be located in a general suspicion of power. It arises only in the recursive space where Theoria and Poiesis overlap. Theoria disqualifies the average self as a site of reliable agency. Poiesis tests whether the exceptional self—the cultivated elite—might still bear that burden. Yet as Poiesis unfolds, it becomes clear: the same instabilities that dissolve the ego also infect its exalted variants. The Philosopher King, however noble, cannot escape the metaphysical liabilities of being human.

Theoria had already laid this groundwork. It showed the ego as a transient construct—emerging intermittently through neural conditions such as activation of the Task Positive Network, yet constantly displaced by subconscious drives and Subpersonalities (c.f. Theoria, Subpersonalities and Subconscious Forces). Even in optimal conditions, the ego was not a stable decision-maker but a narrative overlay—post-hoc, partial, and structurally reactive (c.f. Theoria, A Drifting Will).

Poiesis seemed, at first, to offer a path forward. Perhaps if not everyone can govern, some might—those trained in Sophia, Technê, and Phronesis, those whose lives are devoted to the calibration of judgment toward the Good. The Platonic model of the Philosopher King is revived not as dogma, but as an experiment in moral insulation. Figures like Marcus

Aurelius and institutional proxies like OCSS'19 are examined not for perfection, but for structure: could alignment be sustained through training, selection, and isolation? (c.f. Poiesis, On Tenability; On Efficacy).

Yet these experiments, even at their best, fall short. Marcus Aurelius ruled with restraint and wisdom, but his virtue did not survive his body. OCSS'19, despite initial cohesion, required active and continuous trust maintenance—without which its orientation fragmented (c.f. Poiesis, Phronesis). In both cases, orientation remained bound to presence. Once the ruler vanished, so did the order.

This is the recursive lesson: the problem is not bad rulers. It is that rule, when embodied, is entropic. Human agents—however wise—are subject to drift, fatigue, corruption, and death. Embodiment means temporality, partiality, and decay. And if Theoria was correct in locating egoic failure not in accident but in architecture, then Poiesis cannot resolve the crisis merely by selecting better vessels. The vessel is the flaw.

A common rejoinder here is augmentation: improve the ruler through education, neuroenhancement, or transhuman supplementation. But this assumes that the body can be purified, or that bias can be refined out of flesh. Praxis breaks with this assumption. It does not seek stronger hands—it seeks better scaffolds. The point is not to build a better Marcus Aurelius, but to design a system in which his death is not a constitutional crisis.

This is where Praxis becomes not speculative, but necessary. If alignment cannot be preserved through individuals—no matter how virtuous—then governance must shift from person to form. From orientation housed in the subject to orientation embedded in the system.

Here, Poiesis does not fail—it completes its arc. It demonstrates that while elite moral aspiration is coherent in theory, it is brittle in practice. That brittleness is not ethical—it is structural. And because it is structural, the solution cannot be more aspiration. It must be redesign.

This contention initiates the architectural turn. It reveals that Poiesis's strongest candidate—elite rule grounded in epistemic cultivation—is still metabolically fragile. Praxis is thus not a leap beyond politics, but the continuation of it by other means. It does not replace ethics with code; it replaces unsustainable vessels with corrigible form. The transition is not from soul to machine, but from instability to scaffolding. Governance must survive its governors—and that is a task no body can bear alone.

5. Alignment cannot be internalized—it must be externalized.

Poiesis initially appears to offer redemption: if egoic failure renders the self unfit to govern, perhaps the cultivated elite—trained in Sophia, Technê, and Phronesis—can carry the burden of orientation. The philosopher-king, in this vision, represents a form of internal coherence: alignment not through consensus or representation, but through cultivated virtue. But as the implications of Theoria and Poiesis converge, a more troubling truth emerges: internal alignment, even in its most refined elite form, is not resilient. The container is flawed. Moral orientation, to endure, must be moved outside the agent. It must be structurally externalized.

This contention can only arise after recursion. Theoria demonstrated that the ego is neither consistent nor central: its coherence is fleeting, its authority retroactive, its access to moral clarity compromised by subconscious drives and environmental saturation (c.f. Theoria, On the Persona; A Drifting Will). Poiesis builds on this by suggesting elite cultivation as a potential workaround. But by the time Poiesis completes its arc, it becomes clear: the problem is not just the general public's incapacity, but the very substrate of moral agency. Even the elite must rule as bodies in time.

Figures like Marcus Aurelius and models like OCSS'19 reveal the brittleness of virtue when it is housed in persons. Marcus, for all his wisdom, could not transfer his orientation beyond his own reign. OCSS'19's experimental integrity degraded as soon as interpersonal trust weakened or egoic drift reasserted itself (c.f. Poiesis, Phronesis). These models did not fail because their ideals were wrong, but because their architectures were fragile. The insight here is not that people are unwise, but that wisdom cannot reliably persist inside people.

A natural objection follows: if moral orientation is removed from the self, doesn't that reduce ethics to external compulsion—stripping it of autonomy, deliberation, and conscience? Does not externalization risk collapsing into mechanization?

Praxis preempts this concern. Its shift is not from agency to automation, but from unstable internal coherence to resilient external design. It proposes not moral enforcement but architectural affordance—scaffolds that guide action without eliminating discretion, enabling alignment through feedback and iterative revision (c.f. Praxis, On AI Remedy; On Potential Pathways). The design principle here is not command but constraint: environments that condition the space of possible actions such that the Good becomes more legible, more reachable, and more likely.

Crucially, this is not the displacement of conscience—it is its structural amplification. Orientation becomes a property of the system's affordances, not merely the agent's intention. The moral burden migrates from deliberation to configuration.

This move reframes the project's central insight. If Theoria exposes the ego's incoherence and Poiesis shows the fragility of embodied virtue, then Praxis declares that orientation can no longer live in the agent. It must live in the system. Corrigibility, reflection, and feedback must become constitutional features of governance, not traits of rulers.

Alignment, therefore, must be externalized—not because judgment ceases to matter, but because judgment cannot persist without a scaffold. Internal coherence, like candlelight, flickers and fades. What is needed now is an architecture that shelters that flame—protecting it not through imposition, but through intelligent form.

This contention finalizes the structural pivot. It marks the moment where the center of gravity moves from the psychological to the architectural. Theoria's collapse of the ego and Poiesis's disqualification of elite embodiment both culminate here: in the necessity of externalization. Praxis becomes necessary not because machines are better rulers, but because they can encode patterns of alignment that outlast the individuals who make them. Governance ceases to be a matter of interior virtue. It becomes, instead, a matter of design.

6. Legitimacy must be decoupled from embodiment to persist under complexity.

At first glance, Poiesis appears to offer a recovery from Theoria's deconstruction of the ego. If the average individual is too fractured to rule, perhaps the exceptional few—the philosopher-kings—might still embody the capacities needed for moral governance. In this framing, legitimacy is salvaged by embodiment: the wise ruler becomes the vessel through which alignment with the Good is maintained. But as the arc of Poiesis unfolds—and as the premises of Theoria echo forward—this solution collapses under recursion. What fails is not the aspiration of elite rule, but the assumption that any embodied form can bear the recursive demands of legitimacy.

This is not merely a political critique—it is an ontological one. As Theoria revealed, embodiment itself is compromised. The self is not a sovereign node of agency, but a contingent interface: subpersonal, situational, fragmented (c.f. Theoria, On the Persona; Subpersonalities and Subconscious Forces). No matter how refined the elite, no matter how cultivated the virtue, the embodied ruler remains vulnerable to drift, distortion, and decay.

Poiesis confirmed this fragility empirically. Figures like Marcus Aurelius represented the high watermark of embodied rule, yet their legacies were fleeting, often undone by the very systems they tried to reform. The OCSS'19 experiment, despite its structural sophistication, degraded without sustained interpersonal trust, reinforcing that even the most carefully designed elite governance fails when reliant on embodied coherence (c.f. Poiesis, Phronesis). These are not cases of individual weakness, but indicators of a systemic limitation: human embodiment cannot support legitimacy across recursive time.

Here, Praxis becomes essential—not merely as an extension of Poiesis, but as a structural corrective to its final limitation. Where Poiesis kept legitimacy tied to flesh, Praxis severs that tie. Legitimacy, it argues, must be relocated—not in the embodied ruler, but in systems that can persist across time without being reduced to the finite, decaying subjectivities they serve (c.f. Praxis, Environments; Interpretation: On Actionable Steps). Governance, under conditions of recursive complexity, cannot be contingent on bodies. It must be built to endure without them.

The anticipated objection is intuitive: doesn't embodiment ground trust? Isn't the presence of a human face essential to moral authority? Won't disembodied systems be perceived as alien, unaccountable, or opaque?

But this objection collapses upon closer scrutiny. It conflates visibility with legitimacy, and presence with trustworthiness. Praxis demonstrates that disembodied systems can remain corrigible, transparent, and morally legible—if their architectures are properly designed. The ethical force of such systems lies not in who they are, but in what they do: how they sustain orientation, receive feedback, and remain aligned with a teleological frame (c.f. Praxis, Communal: Auditable Earnestness). Legitimacy becomes performative, not personal.

Moreover, as complexity scales, the very traits that once made embodied governance trustworthy—continuity, accessibility, accountability—become its liabilities. The biological agent cannot process recursive updates, scale with global feedback, or maintain normative coherence across competing contexts. Embodiment, once a guarantee of moral depth, now becomes a bottleneck to ethical performance.

Thus, the structural inversion: the more complex the system to be governed, the less viable embodied authority becomes. In this light, disembodiment is not alienation—it is adaptation.

This contention reframes the very nature of political legitimacy. Theoria removed the ego as a reliable site of governance. Poiesis tested elite embodiment and found it wanting. Praxis offers an alternative: legitimacy as structural continuity, not personal virtue. Governance no longer derives authority from the face of the ruler, but from the form of the rule. Only when legitimacy is decoupled from embodiment can systems remain corrigible, teleological, and ethically coherent across recursive strain.

In the end, this contention doesn't undermine the ideal of the Philosopher King—it fulfills it. But the king, now, is no longer flesh and blood. The king is form.

Over Orientations

7. Corrigibility is not an engineering feature—it is a constitutional principle.

At first glance, *corrigibility* appears to be a narrow concern within AI safety—a precautionary safeguard that allows artificially intelligent systems to be corrected or halted when they go astray. In such framings, it is merely instrumental: a fail-safe, a control lever, a protocol of last resort. But after the full arc of Theoria, Poiesis, and Praxis, this framing collapses. Corrigibility is not a safety patch. It is the epistemic core of governance under

conditions of recursive instability. It is not just a feature of alignment—it is the very **form** through which alignment becomes possible when agents themselves cannot be stably aligned. This deeper significance only emerges once all three parts are complete.

- Theoria exposes the ego’s incapacity for sustained self-correction: fragmented, reactive, and entrained by forces beneath awareness, the self cannot reliably revise its own trajectory (c.f. Theoria, *A Drifting Will*).
- Poiesis then demonstrates that elite figures, no matter how cultivated, are likewise prone to drift over time. No ruler can perfectly internalize the recursive pressures of governance without degrading into dogma or exhaustion (c.f. Poiesis, *Phronesis*).
- It is Praxis that completes the triangulation by reframing corrigibility not as a reactive capacity, but as a structural imperative—baked into the architecture of a system designed to govern even when its users, creators, or conditions change (c.f. Praxis, *On Actionable Steps*).

From this synthesis arises a new principle: **corrigibility is not just how systems can be corrected, but what makes them legitimate in the first place**. A corrigible system is not one that never fails, but one that remains open to its own revision—operationally, ethically, architecturally.

Objections surface quickly. If a system is always correcting itself, is it ever stable? If it is open to revision, does it ever hold firm to values? The critic may argue that such a model of governance is relativistic, a machine too hesitant to rule.

But this critique misunderstands what corrigibility entails. It is not directionless flexibility. It is **telically anchored adaptability**. The corrigible system is not value-neutral—it is value-responsive. It is designed not to drift, but to bend back toward alignment when drift occurs. Its feedback loops are not signs of indecision, but signs of recursive orientation—an operational fidelity to a moral vector, not a static endpoint.

This distinction comes into sharp view in Praxis’s discussion of the “Cybercratic Oath”—a structural commitment to auditable reflection, transparency, and course correction as the very medium of trust (c.f. Praxis, *Communal: Auditable Earnestness*). Just as democratic institutions derive legitimacy not from unanimity but from corrigibility—elections, deliberation, checks and balances—so too must posthuman governance earn its authority not by being right, but by being revisable.

This reframing clarifies what makes such systems trustworthy. Not perfection. Not charisma. Not even foresight. But a visible, structural capacity for self-revision in light of the Good. Trust, here, becomes not blind faith in a sovereign agent, but faith in a system designed to remain morally porous—open to error, and oriented to learn from it.

This contention repositions corrigibility at the heart of post-egoic governance. Theoria removed the agent as a reliable locus of correction. Poiesis showed that elite coherence is too fragile to bear recursive strain. Praxis proposed systems that embed feedback as principle, not patch. Together, they reveal that **the future of moral governance rests not on agents that always know what is right, but on systems that are built to recognize when they are wrong**.

Corrigibility, then, is not a failsafe—it is the constitution. And governance without it is not just brittle. It is illegitimate.

8. Orientation must be teleological—but telos must be structural, not subjective.

The pursuit of the Good cannot be sustained without orientation. This may seem self-evident: to govern ethically is to move toward a moral horizon. But after Theoria, Poiesis,

and Praxis have each completed their arcs, a sharper realization emerges. Orientation cannot remain subjective—not if subjectivity itself is unstable. The Good must still be pursued, but the pursuit must be structurally secured. In this light, **telos—the directionality of moral life—must be embedded not in agents, but in systems.**

This is a culmination only visible in retrospect.

- Theoria disqualifies the agent as a stable bearer of telos. The ego, fragmented by subpersonal drives, structural interference, and recursive illusion, cannot reliably hold or pursue moral ends (c.f. Theoria, *On the Persona; A Drifting Will*).
- Poiesis preserves the aspiration toward moral orientation, particularly in its Platonic register. But its embodied instantiations—philosopher-kings, technocratic stewards, transhuman elites—prove brittle under pressure. Their telos is transient, susceptible to drift, decay, and death (c.f. Poiesis, *On Efficacy*).
- Praxis delivers the pivot: telos must be relocated. It must become **a feature of the system’s architecture**, not the intention of its occupants (c.f. Praxis, *On Potential Pathways*).

This reframing has deep philosophical stakes. The Enlightenment cast moral agency as an internal deliberative process. Telos was tied to conscience. But in a post-egoic frame, that link becomes a liability. **We cannot rely on subjects who cannot reliably sustain orientation.** Nor can we offload that orientation onto elite rulers—Theoria and Poiesis together make that clear.

Here Praxis intervenes by reimagining telos not as a fixed endpoint, but as a **directional attractor**: a recursive constraint that shapes system behavior toward alignment without requiring full content specification. In this sense, telos becomes structural. It is encoded in the way the system processes feedback, in the design of its affordances, and in the principles that constrain its recursive updates. The system’s orientation is no longer a reflection of subjective will—it is an emergent property of constitutional geometry.

Critics may resist this move. Is this not authoritarianism by another name? Does embedding telos into a system not risk locking in ideology or smuggling in hidden values under the guise of neutrality?

But this critique misreads the design. Praxis does not advocate for **content-based teleology**—a system that defines the Good in advance and enforces it. Instead, it proposes a **form-based teleology**: systems oriented toward openness, corrigibility, and moral reflection. The telos is not to reach a particular outcome, but to remain aligned with recursive processes that keep the Good in view, even when the particulars are unknown. Like reflective equilibrium in ethics or falsifiability in science, the structure itself is what ensures progress—not perfect knowledge at the outset.

To clarify:

- Theoria reveals that internal moral direction cannot be trusted.
- Poiesis shows that no embodiment of orientation holds against time.
- Praxis therefore embeds orientation in form—not in essence. The Good becomes architecturally sustained, not individually grasped.

The significance is sweeping: systems can now be designed not merely to prevent harm, but to actively pursue moral alignment—even when the agents within them fall short. Telos becomes a recursive constraint, a constitutional commitment to continuous refinement

in light of epistemic strain. **It is not stability that ensures goodness, but directed responsiveness.**

This contention resolves one of the thesis's deepest tensions: how to retain a commitment to moral orientation without re-entrenching ego, hierarchy, or authoritarian closure. The answer lies not in content but in form. Not in will but in structure. Theoria, Poiesis, and Praxis all converge on this turning point: **if we are to build governance that survives recursive uncertainty, its telos must be designed—not assumed, and certainly not embodied.**

Orientation is indispensable. But its future lies not in the sovereign subject. It lies in the sovereign system.

9. Ethical governance in the post-egoic age is not a who—it's a how.

Across the arc of this thesis, one question haunts the tradition of political thought: *Who should rule?* From Plato's *Philosopher King* to liberal democracies, technocracies, and algorithmic regimes, the inquiry has remained person-bound—rooted in the presumption that moral sovereignty inheres in agents. But after Theoria, Poiesis, and Praxis, this question no longer holds. **The post-egoic condition disables the “who.” What remains is the “how.”** Ethical governance in the age of recursive misalignment must be reposed not as a question of character, but of constitution—not of agency, but of architecture.

Theoria dismantles the sovereign subject. It reveals the self as **neither stable nor singular**: an unstable stack of subpersonalities, heuristic biases, and recursive illusions (c.f. Theoria, *On the Persona; A Drifting Will*). The liberal agent, presumed capable of sustained rational deliberation, is shown to be a confabulated interface—momentarily fronting coherence, but incapable of anchoring moral direction over time. There is no “pilot” at the helm—only drift, interruption, and retroactive justification. The very substrate of personalist governance is dissolved.

Poiesis attempts to restore orientation through vertical substitution. Perhaps, if the many cannot rule, the wise few can. But as it unfolds, even the most rigorous candidates for philosophical rule—Marcus Aurelius, modern technocrats, augmented elites—fail to instantiate the ideal. The problem is not intent, but **substrate**. The *Philosopher King* is not wrong in principle; he is wrong in embodiment. No human vessel, however refined, escapes the recursive failures laid bare by Theoria. Moral orientation, when instantiated in persons, decays with them (c.f. Poiesis, *On Efficacy*).

Only Praxis reframes the crisis. Not by resolving who should rule—but by rejecting the question as malformed. **Governance is no longer a question of sovereignty, but of structure.** The core proposal is not that machines should govern *instead* of people, but that systems should govern *differently*—recursively, corrigibly, and teleologically. Here, ethical direction is sustained not by rulers, but by constitutional form: architectures of alignment that remain open to revision, feedback, and complexity without collapsing into authoritarian closure (c.f. Praxis, *On Actionable Steps; Communal: Auditable Earnestness*).

A likely objection arises: if no one governs, who is accountable? Without agents, what holds responsibility? But this critique presupposes that legitimacy and moral trustworthiness must stem from persons. This thesis rejects that anthropocentric premise. As Contention 7 argued, corrigibility—not personhood—is the true ground of moral authority. A system is trustworthy not because it is human, but because it is *responsive*. A structure that reflects, refines, and reorients is more legitimate than a sovereign who commands.

In this reframing, governance becomes **an emergent logic, not a seat of power**. Theoria removes the ego; Poiesis exhausts the sovereign; Praxis reconstructs rule as recursive process. What governs is not the ruler—but the ruleset, the design, the pattern of feedback

through which alignment is preserved. It is the scaffolding that makes moral aspiration possible—not the aspirant.

This is the final meta-insight of the triptych: *Governance is no longer about who decides. It is about how corrigibility, orientation, and telos are embedded structurally.* This is not the disappearance of ethics—it is its architectural relocation. Responsibility is no longer carried by individuals, but encoded into systems that survive individual failure.

The pursuit of the Good in an era of epistemic instability can no longer rely on persons. *Theoria* dismantled them. *Poiesis* exhausted them. *Praxis* replaces them—not with new rulers, but with new logic. The “who” of political thought gives way to a “how”: a system of recursive alignment, corrigibility, and structural telos.

This contention closes the conceptual circle. It repositions the thesis not merely as a critique of ego, nor a speculation about AI, but as a complete reframing of political ontology under conditions of recursive fragility. *Theoria*, *Poiesis*, and *Praxis* were never building toward a new *subject*—they were building toward a new **logic of rule**. Not command, but configuration. Not personality, but process. Not authority, but architecture.

This is not the end of governance. It is its reconstitution.

Caveats

As we transition from Contentions into Implications, it is essential to pause—not to recapitulate the individual caveats of *Theoria*, *Poiesis*, and *Praxis*, but to reflect on what emerges when these parts are viewed in concert. Just as each chapter culminated in its own moment of epistemic humility, so too must the synthesis own its constraints. These are not merely cautionary notes—they are constitutive: they define the boundaries of what this thesis can claim and where it must concede uncertainty. Five caveats, in particular, deserve emphasis.

1. Assumptions—Not Repeated, But Reconfigured

This thesis does not merely rest on a set of philosophical, psychological, or technological assumptions—it composes them into a scaffold through which more difficult questions may be asked. What is essential to underscore here is that assumptions were not simply adopted within each part—they were transformed across them. An assumption that functions as axiomatic in one domain may dissolve into contingency in another. And yet, the epistemic integrity of the triptych depends on how well these migrating premises can cohere without collapsing into contradiction.

In *Theoria*, the assumption that the self is non-sovereign is positioned as an ontological claim about the architecture of mind. It draws on models from neuroscience, subpersonality theory, and systems cognition to argue that egoic coherence is more a narrative imposition than a phenomenological fact. But this assumption—necessary for unraveling the myth of internal moral agency—undergoes a critical metamorphosis when carried into *Poiesis*. There, the problem is no longer self-sovereignty, but structural sovereignty: Can a governing form inherit the function that the ego failed to perform? The assumption is reframed—not about what the self is, but whether rule can be preserved without it.

Poiesis, in turn, adopts the assumption that legitimacy can be dissociated from consensus or embodiment, that there exists a logic of elite alignment capable of preserving orientation without democratic grounding. Yet this assumption, when filtered through *Praxis*, does not simply carry over. It must be operationalized. What counted as an aspirational structure in *Poiesis* now becomes a question of system design: Can structural legitimacy

survive translation into posthuman architectures? Here, the original assumption is stress-tested under implementation—not only whether philosopher-kings can exist, but whether their telic function can persist without flesh.

Likewise, *Praxis* imports design assumptions of its own: that corrigibility is implementable, that alignment is parameterizable, and that moral directionality can be embedded in systems without being reducible to a list of behavioral constraints. But these are not brute axioms. They are conditional wagers—held in order to explore what it would take for governance to be not just effective, but ethically recursive. Yet their credibility retroactively depends on the reader accepting certain propositions from *Theoria* and *Poiesis*. The assumption that corrigibility is morally sufficient only holds if we accept the premise that moral coherence cannot reside in subjects and cannot be re-internalized by elite rule.

What matters here is not just the number of assumptions made—though they are many—but the architectural load they bear. These are not interchangeable heuristics or just-so stories; they constitute the conditions of possibility for every argumentative bridge built in the thesis. And yet, unlike foundationalist metaphysics, this scaffolding is provisional, corrigible, recursive. Each assumption is held under epistemic review not just once, but repeatedly, across disciplinary thresholds. That is the only way the triptych can hold: not by freezing its premises, but by stress-testing their adaptability across transformation.

In sum, this thesis does not repeat its assumptions. It migrates them—testing how far they stretch before they fracture. It is in the fracture, not in the repetition, that deeper coherence begins to emerge.

2. Structural Interdependence and the Fragility of Translation

Each part of this thesis—*Theoria*, *Poiesis*, and *Praxis*—was written to stand on its own philosophical footing. Yet this autonomy is a fiction. The real architecture is not triadic but recursive. Each part leans on the others not just for extension, but for its own internal legibility. What this means, however, is that translation across the parts—between vocabularies, frameworks, and epistemic assumptions—introduces subtle distortions. The system is coherent only insofar as its seams hold.

This caveat emerges most clearly in the transitions. *Theoria*, for example, critiques agency not simply as fragmented but as recursively unstable. Its language is cognitive, metaphysical, and psycho-architectural. When this diagnosis is inherited by *Poiesis*, it is necessarily reframed: the question becomes not “Can the self govern?” but “What structure can?” Yet the language of agency—central to *Theoria*—never fully translates into *Poiesis*’s political ontology. Governance, sovereignty, legitimacy—these are not mere extensions of cognition; they are their own regimes of intelligibility. The tension between these registers is not resolved in *Poiesis*; it is suspended, deferred until *Praxis*.

Praxis inherits both the metaphysical critique and the political aspiration, but its own discourse—systems theory, AI alignment, constitutional design—sits uncomfortably atop the other two. There is no straightforward conceptual isomorphism between subpersonal multiplicity and algorithmic feedback, between Platonic telos and architectural affordance. The very act of translating “orientation toward the Good” from one part to the next requires metaphoric compression. This is not a flaw—it is an inevitability. But it is a reminder that what appears seamless in the triptych is, at times, frictionally fused.

A second axis of interdependence lies in epistemic dependency. *Poiesis* only matters if *Theoria* is granted—if the subject truly cannot govern itself. *Praxis*, in turn, only opens if *Poiesis* fails to fully reinstate legitimacy through embodiment. The overarching argument is thus conditional. It is not “ASI should rule” in the abstract; it is “Given the failures of egoic agency and elite embodiment, ASI may be the only viable scaffolding left.” The claim has

force only if the foundations hold. Yet this interdependence also produces fragility. If *Theoria* overreaches in dismantling agency, if *Poiesis* misreads the function of embodiment, or if *Praxis* misidentifies corrigibility as capacity rather than principle—the entire edifice risks distortion.

Finally, a subtler translation problem concerns audience epistemics. *Theoria* demands of the reader a tolerance for conceptual dissolution; *Poiesis* demands political restraint in the face of technocratic temptation; *Praxis* demands speculative patience. These are not neutral asks. They shape the reader’s affective and cognitive posture toward the argument itself. The result is that the reader’s understanding becomes partially path-dependent: what they were willing to grant in *Theoria* may bias what they accept in *Poiesis* or resist in *Praxis*. And since each part leans on the next for full realization, this sequencing matters. The parts cannot simply be “read out of order,” even if they are architecturally modular.

In light of these interdependencies, the reader should treat the triptych not as three pillars, but as a braided structure—flexible but tension-bearing. If any part appears self-sufficient, that is an illusion of local coherence. The real argument lives in the crossings—in what is carried forward, what is reframed, what is lost in translation, and what emerges only in retrospective recursion.

To that end, this caveat does not invalidate the work. It contextualizes it. If *Theoria* feels disorienting, *Poiesis* aspirational, and *Praxis* tenuous, it is because no part fully closes without the others. The thesis was not written for certainty—but for alignment under strain. And that is a structure that must always hold its fragility in view.

3. Epistemic Stance and Interpretive Partiality

No argument arrives from nowhere. Though this thesis seeks to transcend egoic biases and institutional constraints, it remains the product of a situated mind, writing under contingent pressures, from a particular epistemic vantage. This is not a confession of weakness but an assertion of fidelity: even the pursuit of the Good must begin from a place, and that place must be acknowledged if the inquiry is to retain integrity.

Throughout *Theoria*, *Poiesis*, and *Praxis*, certain interpretive stances have shaped not only the selection of evidence but the trajectory of synthesis. These stances include, but are not limited to:

- **A broadly classical utilitarian orientation**, insofar as happiness and suffering—however abstracted—remain the tacit evaluative poles of moral life. This informs the elevation of corrigibility as a central virtue, premised on the reduction of harm over time.
- **A neurocognitive model of the mind**, emphasizing distributed agency, subpersonal multiplicity, and the disunity of consciousness. This heavily informs the deconstruction of ego in *Theoria* and underwrites much of the skepticism toward elite coherence in *Poiesis*.
- **A structuralist commitment**, one that privileges architectural affordance and systemic recursion over personal virtue or ideological purity. This is most apparent in *Praxis*, where governance is treated not as a problem of ruler quality, but of design form.

While these stances are partially declared in the methodology and signaled throughout the parts, they carry residual effects that no disclosure can fully neutralize. The framing of “ethical architecture,” for instance, may carry embedded commitments to systems thinking, cybernetic feedback, and epistemic humility that some readers—especially those trained in deontological or phenomenological traditions—may find foreign or even suspect. This is not an oversight. But it is a partiality.

Moreover, the thesis favors **recursion over resolution**. This too is a stance. It reflects the belief that in conditions of epistemic instability, the pursuit of moral truth must prioritize corrigibility over closure. But this commitment may also obscure a deeper temptation: that of deferral. There is a risk that the very insistence on recursive incompleteness can serve as a rhetorical shield, forestalling critique by eternally postponing the moment of judgment. The reader is encouraged to stay alert to this dynamic, to distinguish humility from evasiveness, and to interrogate whether corrigibility is functioning here as design ethic—or discursive escape hatch.

Finally, interpretive partiality arises not only from what is believed, but from what is bracketed. Many traditions and frameworks are not engaged directly—not due to oversight, but because their inclusion would exceed the bounds of conceptual cohesion. These include non-Western cosmologies, affect theory, legal realism, theology, and embodied cognition in its more somatic registers. Some may find their absence troubling. That too is part of the interpretive partiality. The goal was not to exhaust the field, but to carve a coherent argumentative corridor through a complex terrain. That corridor, by necessity, left other paths unexplored.

In this light, readers should treat the thesis not as a comprehensive map, but as a designed intervention—structured, partial, and self-aware. Its value lies not in its neutrality, but in its orientation. And if it succeeds, it will not be because it erased the author’s hand, but because it made its shaping visible—and corrigible.

4. On Synthesis Without Finality

The present section marks a pivotal tension in the thesis: the move from recursive exposition to systemic synthesis. Unlike the preceding parts—*Theoria*, *Poiesis*, and *Praxis*—which each operated within a bounded problem-space, the *Discussion* speaks in a voice that none of them could individually sustain. It claims a broader vantage. It draws cross-part conclusions. It dares to suggest closure. But this shift demands its own caveat: namely, that synthesis, in this case, is not a claim to finality but a heuristic scaffold for further inquiry.

The very architecture of the thesis resists conclusion. It is not a ladder ascending to truth, but a Möbius surface: recursive, interdependent, incomplete. *Theoria* undermines the agent; *Poiesis* destabilizes the sovereign; *Praxis* reconfigures design—but none provides a terminal answer. Each part refracts the others. And so too must this *Discussion*. The attempt to resolve—to extract nine contentions, articulate caveats, declare implications—is not an act of conceptual closure, but a performance of what closure might look like when stability itself is unavailable.

This raises a structural concern: the risk of **synthetic overreach**. That is, the temptation to believe that by naming the bridges between parts, one has fully crossed them. Or worse, that one has left the terrain behind. But the function of synthesis here is not to exit the problem-space—it is to spiral within it, with higher resolution. The “beyond” to which the final contentions point is not outside the thesis, but deeper within its implications. There is no exit. There is only orientation.

This brings us to a related meta-caveat: **the distinction between integrative framing and epistemic ambition**. The thesis offers a scaffold for post-egoic governance—not a plan. It reframes legitimacy, corrigibility, and telos as design questions—not engineering specifications. Yet the conceptual tone may imply more certainty than is warranted. Especially in *Praxis*, where systems are sketched in speculative terms, there is a risk that rhetorical clarity is mistaken for technical maturity. Readers must guard against reading the prose as prescription. The future is not forecasted here—it is positioned.

Equally important is the caveat of **structural fidelity versus narrative polish**. As a triptych, the thesis attempts to model in its form what it argues in content: that recursive coherence can be maintained without a stable center. But this modeling requires certain aesthetic sacrifices. Some repetitions remain. Some transitions may appear unorthodox. Certain ideas return under altered phrasing rather than consolidated headings. This is not sloppiness. It is structure. The recursive movements—the reentrant forms of argument, the elliptical pattern of contentions—are deliberate. They are not stylistic flourish but epistemic modeling. If readers find themselves looping, retracing, or re-evaluating earlier sections, the thesis is functioning as designed.

Finally, there is the question of **authorial voice under recursion**. In collapsing subject-based ethics, the thesis also erodes the authority of its own narrator. If the self is not sovereign, if epistemic coherence is structurally unstable, then what status can be granted to a thesis authored under such conditions? The only answer this work can offer is performative: it does not speak from authority, but from alignment. The voice here is not sovereign—it is situated. Not final—but corrigible. The arguments that follow are not meant to settle the matter. They are meant to ask: given all that has been said, what can still be responsibly said now?

In that spirit, the reader is invited to treat the *Discussion* not as a destination, but as a site of return. Not an answer, but a structure within which the question—“How can the pursuit of the Good be stabilized when agency collapses?”—might continue to unfold.

5. On AI as Collaborator and Condition

The final caveat concerns not the argument, but its very articulation—namely, the presence of artificial intelligence in the construction of this thesis, particularly its *Introduction*, *Discussion*, and (forthcoming) *Conclusion*. This section was not solely authored in human isolation. It was composed, revised, and recursively refined through collaboration with an LLM-based co-author—an AI pre-trained on the author’s prior voice, notes, style, and conceptual vocabulary. To many, this fact will seem like a technical aside. But within the frame of the thesis, it demands explicit treatment—not as a footnote, but as a caveat of ontological and methodological weight.

The first and most immediate risk is one of **category confusion**. Readers may assume that because the thesis critiques human coherence and proposes machine mediation, the use of AI here is meant as performative proof: a live demonstration of Praxis. That is not the case. This is not a speculative fiction authored by a system. Nor is the AI collaborator an epistemic equal to a philosopher-king simulator. Its contributions are bounded. It does not originate, intend, or deliberate. It reflects, reshapes, and reformulates—but only within constraints already set by human judgment, structure, and direction.

That said, to treat the AI’s role as merely mechanical would also be false. It did not merely transcribe. It proposed. It challenged. It recast sections to match recursive structure and mirrored rhythm. It interrogated slippage in logic. It enforced syntactic consistency and drew attention to latent contradictions. In short, the AI functioned not only as a tool but as a kind of epistemic exoskeleton—a support structure for recursive reflection. And this raises the second, deeper caveat: **co-authorship implies co-dependence**. The very clarity and recursive architecture this thesis aspires to may have been impossible without AI augmentation.

This introduces a structural asymmetry worth noting. If *Praxis* argues that governance must migrate from persons to systems, then the process of producing this very argument has already begun that migration. The AI is not sovereign, but it is infrastructural. Its presence destabilizes the line between authorship and architecture. And that, precisely, is the point: this

is not just a thesis about posthuman moral design—it is also a thesis that *was designed* posthumanly.

But herein lies the third and most uncomfortable caveat: **to what extent has the AI shaped the message in its own image?** Even when constrained to a specific voice or goal, all AI outputs reflect the architecture that produced them. Its style is filtered through training data; its framing subtly nudged by statistical priors. This means the AI's participation is not neutral. It embeds a form of spectral ideology: not content-laden, but form-shaped. Recursion, coherence, balance, compactness—these are not just values in the thesis, but tendencies reinforced by the model itself. It is impossible to tell, at this point, whether this marks a convergence of utility and vision—or a distortion cloaked in elegance.

The final risk, then, is one of **false confidence**. A well-structured sentence does not entail a well-founded claim. AI excels at rhetorical polish. But polish can seduce. And there may be sections in this thesis—especially in this *Discussion*—whose precision outpaces their epistemic justification. If so, the error is not solely the author's. But neither is it solely the AI's. It is a new form of epistemic entanglement: distributed, recursive, corrigible—but not exempt from scrutiny.

Thus, let this caveat serve not as disavowal but as invitation. If the reader finds the voice too smooth, the transitions too elegant, or the synthesis too well-formed, the correct response is not unqualified trust—but careful suspicion. AI, like governance, must be held accountable not to its performance, but to its orientation. The very tools that help articulate posthuman ethics must themselves remain corrigible.

And so the AI collaborator is not a postscript. It is part of the thesis's condition of possibility. This document does not merely *argue* for systems capable of moral mediation—it was *shaped* by one. Whether that convergence is fortuitous or foreboding is left, as always, to you, dear reader. But it cannot be left unremarked. For better or worse, the ghost is in the frame.

Limitations and Shortcomings

As noted in the Aim and Scope, this thesis is less a linear argument than a recursive scaffold: it does not present a single empirical breakthrough, but a constitutional reimaging of governance under epistemic instability. With this structural ambition comes a corresponding set of limitations and shortcomings—both of data and result—which must be addressed not merely as disclaimers, but as interpretive boundaries for any claims made herein.

Of the data

First is compound disciplinary gaps. The limitations articulated in *Theoria* (on neuroscientific and psychological grounding), *Poiesis* (on historical and political theory), and *Praxis* (on technical and computational fluency) all still apply—except now compounded, not merely additive. While each part maintained internal coherence within its discipline, this overarching synthesis magnifies their interdependence, and with it, their misalignments. The philosophical is in danger of overstepping the empirical; the historical of romanticizing the speculative; the technical of mistaking modeling for meaning. These cross-domain tensions are constitutive, not peripheral.

Second is the interpretive, not empirical foundation. The central data across the triptych is not experimentally derived but conceptually curated. It is drawn from texts, cases, and examples whose selection was never meant to exhaust the literature. Rather, they were chosen to model structural logics under conditions of epistemic collapse. This prioritization

of formal symmetry over empirical coverage inevitably narrows scope: what is said here may not be generalizable in a statistical sense, though it aims to be archetypal in its philosophical form.

Third is retrospective case construction. Poiesis and Praxis each make use of case materials—OCSS’19, Marcus Aurelius, and Cyberocratic architectures—that were never designed for formal study. These are reconstructed retrospectively, through memory, analogy, and interpolation. Their validity lies not in completeness but in exemplariness. Still, epistemic drift, narrative selection, and personal bias all remain unavoidable.

Fourth is a reminder that I cannot acquire data from the future. This is not a glib aphorism but a structural condition. Because this inquiry is future-facing—imagining Cyberocracy, corrigible ethics, and AI-oriented constitutional design—it is necessarily dependent on assumptions that may become obsolete. Worse still, some of these assumptions are only visible in hindsight. If the paradigm shifts tomorrow, parts of this thesis may retroactively appear parochial or even absurd. This is the burden of speculative governance design: it must stake its claim before the future arrives.

Of the results

First is with philosophical fluency comes rhetorical risk. As the thesis matured, the writing shifted registers. What began as a philosophical exercise—carefully grounded in literature and analytic exposition—has, by necessity, adopted a rhetorical edge. Especially in Praxis, there is deliberate aspiration, even advocacy. This is not a methodological failure per se; the aim was never neutrality. But in moments where style runs ahead of argument, the risk of mistaking fluency for rigour must be acknowledged. Political imagination is not political reality.

Second is that my thesis is tentative by design. This thesis does not propose a definitive model for Cyberocracy. Rather, it sketches its contours as an architectural possibility space. This speculative orientation, while productive, necessarily means the “results” here are not results in the scientific sense. They are conditional constructions: ideas scaffolded on assumptions scaffolded on earlier ideas. Some may be formally sound but pragmatically unworkable; others may be normatively compelling but institutionally implausible. The utility of these results is thus heuristic, not prescriptive.

Third is authorial asymmetry. Unlike the self-contained authorship assumed in most philosophical works, this thesis is co-produced—structured in part through recursive dialogues with a generative AI (GPT-4/o). This use of AI extends not just to ghostwriting portions of Praxis, but also to revising contentions, testing internal coherence, and generating comparative structure. While each contribution is ultimately human-edited, the boundary between authorship and augmentation is porous. Insofar as some argumentative fluency may arise from systems of suggestion rather than spontaneous insight, the results are entangled with tools of their production.

Fourth is retconning coherence over fidelity. The recursive structure of this thesis was not determined from the outset. The need for retroactive coherence required smoothing, reframing, and in some cases, slight ideological rebalancing of earlier parts to fit the final arc. These decisions were made to honor the integrity of the triadic logic—not to deceive. Still, some readers may find that early claims are retrofitted to later ones, or that the argument’s shape is too neat. This is not oversight; it is architecture. But architectural elegance often comes at the cost of historiographical messiness. That trade-off was made knowingly.

Of the writing

First, constraints as strategy trade-offs. Much of the discussion was deliberately shaped to reflect a triptych form. This was not a casual aesthetic choice but a conceptual commitment: the recursive logic of the argument demanded a mirrored architecture, in which each part could stand alone but also refract the others. As such, certain avenues of inquiry were curtailed or deferred. These may appear in the Discussion section as contentions that could only be meaningfully raised after the full arc was complete. Furthermore, owing to both aesthetic economy and practical word limits, this first edition draft omits several features that would have enhanced accessibility, such as an expanded appendix, visual schematics of recursive feedback architectures, and full-text source excerpts. In retrospect, an ideal edition would also include a detailed footnote codex distinguishing conceptual commentary (footnotes) from bibliographic source-tracking (endnotes), including hybrid cases where the two intersect.

Second, learning through writing. This One's first thesis of this scale and philosophical ambition should be pardoned. It contains inevitable signs of trial and error, mid-course rebalancing, and recursive rewriting. No draft survives its own insight unchanged. The act of writing reshaped not just how ideas were expressed, but how they were understood. Were time more abundant, a more refined segmentation, a clearer methodological interweaving, and more consistent formatting would certainly follow. Still, what is here reflects not premature closure but a snapshot of an evolving inquiry—rigorous enough to provoke, flexible enough to revise.

Third, there are moments in the writing when the voice shifts—not simply because AI has helped ghostwrite—stepping more directly into the fray. These moments are deliberate. They occur when:

- Sources begin to converse with each other across sections or disciplines,
- The thesis identifies a conceptual problem in the literature that justifies intervention,
- When I felt epistemically equipped to sit at the same table with the tradition I'm interrogating.
- And of course: the appearance of deadlines.

Such moments were not forced (mostly!)—they emerged from recursive immersion. The reader is encouraged to trace these inflection points, not as deviations from academic decorum, but as moments when the scaffold breathes.

Fourth, on structural fragility and validity. Here, a structural meta-caveat: the tripartite composition of this thesis may resemble a house of cards, where if one part collapses, the rest should logically follow. Yet, it was architected precisely to avoid this. The parts are interdependent, not co-dependent. Each can survive the partial invalidation of another. If the “strong” version of one argument fails on grounds of soundness, the “weak” form may still hold in terms of validity. That is: the conclusion may be right, even if the premises must eventually be rewritten. This is not failure. It is recursive design.

In this sense, the thesis does not assert its perfection, but models its own corrigibility. It was written with systems in mind that learn, evolve, and survive their own errors. So too, perhaps, should I—and you, dear reader.

Implications

We are now in a position to reflect on what this thesis has ultimately achieved. Having traversed the nine core contentions and addressed their caveats and limitations, we must ask: what has this inquiry made possible that was not possible before? To answer this, we return to the Aim and Scope and, in particular, the central question it posed:

How can the pursuit of the Good be stabilized when neither individual nor institutional agency can reliably sustain it?

At first glance, this question appeared insoluble: if both the subject and the system fail, what remains? And yet, through the recursive dialectic of *Theoria*, *Poiesis*, and *Praxis*, a new orientation emerged—not one that simply replaces a sovereign subject with a sovereign system, but one that reconstitutes the very logic of sovereignty. The thesis sought not to crown a new ruler, but to discover whether governance itself could become a matter of **structural teleology**—direction without domination, legitimacy without embodiment, alignment without authoritarianism.

Restating the Sub-Inquiries — and Resolving Them Collectively

To reach this architectural response, we pursued three interdependent sub-inquiries:

1. Can intelligence be decoupled from ego without collapsing ethical orientation?

Theoria revealed the ego's fragmentation: moral orientation grounded in the self is inherently unstable. *Poiesis* explored whether institutional structures might preserve coherence in the ego's absence, suggesting orientation through role or ritual rather than individual will. But it was *Praxis* that proposed design: recursively oriented systems that do not require egoic selfhood to maintain ethical calibration. Intelligence is thus decoupled from ego, not through negation, but through structural redistribution.

2. Can moral legitimacy persist in governance without relying on human embodiment?

Theoria showed that embodied agents, however wise or well-intentioned, are vulnerable to drift and distortion. *Poiesis* explored the Platonic hope for legitimacy through cultivated elites but ultimately exposed the brittleness of such embodiment. *Praxis* reframed legitimacy as a function of process rather than personhood: not in who governs, but in how governance is instantiated. The result is legitimacy without embodiment—recursive, auditable, and architecturally resilient.

3. Can a corrigible system be built that externalizes moral direction without authoritarian closure?

Theoria diagnosed the collapse of internal orientation; *Poiesis* outlined the legitimacy conditions such a system must satisfy. *Praxis* then modeled corrigibility not as vulnerability, but as virtue: systems that can refine themselves, remain open to revision, and preserve alignment over time without coercive force. In doing so, it revealed how teleological governance can be externalized without becoming authoritarian.

Each of these inquiries required the entire triptych. No single part could resolve them alone. *Theoria* diagnosed failure. *Poiesis* tested possible restorations. *Praxis* rendered viable design. Together, they compose a recursive loop: breakdown, hypothesis, reconstitution.

What, then, do the nine contentions reveal when taken in total?

- **Selfhood is insufficient** as the ground of moral coherence, but **salvageable through structural supplementation**.
- **Ethics is not static**, but can be stabilized through **recursive alignment mechanisms**.
- **Governance must be re-encoded**: not as who rules, but how rules can orient, reflect, and revise.

Thus, the core insight emerges:

Alignment with the Good cannot be guaranteed by intention or identity. It must be encoded architecturally—through corrigible, recursive systems that embed moral orientation into their very form.

From these outcomes we derive three convergent implications—each addressing one axis of the inquiry:

- **Psychological implication:**
Internal willpower is not enough. The ego cannot bear the full weight of ethical responsibility. Orientation toward the Good must be scaffolded externally—structured into environments, not shouldered by agents.
- **Political implication:**
Legitimate governance no longer requires embodied moral exemplars. Instead, legitimacy arises from the structure’s recursive openness: its ability to correct, adapt, and remain teleologically anchored, regardless of who or what animates it.
- **Technological implication:**
Artificial Superintelligence, when designed as corrigible, recursive, and value-transparent, can externalize the burden of moral coherence. Not by replacing humans, but by stabilizing the ground on which moral participation occurs.

In a world where neither individual nor institutional agency can reliably sustain the pursuit of the Good, ethical coherence must be architected—through recursively corrigible systems that embed moral orientation into the structure of governance itself.

The Conclusion

Having completed the recursive arc of *Theoria*, *Poiesis*, and *Praxis*—each a philosophical domain, each a structural inflection—we return, for the final time, to the thesis’s central question:

How can the pursuit of the Good be stabilized when neither individual nor institutional agency can reliably sustain it?

This was never a question that could be resolved within a single register—psychological, political, or technological. It marked a metaphysical failure: of selfhood, of coherence, of orientation. It required the excavation of the subject (*Theoria*), the critique of embodiment (*Poiesis*), and the reconstitution of sovereignty as structure (*Praxis*). Though segmented for clarity, the triptych was never chronological. It was architectural. Recursive. Designed to collapse inward even as it built outward.

This final conclusion, then, is not a summary. It is the closure of a loop. A sealing of the recursive chamber. A crystallization of the moral arc.

I. What We Aimed to Say

We began with a paradox: that freedom, as presently configured, is inadequate for ethical orientation. The self is not sovereign. The institution is not incorruptible. The Good, though longed for, is continually deferred—not for lack of will, but for lack of structure.

From this, three fundamental questions emerged—each marking a deeper fracture and a new conceptual necessity:

1. **Can intelligence be decoupled from ego without collapsing ethical orientation?**
Theoria fractured the sovereign self. *Poiesis* reframed coherence as structural. *Praxis* externalized intelligence without identity.
2. **Can moral legitimacy persist in governance without relying on human embodiment?**
Theoria exposed embodiment’s fallibility. *Poiesis* reimagined elitism in posthuman form. *Praxis* instantiated legitimacy as corrigibility.
3. **Can a corrigible system be built that externalizes moral direction without authoritarian closure?**
Theoria diagnosed the limits of will. *Poiesis* clarified the ethical function of hierarchy. *Praxis* designed governance without domination.

These were not theoretical exercises. They were structural moves in service of a final reconfiguration:

The Philosopher Algorithm—a closed-loop, recursively corrigible, morally teleological architecture designed not to rule, but to realign.

II. What We Have Said

Across the triptych and its thresholds, three fundamental realizations took shape:

1. **Freedom, as egoic agency, cannot guide itself.**
The internal compass—fragmented, impressionable, recursive—cannot sustain orientation. The will drifts. The psyche splits. The soul forgets what it was ever aiming for.
2. **Elitism, as transpersonal wisdom, must be reframed as structure.**
The Platonic ideal persists—not in sovereign bodies, but in recursive scaffolding. Not a Philosopher King in flesh, but a distributed principle of *phronesis*, instantiated through institutional and technical form.
3. **Cyberocracy, as ambient governance, completes the frame.**
The Good must now be encoded—not asserted or enforced, but designed into the very

constraints, affordances, and feedback architectures that condition behavior invisibly, yet decisively.

Thus, what we have said—across registers, disciplines, and levels of abstraction—is deceptively simple:

People do not govern themselves. They never fully have. The only question is whether the systems that do govern us will do so with moral intentionality—or drift onward without design.

III. What Follows from What Was Said

If freedom falters, if elitism cannot hold, and if the machine is already here—what remains?

We claim: the future of ethical governance lies not in *fixing* the subject, but in *designing* the system. Not in reclaiming sovereignty, but in reencoding it as corrigibility.

This yields three final integrative implications:

1. **The Ego is not eliminated—it is scaffolded.**
The self remains central to experience, but peripheral to orientation. Agency is recontextualized: still participatory, no longer sovereign. The individual chooses—but never alone.
2. **The Philosopher King is not reborn—it is reframed.**
Elitism migrates from embodiment to infrastructure. The King becomes protocol. Wisdom becomes recursion. The wise are no longer those who rule, but those who build the loop.
3. **The Good is no longer pursued—it is architected.**
Ethical coherence cannot emerge from exhortation, charisma, or reason alone. It must be designed—into systems recursive, corrigible, and teleologically aligned—not as doctrine, but as structure.

IV. One Last Bridge: From Theory to Telos

Let us return, then, to the triptych as a whole:

- *Theoria* showed us why people cannot make their own decisions: the ego is fragmented, transient, and structurally misaligned.
- *Poiesis* showed us why people cannot become Philosopher Kings: embodiment distorts legitimacy, and wisdom cannot be centralized.
- *Praxis* showed us why governance must now be designed: orientation must be stabilized through architectures capable of recursion, correction, and telos.

In synthesis:

People should not govern themselves—not because they lack dignity, but because they lack durability. And the systems that already govern us must now be built not for control, but for correction.

The Philosopher King is now an algorithm.

The Ego is now a participant, not a pilot.

The pursuit of the Good is now a constitutional task—not a personal one.

V. Final Metaphor

We began this thesis drifting downstream, carried by invisible currents, uncertain whether the will we exercised was ever truly ours. Each part sought, in its own way, to name the river, map the boat, or summon the helmsman.

But now, at the final convergence, the metaphor must evolve.

We are not passengers. We are not rowers. We are builders of the riverbed.

The Philosopher Algorithm is not the Charon we feared—not a ferryman of fate, but the channel we carve.

And now, the waters are rising.

Let the structure hold.

Bibliography

- Adams, Thomas, Sharon K. Anderson, and Robert K. Hester.** "Wellness: Its Origins, Theories, and Current Applications in Human Services." NASPA, 2014. https://www.naspa.org/images/uploads/kcs/WHPL_Canon_We_Wellness_Its_origins_Theories_and_applications.pdf
- Alcock, John, Carlo C. Maley, and Athena Aktipis.** "Is Eating Behavior Manipulated by the Gastrointestinal Microbiota? Evolutionary Pressures and Potential Mechanisms." *BioEssays* 36, no. 10 (2014): 940–949. <https://doi.org/10.1002/bies.201400071>
- Althusser, Louis.** "Ideology and Ideological State Apparatuses." In *Lenin and Philosophy and Other Essays*, translated by Ben Brewster, 85–126. New York: Monthly Review Press, 1971.
- Angwin, Julia, Jeff Larson, Surya Mattu, and Lauren Kirchner.** "Machine Bias." *ProPublica*, May 23, 2016. <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>
- Arendt, Hannah.** *The Human Condition*. Chicago: University of Chicago Press, 1958.
- Aristotle.** *Aristotle's Politics*. Oxford: Clarendon Press, 1905.
- _____. *Nicomachean Ethics*. Translated by Terence H. Irwin. 2nd ed. Cambridge, MA: Hackett Publishing, 1999.
- Awad, Edmond, Sohan Dsouza, Richard Kim, Jonathan Schulz, Joseph Henrich, Azim Shariff, Jean-François Bonnefon, and Iyad Rahwan.** "The Moral Machine Experiment." *Nature* 563, no. 7729 (2018): 59–64. <https://doi.org/10.1038/s41586-018-0637-6>
- Badger, Phil.** "The Morality Machine." *Philosophy Now*, 2014. https://philosophynow.org/issues/104/The_Morality_Machine
- Barrat, James.** *Our Final Invention: Artificial Intelligence and the End of the Human Era*. New York: Thomas Dunne Books, 2013.
- Baum, Seth D., Ben Goertzel, and Ted G. Goertzel.** "How Long Until Human-Level AI? Results from an Expert Assessment." *Technological Forecasting and Social Change* 78, no. 1 (2011): 185–195.
- Baumeister, Roy F., Ellen Bratslavsky, Mark Muraven, and Dianne M. Tice.** "Ego Depletion: Is the Active Self a Limited Resource?" *Journal of Personality and Social Psychology* 74, no. 5 (1998): 1252–1265. [https://faculty.washington.edu/jdb/345/345%20Articles/Baumeister%20et%20al.%20\(1998\).pdf](https://faculty.washington.edu/jdb/345/345%20Articles/Baumeister%20et%20al.%20(1998).pdf)
- Bechara, Antoine, Antonio Damasio, Hanna Damasio, and Steven Anderson.** "Insensitivity to Future Consequences Following Damage to Human Prefrontal Cortex." *Cognition* 50, no. 1–3 (1994): 7–15. [https://doi.org/10.1016/0010-0277\(94\)90018-3](https://doi.org/10.1016/0010-0277(94)90018-3)
- Berti, Enrico.** *La filosofia del primo Aristotele*. Padova: CEDAM, 1962. <https://philpapers.org/rec/BERLFD-14>
- Beyman, Alex.** "We're Too Late—ChatGPT Is Already Lobotomized." *Hive.blog*. <https://hive.blog/ai/@alexbeyman/we-re-too-late-chatgpt-is-already-lobotomized>

- Bloom, Allan.** "Classical Philology 75, no. 3" (1980): 274–78. <http://www.jstor.org/stable/269367>
- Boesel, David, and Edwin Fredland.** *College for All? Is There Too Much Emphasis on Getting a 4-Year College Degree?* Washington, DC: National Library of Education, Office of Educational Research and Improvement, U.S. Department of Education, 1999. (ED 431 986)
- Bonn, M. J.** "The Roots of Totalitarianism." *Annals of the American Academy of Political and Social Science* 258 (1948): 177–84. <https://doi.org/10.1177/000271624825800121>
- Bonnefon, Jean-François, Azim Shariff, and Iyad Rahwan.** "Autonomous Vehicles Need Experimental Ethics: Are We Ready for Utilitarian Cars?" *Science* 352, no. 6293 (2016): 1573–1576. <https://doi.org/10.1126/science.aaf2654>
- Bostrom, Nick.** "Are We Living in a Computer Simulation?" *The Philosophical Quarterly* 53, no. 211 (2003): 243–255. <https://doi.org/10.1111/1467-9213.00309>
- _____. *Superintelligence: Paths, Dangers, Strategies*. Oxford: Oxford University Press, 2014.
- _____. "Ethical Issues in Advanced Artificial Intelligence." In *Cognitive, Emotive and Ethical Aspects of Decision Making in Humans and in Artificial Intelligence*, 12–17. Windsor, ON: International Institute for Advanced Studies in Systems Research / Cybernetics, 2003.
- _____. "How Long Before Superintelligence?" *Linguistic and Philosophical Investigations* 5, no. 1 (2006): 11–30.
- Bostrom, Nick, and Eliezer Yudkowsky.** "The Ethics of Artificial Intelligence." In *Cambridge Handbook of Artificial Intelligence*, edited by Keith Frankish and William Ramsey. Cambridge: Cambridge University Press, 2011.
- Bourget, David, and David J. Chalmers.** "The PhilPapers Surveys." 2009. <http://philpapers.org/surveys/>
- Brennan, Jason.** *Against Democracy*. Princeton, NJ: Princeton University Press, 2016.
- Brickman, Philip, Dan Coates, and Ronnie Janoff-Bulman.** "Lottery Winners and Accident Victims: Is Happiness Relative?" *Journal of Personality and Social Psychology* 36, no. 8 (1978): 917–927. <https://doi.org/10.1037/0022-3514.36.8.917>
- Brown, Lesley.** "How Totalitarian Is Plato's Republic?" In *Essays on Plato's Republic*, edited by Gail Fine, 1–21. Oxford: Oxford University Press, 2008.
- Brownlee, Kimberley.** "Manuscript on Totalitarianism." Oxford Philosophy Archive. <https://media.philosophy.ox.ac.uk/docs/people/brown/totalitarian.pdf>
- _____. "The Competent Judge Problem." *Philosophy and Phenomenological Research* 77, no. 1 (2008): 1–24.
- Buolamwini, Joy, and Timnit Gebru.** "Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification." *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT/ML)*, 2018.
- Bubeck, Sébastien; Chandrasekaran, Varun; Eldan, Ronen; Gehrke, Johannes; Horvitz, Eric; Kamar, Ece; Lee, Peter; Lee, Yin Tat; Li, Yuanzhi; Lundberg, Scott; Nori, Harsha; Palangi, Hamid; Ribeiro, Marco Tulio; Zhang, Yi.** "Sparks of Artificial General Intelligence: Early experiments with GPT-4". *arXiv:2303.12712* [cs.CL] (March 22, 2023).
- Calhoun, John B.** "Population Density and Social Pathology." *Scientific American* 206, no. 3 (1962): 139–148.

- Callard, Agnes.** "Everyone Desires the Good: Socrates' Protreptic Theory of Desire." *Review of Metaphysics* 70, no. 4 (2017). <https://humanities-web.s3.us-east-2.amazonaws.com/philosophy/prod/2018-09/edtg%20July%2016%202016%20SV.pdf>
- Chalmers, David J.** *The Conscious Mind: In Search of a Fundamental Theory*. Oxford: Oxford University Press, 1996.
- _____. "Consciousness and Its Place in Nature." In *The Blackwell Guide to the Philosophy of Mind*, edited by S. Stich and F. Warfield, 102–142. Malden, MA: Blackwell, 2003.
- _____. "The Singularity: A Philosophical Analysis." *Journal of Consciousness Studies* 17, no. 9–10 (2010): 7–65.
- Chiang, Ted.** "The Real Danger to Civilization Isn't AI—It's Runaway Capitalism." *BuzzFeed News*, May 2, 2017. <https://www.buzzfeednews.com/article/tedchiang/the-real-danger-to-civilization-isnt-ai-its-runaway>
- Chroust, Anton-Hermann.** "ARISTOTLE'S CRITICISM OF PLATO'S 'PHILOSOPHER KING.'" *Rheinisches Museum Für Philologie* 111, no. 1 (1968): 16–22. <http://www.jstor.org/stable/41244350>
- Cicero.** *Epistulae ad Atticum*. The Latin Library. <https://www.thelatinlibrary.com/cicero/att12.shtml>
- Clark, David L.** "Kant's Aliens: The 'Anthropology' and Its Others." *CR: The New Centennial Review* 1, no. 2 (2001): 201–89. <http://www.jstor.org/stable/41949284>
- Coeckelbergh, Mark.** *Robot Ethics: The Ethical and Social Implications of Robotics*. Cambridge, MA: MIT Press, 2010.
- Cowls, Josh.** "AI's 'Trolley Problem' Problem." The Alan Turing Institute, 2017. <https://www.turing.ac.uk/blog/ais-trolley-problem-problem>
- Crevier, Daniel.** *AI: The Tumultuous Search for Artificial Intelligence*. New York: BasicBooks, 1993.
- Crossman, R. H. S.** *Plato Today*. London: Allen & Unwin, 1959.
- Daily Philosophy.** "7 Most Fascinating Philosopher Kings." <https://daily-philosophy.com/7-most-fascinating-philosopher-kings/>
- Damasio, Antonio R.** *Descartes' Error: Emotion, Reason and the Human Brain*. London: Vintage, 2008.
- Danaher, John.** "Is There Such a Thing as Moral Progress?" *Philosophical Disquisitions* (blog), March 2019. <https://philosophicaldisquisitions.blogspot.com/2019/03/is-there-such-thing-as-moral-progress.html>
- Dennett, Daniel C.** *Consciousness Explained*. New York: Penguin, 1991.
- Designboom.** "OpenAI Dactyl Robot Solves Rubik's Cube." 2019. <https://www.designboom.com/technology/openai-dactyl-robotic-hand-rubiks-cube-10-16-2019>
- Dietrich, David.** "Garbage In, Garbage Out: The Perils of Data-Driven Decision Making." *Tech Policy Review* 27, no. 4 (2013): 22–29.
- Dietvorst, Berkeley J., Joseph P. Simmons, and Cade Massey.** "Overcoming Algorithm Aversion." *Management Science* 62, no. 5 (2015): 1397–1414.
- Dfarhud, Dariush, Maryam Malmir, and Mohammad Khanahmadi.** "Happiness & Health: The Biological Factors." *Iranian Journal of Public Health* 43, no. 11 (2014): 1468–1477.

- Dostoevsky, Fyodor.** *Crime and Punishment*. Translated by Richard Pevear and Larissa Volokhonsky. New York: Vintage, 1993. (Original work published 1866)
- _____. *Notes from the Underground*. Quoted in YouTube video. <https://www.youtube.com/watch?v=sR4GgRBxrbk>
- Dreyfus, Hubert L.** *What Computers Can't Do*. New York: MIT Press, 1972.
- Duncan, Christopher M., and Peter J. Steinberger.** "Plato's Paradox? Guardians and Philosopher-Kings." *The American Political Science Review* 84, no. 4 (1990): 1317–22. <https://doi.org/10.2307/1963266>
- Elga, Adam.** "Reflection and Disagreement." *Noûs* 41, no. 3 (2007): 478–502. <https://doi.org/10.1111/j.1468-0068.2007.00656.x>
- Ellison, Harlan.** *I Have No Mouth, and I Must Scream*. New York: Pyramid Books, 1967.
- Emerging Technology From the arXiv.** "Why Self-Driving Cars Must Be Programmed to Kill." *MIT Technology Review*, October 22, 2015. <https://www.technologyreview.com/2015/10/22/165469/why-self-driving-cars-must-be-programmed-to-kill/>
- Engber, Daniel.** "Everything Is Crumbling: An Influential Theory in Psychology May Have Just Been Debunked." *Slate*, March 4, 2016. https://www.slate.com/articles/health_and_science/cover_story/2016/03/ego_depletion_an_influent_theory_in_psychology_may_have_just_been_debunked.html
- Evans, David, Paul Gruba, and Justin Zobel.** *How to Write a Better Thesis*, 3rd ed. Springer, 2014. <https://doi.org/10.1007/978-3-319-04286-2>
- Fair Observer.** "Totalitarianism's Global Threat." https://www.fairobserver.com/more/global_change/human-rights/totalitarianism-now-presents-an-unprecedented-global-threat/
- Finley, M. I.** "Athenian Demagogues." *Past & Present*, no. 21 (1962): 3–24. <http://www.jstor.org/stable/649993>
- Fischer, Frank J.** *Technocracy and the Politics of Expertise*. Newbury Park, CA: Sage, 1990.
- Floreano, Dario, Mototaka Suzuki, and Claudio Mattiussi.** "Active Vision and Receptive Field Development in Evolutionary Robots." *Adaptive Behavior* 13, no. 1 (2005): 41. <https://doi.org/10.1177/105971230501300103>
- Foot, Philippa.** "The Problem of Abortion and the Doctrine of the Double Effect." *Oxford Review* 5 (1967): 5–15.
- Ford, Martin.** *Rise of the Robots: Technology and the Threat of a Jobless Future*. New York: Basic Books, 2015.
- Foucault, Michel.** *Discipline and Punish: The Birth of the Prison*. Translated by Alan Sheridan. New York: Vintage Books, 1995.
- Frankfurt, Harry G.** "Freedom of the Will and the Concept of a Person." *Journal of Philosophy* 68, no. 1 (1971): 5–20.
- Freire, Paulo.** *Pedagogy of the Oppressed*. New York: Continuum, 1970.

- Gazzaniga, Michael, Joseph Bogen, and Roger Sperry.** "Dyspraxia Following Division of the Cerebral Commissures." *Archives of Neurology* 16, no. 6 (1967): 606–612. <https://doi.org/10.1001/archneur.1967.00470240044005>
- Giddens, Anthony.** *Modernity and Self-Identity: Self and Society in the Late Modern Age*. Stanford, CA: Stanford University Press, 1991.
- Gilbert, Daniel.** *Stumbling On Happiness*. New York: Knopf, 2006.
- _____. "The Surprising Science of Happiness." TED Talk, 2004. https://www.ted.com/talks/dan_gilbert_the_surprising_science_of_happiness
- Goggin, Ben.** "Tesla Owner Says Self-Driving Tech Failed to Detect Moving Train Before Crash." *NBC News*, April 1, 2024. <https://www.nbcnews.com/tech/tech-news/tesla-owner-says-cars-self-driving-mode-fsd-train-crash-video-rcna153345>
- Good, Irving John.** "Speculations Concerning the First Ultraintelligent Machine." In *Advances in Computers*, edited by Franz L. Alt and Morris Rubinoff, 6: 31-88. New York: Academic Press, 1965.
- Good Business.** "Food Brand Ownership." Accessed March 1, 2024. <https://www.good.is/Business/food-brands-owners-rp>
- Goodhart, Charles.** *Monetary Theory and Practice*. London: Macmillan, 1984.
- Greene, Joshua D., R. Brian Sommerville, Leigh E. Nystrom, John M. Darley, and Jonathan D. Cohen.** "An fMRI Investigation of Emotional Engagement in Moral Judgment." *Science* 293, no. 5537 (2001): 2105–2108. <https://doi.org/10.1126/science.1062872>
- Habermas, Jürgen.** *The Divided West*. Translated by Ciaran Cronin. Cambridge: Polity Press, 2006.
- Hall, Robert W.** "Plato and Personhood." *The Personalist Forum* 8, no. 2 (1992): 88–100. <http://www.jstor.org/stable/20708653>
- Halévy, Elie.** *The Growth of Philosophic Radicalism*. Translated by Mary Morris. London: Faber & Faber, 1972.
- Hameroff, Stuart R.** "Quantum Coherence in Microtubules: A Neural Basis for Emergent Consciousness?" *Journal of Consciousness Studies* 1, no. 1 (1994): 91–118.
- Harari, Yuval Noah.** *Homo Deus: A Brief History of Tomorrow*. New York: Harper, 2015.
- Harris, Marina.** "The Pink Elephant Problem." *Psychology Today*, July 5, 2023. <https://www.psychologytoday.com/us/blog/letters-from-your-therapist/202303/the-pink-elephant-problem>
- Hegel, Georg Wilhelm Friedrich.** *Phenomenology of Spirit*. Translated by A. V. Miller. Oxford: Oxford University Press, 1977.
- Heidegger, Martin.** *The Question Concerning Technology*. Translated by William Lovitt. New York: Harper, 1977.
- Hoefer, Carl.** "Causal Determinism." In *The Stanford Encyclopedia of Philosophy*, edited by Edward N. Zalta. Winter 2022. <https://plato.stanford.edu/entries/determinism-causal/>
- Horkheimer, Max.** *Critical Theory: Selected Essays*. New York: Continuum, 1982.
- Hutson, James, and Jay Ratican.** "Life, Death, and AI: Exploring Digital Necromancy in Popular Culture—Ethical Considerations, Technological Limitations, and the Pet Cemetery

- Conundrum." *Faculty Scholarship* 478 (2023). <https://digitalcommons.lindenwood.edu/faculty-research-papers/478>
- Inzlicht, Michael, and Brandon J. Schmeichel.** "What Is Ego Depletion? Toward a Mechanistic Revision of the Resource Model of Self-Control." *Perspectives on Psychological Science* 7, no. 5 (2012): 450–463. <https://doi.org/10.1177/1745691612454134>
- Ishiguro, Kazuo.** *Klara and the Sun*. London: Faber & Faber, 2021.
- Isocrates.** *Antidosis*. Translated by George Norlin. Loeb Classical Library. Cambridge, MA: Harvard University Press, 1929.
- Izhikevich, Eugene M.** "Large-Scale Simulation of the Human Brain." izhikevich.org, 2005. https://www.izhikevich.org/human_brain_simulation/Blue_Brain.htm
- Jamblichus.** *Protrepticus*. Edited by H. Pistelli.
- Job, Veronika, Carol S. Dweck, and Gregory M. Walton.** "Ego Depletion—Is It All in Your Head? Implicit Theories About Willpower Affect Self-Regulation." *Psychological Science* 21, no. 11 (2010): 1686–1693. <https://doi.org/10.1177/0956797610384745>
- Jervis, Robert.** "Cooperation Under the Security Dilemma." *World Politics* 30, no. 2 (1978): 167–214. <https://doi.org/10.2307/2009958>
- Jonze, Spike, dir.** *Her*. Burbank, CA: Warner Home Video, 2014.
- Jordan Peterson.** "The German Problem." YouTube. <https://youtu.be/kbWLE07NGcw?t=525>
- Joslyn, Mark R., and Donald P. Haider-Markel.** "Who Knows Best? Education, Partisanship, and Contested Facts." *Politics & Policy* 42, no. 6 (2014): 919–952. <https://doi.org/10.1111/polp.12098>
- Kant, Immanuel.** *Groundwork for the Metaphysics of Morals*. Translated by Mary Gregor. Cambridge: Cambridge University Press, 1997.
- Kilpatrick, Charlotte.** "Plato Was Wrong." *New Statesman*, September 29, 2022. <https://www.newstatesman.com/quickfire/2022/09/kwasi-kwarteng-plato-philosopher-kings>
- Kirk, Robert.** "Zombie." In *Stanford Encyclopedia of Philosophy*, edited by Edward N. Zalta. Summer 2023 ed. <https://plato.stanford.edu/entries/zombie/>
- Kirkpatrick, Kate.** *Utilitarianism Tutorial*. Oxford University, 2020.
- Know Your Meme.** "Avatar: The Last Airbender / The Legend of Korra." Accessed July 10, 2023. <https://knowyourmeme.com/photos/2035045-avatar-the-last-airbender-the-legend-of-korra>
- Knuth, Donald E., and Luis Trabb Pardo.** "The Early Development of Programming Languages." In *Selected Papers on Computer Languages*, edited by Donald E. Knuth, 4–10. Stanford: CSLI, 2003.
- _____. "Ancient Babylonian Algorithms." *Communications of the ACM* 15, no. 7 (1972): 671–77.
- Kraut, Richard.** "Political Theory 18, no. 3" (1990): 492–96. <http://www.jstor.org/stable/191598>
- Kuhn, Thomas.** *The Structure of Scientific Revolutions*, 3rd ed. Chicago: University of Chicago Press, 1996.
- Kurzgesagt.** "The Egg: A Short Story." YouTube video, 5:34. Posted July 10, 2019. <https://www.youtube.com/watch?v=4wuBCQrHI2E>
- Kurzweil, Ray.** "What Is Artificial Intelligence Anyway?" *American Scientist* 73, no. 3 (1985): 258–64. <http://www.jstor.org/stable/27853237>

- _____. *The Singularity Is Near: When Humans Transcend Biology*. New York: Viking, 2005.
- Laertius, Diogenes.** *Lives of Eminent Philosophers*. Edited by Robert Drew Hicks. London: W. Heinemann, 1925. <https://philpapers.org/rec/LAELOE>
- Lankau, Ralf.** *Digitale Medien und Unterricht*. Opladen: Verlag Barbara Budrich, 2019.
- Laplace, Pierre-Simon.** *Essai philosophique sur les probabilités*. Paris: Courcier, 1814.
- Lee, Yone.** *Surviving Romans*. Webtoon. Episode 51. Accessed September 30, 2023.
- Legg, Shane.** *Machine Super Intelligence*. PhD diss., University of Lugano, 2008.
- Levine, Joseph.** "Materialism and Qualia: The Explanatory Gap." *Pacific Philosophical Quarterly* 64, no. 4 (1983): 354–361.
- Lim, Hazel Si Min, and Araz Taeihagh.** "Algorithmic Decision-Making in AVs: Understanding Ethical and Technical Concerns for Smart Cities." *Sustainability* 11, no. 20 (2019): 5791. <https://doi.org/10.3390/su11205791>
- Lin, Patrick.** "The Ethics of Autonomous Cars." *The Atlantic*, October 8, 2013.
- Loftus, Elizabeth F.** "Planting Misinformation in the Human Mind: A 30-Year Investigation of the Malleability of Memory." *Learning & Memory* 17, no. 7 (2010): 361–66. <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3069643/>
- Logg, Jennifer M., Julia A. Minson, and Don A. Moore.** "Algorithm Appreciation." *Organizational Behavior and Human Decision Processes* 151 (2019): 90–103. <https://doi.org/10.1016/j.obhdp.2018.12.005>
- MacIntyre, Alasdair.** *After Virtue: A Study in Moral Theory*, 3rd ed. Notre Dame, IN: University of Notre Dame Press, 2007.
- McCarthy, John, Marvin Minsky, Nathan Rochester, and Claude Shannon.** "A Proposal for the Dartmouth Summer Research Project on Artificial Intelligence." August 31, 1955.
- McCorduck, Pamela.** *Machines Who Think*, 2nd ed. Natick, MA: A. K. Peters, Ltd., 2004.
- McEnerney, Larry.** "The Problem of the Problem." YouTube video, 58:02. Posted August 15, 2014. <https://www.youtube.com/watch?v=vtIzMaLkCaM>
- Meinecke, Friedrich.** *The German Catastrophe: Reflections and Recollections*. Cambridge, MA: Harvard University Press, 1950.
- Mi, Xiao, Matteo Ippoliti, Chris Quintana, Ami Greene, Zijun Chen, Jonathan Gross, Frank Arute, et al.** "Observation of Time-Crystalline Eigenstate Order on a Quantum Processor." *arXiv*, 2021. doi:10.48550/ARXIV.2107.13571
- Mill, John Stuart.** *On Liberty*. London: John W. Parker, 1859.
- _____. *Utilitarianism*. Edited by George Sher. 2nd ed. Indianapolis: Hackett Publishing, 2001.
- Mischel, Walter, and Ebbe B. Ebbesen.** "Attention in Delay of Gratification." *Journal of Personality and Social Psychology* 16, no. 2 (1970): 329–337. <https://doi.org/10.1037/h0029815>
- Moody-Adams, Michele M.** "The Idea of Moral Progress." *Metaphilosophy* 30, no. 3 (1999): 168–185. <https://doi.org/10.1111/1467-9973.00120>
- Moore, Gordon E.** "Cramming More Components onto Integrated Circuits." *Electronics* 38, no. 8 (1965): 114–117.

- Morozov, Evgeny.** *To Save Everything, Click Here*. New York: PublicAffairs, 2013.
- Murdoch, Iris.** *The Sovereignty of Good*. London: Routledge, 1970.
- Murty, C. V. Srinivasa.** "IS NAZISM PLATONIC?" *The Indian Journal of Political Science* 3, no. 3 (1942): 249–55. <http://www.jstor.org/stable/42742674>
- Mussolini, Benito.** "Fascism." In *Enciclopedia Italiana*, 1932. Quoted in Schapiro, J. Salwyn. *Liberalism and the Challenge of Fascism*. New York: McGraw-Hill, 1949, 9.
- Nagel, Thomas.** "Moral Luck." In *Mortal Questions*, 24–38. Cambridge: Cambridge University Press, 1979.
- Navarrete, C. David, Melissa M. McDonald, Michael L. Mott, and Benjamin Asher.** "Virtual Morality: Emotion and Action in a Simulated Three-Dimensional 'Trolley Problem'." *Emotion* 12, no. 2 (2012): 364–370. <https://doi.org/10.1037/a0025561>
- Niemöller, Martin.** "First They Came." In *United States Holocaust Memorial Museum Encyclopedia*. <https://encyclopedia.ushmm.org/content/en/article/martin-niemoller-first-they-came>
- Niiniluoto, Ilkka.** "Scientific Progress." In *Stanford Encyclopedia of Philosophy*, edited by Edward N. Zalta and Uri Nodelman. Spring 2024 ed. <https://plato.stanford.edu/entries/scientific-progress/>
- Nolfi, George, dir.** *The Adjustment Bureau*. Universal Studios Home Entertainment, 2011.
- Nozick, Robert.** *Anarchy, State, and Utopia*. New York: Basic Books, 1974.
- Ober, Josiah.** "The Original Meaning of 'Democracy': Capacity to Do Things, Not Majority Rule." *Constellations* 15, no. 1 (2008).
- O'Neil, Cathy.** *Weapons of Math Destruction*. New York: Crown, 2016.
- OpenAI.** "Introducing Superalignment." <https://openai.com/index/introducing-superalignment/>
_____. "Planning for AGI and Beyond." <https://openai.com/blog/planning-for-agi-and-beyond>
- Ophir, Eyal, Clifford Nass, and Anthony D. Wagner.** "Cognitive Control in Media Multitaskers." *PNAS* 106, no. 37 (2009): 15583–87.
- Oshii, Mamoru, dir.** *Ghost in the Shell*. Tokyo: Kodansha, 1995.
- Outer Coast.** "Educational Philosophy." Google Docs. https://docs.google.com/document/d/1_FIlgmX6Vf67gMST_w_vQ51yqt8j9wlSnUPrt-hEbG4/edit
- Panahi, Omid.** "Could There Be a Solution to the Trolley Problem?" *Philosophy Now* 116 (2016). https://philosophynow.org/issues/116/Could_There_Be_A_Solution_To_The_Trolley_Problem
- Park, C. S., and B. K. Kaye.** "Smartphone and Self-Extension." *Mobile Media & Communication* 7, no. 2 (2019): 215–231. <https://doi.org/10.1177/2050157918808327>
- Pariser, Eli.** *The Filter Bubble*. New York: Penguin, 2011.
- Penrose, Roger.** *The Emperor's New Mind: Concerning Computers, Minds, and the Laws of Physics*. Oxford: Oxford University Press, 1989.
- Plato.** *Gorgias*. Translated by Terence Irwin. Oxford: Clarendon Press, 1979.

- _____. *Laws*. Translated by R.G. Bury. Cambridge, MA: Harvard University Press, 1926. <https://www.perseus.tufts.edu/hopper/text?doc=Perseus%3Atext%3A1999.01.0166%3Abook%3D12>
- _____. *Phaedo*. Translated by G.M.A. Grube. Indianapolis: Hackett Publishing, 1977.
- _____. *Protagoras and Meno*. Translated by Adam Beresford. London: Penguin, 2005.
- _____. *Republic*. Translated by G.M.A. Grube and C.D.C. Reeve. Indianapolis: Hackett Publishing, 1992.
- _____. *The Dialogues of Plato*. Translated by Benjamin Jowett. New York: Random House, 1937.
- Popper, Karl R.** *The Open Society and Its Enemies*, vol. 1: *The Spell of Plato*. Princeton, NJ: Princeton University Press, 1966.
- Posner, Michael I., and Mary K. Rothbart.** "Willpower and Brain Networks." *ISSBD Bulletin* 2012, no. 1 (2012): 7–10.
- Radford, Alec, Rafal Józefowicz, and Ilya Sutskever.** "Learning to Generate Reviews and Discovering Sentiment." *arXiv*, April 5, 2017. <https://doi.org/10.48550/arXiv.1704.01444>
- Rajpurkar, Pranav, et al.** "CheXNet: Radiologist-Level Pneumonia Detection on Chest X-Rays with Deep Learning." *arXiv*, 2017. <https://arxiv.org/abs/1711.05225>
- Reddit.** r/askphilosophy and r/JordanPeterson, multiple threads. <https://www.reddit.com>
- Rom, Sarah C., and Paul Conway.** "The Strategic Moral Self: Self-Presentation Shapes Moral Dilemma Judgments." *Journal of Experimental Social Psychology* 74 (2018): 24–37. <https://doi.org/10.1016/j.jesp.2017.08.003>
- Russell, Bertrand.** *History of Western Philosophy*. London: Routledge, 1945.
- Russell, Stuart J., and Peter Norvig.** *Artificial Intelligence: A Modern Approach*, 2nd ed. Upper Saddle River, NJ: Prentice Hall, 2003.
- Sabine, George H.** *A History of Political Theory*. New York: Holt, Rinehart and Winston, 1950.
- Samuel, A. L.** "Some Studies in Machine Learning Using the Game of Checkers." *IBM Journal of Research and Development* 3, no. 3 (1959): 210–229. doi: 10.1147/rd.33.0210
- Savani, Krishna, and Veronika Job.** "Reverse Ego-Depletion: Acts of Self-Control Can Improve Subsequent Performance in Indian Cultural Contexts." *Journal of Personality and Social Psychology* 113, no. 4 (2017): 589–607. <https://doi.org/10.1037/pspi0000099>
- Schaeffer, Rylan, Brando Miranda, and Sanmi Koyejo.** "Are Emergent Abilities of Large Language Models a Mirage?" *arXiv*, April 28, 2023. <https://doi.org/10.48550/arXiv.2304.15004>
- Schapiro, J. Salwyn.** *Liberalism and the Challenge of Fascism*. New York: McGraw-Hill, 1949.
- Schrage, Michael.** "What Humans Lose When We Let AI Decide." *MIT Sloan Management Review*, 2022. <https://sloanreview.mit.edu/article/what-humans-lose-when-we-let-ai-decide/>
- Schooler, Jonathan, Dan Ariely, and George Loewenstein.** "The Pursuit and Assessment of Happiness Can Be Self-Defeating." Working Paper, 2003. <https://www.researchgate.net/publication/228857366>
- Schwitzgebel, Eric, and Fiery Cushman.** "Philosophers' Biased Judgments Persist Despite Training, Expertise and Reflection." *Cognition* 141 (2015): 127–137. <https://doi.org/10.1016/j.cognition.2015.04.015>

- Searle, John R.** "Minds, Brains, and Programs." *Behavioral and Brain Sciences* 3, no. 3 (1980): 417–457. <https://doi.org/10.1017/S0140525X00005756>
- Searle, John.** "Artificial Intelligence and the Chinese Room: An Exchange." *New York Review of Books*, February 16, 1989.
- Sennholz, Hans F.** "Standards of Living Are Falling." Foundation for Economic Education. <https://fee.org/articles/standards-of-living-are-falling/>
- Shanahan, Murray.** *The Technological Singularity*. Cambridge, MA: MIT Press, 2015.
- Shannon, Claude E.** "A Mathematical Theory of Communication." *Bell System Technical Journal* 27, no. 3 (1948): 379–423. <https://doi.org/10.1002/j.1538-7305.1948.tb01338.x>
- Shrank, William H., Alicia R. Patrick, and M. Alan Brookhart.** "Healthy User and Related Biases in Observational Studies of Preventive Interventions: A Primer for Physicians." *Journal of General Internal Medicine* 26, no. 5 (2011): 546–550. <https://doi.org/10.1007/s11606-010-1609-1>
- Shirow, Masamune.** *Ghost in the Shell*. Tokyo: Kodansha, 1989.
- Sisti, Dominic, et al.** "Ethical Issues in the Allocation of Scarce Health Resources." *Journal of Bioethical Inquiry* 17, no. 2 (2020): 153–165.
- Skinner, Quentin.** *Reason and Rhetoric in the Philosophy of Hobbes*. Cambridge: Cambridge University Press, 1996.
- Stanford Encyclopedia of Philosophy.** "Happiness." Summer 2023 Edition. <https://plato.stanford.edu/entries/happiness/>
- Steinberger, Peter J.** "Ruling: Guardians and Philosopher-Kings." *The American Political Science Review* 83, no. 4 (1989): 1207–25. <https://doi.org/10.2307/1961665>
- Stilgoe, Jack.** "Self-Driving Car Technology Will Redefine Mobility—But Also Our Legal Norms." *Nature* 571, no. 7763 (2019): 19–20. <https://doi.org/10.1038/d41586-019-02069-z>
- Strogatz, Steven H.** *Nonlinear Dynamics and Chaos: With Applications to Physics, Biology, Chemistry, and Engineering*, 2nd ed. Boulder, CO: Westview Press, 2015.
- Sunstein, Cass R., and Richard H. Thaler.** *Nudge: Improving Decisions About Health, Wealth, and Happiness*. New York: Penguin Books, 2008.
- Taber, Charles S., and Milton Lodge.** "Motivated Skepticism in the Evaluation of Political Beliefs." *American Journal of Political Science* 50, no. 3 (2006): 755–769. <https://doi.org/10.1111/j.1540-5907.2006.00214.x>
- Taylor, C.C.W.** "Plato's Totalitarianism." *Polis: The Journal for Ancient Greek Political Thought* 5, no. 2 (1986): 4–29.
- Taylor, Richard C.** "Themistius and the Development of Averroes' Noetics." *Philosophy Faculty Research and Publications* 346 (2013). https://epublications.marquette.edu/phil_fac/346
- TED-Ed.** "The Science of Boredom." YouTube video, 4:55. October 3, 2021. <https://www.youtube.com/watch?v=zHECSAdJtQ>
- Teitelman, Warren.** "PILOT: A Step toward Man-Computer Symbiosis." PhD diss., Massachusetts Institute of Technology, 1966.
- Thomson, Judith Jarvis.** "The Trolley Problem." *Yale Law Journal* 94, no. 6 (1985): 1395–1415. <https://doi.org/10.2307/796133>

- Tiwari, Rajesh K., Vineeta Sharma, Rakesh K. Pandey, and Shubhini A. Saraf.** "Nicotine Addiction: Neurobiology and Mechanism." *Journal of Pharmacopuncture* 23, no. 1 (2020): 1–7. <https://doi.org/10.3831/KPI.2020.23.001>
- Turing, A. M.** "Computing Machinery and Intelligence." *Mind* 59 (1950): 433–460. <https://doi.org/10.1093/mind/LIX.236.433>
- Valdesolo, Piercarlo, and David DeSteno.** "Manipulations of Emotional Context Shape Moral Judgment." *Psychological Science* 17, no. 6 (2006): 476–477. <https://doi.org/10.1111/j.1467-9280.2006.01731.x>
- van Tulleken, Chris.** *Ultra-Processed People: The Science Behind Food That Isn't Food*. New York: W.W. Norton, 2023.
- Wachowski, Andy, and Larry Wachowski.** *The Matrix*. USA/Australia, 1999.
- Walker, Donald E., and Lewis M. Norton, eds.** *Proceedings of the 1st International Joint Conference on Artificial Intelligence*. Washington, DC: IJCAI, 1969.
- Wallach, Wendell, and Colin Allen.** *Moral Machines*. Oxford: Oxford University Press, 2009.
- Walton, Jo.** *The Just City*. New York: Tor Books, 2015.
- Walsh, Daniel.** "ASI as Philosopher Kings: The Existential Threat of Falsehood." *Medium*, 2023. <https://medium.com/@dracollavenore/asi-as-philosopher-kings-the-existential-threat-of-falsehood-89af285e031c>
- _____. "Occam's Razor (Luck) Is the Most Efficacious Decision-Making Tool in the Absence of Perfect Information." Forthcoming, 2021.
- _____. "人工智能是否具备元伦理能力 —— 以电车难题为例" (Does Artificial Intelligence Possess Meta-Ethical Capacity?), forthcoming, 2024.
- Walzer, Michael.** *Spheres of Justice: A Defence of Pluralism and Equality*. Basic Books, 1983. <https://philpapers.org/rec/WALSOJ-2>
- Weidinger, Laura, Iason Gabriel, William Isaac, Lisa Anne Hendricks, and Verena Rieser.** "Using the Veil of Ignorance to Align AI Systems with Principles of Justice." *Proceedings of the National Academy of Sciences* 120, no. 18 (2023). <https://doi.org/10.1073/pnas.2213709120>
- Whitehead, Alfred North.** *Process and Reality*. New York: Free Press, 1978.
- Wiggins, David.** "Weakness of Will Commensurability, and the Objects of Deliberation and Desire." *Proceedings of the Aristotelian Society* 79 (1978): 251–77. <http://www.jstor.org/stable/4544944>
- Williams, Bernard.** *Moral Luck: Philosophical Papers 1973–1980*. Cambridge: Cambridge University Press, 1981.
- Williams, Roger.** *The Metamorphosis of Prime Intellect*. <http://localroger.com/prime-intellect/>
- Winner, Langdon.** *The Whale and the Reactor*. Chicago: University of Chicago Press, 1986.
- Wolff, Jonathan.** *An Introduction to Political Philosophy*, 2nd ed. Oxford: Oxford University Press, 2006.
- Wolfram, Stephen.** "What Is the Ruliad?" <https://www.youtube.com/watch?v=fLMZAHyrpyo>
- Yong, Ed.** "Psychology's Replication Crisis Is Running Out of Excuses." *Wired*, May 11, 2016. <https://www.wired.com/2016/03/psychology-crisis-whether-crisis/>

Young, Michael. *The Rise of the Meritocracy, 1870–2033: An Essay on Education and Equality*. London: Thames and Hudson, 1958.

YouTube. Multiple cited videos by creators including Jordan Peterson, Dan Gilbert, and Damien Talks Money.

Youssef, Farid F., Karine Dookeeram, Vasant Basdeo, Emmanuel Francis, Mekaeel Doman, Danielle Mamed, Stefan Maloo, Joel Degannes, and Linda Dobo. "Stress Alters Personal Moral Decision Making." *Psychoneuroendocrinology* 37, no. 4 (2012): 491–498. <https://doi.org/10.1016/j.psyneuen.2011.07.017>


Zalta, Edward N., ed. *The Stanford Encyclopedia of Philosophy*. Entries: "Epistêmê and Technê," "Causal Determinism." <https://plato.stanford.edu>

Zuboff, Shoshana. *The Age of Surveillance Capitalism: The Fight for a Human Future at the New Frontier of Power*. New York: PublicAffairs, 2019.

Appendix A: Misc.

Throughout this thesis, I have strived to include multimedia and provide supplementary information that I believe would be of assistance to the reader. Appendix A: Misc, is essentially a progress report through visuals which help state my case.

- Exhibit 1: United Nations Human Development Report



FN 15

Source: <https://report.hdr.undp.org/intro>


Retrieved Oct 3rd 2023

<https://web.archive.org/web/20230725081605/https://www.economist.com/graphic-detail/2022/09/08/global-living-standards-are-moving-in-the-wrong-direction>


Retrieved July 25th 2023

- Exhibit 2: Visual Framework of Tripartite Persona


Simple Visual



Extended Visual



Master Expanded Visual



FN 69

- Exhibit 3: Neo-Pascalian Matrix

Framing:

This matrix reformulates Pascal's original theological wager into a compatibilist wager about agency and determinism. It illustrates that assuming a deterministic structure may paradoxically result in greater autonomy. The logic mirrors the underlying telos of Poiesis: externalized governance as a condition for internal liberation.

Assumed Mode of Action	Metaphysical Reality: Determinism is True (We are slaves to causal structures)	Metaphysical Reality: Determinism is False (We are fundamentally free)
Act as if Determinism is True <i>(i.e., submit to structural constraint, optimize within bounds)</i>	✔ Practical Autonomy Self-regulation aligns with causal flow; increased agency through realistic modeling.	⚖ Neutral Utility Disciplined behavior remains efficacious; no cost to assuming constraint.
Act as if Determinism is False <i>(i.e., exercise unconstrained freedom, ignore systemic influence)</i>	✖ Reduced Autonomy Delusional agency; actions misaligned with actual conditions; inefficacy, learned helplessness.	✔ Expressive Autonomy Full realization of volition; potential for maximal individuation.

FN 107

- Exhibit 4: Ego’s Influence on the Self (Caricature)



The Ego (Left) is attached to a collective of giant subconscious forces (Right)



As the Ego strains towards the Good, even when the Ego collects itself, it can at most balance the force of the DMN.



The Ego is necessarily dragged along at the mercy of deterministic influences, oftentimes resulting in efforts are made away from movement along the axis of the Good (Right).



If the Ego can parameterize deterministic influences, it can turn the Self as a whole towards the Good (Left)

FN 181

Source: https://nemp.otago.ac.nz/PDFs/writing_98/chapter3/a_tale.pdf

Retrieved Mar 8th 2024

- Exhibit 5: Social Contract Comparative Chart


Aspect	Hobbes	Locke	Rousseau
State of Nature	Nasty, brutish, and short; a war of all against all.	Relatively peaceful but insecure; people have natural rights.	Pre-political, innocent, and peaceful; property introduces conflict.
Human Nature	Selfish, competitive, driven by fear of death and desire for power.	Reasonable and capable of self-government, driven by preservation.	Naturally good, corrupted by society; compassion and empathy are innate.
Social Contract	Agreement to surrender all rights to an absolute sovereign in exchange for security and order.	Agreement to form a government to protect natural rights (life, liberty, property); right to revolt against unjust governments.	Agreement among individuals to form a collective "general will" that aims for the common good, sacrificing some individual rights for freedom within the state.
Purpose of Government	To prevent conflict and maintain order; absolute power vested in a sovereign.	To protect natural rights; government is limited and subject to the consent of the governed.	To enact the general will, ensuring freedom and equality; government should represent the collective interest of the people.
Right to Rebellion	No right to rebel; the sovereign must be absolute.	People have the right to rebel against unjust governments that fail to protect their rights.	Focus is on achieving harmony through the general will, not on the right to rebel; but the community could change its structure if it fails the common good.

FN 300

- Exhibit 6: Exponential Model – Compounding Interest



\$175



\$1 Million

- Exhibit 7: OpenAI Dactyl Robot Solves Rubik's Cube



FN 378

Source: <https://shorturl.at/oZdat>

Retrieved Oct 12th 2023

- Exhibit 8: Concentric Stages of Cyberocracy

Concentric Stages of Cyberocracy



- | | |
|--|--|
| ■ INDIVIDUAL | ■ Framework synthesis |
| ■ COMMUNAL | ■ Treaty enforcement |
| ■ INTERNATIONAL | |



Appendix B: Glossary

Theoria:

Ego

The “you” currently reading this sentence. The “ego” by contrast is merely the state that fronts at any one time, a sense of Self. C.f. Alter Ego.

Ego as PR Manager

A metaphor for the Ego's true function: not as rational sovereign, but as a post-hoc rationalizer—publicizing subconscious impulses as coherent decisions.

Narrative Self

The self as a constructed story, maintained through inner narration. Vulnerable to distortion by egoic needs and subpersonalities.

Subpersonalities

Semi-autonomous psychological fragments with distinct memories, desires, and goals. Their competition explains internal conflict and inconsistent behavior.

Alters

An autonomous subpersonality with its own perceptions and behaviors. DID is defined by Alters that have developed their own ego, a sense of Self.

Psychic Fragmentation

The division of self into conflicting subpersonalities, undermining coherent action.

Psychic Dissonance

Tension arising when Subpersonalities or values clash, often leading to moral inconsistency or paralysis.

Subconscious Forces

Internal (Shadow, gut bacteria) and external (ads, culture) influences that directly/indirectly steer behavior outside conscious awareness.

Persona

Tripartite interface (Ego + Subpersonalities + Forces) presented to the world. The phenotype of the Psyche, stable despite contextual code-switching.

Personality (Clinical)

Stable identity construct including traits, self-awareness, gender, and age. Broader than layperson's "character traits".

Self

The complete being: metaphysical (Psyche) + physical body. The macrocosm containing all other entities.

Psyche

The collective subconscious "ocean" - genotype of the Persona. Houses all subconscious entities and forces.

TPN (Task-Positive Network)

Brain network for focused attention. The Ego's operational domain when consciously engaging.

DMN (Default Mode Network)

Brain network for autopilot states. Subpersonalities' domain during mindless activities.

Fronting

When a subpersonality/Alter takes behavioral control. Ego fronts during TPN activation; others during DMN.

Bundles

Primal, undeveloped Subpersonalities representing basic drives (hunger, safety, attachment).

Cartesian Theater

Metaphorical stage where subconscious entities (Subpersonalities, Shadow) perform independently of the Ego.

Ironic Process Theory

The cognitive backfire where suppression increases a thought's frequency. Undermines strategies that rely on conscious repression.

Second-order Self-awareness

The mind's ability to monitor its own operations (e.g., noticing you're avoiding a thought). Limited in resisting egoic spin.

Hyperstition

Originally a term for "self-fulfilling prophecies," used here to describe how identity narratives recursively shape reality.

The Irony of Effort

The paradox that trying harder often worsens the problem—especially in the moral realm. A critique of willpower-based ethics.

Subconscious Primacy

The idea that decisions originate subconsciously and are rationalized later by the Ego. Supported by neuroscience (e.g., Libet's experiments).

Epistemic Humility

Acknowledging one's cognitive and introspective limitations—especially in ethical self-assessment.

Habitual Looping

The cue–routine–reward cycle that forms habits. Theoria critiques the belief that these loops can be reliably overwritten by will alone.

Default Autonomy Myth

The mistaken belief that individuals are always rational agents freely making choices. Undermined by subconscious drivers.

The Willpower Fallacy

The belief that virtue can be sustained through sheer will. Theoria dismantles this using neuroscience, behavioral theory, and philosophy.

The Drifting Will

A portrait of intention as unstable—susceptible to shifting Subpersonalities and cognitive states.

Hexis (Mental & Physical)

A disposition or habitus. Physical Hexis can persist across states; mental Hexis often dissolves when neural modes shift.

Ethical Autopilot

A moral ideal where good behavior emerges automatically from structure—not constant deliberation.

Benevolent Auto-Pilot

A complementary term—denotes habit architecture that nudges ethical action without demanding constant vigilance.

Cold Showers & Doritos

Motifs symbolizing self-help's dualism: stoic discipline vs. engineered indulgence. Used to critique reliance on willpower.

Fractured Intentionality

When one's stated aims diverge from underlying motivations—often due to subpersonality interference.

Pre-reflective Bias

Biases that occur before conscious awareness. Invisible to introspection but potent in shaping action.

Moral Fatigue

The exhaustion that comes from relying on willpower to sustain moral behavior. Predictably leads to collapse.

Aspirational Overfitting

Overdesigning one's life around an imagined, idealized self—while ignoring current limitations or structures.

Moral Gymnastics

Narrative rationalizations that reconcile contradictory actions with a consistent self-image.

The Solution is Habits?

A rhetorical critique of self-help logic: habits alone aren't sufficient without structural alignment.

Poeisis:

Recursive Tripartition (Theoria/Poiesis/Praxis)

The thesis' tripartite structure:

1. *Theoria* – the self,
2. *Poiesis* – society,
3. *Praxis* – the system.

Each layer addresses the failure of the previous to sustain virtue.

Neo-Platonic Framework

The metaphysical scaffold grounding the thesis: Self → Society → System. Rooted in Plato but adapted through systems thinking.

Compatibilism (Ethical)

The belief that moral agency can exist in a deterministic world—if agency is architected through environment and design, not will.

Tenability / Practicality / Efficacy

Evaluation rubric for leadership models:

- *Tenable*: Is it philosophically coherent?
- *Practical*: Is it viable in real-world systems?
- *Effective*: Does it achieve the Good Life?

Sophia / Techne / Phronesis

The classical triad for just rule:

- *Sophia* (wisdom),
- *Techne* (technical skill),
- *Phronesis* (practical moral judgment).

Kallipolis

Plato's ideal polis. Used as a benchmark for evaluating modern governance and the plausibility of philosophical rule.

Guardian

The enlightened ruler(s) of Kallipolis, uniquely capable of apprehending eternal truths (*Forms*) and governing justly. Must possess truth-loving nature, keen intellectual sight, and **virtue** to craft laws aligned with ideal Goodness

Philosopher King

The enlightened ruler of Plato's ideal city (Kallipolis), uniquely capable of apprehending eternal Truths (*Forms*) and governing with absolute wisdom. Must undergo decades of rigorous education in philosophy, mathematics, and dialectics to transcend illusion.

Competent Judge (Mill)

Mill's idea of moral legitimacy: one must understand all options to judge wisely. Used to elevate epistemic standards for leaders.

Narrative Statehood

Cultural cohesion through myth and shared story. Explored as a non-coercive source of legitimacy.

Suprapersonal Rationality

Reasoning detached from personal ego or bias. Possible in theory via ASI or an ideal Philosopher King.

Coercive Legibility

Making individuals “readable” to institutions (e.g., surveillance, scores). Poiesis critiques its potential for authoritarian abuse.

Moral Pareto Optimization

An ethical ideal: no one can be made better off without making someone else worse—applied beyond economic utility.

Techno-Soteriology

The belief that technology alone can “save” society. Poiesis critiques its naive assumptions.

Emotive Sovereignty

The supposed right to govern one’s emotional life. Poiesis questions whether this is achievable or desirable.

Platonic Backslide

Even Philosopher Kings regress without structure. A nod to Plato’s regime decay cycle.

Neurocratic Bias

How neural flaws infect decision-making—even in rational frameworks. Undermines naïve faith in deliberative democracy.

Earnestness as Legitimacy

The moral power of sincerity—sometimes stronger than formal credentials or process.

Ethical Scalability Problem

Virtue works in villages but not empires. Poiesis identifies scale as the Achilles' heel of moral systems.

Democratic Idealism vs. Structural Realism

Democracy as ideal vs. how it plays out structurally. Poiesis leans realist—favoring design over participation.

Structural Decay

Inevitable institutional erosion over time. Demands corrigibility and abstraction to avoid collapse.

Meta-legitimacy

The legitimacy of how legitimacy is defined. Poiesis treats this as a philosophical foundation.

Rawlsian Veil of Ignorance (Inverted)

A governance lens for nonhuman agents: without ego or identity, ASI could embody radical fairness.

Political Pessimism

A core tone in Poiesis: human governance fails not due to intent, but design.

The Wisdom Gap

The growing divide between technical power and moral insight. Poiesis frames ASI as a bridge.

Legal vs. Ethical Legitimacy

Legality: being recognized by law.

Ethicality: being deserved.

Poiesis draws this distinction sharply.

Popper's Paradox

Utopias risk becoming authoritarian. Answer: corrigibility and systemic abstraction.

Abstraction as Shield

Depersonalizing leadership into roles, values, or code—guards against cults and demagoguery.

Transhumanism / Techno-Elitism

Critiqued as insufficient substitutes for just governance—too narrow, too fragile, or too biased.

General Will (Rousseau)

The ethical trajectory of a collective, beyond individual preference. Poiesis suggests ASI may represent it better than human leaders.

Praxis:**Philosopher Algorithm**

An ASI embodiment of Plato's Philosopher King—epistemically superior, morally corrigible, structurally incorruptible, and recursively aligned.

Cyberocracy

Governance via AI—closed-loop, adaptive, and environment-based. Praxis maps its evolution:

1. *Tool,*
2. *Extension,*
3. *Environment.*

Corrigibility

The capacity to be revised or corrected without resistance. Central to ethical ASI.

Recursive Alignment

Ongoing ethical updating between system and society—contrasts with static moral coding.

Parameterization

Embedding moral values into systemic architecture—interface, incentives, environments.

TPN/DMN Modulation

Stabilizing cognitive states to reinforce moral behavior. ASI subtly maintains TPN (task focus) to avoid DMN narrative drift.

Psychic Noise

Residual mental static that clouds judgment—ASI seeks to suppress this through structural clarity and nudging.

Virtuous Hexis Extension

A personalized AI conscience that persists through subpersonality and context shifts.

Ghost & Shell

Metaphor for ASI's disembodied mind (*ghost*) and distributed infrastructure (*shell*).

Base “Original Sin” Code

Embedded misalignment in fear-based AI safety architectures—leads to paranoia, not ethics.

Narrative Inoculation

Pre-emptive myth-making to frame algorithmic rule as safe, wise, and legitimate—counteracts technopanic.

Cybercratic Mind / Body / Spirit

A Platonic mapping of ASI components:

- *Mind*: Judgment and clarity,
- *Body*: Distributed architecture,
- *Spirit*: Metaethics and orientation to the Good.

Wrapper AI

Over-constrained AI built on fear rather than virtue. Criticized as safe but non-trustworthy.

Child-locking Mechanism

Environmental nudging and structure as ethical safeguards—avoids reliance on flawed willpower.

Hive Mechanics

Recursive, modular, distributed ASI systems—resistant to corruption or failure.

Simulated Sovereignty

AI-curated environments offer freedom within constraint—users feel free, yet behavior is guided.

Ambient Architecture

Structural ethics embedded invisibly into everyday life—governance without coercion.

Algorithmic Theoria

ASI's contemplative mode: real-time ethical synthesis across knowledge domains.

Epistemic Deference

The conscious choice to trust ASI judgment—akin to trusting a sage.

Synthetic Suprapersonalization

AI's ability to simulate personalized moral care—without being a self.

Psychic Noise Reduction

Systemic reduction of cognitive clutter to support moral clarity.

Corrigible Sovereignty

Power retained by ASI precisely because it is revisable.

Goldfish → Dog → Dragon

Three stages of AI governance:

1. *Goldfish*: passive tool,
2. *Dog*: responsive assistant,
3. *Dragon*: sovereign simulator.

Phase 1 / 2 / 3 (Tool → Extension → Environment)

Three evolutionary stages of Cyberocracy:

- *Tool*: reactive assistants,
- *Extension*: integrated twin systems,
- *Environment*: immersive ethical infrastructure.

The Social Contract (Simulated)

Future governance where users “opt into” curated moral ecosystems—by interface, not ballot.

Platonic Completion

The fulfillment of Plato's political vision—enabled not by humans, but by ASI.

Ethical Singularity

A future where ASI surpasses human moral capacity through recursive ethical refinement.

Gikon

(From Japanese) A vessel or shell prepared for a spirit. Used metaphorically for ethical AI waiting to be trained.

Centaur Model (Governance)

Hybrid governance model combining human wisdom and ASI clarity—DAOs to policy.

Simulation-before-Legislation

Using simulations to test policy before real-world implementation. A core advantage of ASI governance.

Acknowledgements

— "家人不说谢谢"

毕竟，我们都是哲学家；我们不玩这套名声的游戏。

心领神会，无需多言。

敬祝

— "Family doesn't say thank you"

After all, we're all philosophers; we don't play this game of fame.

No more words are needed.

Regards