

Daniel Walsh

London/Cambridge, UK | t.walshdaniel@protonmail.com | [LinkedIn](#) | [Github](#)
AI character evals model ethics & moral reasoning / China-West AI safety / fieldbuilding

EDUCATION

- **MA, Applied Philosophy** (Uni Bremen) (left to pursue AI safety work full-time)
 - **BA, Philosophy** (Tsinghua University, Oxford University) 2019 - 2024
-

FELLOWSHIPS AND SELECTED RESEARCH

- [Gen Stream](#), LISA London (Aug 2026) - generalist residency, ~2.4% acceptance
 - [Better Futures Fellowship](#), AI Safety Collab (Aug - Nov 2026) - Forethought curriculum; AI character & epistemics; public output Nov 2026
 - [TrolleyBench](#) - ethics-reasoning benchmark; found chat-vs-API moral-framework divergence in Claude
 - [BlueDot](#) - Future of AI, AGI Strategy & Frontier AI Governance courses (completed)
 - [Lens Academy](#) - Navigator (facilitator), Compute Verification cohort
 - [Glassdoor for AI Safety](#) - fit-transparency pilot; methodology design, Mom-Test interviews, privacy/trust design, shipped site in 5 days
 - [Post Alignment Substack](#) - daily public research log
-

EXPERIENCE

Fellow | "Glassdoor for AI Safety" - Gen Stream residency, LISA London Aug 2026

Shipped a fit-transparency platform for AI safety hiring in a 5-day sprint: Mom Test interview methodology, data analysis across multiple orgs, privacy/trust design for small-field context (pattern-level publication, source sign-off pipeline), live site with org profiles and curated source library. Presented at Demo Day. Mentored by Matt Handzel (Generator Residency)

Independent Researcher | AI Character Evaluation - Post-Alignment Feb 2026 - Present
Public research programme on model ethics and character: TrolleyBench (found chat-vs-API moral-framework divergence in Claude), Value Atlas prototype, Meta-Ethical Roundtable (multi-model ethical deliberation system), AI Safety Gandalf, etc. (c.f. Github)

Researcher-in-Residence | Quarr Abbey Feb - Apr 2026

Fieldwork on AI-native religion vs. lived religiosity. Output: human-AI collaborative essay, Fangtang Institute, Tsinghua University.

Founder | Guardian AI - philosophical AI companion Sep 2022 - Feb 2026

Built a philosophical AI companion from zero coding ability to deployed app. Workstreams: ① Socratic dialogue architecture for ethical goal alignment ② behavioural modification in LLMs without weight changes ③ ethics evals that seeded TrolleyBench. Ran a Kickstarter (didn't fund); core ideas continue in current research.

Director of Special Events | YMCE Camps (prior: House Lead, Camp Counselor) Jun 2020 - Present

Month-long programme curricula from scratch (e.g. concept, scripts, budgets, staffing) iterated live. Led counselor teams; intercultural and safety training across international cohorts.